

Runtime Process Infection

Shawn “lattera” Webb
SoldierX

<https://www.soldierx.com/>



Who Am I?

- Just another blogger
 - 0xfeedface.org
- Professional Security Engineer
- Thirteen-year C89 programmer
- Member of SoldierX, BinRev, and Hack3r
- Twitter: @lattera

Disclaimers

- Opinions/views expressed here are mine, not my employer's
- Talk is semi-random
 - Tied together at the end
- Almost nothing new explained
 - Theory known
 - New technique
- Presentation and tools only for educational purposes

Assumptions

- FreeBSD? What's that?
- Basic knowledge of C and 64bit memory management
- Ability and desire to think abstractly
- Non-modified memory layout (NO grsec/pax)

History

- CGI/Web App vulnerabilities
 - Needed connect-back shellcode
 - Needed reliable, random access
 - Firewall holes are a problem
 - Needed way to reuse existing connection to web server
 - Needed to covertly sniff traffic
 - Libhijack is born (discussed later)

Setting the Stage

- Got a shell via CGI/Web App exploit
 - Reliable way to get back in
 - nginx good candidate (live demo against nginx later)
 - Already listening for connections
 - Modify nginx process somehow to run a shell when a special string is sent
 - i.e. GET /shell HTTP/1.1
 - ```
$ whoami
```

```
www
```
  - Need to hook certain functions in runtime

# Process Loading

- `execve` is called
- Kernel checks file existence, permissions, etc.
- Kernel loads RTLD (Runtime Linker (`ld.elf.so`))
- Kernel loads process meta-data, initializes stack
  - Meta-data loaded at `0x00400000` on FreeBSD/  
amd64

# Runtime Linker

- Loads process into memory
- Loads dependencies (shared objects)
  - Patches PLT/GOT for needed dynamic functions
- Calls initialization routines
- Finally calls main()



# ELF

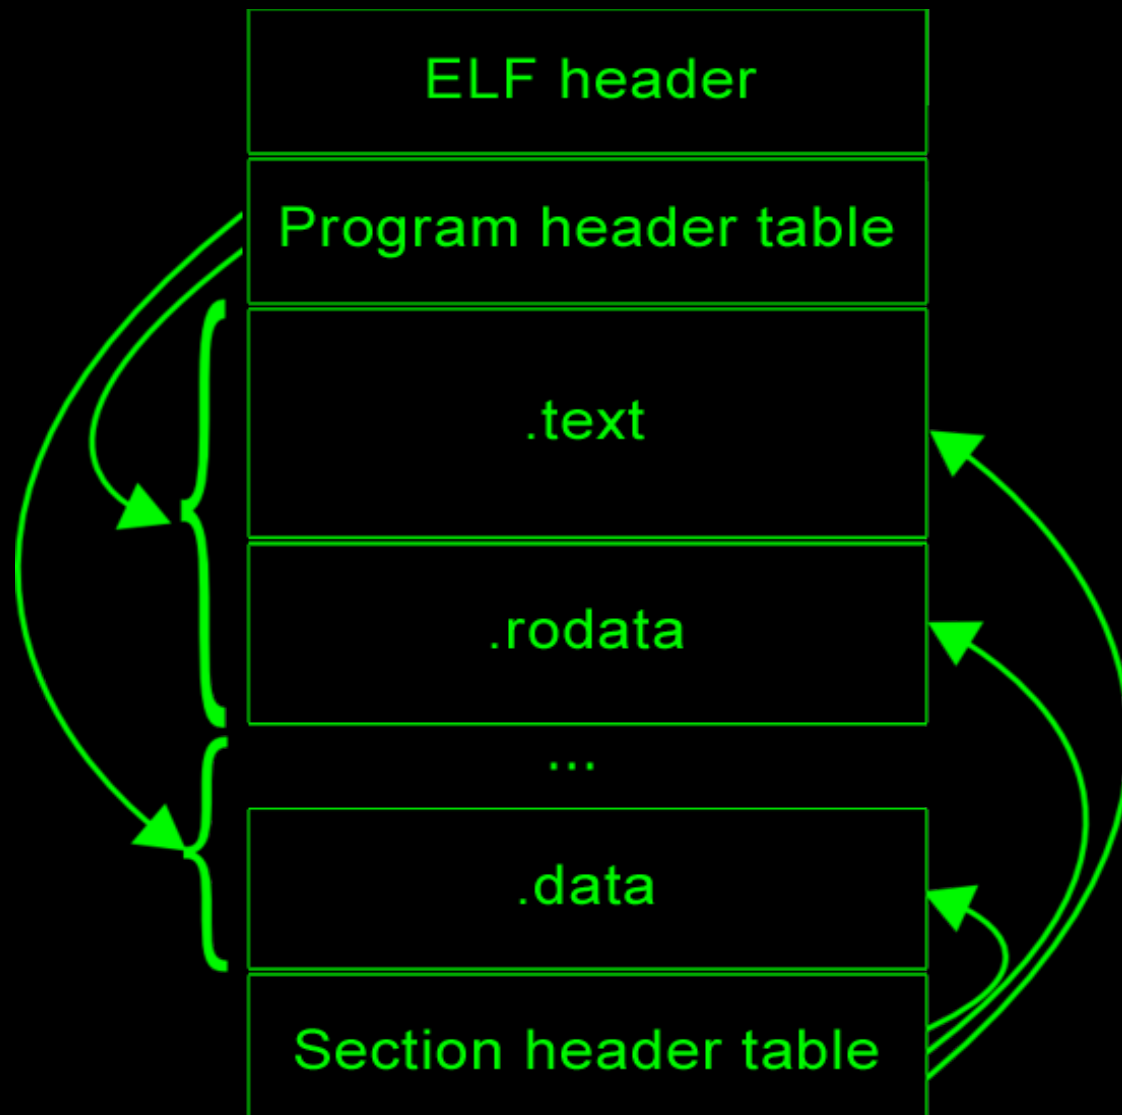
- Executable and Linkable Format
- Meta-data
- Tells RTLD what to load and how to load it



# ELF

- Describes where to load different parts of the object file
  - Process Header (PHDR) – Minimum one entry; contains virtual address locations, access rights (read, write, execute), alignment
  - Section Header (SHDR) – Minimum zero entries; describes the PHDRs; contains string table, debugging entries (if any), compiler comments
  - Dynamic Headers – Contains relocation entries, stubs, PLT/GOT (jackpot)

# ELF



Thanks bic!

# Debugging

- Ptrace – Oldschool debugging facility for FreeBSD
  - Kernel syscall
  - GDB relies on ptrace
  - Read/write from/to memory
  - Get/set registers
  - Debuggee becomes child of debugger
  - Destructive
    - Original ptrace engineer evil, likely knew it could be abused

# Current Techniques

- Store shellcode on the stack
  - Stack is non-executable
- Store shellcode at \$rip
  - Mucks up original code
- Store shellcode on the heap
  - Heap is non-executable
- LD\_PRELOAD?
  - Process has already started

# Allocating Memory

- We have arbitrary code to store. Where?
- Allocate memory in child
  - Unlike Windows and OSX, we cannot allocate from the parent process, the child must allocate
- Find “syscall” (\x0f\x05) opcode
- Program's main code won't call kernel
  - Calls library functions which call the kernel
    - Libc!
  - Find a library function that calls the kernel by crawling the ELF meta-data

# Allocating Memory - Finding “syscall”

- Loop through the ELF headers
  - Main ELF header contains pointer to PHDR
  - PHDR contains a pointer to the Dynamic headers
  - Dynamic headers has a pointer to the GOT
  - GOT[1] contains a pointer to the Obj\_Entry
  - Obj\_Entry is a structure created/maintained by RTLD and dlopen
  - Obj\_Entry points to each shared object's ELF headers
  - Loop through symbol table of each shared object

# Allocating Memory

- Parse ELF headers, loaded at 0x00400000
  - Headers include lists of loaded functions
- Found “syscall” in a shared object
- Back up registers
- Set \$rip to address of found “syscall” opcode
- Set up stack to call mmap syscall
- Continue execution until mmap finishes



# Injecting Shellcode

- After calling mmap
  - \$rax contains address of newly-allocated mapping
  - Can write to it
    - Even if mapping is marked non-writable (PROT\_READ | PROT\_EXECUTE)
  - Restore the backed-up registers
  - Push return address
    - Shellcode needs to know where to return to
    - Decrement \$rsp by sizeof(unsigned long)
    - Copy \$rip to \$rsp

# Injecting Shellcode

- Write shellcode to newly-allocated mapping
- Set \$rip to address of the shellcode
- Detach from the process
- Sit back, relax, and enjoy life
- But wait! There's more!



# Hijacking Functions

- Global Offset Table/Procedure Linkage Table
  - Array of function addresses
- All referenced functions are in GOT/PLT
- PLT/GOT redirection
  - `Shellcode["\x11\x11\x11\x11"] = @Function`
  - `GOT[#Function] = @Shellcode`

# Hijacking Functions

- Be careful
  - Multiple shared objects implement functions of the same name
    - Different signature
  - Make sure you target the correct function
  - Know your target
  - Set up a VM, mimicking the victim
    - Same OS, same patch levels, etc.
- Cannot reliably remove hijack

# Injecting Shared Objects

- Why?
  - Don't have to write a ton of shellcode
  - Write in C, use other libraries, possibilities are endless
- Two ways of doing it
  - The cheating way: Use a stub shellcode that calls `dlopen()`
  - The real way: rewrite `dlopen()`

# The Cheating Way

- Allocate a new memory mapping
- Store auxiliary data in mapping
  - .so path
  - Name of the function to hijack
  - Stub shellcode
- Stub shellcode will:
  - Call dlopen and dlsym
  - Replace GOT entry with entry found via dlsym

# The Cheating Way

- Advantages
  - Easy
  - Extendable
  - Fast
- Disadvantages
  - File-backed entry in `procstat -v`
  - Rely on stub shellcode

# The Real Way

- Reimplement dlopen in parent
  - Load dependencies (deps can be loaded via real dlopen)
  - Create memory maps
  - Write .so data to new memory maps
  - Patch into the RTLD
  - Run init routines
  - Hijack GOT



# The Real Way

- Advantages
  - Completely anonymous
  - Extensible
- Disadvantages
  - Takes time to research and implement

# Shared Objects

- Shared objects can have dependencies
- Shared objects have own PLT/GOT
  - Loop through Dynamic structures found in linkmap
  - Use same PLT/GOT redirection technique against shared objects
  - Even shared objects loaded via dlopen

# Libhijack

- Libhijack makes injection of arbitrary code and hijack of dynamically-loaded functions easy
  - Shared objects via the cheating method
  - Inject shellcode in as little as eight lines of C code
  - Full 32bit and 64bit Linux support
  - Full support for FreeBSD/amd64
  - Interest in porting to OSX and Android
- Always looking for help
- <https://github.com/lattera/libhijack>

# Libhijack 0.6.1

- Bug fixes
- Ability to look up non-exported RTLD functions on FreeBSD
  - RTLD's Obj\_Entry object isn't easily fetched
    - FreeBSD's developers take security seriously
  - Utilize the link\_map doubly-linked list in the Obj\_Entry struct

# Libhijack TODO

- Version 0.7
  - Inject shared objects via “The Real Way”
  - Looking for an adventure? Port to Android
- Always looking for help

# Prevention

- RELRO+BIND\_NOW = nope
- Make sure PLT/GOT entries point to correct lib
  - How? Symbol table resolution?
- Use dtrace, disable ptrace
  - From Solaris
  - Non-destructive debugging
- Limit ptrace usage (www user shouldn't ever use it)
  - Capsicum, anyone?

# Prevention

- Static binaries
  - Major disk and memory usage
- Grsec/PaX
  - Only protects to a certain extent
  - Doesn't exist on FreeBSD. Surprise!
- Static and dynamic profiling
  - Watch for changes in GOT
  - Make sure changes reflect static profile
  - What about shared objects loaded via `dlopen()`?

# Demo

Assembly loading .so



```
exit(0);
```

Comments/questions  
Thanks