

Quick introduction to reverse engineering for beginners

Dennis Yurichev
<dennis@yurichev.com>



©2013, Dennis Yurichev.

This work is licensed under the Creative Commons Attribution-NonCommercial-NoDerivs 3.0 Unported License. To view a copy of this license, visit

<http://creativecommons.org/licenses/by-nc-nd/3.0/>.

Text version (August 15, 2013).

Probably, newer version of this text, and also Russian language version is also accessible at

<http://yurichev.com/RE-book.html>

You may also subscribe to my twitter, to get information about updates of this text, etc: [@yurichev](#)

Contents

Preface	v
About author	vi
Thanks	vii
1 Compiler's patterns	1
1.1 Hello, world!	2
1.1.1 x86	2
1.1.2 ARM	4
1.2 Stack	8
1.2.1 Save return address where function should return control after execution	8
1.2.2 Function arguments passing	9
1.2.3 Local variable storage	9
1.2.4 (Windows) SEH	11
1.2.5 Buffer overflow protection	11
1.3 printf() with several arguments	12
1.3.1 x86	12
1.3.2 ARM: 3 printf() arguments	13
1.3.3 ARM: 8 printf() arguments	14
1.3.4 By the way	17
1.4 scanf()	18
1.4.1 About pointers	18
1.4.2 x86	18
1.4.3 ARM	20
1.4.4 Global variables	20
1.4.5 scanf() result checking	23
1.5 Passing arguments via stack	25
1.5.1 x86	25
1.5.2 ARM	26
1.6 One more word about results returning.	28
1.7 Pointers	29
1.7.1 C++ references	30
1.8 Conditional jumps	31
1.8.1 x86	31
1.8.2 ARM	33
1.9 switch()/case/default	35
1.9.1 Few number of cases	35
1.9.2 A lot of cases	38
1.10 Loops	44
1.10.1 x86	44
1.10.2 ARM	46
1.10.3 One more thing	48

1.11	strlen()	49
1.11.1	x86	49
1.11.2	ARM	51
1.12	Division by 9	54
1.12.1	ARM	55
1.13	Work with FPU	57
1.13.1	Simple example	57
1.13.2	Passing floating point number via arguments	60
1.13.3	Comparison example	62
1.14	Arrays	70
1.14.1	Simple example	70
1.14.2	Buffer overflow	73
1.14.3	Buffer overflow protection methods	76
1.14.4	One more word about arrays	79
1.14.5	Multidimensional arrays	79
1.15	Bit fields	82
1.15.1	Specific bit checking	82
1.15.2	Specific bit setting/clearing	85
1.15.3	Shifts	88
1.15.4	CRC32 calculation example	91
1.16	Structures	95
1.16.1	SYSTEMTIME example	95
1.16.2	Let's allocate place for structure using malloc()	97
1.16.3	struct tm	99
1.16.4	Fields packing in structure	104
1.16.5	Nested structures	106
1.16.6	Bit fields in structure	107
1.17	C++ classes	114
1.17.1	Simple example	114
1.17.2	Class inheritance in C++	119
1.17.3	Encapsulation in C++	122
1.17.4	Multiple inheritance in C++	124
1.17.5	C++ virtual methods	128
1.18	Unions	131
1.18.1	Pseudo-random number generator example	131
1.19	Pointers to functions	133
1.19.1	GCC	135
1.20	SIMD	137
1.20.1	Vectorization	137
1.20.2	SIMD strlen() implementation	143
1.21	64 bits	147
1.21.1	x86-64	147
1.21.2	ARM	153
1.22	C99 restrict	154
2	Couple things to add	157
2.1	LEA instruction	158
2.2	Function prologue and epilogue	159
2.3	npad	160
2.4	Signed number representations	162
2.4.1	Integer overflow	162
2.5	Arguments passing methods (calling conventions)	163
2.5.1	cdecl	163
2.5.2	stdcall	163

2.5.3	fastcall	163
2.5.4	thiscall	164
2.5.5	x86-64	164
2.5.6	Returning values of <i>float</i> and <i>double</i> type	164
2.6	position-independent code	165
3	Finding important/interesting stuff in the code	168
3.1	Communication with the outer world	168
3.2	String	169
3.3	Constants	169
3.3.1	Magic numbers	169
3.4	Finding the right instructions	170
3.5	Suspicious code patterns	171
3.6	Magic numbers usage while tracing	172
3.7	Old-school methods, nevertheless, interesting to know	172
3.7.1	Memory “snapshots” comparing	172
4	Tasks	173
4.1	Easy level	173
4.1.1	Task 1.1	173
4.1.2	Task 1.2	174
4.1.3	Task 1.3	176
4.1.4	Task 1.4	177
4.1.5	Task 1.5	179
4.1.6	Task 1.6	180
4.1.7	Task 1.7	181
4.1.8	Task 1.8	182
4.1.9	Task 1.9	182
4.1.10	Task 1.10	183
4.2	Middle level	183
4.2.1	Task 2.1	183
4.3	crackme / keygenme	190
5	Tools	191
5.0.1	Debugger	191
6	Books/blogs worth reading	192
6.1	Books	192
6.1.1	Windows	192
6.1.2	C/C++	192
6.1.3	x86 / x86-64	192
6.1.4	ARM	192
6.2	Blogs	192
6.2.1	Windows	192
7	More examples	193
7.1	“QR9”: Rubik’s cube inspired amateur crypto-algorithm	193
7.2	SAP	221
7.2.1	About SAP client network traffic compression	221
7.2.2	SAP 6.0 password checking functions	232
7.3	Oracle RDBMS	235
7.3.1	V\$VERSION table in Oracle RDBMS	235
7.3.2	X\$KSMLRU table in Oracle RDBMS	243
7.3.3	V\$TIMER table in Oracle RDBMS	245

8 Other things	249
8.1 Compiler's anomalies	249
9 Tasks solutions	250
9.1 Easy level	250
9.1.1 Task 1.1	250
9.1.2 Task 1.2	250
9.1.3 Task 1.3	250
9.1.4 Task 1.4	251
9.1.5 Task 1.5	251
9.1.6 Task 1.6	252
9.1.7 Task 1.7	252
9.1.8 Task 1.8	253
9.1.9 Task 1.9	253
9.2 Middle level	254
9.2.1 Task 2.1	254
Afterword	259
9.3 Donate	259
9.4 Questions?	259
Bibliography	260
Index	261

Preface

Here (will be) some of my notes about reverse engineering in English language for those beginners who like to learn to understand x86 (which is a most large mass of all executable software in the world) and ARM code created by C/C++ compilers.

There are two most used assembly language syntax: Intel (most used in DOS/Windows) and AT&T (used in *NIX)¹. Here we use Intel syntax. IDA 5 produce Intel syntax listings too.

¹http://en.wikipedia.org/wiki/X86_assembly_language#Syntax

About author

Dennis Yurichev is experienced reverse engineering, available for hire as reverse engineer, consultant, trainer. His CV is available [here](#).

Thanks

Andrey "herm1t" Baranovich, Slava "Avid" Kazakov, Stanislav "Beaver" Bobrytskyy, Alexander Lysenko, Alexander "Lstar" Chernenkiy, Arnaud Patard (rtp on #debian-arm IRC).

Chapter 1

Compiler's patterns

When I first learn C and then C++, I was just writing small pieces of code, compiling it and see what is producing in assembly language, that was an easy way to me. I did it a lot times and relation between C/C++ code and what compiler produce was imprinted in my mind so deep so that I can quickly understand what was in C code when I look at produced x86 code. Perhaps, this method may be helpful for anyone else, so I'll try to describe here some examples.

1.1 Hello, world!

Let's start with famous example from the book "The C programming Language" [\[Ker88\]](#):

```
#include <stdio.h>

int main()
{
    printf("hello, world");
    return 0;
};
```

1.1.1 x86

MSVC

Let's compile it in MSVC 2010: `cl 1.cpp /Fa1.asm`
(/Fa option mean generate assembly listing file)

Listing 1.1: MSVC 2010

```
CONST    SEGMENT
$SG3830 DB      'hello, world', 00H
CONST    ENDS
PUBLIC   _main
EXTRN    _printf:PROC
; Function compile flags: /Odtp
_TEXT    SEGMENT
_main    PROC
    push    ebp
    mov     ebp, esp
    push    OFFSET $SG3830
    call    _printf
    add     esp, 4
    xor     eax, eax
    pop     ebp
    ret     0
_main    ENDP
_TEXT    ENDS
```

Compiler generated `1.obj` file which will be linked into `1.exe`.

In our case, the file contain two segments: `CONST` (for data constants) and `_TEXT` (for code).

The string `'hello, world'` in C/C++ has type `const char*`, however hasn't its own name.

But compiler need to work with the string somehow, so it define internal name `$SG3830` for it.

As we can see, the string is terminated by zero byte — it's C/C++ standard for strings.

In the code segment `_TEXT` there are only one function so far — `main()`.

Function `main()` starting with prologue code and ending with epilogue code, like almost any function.

Read more about it in section about function prolog and epilog [2.2](#).

After function prologue we see a function `printf()` call: `CALL _printf`.

Before the call, string address (or pointer to it) containing our greeting is placed into stack with help of `PUSH` instruction.

When `printf()` function returning control flow to `main()` function, string address (or pointer to it) is still in stack.

Because we do not need it anymore, stack pointer (ESP register) is to be corrected.

`ADD ESP, 4` mean add 4 to the value in ESP register.

Why 4? Since it is 32-bit code, we need exactly 4 bytes for address passing through the stack. It's 8 bytes in x64-code

`"ADD ESP, 4"` is equivalent to `"POP register"` but without any register usage¹.

¹CPU flags, however, modified

Some compilers like Intel C++ Compiler, at the same point, could emit `POP ECX` instead of `ADD` (for example, such pattern can be observed in Oracle RDBMS code, compiled by Intel C++ compiler), and this instruction has almost the same effect, but `ECX` register contents will be rewritten.

Probably, Intel C++ compiler using `POP ECX` because this instruction's opcode is shorter than `ADD ESP, x` (1 byte against 3).

Read more about stack in relevant section [1.2](#).

After `printf()` call, in original C/C++ code was `return 0` — return zero as a `main()` function result.

In the generated code this is implemented by instruction `XOR EAX, EAX`

`XOR`, in fact, just “eXclusive OR”², but compilers using it often instead of `MOV EAX, 0` — slightly shorter opcode again (2 bytes against 5).

Some compilers emitting `SUB EAX, EAX`, which mean *SUBtract EAX value from EAX*, which is in any case will result zero.

Last instruction `RET` returning control flow to calling function. Usually, it's C/C++ CRT³ code, which, in turn, return control to operation system.

GCC

Now let's try to compile the same C/C++ code in GCC 4.4.1 compiler in Linux: `gcc 1.c -o 1`

After, with the IDA [5](#) disassembler assistance, let's see how `main()` function was created.

Note: we could also see GCC assembler result applying option `-S -masm=intel`

Listing 1.2: GCC

```
main      proc near
var_10    = dword ptr -10h

          push    ebp
          mov     ebp, esp
          and     esp, 0FFFFFFF0h
          sub     esp, 10h
          mov     eax, offset aHelloWorld ; "hello, world"
          mov     [esp+10h+var_10], eax
          call    _printf
          mov     eax, 0
          leave
          retn
main      endp
```

Almost the same. Address of “hello world” string (stored in data segment) is saved in `EAX` register first, then it stored into stack. Also, in function prologue we see `AND ESP, 0FFFFFFF0h` — this instruction aligning `ESP` value on 16-byte border, resulting all values in stack aligned too (CPU performing better if values it working with are located in memory at addresses aligned by 4 or 16 byte border)⁴.

`SUB ESP, 10h` allocate 16 bytes in stack, although, as we could see below, only 4 need here.

This is because the size of allocated stack is also aligned on 16-byte border.

String address (or pointer to string) is then writing directly into stack space without `PUSH` instruction use. `var_10` — is local variable, but also argument for `printf()`. Read below about it.

Then `printf()` function is called.

Unlike MSVC, GCC while compiling without optimization turned on, emitting `MOV EAX, 0` instead of shorter opcode.

The last instruction `LEAVE` — is `MOV ESP, EBP` and `POP EBP` instructions pair equivalent — in other words, this instruction setting back stack pointer (`ESP`) and `EBP` register to its initial state.

This is necessary because we modified these register values (`ESP` and `EBP`) at the function start (executing `MOV EBP, ESP` / `AND ESP, ...`).

²http://en.wikipedia.org/wiki/Exclusive_or

³C Run-Time Code

⁴Wikipedia: Data structure alignment

1.1.2 ARM

For my experiments with ARM CPU I choose two compilers: popular in embedded area Keil Release 6/2013 and Apple Xcode 4.6.3 IDE (with LLVM-GCC 4.2 compiler), producing code for ARM-comptaible processors and SoCs⁵ in iPod/iPhone/iPad, Windows 8 and Window RT tables⁶ and also such devices as Raspberry Pi.

Non-optimizing Keil + ARM mode

Let's start by compiling our example in Keil:

```
armcc.exe --arm --c90 -O0 1.c
```

armcc compiler producing assembly listing, but it has some high-level ARM-processor related macros⁷, but it's more important for us to see instructions "as is", so let's see compiled results in IDA [5](#).

Listing 1.3: Non-optimizing Keil + ARM mode + IDA [5](#)

```
.text:00000000          main
.text:00000000 10 40 2D E9      STMFD    SP!, {R4,LR}
.text:00000004 1E 0E 8F E2      ADR     R0, aHelloWorld ; "hello, world"
.text:00000008 15 19 00 EB      BL      __2printf
.text:0000000C 00 00 A0 E3      MOV     R0, #0
.text:00000010 10 80 BD E8      LDMFD    SP!, {R4,PC}

.text:000001EC 68 65 6C 6C+aHelloWorld DCB "hello, world",0 ; DATA XREF: main+4
```

Here is couple of ARM-related facts we should know in order to proceed. ARM processor has at least two major modes: ARM mode and thumb. In first (ARM) mode all instructions are enabled and each has 32-bit (4 bytes) size. In second (thumb) mode each instruction has 16-bit (or 2 bytes) size⁸. Thumb mode may look attractive because program in it may be 1) compact; 2) executing faster on microcontroller having 16-bit memory datapath. Nothing come for free of charge, so, in thumb mode there are reduced instruction set, only 8 registers are accessible and one need several thumb instructions for doing some operations when in ARM mode you'll need just one. Starting at ARMv7, there are also thumb-2 instructions set present, this is a thumb extended to support much bigger instructions set. There is a common misconception that thumb-2 is a mix of ARM and thumb. It's not correct. But rather thumb-2 was extended to support processor features so fully, so now it can compete with ARM mode. A program for ARM processor may be mix of procedures compiled for both modes. Majority of iPod/iPhone/iPad applications are compiled for thumb-2 instructions set, because Xcode do this by default.

In example we see here we can easily see that each instruction has size of 4 bytes. Indeed, we compiled our code for ARM mode, but for thumb.

The very first instruction "STMFD SP!, {R4,LR}"⁹ works here as PUSH in x86, instruction, writing values of two (R4 and LR) registers into stack. Indeed, in output listing, *armcc* compiler, for the sake of simplification, showing here "PUSH {r4,lr}" instruction. But it's not quite correct, PUSH instruction available only in thumb mode, so, to make things less messy, I offered to work in IDA [5](#).

So this instruction writes values of R4 and LR registers at the address in memory to which SP¹⁰ pointing, then decrements SP so it will points to a place in stack free for new entries.

This instruction, like PUSH instruction in thumb mode, is able save several register values at once and this may be useful. By the way, there is no such thing in x86. It's also can be noted that STMFD — generalization of PUSH instruction (extending its features), because it can work with any register, not just with SP and this can be very useful.

"ADR R0, aHelloWorld" instruction adding PC register value to the offset, where the "hello, world" string is located. How PC register used here, one might ask? This is so called "position-independent code" ¹¹, it is intended to be executed not to be fixed to any addresses in memory. In the opcode of ADR instruction, here is encoded a

⁵system on chip

⁶http://en.wikipedia.org/wiki/List_of_Windows_8_and_RT_tablet_devices

⁷for example, ARM mode lacks PUSH/POP instructions

⁸NOTTRANSLATED

⁹Store Multiple Full Descending

¹⁰stack pointer

¹¹Read more about it in relevant section [2.6](#)

difference between address of this instruction and the place where the string is located. Difference will always be constant, without any dependence to the address where that code being loaded, by operation system, presumably. That's why all we need is to add address of current instruction (from PC) in order to get absolute address of our C-string in memory.

'BL __2printf'¹² instruction calling printf() function. That's how this instruction works:

- write address after BL instruction (0xC) into LR¹³ register;
- then pass control flow into printf() by writing its address into PC¹⁴ register.

Because, when printf() finishes its work, it should have information, where it should return control, that's why each function passes control to the address stored in LR register.

That is the difference between "pure" RISC-processors like ARM and x86, where address of return is stored in stack¹⁵.

By the way, absolute 32-bit address or offset cannot be encoded in 32-bit BL instruction, because it has space only for 24 bits. It's also worth to note that all ARM mode instructions has size 4 bytes (32 bits), hence they all can be located only on 4-byte boundary addresses. This mean, last 2 bits of instruction address (always zero bits) may be omitted. In summary, we have 26 bit for offset encoding, this is enough to represent offset $\pm \approx 32M$.

Next 'MOV R0, #0'¹⁶ instruction just writes 0 into R0 register. That's because our C-function returning 0 and returning value is to be placed in R0.

The last instruction 'LDMFD SP!, R4, PC'¹⁷ is an inversive instruction of STMFD, it loads values from stack for saving them into R4 and PC, incrementing stack pointer SP. It can be said, it is similar to POP. Note: the very first instruction STMFD saved R4 and LR into stack, but R4 and PC are *restored*. As I wrote before, in LR¹⁸ register address of place saved, to where each function should return control. The very first function saving its value in stack because our main() function will use that register in order to call printf(). And then, in the function end this value can be written to PC, thus, by passing control to where our function was called. Since our main() function is usually primary function in C/C++, apparently, control will be returned to operation system loader or to some place in runtime C functions, or something like that.

DCB — assembly language directive, defining array of bytes or ASCII-strings, similar to DB directive in x86-assembly language.

Non-optimizing Keil: thumb mode

Let's compile the same example in Keil in thumb mode:

```
armcc.exe --thumb --c90 -O0 1.c
```

We will get (in IDA 5):

Listing 1.4: Non-optimizing Keil + thumb mode + IDA 5

```
.text:00000000      main
.text:00000000 10 B5      PUSH    {R4,LR}
.text:00000002 C0 A0      ADR     R0, aHelloWorld ; "hello, world"
.text:00000004 06 F0 2E F9    BL      __2printf
.text:00000008 00 20      MOVS   R0, #0
.text:0000000A 10 BD      POP     {R4,PC}

.text:00000304 68 65 6C 6C+aHelloWorld    DCB "hello, world",0 ; DATA XREF: main+2
```

We can easily spot 2-byte (16-bit) opcodes, this is, as I mentioned, thumb. Except BL instruction. In fact, it consisted in two 16-bit instructions. That's because it's not possible to load offset to printf() function into PC when using so small space in one 16-bit opcode, obviously. That's why first 16-bit instruction loads higher 10 bits

¹²Branch with Link

¹³link register

¹⁴program counter

¹⁵Read more about this in next section 1.2

¹⁶MOVE

¹⁷Load Multiple Full Descending

¹⁸link register

of offset and second — loads 11 lower bits of offset. As I mentioned, all instructions in thumb mode has size of 2 bytes or 16 bits. This mean, it's not possible for thumb-instruction to be on odd address whatsoever. Considering this, last address bit may be omitted while instruction encoding. Summarizing, in BL thumb-instruction, $\pm \approx 2M$ can be encoded as offset from current address.

Other instructions in functions are: PUSH and POP works just like described STMFD/LDMFD, but SP register not mentioned explicitly here. ADR works just like in previous example. MOVS writes 0 in R0 register to zero returning.

Optimizing Xcode (LLVM) + ARM mode

Xcode 4.6.3 without optimization turned on, produces a lot of redundant code, so we'll study that version where instruction count as small as possible: -O3.

Listing 1.5: Optimizing Xcode (LLVM) + ARM mode

```
__text:000028C4      _hello_world
__text:000028C4 80 40 2D E9      STMFD      SP!, {R7,LR}
__text:000028C8 86 06 01 E3      MOV        R0, #0x1686
__text:000028CC 0D 70 A0 E1      MOV        R7, SP
__text:000028D0 00 00 40 E3      MOVT       R0, #0
__text:000028D4 00 00 8F E0      ADD        R0, PC, R0
__text:000028D8 C3 05 00 EB      BL         _puts
__text:000028DC 00 00 A0 E3      MOV        R0, #0
__text:000028E0 80 80 BD E8      LDMFD      SP!, {R7,PC}

__cstring:00003F62 48 65 6C 6C+aHelloWorld_0 DCB "Hello world!",0
```

Instructions STMFD and LDMFD are familiar to us.

MOV instruction just writes 0x1686 number into R0 register, this is offset pointing to the “Hello world!” string.

R7 register, as it is standardized in [App10] is frame pointer, more on it below.

MOVT R0, #0 instruction writes 0 into higher 16 bit of register. The issue is here in that generic MOV instruction in ARM mode may writes only lower 16 bit of register. Remember, all instruction's opcodes in ARM mode are limited in size to 32 bits. Of course, this limitation is not related to moving between registers. So that's why additional instruction MOVT exist for writing into higher bits (from 16 to 31 inclusive). However, its usage here is redundant, because “MOV R0, #0x1686” instruction above cleared higher part of register. Probably, it's compiler's shortcoming.

“ADD R0, PC, R0” instruction adding PC to R0, for calculating absolute address of “Hello world!” string, and as we already know that, it's “position-independent code”, so this corrective is essential here.

BL instruction calling puts() instead of printf().

GCC replaced first printf() call to puts(). Indeed: printf() with sole argument is almost analogous to puts().

Almost, because we need to be sure that this string will not contain printf-control statements starting with %: then effect of these two functions will be different.

Why compiler replaced printf() to puts()? Because puts() () work faster ¹⁹.

puts() working faster because it just passes characters to stdout not comparing each with % symbol.

Next, we see familiar to us “MOV R0, #0” instruction, intended to set 0 to R0 register.

Optimizing Xcode (LLVM) + thumb-2 mode

By default, Xcode 4.6.3 generating code for thumb-2 in such manner:

Listing 1.6: Optimizing Xcode (LLVM) + thumb-2 mode

```
__text:00002B6C      _hello_world
__text:00002B6C 80 B5      PUSH      {R7,LR}
__text:00002B6E 41 F2 D8 30      MOVW      R0, #0x13D8
__text:00002B72 6F 46      MOV        R7, SP
__text:00002B74 C0 F2 00 00      MOVT.W    R0, #0
__text:00002B78 78 44      ADD        R0, PC
```

¹⁹http://www.ciselant.de/projects/gcc_printf/gcc_printf.html

```

__text:00002B7A 01 F0 38 EA      BLX      _puts
__text:00002B7E 00 20      MOVS     R0, #0
__text:00002B80 80 BD      POP     {R7,PC}

...

__cstring:00003E70 48 65 6C 6C 6F 20+aHelloWorld  DCB "Hello world!",0xA,0

```

BL and BLX instructions in thumb mode, as we remember, encoded as pair of 16-bit instructions and in thumb-2, these *surrogate* opcodes extended in such way so that new instruction may be encoded here as 32-bit instructions. That's easily observable — opcodes of thumb-2 instructions are also beginning with *0xFx* or *0xE_x*. But in IDA 5 listings, first byte of opcode is at the place of second, that's because instructions here encoded as follows: last byte and then first one (for thumb and thumb-2 modes), or, (for instructions in ARM mode): fourth byte, then third, then second and first. So as we see, MOVW, MOVT.W and BLX instructions are beginning with *0xFx*.

One of thumb-2 instructions is 'MOVW R0, #0x13D8' — it writes 16-bit value into lower part of R0 register.

Also 'MOVT.W R0, #0' — this instruction works just like MOVT from previous example, but it works in thumb-2.

Among other differences, here is BLX instruction used instead of BL. Difference in that way that beside saving of return address in LR register and passing control to puts() function, processor is switching from thumb mode to ARM or back. This instruction in place here because the instruction to which control is passed looks like (it's encoded in ARM mode):

```

__symbolstub1:00003FEC _puts      ; CODE XREF: _hello_world+E
__symbolstub1:00003FEC 44 F0 9F E5  LDR PC, =__imp__puts

```

So, observant reader may ask: why not to call puts() right at the place of code where it needed?

But that's not very space-efficient, and that's why.

Almost any program uses external dynamic libraries, like DLL in Windows, .so in *NIX or .dylib in Mac OS X. Often used library functions are stored in dynamic libraries, including standard C-function puts().

In executable binary file (Windows PE .exe, ELF or Mach-O) a section of imports is present, that is list of symbols (functions or global variables) being imported from external modules and also names of these modules.

Operation system loader loads all modules need and, while enumerating importing symbols in primary module, sets correct addresses of each symbol.

In our case, __imp__puts is 32-bit variable where OS loader will write correct address of that function in external library. So that LDR instruction just takes 32-bit value from this variable and, writing it into PC register, just passing control to it.

So to reduce a time OS loader needs for doing this procedure, it's good idea for it to write address of each symbol only once, to special place for it.

Besides, as we already figured out, it's not possible to load 32-bit value into register using only one instruction, without memory access. So, it is optimal to allocate separate function working in ARM mode with only one goal — to pass control to dynamic library. And then to jump to this short one-instruction function (so called thunk-function) from thumb-code.

By the way, in previous example (compiled for ARM mode) control passing by BL instruction is going to the same thunk-function, however, processor mode is not switched (hence, absence of "X" in instruction mnemonic).

1.2 Stack

Stack — is one of the most fundamental things in computer science.²⁰

Technically, it's just a memory block in process memory + ESP or RSP register in x86, or SP register in ARM, as a pointer within this block.

Most frequently used stack access instructions are PUSH and POP (both in x86 and ARM thumb-mode). PUSH subtracting ESP/RSP/SP by 4 and then writing contents of its sole operand to the memory address pointing by ESP/RSP/SP.

POP is reverse operation: get a data from memory pointing by ESP/RSP/SP, put it to operand (often register) and then add 4 to ESP/RSP/SP. Of course, this is for 32-bit environment. 8 will be here instead of 4 in x64 environment.

After stack allocation, stack pointer pointing to the end of stack. PUSH increasing stack pointer, and POP decreasing. The end of stack is actually at the beginning of allocated for stack memory block. It seems strange, but it is so.

Nevertheless, ARM has instructions supporting ascending stacks, but also descending stacks. For example, STMFD²¹/LDMFD²², STMED²³/LDMED²⁴ instructions are intended for work with descending stack. STMFA²⁵/LMDFA²⁶, STMEA²⁷/LDMEA²⁸ instructions are intended for work with ascending stack.

What stack is used for?

1.2.1 Save return address where function should return control after execution

x86

While calling another function by CALL instruction, the address of point exactly after CALL instruction is saved to stack, and then unconditional jump to the address from CALL operand is executed.

CALL is PUSH address_after_call / JMP operand instructions pair equivalent.

RET is fetching value from stack and jump to it — it is POP tmp / JMP tmp instructions pair equivalent.

Stack overflow is simple, just run eternal recursion:

```
void f()
{
    f();
};
```

MSVC 2008 reporting about problem:

```
c:\tmp6>cl ss.cpp /Fass.asm
Microsoft (R) 32-bit C/C++ Optimizing Compiler Version 15.00.21022.08 for 80x86
Copyright (C) Microsoft Corporation. All rights reserved.

ss.cpp
c:\tmp6\ss.cpp(4) : warning C4717: 'f' : recursive on all control paths, function will cause runtime stack
overflow
```

...but generates right code anyway:

```
?f@@YAXXZ PROC                                ; f
; File c:\tmp6\ss.cpp
; Line 2
    push    ebp
    mov     ebp, esp
; Line 3
    call    ?f@@YAXXZ                          ; f
```

²⁰http://en.wikipedia.org/wiki/Call_stack

²¹Store Multiple Full Descending

²²Load Multiple Full Descending

²³Store Multiple Empty Descending

²⁴Load Multiple Empty Descending

²⁵Store Multiple Full Ascending

²⁶Load Multiple Full Ascending

²⁷Store Multiple Empty Ascending

²⁸Load Multiple Empty Ascending


```
; Line 4
    pop     ebp
    ret     0
?f@@YAXXZ ENDP                ; f
```

...Also, if we turn on optimization (/Ox option), the optimized code will not overflow stack, but will work *correctly*²⁹:

```
?f@@YAXXZ PROC                ; f
; File c:\tmp6\ss.cpp
; Line 2
$LL3@f:
; Line 3
    jmp     SHORT $LL3@f
?f@@YAXXZ ENDP                ; f
```

GCC 4.4.1 generating the same code in both cases, although not warning about problem.

ARM

ARM programs also use stack for return address savings, but differently. As it was mentioned in “Hello, world!” 1.1.2, return address is saved to LR (*link register*). However, if one need to call some another function and use LR register one more time, its value should be saved. Usually, it is saved in function prologue, often, we see there instruction like “PUSH R4-R7, LR”, and also this instruction in epilogue “POP R4-R7, PC” — thus register values to be used in function are saved in stack, including LR.

Nevertheless, if some function never call any other function, in ARM terminology it’s called *leaf function*³⁰. As a consequence, “leaf” function don’t use LR register. And if this function is small, if it use small number of registers, it may not use stack at all. Thus, it is possible to call “leaf” function without stack usage. This could be faster than in x86, because external RAM is not used for stack³¹. Or, it can be useful for such situations, when memory for stack is not yet allocated or not available.

1.2.2 Function arguments passing

```
push arg3
push arg2
push arg1
call f
add esp, 4*3
```

Callee³² function get its arguments via stack pointer.

See also section about calling conventions 2.5.

It is important to note that no one oblige programmers to pass arguments through stack, it is not prerequisite. One could implement any other method not using stack.

For example, it is possible to allocate a place for arguments in heap, fill it and pass to a function via pointer to this pack in EAX register. And this will work³³.

However, it is convenient tradition in x86 and ARM to use stack for this.

1.2.3 Local variable storage

A function could allocate some space in stack for its local variables just shifting stack pointer deeply enough to stack bottom.

It is also not prerequisite. You could store local variables wherever you like. But traditionally it is so.

²⁹irony here

³⁰<http://infocenter.arm.com/help/index.jsp?topic=/com.arm.doc.faq/ka13785.html>

³¹Some time ago, on PDP-11 and VAX, CALL instruction (calling other functions) was expensive, up to 50% of execution time might be spent on it, so it was common sense that big number of small function is anti-pattern [Ray03, Chapter 4, Part II].

³²Function being called

³³For example, in “The Art of Computer Programming” book by Donald Knuth, in section 1.4.1 dedicated to subroutines [Knu98, section 1.4.1], we can read about one way to supply arguments to subroutine is simply to list them after the JMP instruction passing control to subroutine. Knuth writes that this method was particularly convenient on System/360.

x86: `alloca()` function

It is worth noting `alloca()` function.³⁴

This function works like `malloc()`, but allocate memory just in stack.

Allocated memory chunk is not needed to be freed via `free()` function call because function epilogue 2.2 will return ESP back to initial state and allocated memory will be just annuled.

It is worth noting how `alloca()` implemented.

This function, if to simplify, just shifting ESP deeply to stack bottom so much bytes you need and set ESP as a pointer to that *allocated* block. Let's try:

```
#include <malloc.h>
#include <stdio.h>

void f()
{
    char *buf=(char*)alloca (600);
    _snprintf (buf, 600, "hi! %d, %d, %d\n", 1, 2, 3);

    puts (buf);
};
```

(`_snprintf()` function works just like `printf()`, but instead dumping result into stdout (e.g., to terminal or console), write it to `buf` buffer. `puts()` copies `buf` contents to stdout. Of course, these two function calls might be replaced by one `printf()` call, but I would like to illustrate small buffer usage.)

Let's compile (MSVC 2010):

Listing 1.7: MSVC 2010

```
...

mov     eax, 600           ; 00000258H
call    __alloca_probe_16
mov     esi, esp

push    3
push    2
push    1
push    OFFSET $SG2672
push    600                ; 00000258H
push    esi
call    __snprintf

push    esi
call    _puts
add     esp, 28            ; 0000001cH

...
```

The sole `alloca()` argument passed via EAX (instead of pushing into stack)³⁵. After `alloca()` call, ESP is now pointing to the block of 600 bytes and we can use it as memory for `buf` array.

GCC 4.4.1 can do the same without calling external functions:

Listing 1.8: GCC 4.4.1

```
f      public f
      proc near                ; CODE XREF: main+6

s      = dword ptr -10h
var_C  = dword ptr -0Ch

      push    ebp
```

³⁴As of MSVC, function implementation can be found in `alloca16.asm` and `chkstk.asm` in `C:\Program Files (x86)\Microsoft Visual Studio 10.0\VC\crt\src\intel`

³⁵It's because `alloca()` is rather compiler intrinsic than usual function

```

mov     ebp, esp
sub     esp, 38h
mov     eax, large gs:14h
mov     [ebp+var_C], eax
xor     eax, eax
sub     esp, 624
lea     eax, [esp+18h]
add     eax, 0Fh
shr     eax, 4           ; align pointer
shl     eax, 4           ; by 16-byte border
mov     [ebp+s], eax
mov     eax, offset format ; "hi! %d, %d, %d\n"
mov     dword ptr [esp+14h], 3
mov     dword ptr [esp+10h], 2
mov     dword ptr [esp+0Ch], 1
mov     [esp+8], eax     ; format
mov     dword ptr [esp+4], 600 ; maxlen
mov     eax, [ebp+s]
mov     [esp], eax       ; s
call    _snprintf
mov     eax, [ebp+s]
mov     [esp], eax       ; s
call    _puts
mov     eax, [ebp+var_C]
xor     eax, large gs:14h
jz      short locret_80484EB
call    ___stack_chk_fail

```

```

locret_80484EB:                                ; CODE XREF: f+70

```

```

    leave
    retn
f      endp

```

1.2.4 (Windows) SEH

SEH (*Structured Exception Handling*) records are also stored in stack (if needed). ³⁶.

1.2.5 Buffer overflow protection

More about it here [1.14.2](#).

³⁶Classic Matt Pietrek article about SEH: <http://www.microsoft.com/msj/0197/Exception/Exception.aspx>

1.3 printf() with several arguments

Now let's extend *Hello, world!* [1.1](#) example, replacing `printf()` in `main()` function body by this:

```
printf("a=%d; b=%d; c=%d", 1, 2, 3);
```

1.3.1 x86

Let's compile it by MSVC 2010 and we got:

```
$SG3830 DB      'a=%d; b=%d; c=%d', 00H
...
        push    3
        push    2
        push    1
        push    OFFSET $SG3830
        call    _printf
        add     esp, 16                ; 00000010H
```

Almost the same, but now we can see that `printf()` arguments are pushing into stack in reverse order: and the first argument is pushing in as the last one.

By the way, variables of *int* type in 32-bit environment has 32-bit width, that's 4 bytes.

So, we got here 4 arguments. $4 * 4 = 16$ — they occupy in stack exactly 16 bytes: 32-bit pointer to string and 3 number of *int* type.

When stack pointer (ESP register) is corrected by “ADD ESP, X” instruction after some function call, often, the number of function arguments could be deduced here: just divide X by 4.

Of course, this is related only to *cdecl* calling convention.

See also section about calling conventions [2.5](#).

It is also possible for compiler to merge several “ADD ESP, X” instructions into one, after last call:

```
push a1
push a2
call ...
...
push a1
call ...
...
push a1
push a2
push a3
call ...
add esp, 24
```

Now let's compile the same in Linux by GCC 4.4.1 and take a look in IDA [5](#) what we got:

```
main      proc near
var_10     = dword ptr -10h
var_C      = dword ptr -0Ch
var_8      = dword ptr -8
var_4      = dword ptr -4

        push    ebp
        mov     ebp, esp
        and     esp, 0FFFFFF0h
        sub     esp, 10h
        mov     eax, offset aADBD CD ; "a=%d; b=%d; c=%d"
        mov     [esp+10h+var_4], 3
        mov     [esp+10h+var_8], 2
        mov     [esp+10h+var_C], 1
        mov     [esp+10h+var_10], eax
        call    _printf
```

```

mov     eax, 0
leave
retn
main    endp

```

It can be said, the difference between code by MSVC and GCC is only in method of placing arguments into stack. Here GCC working directly with stack without PUSH/POP.

1.3.2 ARM: 3 printf() arguments

Traditionally, ARM arguments passing scheme (calling convention) is as follows: the 4 first arguments are passed in R0-R3 registers and remaining arguments — via stack. This resembling arguments passing scheme in fastcall 2.5.3 or win64 2.5.5.

Non-optimizing Keil + ARM mode

Listing 1.9: Non-optimizing Keil + ARM mode

```

.text:00000014      printf_main1
.text:00000014 10 40 2D E9      STMFD    SP!, {R4,LR}
.text:00000018 03 30 A0 E3      MOV     R3, #3
.text:0000001C 02 20 A0 E3      MOV     R2, #2
.text:00000020 01 10 A0 E3      MOV     R1, #1
.text:00000024 1D 0E 8F E2      ADR     R0, aADBDCD      ; "a=%d; b=%d; c=%d\n"
.text:00000028 0D 19 00 EB      BL     __2printf
.text:0000002C 10 80 BD E8      LDMFD    SP!, {R4,PC}

```

So, the first 4 arguments are passing via R0-R3 arguments in this order: pointer to printf() format string in R0, then 1 in R1, 2 in R2 and 3 in R3.

Nothing unusual so far.

Optimizing Keil + ARM mode

Listing 1.10: Optimizing Keil + ARM mode

```

.text:00000014      EXPORT printf_main1
.text:00000014      printf_main1
.text:00000014 03 30 A0 E3      MOV     R3, #3
.text:00000018 02 20 A0 E3      MOV     R2, #2
.text:0000001C 01 10 A0 E3      MOV     R1, #1
.text:00000020 1E 0E 8F E2      ADR     R0, aADBDCD      ; "a=%d; b=%d; c=%d\n"
.text:00000024 CB 18 00 EA      B       __2printf

```

This is optimized (-O3) version for ARM mode and here we see B as the last instruction instead of familiar BL. Another difference between this optimized version and previous one, compiled without optimization, is also in the fact that there are no function prologue and epilogue (instructions saving R0 and LR registers values). B instruction just jumping to another address, without any LR register manipulation, that is, it's JMP analogue in x86. Why it works fine? Because this code is in fact equivalent to the previous. There are two main reasons: 1) stack is not modified, as well as SP stack pointer; 2) printf() call is the last one, there are nothing going on after it. After finishing, printf() function will just return control to the address stored in LR. But the address of the place from where our function was called is now in LR! And consequently, control from printf() will returned to that place. As a consequent, we don't need to save LR, because we don't need to modify LR. And we don't need to modify LR because there are no other functions calls except printf(), furthermore, after this call we are not planning to do anything! That's why this optimization is possible.

Another similar example was described in "switch()/case/default" section, here 1.9.1.

Optimizing Keil + thumb mode

Listing 1.11: Optimizing Keil + thumb mode

```
.text:0000000C      printf_main1
.text:0000000C 10 B5      PUSH    {R4,LR}
.text:0000000E 03 23      MOVS    R3, #3
.text:00000010 02 22      MOVS    R2, #2
.text:00000012 01 21      MOVS    R1, #1
.text:00000014 A4 A0      ADR     R0, aADBDCD      ; "a=%d; b=%d; c=%d\n"
.text:00000016 06 F0 EB F8    BL     __2printf
.text:0000001A 10 BD      POP     {R4,PC}
```

There are no significant difference from non-optimized code for ARM mode.

1.3.3 ARM: 8 printf() arguments

To see, how other arguments will be passed via stack, let's change our example again by increasing number of arguments to be passed to 9 (printf() format string + 8 int variables):

```
void printf_main2()
{
    printf("a=%d; b=%d; c=%d; d=%d; e=%d; f=%d; g=%d; h=%d\n", 1, 2, 3, 4, 5, 6, 7, 8);
};
```

Optimizing Keil: ARM mode

```
.text:00000028      printf_main2
.text:00000028
.text:00000028      var_18      = -0x18
.text:00000028      var_14      = -0x14
.text:00000028      var_4       = -4
.text:00000028
.text:00000028 04 E0 2D E5      STR     LR, [SP,#var_4]!
.text:0000002C 14 D0 4D E2      SUB     SP, SP, #0x14
.text:00000030 08 30 A0 E3      MOV     R3, #8
.text:00000034 07 20 A0 E3      MOV     R2, #7
.text:00000038 06 10 A0 E3      MOV     R1, #6
.text:0000003C 05 00 A0 E3      MOV     R0, #5
.text:00000040 04 C0 8D E2      ADD     R12, SP, #0x18+var_14
.text:00000044 0F 00 8C E8      STMIA   R12, {R0-R3}
.text:00000048 04 00 A0 E3      MOV     R0, #4
.text:0000004C 00 00 8D E5      STR     R0, [SP,#0x18+var_18]
.text:00000050 03 30 A0 E3      MOV     R3, #3
.text:00000054 02 20 A0 E3      MOV     R2, #2
.text:00000058 01 10 A0 E3      MOV     R1, #1
.text:0000005C 6E 0F 8F E2      ADR     R0, aADBDCDDDEDFDGD ; "a=%d; b=%d; c=%d; d=%d; e=%d; f=%d
; g=%d" ...
.text:00000060 BC 18 00 EB      BL     __2printf
.text:00000064 14 D0 8D E2      ADD     SP, SP, #0x14
.text:00000068 04 F0 9D E4      LDR     PC, [SP+4+var_4],#4
```

This code can be divided into several parts:

- Function prologue:

The very first “STR LR, [SP,#var_4]!” instruction save LR in stack, because we will use this register for printf() call.

The second “SUB SP, SP, #0x14” instruction decreasing SP stack pointer, but in fact, this procedure is needed for allocating a space of size 0x14 (20) bytes in the stack. Indeed, we need to pass 5 32-bit values via stack to printf(), and each one occupy 4 bytes, that is $5 * 4 = 20$ — exactly. Other 4 32-bit values will be passed in registers.

- Passing 5, 6, 7 and 8 via stack:

Then values 5, 6, 7 and 8 are written to R0, R1, R2 and R3 registers respectively. Then “ADD R12, SP, #0x18+var_14” instruction write an address of place in stack, where these 4 variables will be written, into R12 register. *var_14* is an assembly macro, equal to $-0x14$, such macros are created by IDA 5 in order to show simply how code accessing stack. *var_?* macros created by IDA 5 reflecting local variables in stack. So, $SP + 4$ will be written into R12 register. The next “STMIA R12, R0-R3” instruction writes R0-R3 registers contents at the place in memory to which R12 pointing. STMIA instruction meaning *Store Multiple Increment After*. *Increment After* meaning that R12 will be increasing by 4 after each register value write.

- Passing 4 via stack: 4 is stored in R0 and then, this value, with the help of “STR R0, [SP, #0x18+var_18]” instruction, is saved in stack. *var_18* is $-0x18$, offset will be 0, so, value from R0 register (4) will be written to a place SP pointing to.

- Passing 1, 2 and 3 via registers:

Values of first 3 numbers (a, b, c) (1, 2, 3 respectively) are passing in R1, R2 and R3 registers right before `printf()` call, and other 5 values are passed via stack and this is how:

- `printf()` call:

- Function epilogue:

“ADD SP, SP, #0x14” instruction returning SP to former place, thus annulling what was written to stack. Of course, what was written in stack will stay there, but it all will be rewritten while execution of following functions.

“LDR PC, [SP+4+var_4], #4” instruction loading saved in stack LR value into PC, providing exit from the function.

Optimizing Keil: thumb mode

```
.text:0000001C      printf_main2
.text:0000001C
.text:0000001C      var_18          = -0x18
.text:0000001C      var_14          = -0x14
.text:0000001C      var_8           = -8
.text:0000001C
.text:0000001C 00 B5          PUSH      {LR}
.text:0000001E 08 23          MOVS     R3, #8
.text:00000020 85 B0          SUB      SP, SP, #0x14
.text:00000022 04 93          STR      R3, [SP, #0x18+var_8]
.text:00000024 07 22          MOVS     R2, #7
.text:00000026 06 21          MOVS     R1, #6
.text:00000028 05 20          MOVS     R0, #5
.text:0000002A 01 AB          ADD      R3, SP, #0x18+var_14
.text:0000002C 07 C3          STMIA   R3!, {R0-R2}
.text:0000002E 04 20          MOVS     R0, #4
.text:00000030 00 90          STR      R0, [SP, #0x18+var_18]
.text:00000032 03 23          MOVS     R3, #3
.text:00000034 02 22          MOVS     R2, #2
.text:00000036 01 21          MOVS     R1, #1
.text:00000038 A0 A0          ADR      R0, aADBDCDDDEDFDGD ; "a=%d; b=%d; c=%d; d=%d; e=%d; f=%d
; g=%d"...
.text:0000003A 06 F0 D9 F8      BL       __2printf
.text:0000003E
.text:0000003E      loc_3E              ; CODE XREF: example13_f+16
.text:0000003E 05 B0          ADD      SP, SP, #0x14
.text:00000040 00 BD          POP      {PC}
```

Almost same as in previous example, however, this is thumb code and values are packed into stack differently: 8 for the first time, then 5, 6, 7 for the second and 4 for the third.

Optimizing Xcode (LLVM): ARM mode

```
__text:0000290C          _printf_main2
__text:0000290C
__text:0000290C          var_1C          = -0x1C
__text:0000290C          var_C          = -0xC
__text:0000290C
__text:0000290C 80 40 2D E9          STMFD          SP!, {R7,LR}
__text:00002910 0D 70 A0 E1          MOV           R7, SP
__text:00002914 14 D0 4D E2          SUB           SP, SP, #0x14
__text:00002918 70 05 01 E3          MOV           R0, #0x1570
__text:0000291C 07 C0 A0 E3          MOV           R12, #7
__text:00002920 00 00 40 E3          MOVT          R0, #0
__text:00002924 04 20 A0 E3          MOV           R2, #4
__text:00002928 00 00 8F E0          ADD           R0, PC, R0
__text:0000292C 06 30 A0 E3          MOV           R3, #6
__text:00002930 05 10 A0 E3          MOV           R1, #5
__text:00002934 00 20 8D E5          STR           R2, [SP,#0x1C+var_1C]
__text:00002938 0A 10 8D E9          STMFA         SP, {R1,R3,R12}
__text:0000293C 08 90 A0 E3          MOV           R9, #8
__text:00002940 01 10 A0 E3          MOV           R1, #1
__text:00002944 02 20 A0 E3          MOV           R2, #2
__text:00002948 03 30 A0 E3          MOV           R3, #3
__text:0000294C 10 90 8D E5          STR           R9, [SP,#0x1C+var_C]
__text:00002950 A4 05 00 EB          BL           _printf
__text:00002954 07 D0 A0 E1          MOV           SP, R7
__text:00002958 80 80 BD E8          LDMFD         SP!, {R7,PC}
```

Almost the same what we already figured out, with the exception of STMFA (Store Multiple Full Ascending) instruction, it is synonym to STMIB (Store Multiple Increment Before) instruction. This instruction increasing SP and only then writing next register value into memory, but not vice versa.

Another thing we easily spot is that instructions ostensibly located randomly. For instance, R0 register value is prepared in three places, at addresses 0x2918, 0x2920 and 0x2928, when it would be possible to do it in one single place. However, optimizing compiler has his own reasons about how to place instructions better. Usually, processor attempts to execute instructions located side-by-side. For example, instructions like ‘MOVT R0, #0’ and ‘ADD R0, PC, R0’ cannot be executed simultaneously, because they both modifying R0 register. On the other hand, ‘MOVT R0, #0’ and ‘MOV R2, #4’ instructions can be executed simultaneously because effects of their execution are not conflicting with each other. Presumably, compiler tries to generate code in such way, where it’s possible, of course.

Optimizing Xcode (LLVM): thumb-2 mode

```
__text:00002BA0          _printf_main2
__text:00002BA0
__text:00002BA0          var_1C          = -0x1C
__text:00002BA0          var_18          = -0x18
__text:00002BA0          var_C          = -0xC
__text:00002BA0
__text:00002BA0 80 B5          PUSH          {R7,LR}
__text:00002BA2 6F 46          MOV           R7, SP
__text:00002BA4 85 B0          SUB           SP, SP, #0x14
__text:00002BA6 41 F2 D8 20          MOVW          R0, #0x12D8
__text:00002BAA 4F F0 07 0C          MOV.W         R12, #7
__text:00002BAE C0 F2 00 00          MOVT.W        R0, #0
__text:00002BB2 04 22          MOVS          R2, #4
__text:00002BB4 78 44          ADD           R0, PC ; char *
__text:00002BB6 06 23          MOVS          R3, #6
__text:00002BB8 05 21          MOVS          R1, #5
__text:00002BBA 0D F1 04 0E          ADD.W         LR, SP, #0x1C+var_18
__text:00002BBE 00 92          STR           R2, [SP,#0x1C+var_1C]
__text:00002BC0 4F F0 08 09          MOV.W         R9, #8
__text:00002BC4 8E E8 0A 10          STMIA.W       LR, {R1,R3,R12}
```


__text:00002BC8	01 21	MOVS	R1, #1
__text:00002BCA	02 22	MOVS	R2, #2
__text:00002BCC	03 23	MOVS	R3, #3
__text:00002BCE	CD F8 10 90	STR.W	R9, [SP,#0x1C+var_C]
__text:00002BD2	01 F0 0A EA	BLX	_printf
__text:00002BD6	05 B0	ADD	SP, SP, #0x14
__text:00002BD8	80 BD	POP	{R7,PC}

Almost the same as in previous example, with the exception that thumb-instructions are used there instead.

1.3.4 By the way

By the way, this difference between passing arguments in x86 and ARM is a good illustration that CPU is not aware of how arguments is passed to functions. It is also possible to create hypothetical compiler which is able to pass arguments via some special structure not using stack at all.

1.4 scanf()

Now let's use `scanf()`.

```
int main()
{
    int x;
    printf ("Enter X:\n");

    scanf ("%d", &x);

    printf ("You entered %d...\n", x);

    return 0;
};
```

OK, I agree, it is not clever to use `scanf()` today. But I wanted to illustrate passing pointer to `int`.

1.4.1 About pointers

It's one of the most fundamental things in computer science. Often, large array, structure or object, it's too costly to pass to another function, while passing its address is much easier. More than that: if calling function should modify something in that large array or structure, to return it as a whole is absurdical as well. So the most simple thing to do is to pass an address of array or structure to function, and let it change it what need.

In C/C++ it's just an address of some place in memory.

In x86, address is represented as 32-bit number (i.e., occupying 4 bytes), while in x86-64 it's 64-bit number (occupying 8 bytes). By the way, that's a reason of some people's indignation related to switching to x86-64 — all pointers on that architecture will require twice as more space.

With some effort, it's possible to work only with untyped pointers, for example, standard function `memcpy()`, copying a block from one place in memory to another, takes 2 pointers of `void*` type on input, because it's not possible to predict block type you would like to copy, and it's not even important to know, only block size is important.

Also, pointers are widely used when function need to return more than one value (we will back to this in future 1.7). `scanf()` is just that case. In addition to the function's need to show, how many values were read successfully, it also should to return all these values.

In C/C++ pointer type is needed only for type checking on compiling stage. Internally, in compiled code, there are no information about pointers types.

1.4.2 x86

What we got after compiling in MSVC 2010:

```
CONST    SEGMENT
$SG3831  DB      'Enter X:', 0aH, 00H
        ORG $+2
$SG3832  DB      '%d', 00H
        ORG $+1
$SG3833  DB      'You entered %d...', 0aH, 00H
CONST    ENDS
PUBLIC   _main
EXTRN    _scanf:PROC
EXTRN    _printf:PROC
; Function compile flags: /Odtp
_TEXT    SEGMENT
_x$ = -4                                ; size = 4
_main    PROC
    push    ebp
    mov     ebp, esp
    push    ecx
    push    OFFSET $SG3831
    call    _printf
```

```

add     esp, 4
lea     eax, DWORD PTR _x$[ebp]
push    eax
push    OFFSET $SG3832
call    _scanf
add     esp, 8
mov     ecx, DWORD PTR _x$[ebp]
push    ecx
push    OFFSET $SG3833
call    _printf
add     esp, 8
xor     eax, eax
mov     esp, ebp
pop     ebp
ret     0
_main   ENDP
_TEXT   ENDS

```

Variable x is local.

C/C++ standard tell us it must be visible only in this function and not from any other place. Traditionally, local variables are placed in the stack. Probably, there could be other ways, but in x86 it is so.

Next after function prologue instruction PUSH ECX is not for saving ECX state (notice absence of corresponding POP ECX at the function end).

In fact, this instruction just allocate 4 bytes in stack for x variable storage.

x will be accessed with the assistance of _x\$ macro (it equals to -4) and EBP register pointing to current frame.

Over a span of function execution, EBP is pointing to current stack frame and it is possible to have an access to local variables and function arguments via EBP+offset.

It is also possible to use ESP, but it's often changing and not very convenient. So it can be said, EBP is *frozen state* of ESP at the moment of function execution start.

Function scanf () in our example has two arguments.

First is pointer to the string containing "%d" and second — address of variable x.

First of all, address of x is placed into EAX register by lea eax, DWORD PTR _x\$[ebp] instruction LEA meaning *load effective address*, but over a time it changed its primary application [2.1](#).

It can be said, LEA here just placing to EAX sum of EBP value and _x\$ macro.

It is the same as lea eax, [ebp-4].

So, 4 subtracting from EBP value and result is placed to EAX. And then value in EAX is pushing into stack and scanf () is called.

After that, printf () is called. First argument is pointer to string: "You entered %d...\n".

Second argument is prepared as: mov ecx, [ebp-4], this instruction placing to ECX not address of x variable but its contents.

After, ECX value is placing into stack and last printf () called.

Let's try to compile this code in GCC 4.4.1 under Linux:

```

main      proc near
var_20    = dword ptr -20h
var_1C    = dword ptr -1Ch
var_4     = dword ptr -4

        push    ebp
        mov     ebp, esp
        and     esp, 0FFFFFFF0h
        sub     esp, 20h
        mov     [esp+20h+var_20], offset aEnterX ; "Enter X:"
        call    _puts
        mov     eax, offset aD ; "%d"
        lea     edx, [esp+20h+var_4]
        mov     [esp+20h+var_1C], edx
        mov     [esp+20h+var_20], eax
        call    ___isoc99_scanf
        mov     edx, [esp+20h+var_4]

```

```

mov     eax, offset aYouEnteredD___ ; "You entered %d...\n"
mov     [esp+20h+var_1C], edx
mov     [esp+20h+var_20], eax
call    _printf
mov     eax, 0
leave
retn
main    endp

```

GCC replaced first `printf()` call to `puts()`, it was already described [1.1.2](#) why it was done. As before — arguments are placed into stack by `MOV` instruction.

1.4.3 ARM

Optimizing Keil + thumb mode

```

.text:00000042      scanf_main
.text:00000042
.text:00000042      var_8          = -8
.text:00000042
.text:00000042 08 B5          PUSH     {R3,LR}
.text:00000044 A9 A0          ADR      R0, aEnterX      ; "Enter X:\n"
.text:00000046 06 F0 D3 F8      BL       __2printf
.text:0000004A 69 46          MOV      R1, SP
.text:0000004C AA A0          ADR      R0, aD          ; "%d"
.text:0000004E 06 F0 CD F8      BL       __0scanf
.text:00000052 00 99          LDR      R1, [SP,#8+var_8]
.text:00000054 A9 A0          ADR      R0, aYouEnteredD___ ; "You entered %d...\n"
.text:00000056 06 F0 CB F8      BL       __2printf
.text:0000005A 00 20          MOVS    R0, #0
.text:0000005C 08 BD          POP      {R3,PC}

```

A pointer to *int*-typed variable should be passed to `scanf()` so it can return value via it. *int* is 32-bit value, so we need 4 bytes for storing it somewhere in memory, and it fits exactly in 32-bit register. A place for local variable *x* is allocated in stack and IDA [5](#) named it `var_8`, however, it's not necessary to allocate it, because, SP stack pointer is already pointing to the place which may be used. So, SP stack pointer value is copied to R1 register and, together with format-string, passed into `scanf()`. Later, with the help of `LDR` instruction, this value is moved from stack into R1 register in order to be passed into `printf()`.

Examples compiled for ARM-mode and also examples compiled with Xcode LLVM are not differ significantly from what we saw here, so they are omitted.

1.4.4 Global variables

x86

What if *x* variable from previous example will not be local but global variable? Then it will be accessible from any place but not only from function body. It is not very good programming practice, but for the sake of experiment we could do this.

```

_DATA    SEGMENT
COMM     _x:DWORD
$SG2456  DB      'Enter X:', 0aH, 00H
         ORG     $+2
$SG2457  DB      '%d', 00H
         ORG     $+1
$SG2458  DB      'You entered %d...', 0aH, 00H
_DATA    ENDS
PUBLIC   _main
EXTRN    _scanf:PROC
EXTRN    _printf:PROC
; Function compile flags: /OdtP
_TEXT    SEGMENT

```

```

_main    PROC
    push    ebp
    mov     ebp, esp
    push    OFFSET $SG2456
    call    _printf
    add     esp, 4
    push    OFFSET _x
    push    OFFSET $SG2457
    call    _scanf
    add     esp, 8
    mov     eax, DWORD PTR _x
    push    eax
    push    OFFSET $SG2458
    call    _printf
    add     esp, 8
    xor     eax, eax
    pop     ebp
    ret     0
_main    ENDP
_TEXT    ENDS

```

Now `x` variable is defined in `_DATA` segment. Memory in local stack is not allocated anymore. All accesses to it are not via stack but directly to process memory. Its value is not defined. This mean that memory will be allocated by operation system, but not compiler, neither operation system will not take care about its initial value at the moment of `main()` function start. As experiment, try to declare large array and see what will it contain after program loading.

Now let's assign value to variable explicitly:

```
int x=10; // default value
```

We got:

```

_DATA    SEGMENT
_x       DD      0aH
...

```

Here we see value `0xA` of `DWORD` type (`DD` meaning `DWORD` = 32 bit).

If you will open compiled `.exe` in [IDA 5](#), you will see `x` placed at the beginning of `_DATA` segment, and after you'll see text strings.

If you will open compiled `.exe` in [IDA 5](#) from previous example where `x` value is not defined, you'll see something like this:

```

.data:0040FA80 _x          dd ?          ; DATA XREF: _main+10
.data:0040FA80          ; _main+22
.data:0040FA84 dword_40FA84 dd ?          ; DATA XREF: _memset+1E
.data:0040FA84          ; unknown_libname_1+28
.data:0040FA88 dword_40FA88 dd ?          ; DATA XREF: ___sbh_find_block+5
.data:0040FA88          ; ___sbh_free_block+2BC
.data:0040FA8C ; LPVOID lpMem
.data:0040FA8C lpMem      dd ?          ; DATA XREF: ___sbh_find_block+B
.data:0040FA8C          ; ___sbh_free_block+2CA
.data:0040FA90 dword_40FA90 dd ?          ; DATA XREF: _V6_HeapAlloc+13
.data:0040FA90          ; __calloc_impl+72
.data:0040FA94 dword_40FA94 dd ?          ; DATA XREF: ___sbh_free_block+2FE

```

`_x` marked as `?` among another variables not required to be initialized. This mean that after loading `.exe` to memory, place for all these variables will be allocated and some random garbage will be here. But in `.exe` file these not initialized variables are not occupy anything. It is suitable for large arrays, for example.

It is almost the same in Linux, except segment names and properties: not initialized variables are located in `_bss` segment. In [ELF³⁷](#) file format this segment has such attributes:

³⁷ Executable file format widely used in *NIX system including Linux

```
; Segment type: Uninitialized
; Segment permissions: Read/Write
```

If to assign some value to variable, e.g. 10, it will be placed in `_data` segment, this is segment with such attributes:

```
; Segment type: Pure data
; Segment permissions: Read/Write
```

ARM: Optimizing Keil + thumb mode

```
.text:00000000 ; Segment type: Pure code
.text:00000000      AREA .text, CODE
...
.text:00000000 main
.text:00000000      PUSH    {R4,LR}
.text:00000002      ADR     R0, aEnterX      ; "Enter X:\n"
.text:00000004      BL      __2printf
.text:00000008      LDR     R1, =x
.text:0000000A      ADR     R0, aD          ; "%d"
.text:0000000C      BL      __0scanf
.text:00000010      LDR     R0, =x
.text:00000012      LDR     R1, [R0]
.text:00000014      ADR     R0, aYouEnteredD___ ; "You entered %d...\n"
.text:00000016      BL      __2printf
.text:0000001A      MOVS    R0, #0
.text:0000001C      POP     {R4,PC}
...
.text:00000020 aEnterX      DCB "Enter X:",0xA,0      ; DATA XREF: main+2
.text:0000002A      DCB     0
.text:0000002B      DCB     0
.text:0000002C off_2C      DCD x                      ; DATA XREF: main+8
.text:0000002C                      ; main+10
.text:00000030 aD          DCB "%d",0                ; DATA XREF: main+A
.text:00000033      DCB     0
.text:00000034 aYouEnteredD___ DCB "You entered %d...",0xA,0 ; DATA XREF: main+14
.text:00000047      DCB     0
.text:00000047 ; .text      ends
.text:00000047
...
.data:00000048 ; Segment type: Pure data
.data:00000048      AREA .data, DATA
.data:00000048      ; ORG 0x48
.data:00000048      EXPORT x
.data:00000048 x          DCD 0xA                    ; DATA XREF: main+8
.data:00000048                      ; main+10
.data:00000048 ; .data      ends
```

So, `x` variable is now global and it's located, and, somehow, it's now located in other segment, namely data segment (`.data`). One could ask, why text strings are located in code segment (`.text`) and `x` can be located right here? Because, this is variable, and by its definition, it can be changed. And probably, can be changed very often.

Segment of code not infrequently can be located in microcontroller ROM (remember, we now deal with embedded microelectronics, and memory scarcity is common here), and changeable variables — in RAM. It's not very economically to store constant variables in RAM when one have ROM.

Onwards, we see, in code segment, a pointer to `x` (`off_2C`) variable, and all operations with variable occurred via this pointer. This is because `x` variable can be located somewhere far from this code fragment, so its address should be saved somewhere in close proximity to the code. `LDR` instruction in thumb mode can address only variable in range of 1020 bytes from the place where it's located. Same instruction in ARM-mode — variables in range ± 4095 bytes, this, address of `x` variable should be located somewhere in close proximity, because, there are no guarantee that linker will be able to place this variable near the code, it could be even in other memory chip!

One more thing: if variable will be declared as *const*, Keil compiler will allocate it in `.constdata` segment. Perhaps, thereafter, linker will be able to place this segment in ROM too, along with code segment.

1.4.5 scanf() result checking

x86

As I noticed before, it is slightly old-fashioned to use `scanf()` today. But if we have to, we need at least check if `scanf()` finished correctly without error.

```
int main()
{
    int x;
    printf ("Enter X:\n");

    if (scanf ("%d", &x)==1)
        printf ("You entered %d...\n", x);
    else
        printf ("What you entered? Huh?\n");

    return 0;
};
```

By standard, `scanf()`³⁸ function returning number of fields it successfully read.

In our case, if everything went fine and user entered some number, `scanf()` will return 1 or 0 or EOF in case of error.

I added C code for `scanf()` result checking and printing error message in case of error.

What we got in assembly language (MSVC 2010):

```
; Line 8
    lea    eax, DWORD PTR _x$[ebp]
    push   eax
    push   OFFSET $SG3833 ; '%d', 00H
    call   _scanf
    add    esp, 8
    cmp    eax, 1
    jne    SHORT $LN2@main
; Line 9
    mov    ecx, DWORD PTR _x$[ebp]
    push   ecx
    push   OFFSET $SG3834 ; 'You entered %d...', 0aH, 00H
    call   _printf
    add    esp, 8
; Line 10
    jmp    SHORT $LN1@main
$LN2@main:
; Line 11
    push   OFFSET $SG3836 ; 'What you entered? Huh?', 0aH, 00H
    call   _printf
    add    esp, 4
$LN1@main:
; Line 13
    xor    eax, eax
```

Caller function (`main()`) should have access to result of callee function (`scanf()`), so callee leave this value in EAX register.

After, we check it using instruction `CMP EAX, 1` (*CoMPare*), in other words, we compare value in EAX with 1.

JNE conditional jump follows `CMP` instruction. JNE mean *Jump if Not Equal*.

So, if EAX value not equals to 1, then the processor will pass execution to the address mentioned in operand of JNE, in our case it is `$LN2@main`. Passing control to this address, microprocessor will execute function `printf()` with argument "What you entered? Huh?". But if everything is fine, conditional jump will not be taken, and another `printf()` call will be executed, with two arguments: 'You entered %d...' and value of variable `x`.

³⁸MSDN: [scanf, wscanf](#)

Because second subsequent `printf()` not needed to be executed, there are `JMP` after (unconditional jump) it will pass control to the place after second `printf()` and before `XOR EAX, EAX` instruction, which implement `return 0`.

So, it can be said that most often, comparing some value with another is implemented by `CMP/Jcc` instructions pair, where *cc* is *condition code*. `CMP` comparing two values and set processor flags³⁹. `Jcc` check flags needed to be checked and pass control to mentioned address (or not pass).

But in fact, this could be perceived paradoxial, but `CMP` instruction is in fact `SUB` (subtract). All arithmetic instructions set processor flags too, not only `CMP`. If we compare 1 and 1, $1 - 1$ will be zero in result, `ZF` flag will be set (meaning that last result was zero). There are no any other circumstances when it's possible except when operands are equal. `JNE` checks only `ZF` flag and jumping only if it is not set. `JNE` is in fact a synonym of `JNZ` (*Jump if Not Zero*) instruction. Assembler translating both `JNE` and `JNZ` instructions into one single opcode. So, `CMP` can be replaced to `SUB` and almost everything will be fine, but the difference is in that `SUB` alter value at first operand. `CMP` is "*SUB without saving result*".

Code generated by GCC 4.4.1 in Linux is almost the same, except differences we already considered.

ARM: Optimizing Keil + thumb mode

Listing 1.12: Optimizing Keil + thumb mode

```
var_8      = -8

        PUSH    {R3,LR}
        ADR     R0, aEnterX      ; "Enter X:\n"
        BL      __2printf
        MOV     R1, SP
        ADR     R0, aD           ; "%d"
        BL      __0scanf
        CMP     R0, #1
        BEQ     loc_1E
        ADR     R0, aWhatYouEntered ; "What you entered? Huh?\n"
        BL      __2printf

loc_1A:
        ; CODE XREF: main+26
        MOVS    R0, #0
        POP     {R3,PC}

loc_1E:
        ; CODE XREF: main+12
        LDR     R1, [SP,#8+var_8]
        ADR     R0, aYouEnteredD___ ; "You entered %d...\n"
        BL      __2printf
        B       loc_1A
```

New instructions here are `CMP` and `BEQ`.

`CMP` is similar to the x86 instruction, it subtracts one argument from another and save flags.

`BEQ` (*Branch Equal*) is jumping to another address if operands while comparing were equal to each other, or, if result of last computation was zero, or if `Z` flag is 1. Same thing as `JZ` in x86.

Everything else is simple: execution flow is forking into two branches, then the branches are converging at the place where 0 is written into `R0`, as a value returned from the function, and then function finishing.

³⁹ About x86 flags, see also: [http://en.wikipedia.org/wiki/FLAGS_register_\(computing\)](http://en.wikipedia.org/wiki/FLAGS_register_(computing)).

1.5 Passing arguments via stack

Now we figured out that caller function passing arguments to callee via stack. But how callee⁴⁰ access them?

```
#include <stdio.h>

int f (int a, int b, int c)
{
    return a*b+c;
};

int main()
{
    printf ("%d\n", f(1, 2, 3));
    return 0;
};
```

1.5.1 x86

What we have after compilation (MSVC 2010 Express):

Listing 1.13: MSVC 2010 Express

```
_TEXT    SEGMENT
_a$ = 8                                     ; size = 4
_b$ = 12                                    ; size = 4
_c$ = 16                                    ; size = 4
_f      PROC
; File c:\...\1.c
; Line 4
    push    ebp
    mov     ebp, esp
; Line 5
    mov     eax, DWORD PTR _a$[ebp]
    imul    eax, DWORD PTR _b$[ebp]
    add     eax, DWORD PTR _c$[ebp]
; Line 6
    pop     ebp
    ret     0
_f      ENDP

_main    PROC
; Line 9
    push    ebp
    mov     ebp, esp
; Line 10
    push    3
    push    2
    push    1
    call    _f
    add     esp, 12                          ; 0000000cH
    push    eax
    push    OFFSET $SG2463 ; '%d', 0aH, 00H
    call    _printf
    add     esp, 8
; Line 11
    xor     eax, eax
; Line 12
    pop     ebp
    ret     0
_main    ENDP
```

What we see is that 3 numbers are pushing to stack in function main() and f(int,int,int) is called then. Argument access inside f() is organized with help of macros like: _a\$ = 8, in the same way as local variables

⁴⁰function being called

accessed, but difference in that these offsets are positive (addressed with *plus* sign). So, adding `_a$` macro to EBP register value, *outer* side of stack frame is addressed.

Then a value is stored into EAX. After IMUL instruction execution, EAX value is product⁴¹ of EAX and what is stored in `_b`. After IMUL execution, ADD is summing EAX and what is stored in `_c`. Value in EAX is not needed to be moved: it is already in place it need. Now return to caller — it will take value from EAX and used it as `printf()` argument.

Let's compile the same in GCC 4.4.1:

Listing 1.14: GCC 4.4.1

```

public f
f      proc near                ; CODE XREF: main+20
      arg_0      = dword ptr  8
      arg_4      = dword ptr  0Ch
      arg_8      = dword ptr  10h

      push      ebp
      mov       ebp, esp
      mov       eax, [ebp+arg_0]
      imul      eax, [ebp+arg_4]
      add       eax, [ebp+arg_8]
      pop       ebp
      retn
f      endp

public main
main   proc near                ; DATA XREF: _start+17
      var_10     = dword ptr -10h
      var_C      = dword ptr -0Ch
      var_8      = dword ptr -8

      push      ebp
      mov       ebp, esp
      and       esp, 0FFFFFFF0h
      sub       esp, 10h        ; char *
      mov       [esp+10h+var_8], 3
      mov       [esp+10h+var_C], 2
      mov       [esp+10h+var_10], 1
      call      f
      mov       edx, offset aD  ; "%d\n"
      mov       [esp+10h+var_C], eax
      mov       [esp+10h+var_10], edx
      call      _printf
      mov       eax, 0
      leave
      retn
main   endp

```

Almost the same result.

1.5.2 ARM

Non-optimizing Keil + ARM mode

```

.text:000000A4 00 30 A0 E1      MOV     R3, R0
.text:000000A8 93 21 20 E0      MLA     R0, R3, R1, R2
.text:000000AC 1E FF 2F E1      BX      LR
...
.text:000000B0                main
.text:000000B0 10 40 2D E9      STMFD   SP!, {R4,LR}

```

⁴¹result of multiplication

```

.text:000000B4 03 20 A0 E3      MOV     R2, #3
.text:000000B8 02 10 A0 E3      MOV     R1, #2
.text:000000BC 01 00 A0 E3      MOV     R0, #1
.text:000000C0 F7 FF FF EB      BL      f
.text:000000C4 00 40 A0 E1      MOV     R4, R0
.text:000000C8 04 10 A0 E1      MOV     R1, R4
.text:000000CC 5A 0F 8F E2      ADR     R0, aD_0          ; "%d\n"
.text:000000D0 E3 18 00 EB      BL      __2printf
.text:000000D4 00 00 A0 E3      MOV     R0, #0
.text:000000D8 10 80 BD E8      LDMFD   SP!, {R4,PC}

```

In `main()` function, two other functions are simply called, and three values are passed to the first one (`f`).

As I mentioned before, in ARM, first 4 values are usually passed in first 4 registers (R0-R3).

`f` function, as it seems, use first 3 registers (R0-R2) as arguments.

MLA (*Multiply Accumulate*) instruction multiplies two first operands (R3 and R1), adds third operand (R2) to product and places result into zeroth operand (R0), via which, by standard, values are returned from functions.

Multiplication and addition at once⁴² (*Fused multiply-add*) is very useful operation, by the way, there are no such instruction in x86, if not to count new FMA-instruction⁴³ in SIMD.

The very first `MOV R3, R0`, instruction, apparently, redundant (single MLA instruction could be used here instead), compiler wasn't optimized it, because, this is non-optimizing compilation.

BX instruction returns control to the address stored in LR and, if need, switches processor mode from thumb to ARM or vice versa. This can be necessary because, as we can see, `f` function is not aware, from which code it may be called, from ARM or thumb. This, if it will be called from thumb code, BX will not only return control to the calling function, but also will switch processor mode to thumb mode. Or not switch, if the function was called from ARM code.

Optimizing Keil + ARM mode

```

.text:00000098          f
.text:00000098 91 20 20 E0      MLA     R0, R1, R0, R2
.text:0000009C 1E FF 2F E1      BX      LR

```

And here is `f` function compiled by Keil compiler in full optimization mode (`-O3`). `MOV` instruction was optimized and now `MLA` uses all input registers and place result into R0, exactly where calling function will read it and use.

Optimizing Keil + thumb mode

```

.text:0000005E 48 43      MULS    R0, R1
.text:00000060 80 18      ADDS    R0, R0, R2
.text:00000062 70 47      BX      LR

```

MLA instruction is not available in thumb mode, so, compiler generates the code doing these two operations separately. First `MULS` instruction multiply R0 by R1 leaving result in R1. Second (`ADDS`) instruction adds result and R2 leaving result in R0.

⁴²[wikipedia: Multiply-accumulate operation](https://en.wikipedia.org/wiki/Multiply-accumulate_operation)

⁴³https://en.wikipedia.org/wiki/FMA_instruction_set

1.6 One more word about results returning.

As of x86, function execution result is usually returned⁴⁴ in EAX register. If it's byte type or character (*char*) — then in lowest register EAX part — AL. If function returning *float* number, FPU register ST(0) will be used instead. In ARM, result is usually returned in R0 register.

That is why old C compilers can't create functions capable of returning something not fitting in one register (usually type *int*), but if one need it, one should return information via pointers passed in function arguments. Now it is possible, to return, let's say, whole structure, but its still not very popular. If function should return a large structure, caller must allocate it and pass pointer to it via first argument, hiddenly and transparently for programmer. That is almost the same as to pass pointer in first argument manually, but compiler hide this.

Small example:

```
struct s
{
    int a;
    int b;
    int c;
};

struct s get_some_values (int a)
{
    struct s rt;

    rt.a=a+1;
    rt.b=a+2;
    rt.c=a+3;

    return rt;
};
```

...what we got (MSVC 2010 /Ox):

```
$T3853 = 8                ; size = 4
_a$ = 12                  ; size = 4
?get_some_values@@YA?AUs@@H@Z PROC                ; get_some_values
    mov     ecx, DWORD PTR _a$[esp-4]
    mov     eax, DWORD PTR $T3853[esp-4]
    lea     edx, DWORD PTR [ecx+1]
    mov     DWORD PTR [eax], edx
    lea     edx, DWORD PTR [ecx+2]
    add     ecx, 3
    mov     DWORD PTR [eax+4], edx
    mov     DWORD PTR [eax+8], ecx
    ret     0
?get_some_values@@YA?AUs@@H@Z ENDP                ; get_some_values
```

Macro name for internal variable passing pointer to structure is \$T3853 here.

⁴⁴ See also: [MSDN: Return Values \(C++\)](#)

1.7 Pointers

Pointers are often used to return values from function (recall `scanf()` case 1.4). For example, when function should return two values:

```
void f1 (int x, int y, int *sum, int *product)
{
    *sum=x+y;
    *product=x*y;
};

void main()
{
    int sum, product;

    f1(123, 456, &sum, &product);
    printf ("sum=%d, product=%d\n", sum, product);
};
```

This compiling into:

Listing 1.15: Optimizing MSVC 2010

```
CONST    SEGMENT
$SG3863 DB      'sum=%d, product=%d', 0aH, 00H
$SG3864 DB      'sum=%d, product=%d', 0aH, 00H
CONST    ENDS
_TEXT    SEGMENT
_x$ = 8                                ; size = 4
_y$ = 12                               ; size = 4
_sum$ = 16                             ; size = 4
_product$ = 20                         ; size = 4
f1 PROC                                ; f1
    mov     ecx, DWORD PTR _y$[esp-4]
    mov     eax, DWORD PTR _x$[esp-4]
    lea     edx, DWORD PTR [eax+ecx]
    imul    eax, ecx
    mov     ecx, DWORD PTR _product$[esp-4]
    push    esi
    mov     esi, DWORD PTR _sum$[esp]
    mov     DWORD PTR [esi], edx
    mov     DWORD PTR [ecx], eax
    pop     esi
    ret     0
f1 ENDP                                ; f1

_product$ = -8                         ; size = 4
_sum$ = -4                             ; size = 4
_main    PROC
    sub     esp, 8
    lea     eax, DWORD PTR _product$[esp+8]
    push    eax
    lea     ecx, DWORD PTR _sum$[esp+12]
    push    ecx
    push    456                        ; 000001c8H
    push    123                        ; 0000007bH
    call    f1                        ; f1
    mov     edx, DWORD PTR _product$[esp+24]
    mov     eax, DWORD PTR _sum$[esp+24]
    push    edx
    push    eax
    push    OFFSET $SG3863
    call    _printf
    ...
```

1.7.1 C++ references

In C++, references are pointers as well, but they are called *safe*, because it's harder to make a mistake while working with them [ISO13, 8.3.2]. For example, reference should be always be pointing to the object of corresponding type and can't be NULL [Cli, 8.6]. Even more, reference cannot be changed, it's not possible to point to to another object (reseat) [Cli, 8.5].

If we will try to change are example to use references instead of pointers:

```
void f2 (int x, int y, int & sum, int & product)
{
    sum=x+y;
    product=x*y;
};
```

Then we'll figure out that compiled code is just the same:

Listing 1.16: Optimizing MSVC 2010

```
_x$ = 8 ; size = 4
_y$ = 12 ; size = 4
_sum$ = 16 ; size = 4
_product$ = 20 ; size = 4
?f2@@YAXHHAH0@Z PROC ; f2
    mov     ecx, DWORD PTR _y$[esp-4]
    mov     eax, DWORD PTR _x$[esp-4]
    lea     edx, DWORD PTR [eax+ecx]
    imul    eax, ecx
    mov     ecx, DWORD PTR _product$[esp-4]
    push    esi
    mov     esi, DWORD PTR _sum$[esp]
    mov     DWORD PTR [esi], edx
    mov     DWORD PTR [ecx], eax
    pop     esi
    ret     0
?f2@@YAXHHAH0@Z ENDP ; f2
```

(A reason why C++ functions has such strange names, will be described later 1.17.1.)

1.8 Conditional jumps

Now about conditional jumps.

```
void f_signed (int a, int b)
{
    if (a>b)
        printf ("a>b\n");
    if (a==b)
        printf ("a==b\n");
    if (a<b)
        printf ("a<b\n");
};

void f_unsigned (unsigned int a, unsigned int b)
{
    if (a>b)
        printf ("a>b\n");
    if (a==b)
        printf ("a==b\n");
    if (a<b)
        printf ("a<b\n");
};

int main()
{
    f_signed(1, 2);
    f_unsigned(1, 2);
    return 0;
};
```

1.8.1 x86

x86 + MSVC

What we have in `f_signed()` function:

Listing 1.17: MSVC

```
_a$ = 8 ; size = 4
_b$ = 12 ; size = 4
_f_signed PROC
    push    ebp
    mov     ebp, esp
    mov     eax, DWORD PTR _a$[ebp]
    cmp     eax, DWORD PTR _b$[ebp]
    jle     SHORT $LN3@f_signed
    push    OFFSET $SG737 ; 'a>b', 0aH, 00H
    call    _printf
    add     esp, 4
$LN3@f_signed:
    mov     ecx, DWORD PTR _a$[ebp]
    cmp     ecx, DWORD PTR _b$[ebp]
    jne     SHORT $LN2@f_signed
    push    OFFSET $SG739 ; 'a==b', 0aH, 00H
    call    _printf
    add     esp, 4
$LN2@f_signed:
    mov     edx, DWORD PTR _a$[ebp]
    cmp     edx, DWORD PTR _b$[ebp]
    jge     SHORT $LN4@f_signed
    push    OFFSET $SG741 ; 'a<b', 0aH, 00H
    call    _printf
    add     esp, 4
$LN4@f_signed:
```

```

    pop    ebp
    ret    0
_f_signed ENDP

```

First instruction JLE mean *Jump if Larger or Equal*. In other words, if second operand is larger than first or equal, control flow will be passed to address or label mentioned in instruction. But if this condition will not trigger (second operand less than first), control flow will not be altered and first `printf()` will be called. The second check is JNE: *Jump if Not Equal*. Control flow will not altered if operands are equals to each other. The third check is JGE: *Jump if Greater or Equal* — jump if second operand is larger than first or they are equals to each other. By the way, if all three conditional jumps are triggered, no `printf()` will be called whatsoever. But, without special intervention, it is nearly impossible.

`f_unsigned()` function is the same, with the exception that JBE and JAE instructions are used here instead of JLE and JGE, see below about it:

GCC

GCC 4.4.1 produce almost the same code, but with `puts()` [1.1.2](#) instead of `printf()`.

Now let's take a look of `f_unsigned()` produced by GCC:

Listing 1.18: GCC

```

.globl f_unsigned
.type    f_unsigned, @function
f_unsigned:
    push    ebp
    mov     ebp, esp
    sub     esp, 24
    mov     eax, DWORD PTR [ebp+8]
    cmp     eax, DWORD PTR [ebp+12]
    jbe     .L7
    mov     DWORD PTR [esp], OFFSET FLAT:.LC0 ; "a>b"
    call    puts
.L7:
    mov     eax, DWORD PTR [ebp+8]
    cmp     eax, DWORD PTR [ebp+12]
    jne     .L8
    mov     DWORD PTR [esp], OFFSET FLAT:.LC1 ; "a==b"
    call    puts
.L8:
    mov     eax, DWORD PTR [ebp+8]
    cmp     eax, DWORD PTR [ebp+12]
    jae     .L10
    mov     DWORD PTR [esp], OFFSET FLAT:.LC2 ; "a<b"
    call    puts
.L10:
    leave
    ret

```

Almost the same, with exception of instructions: JBE — *Jump if Below or Equal* and JAE — *Jump if Above or Equal*. These instructions (JA/JAE/JBE/JBE) are distinct from JG/JGE/JL/JLE in that way, they works with unsigned numbers.

See also section about signed number representations [2.4](#). So, where we see usage of JG/JL instead of JA/JBE or otherwise, we can almost be sure about signed or unsigned type of variable.

Here is also `main()` function, where nothing new to us:

Listing 1.19: `main()`

```

main:
    push    ebp
    mov     ebp, esp
    and     esp, -16
    sub     esp, 16
    mov     DWORD PTR [esp+4], 2

```



```

mov     DWORD PTR [esp], 1
call    f_signed
mov     DWORD PTR [esp+4], 2
mov     DWORD PTR [esp], 1
call    f_unsigned
mov     eax, 0
leave
ret

```

1.8.2 ARM

Optimizing Keil + ARM mode

Listing 1.20: Optimizing Keil + ARM mode

```

.text:000000B8          EXPORT f_signed
.text:000000B8          f_signed          ; CODE XREF: main+C
.text:000000B8 70 40 2D E9          STMFD     SP!, {R4-R6,LR}
.text:000000BC 01 40 A0 E1          MOV      R4, R1
.text:000000C0 04 00 50 E1          CMP      R0, R4
.text:000000C4 00 50 A0 E1          MOV      R5, R0
.text:000000C8 1A 0E 8F C2          ADRGT    R0, aAB          ; "a>b\n"
.text:000000CC A1 18 00 CB          BLGT     __2printf
.text:000000D0 04 00 55 E1          CMP      R5, R4
.text:000000D4 67 0F 8F 02          ADREQ    R0, aAB_0        ; "a==b\n"
.text:000000D8 9E 18 00 0B          BLEQ     __2printf
.text:000000DC 04 00 55 E1          CMP      R5, R4
.text:000000E0 70 80 BD A8          LDMGEFD  SP!, {R4-R6,PC}
.text:000000E4 70 40 BD E8          LDMFD     SP!, {R4-R6,LR}
.text:000000E8 19 0E 8F E2          ADR      R0, aAB_1        ; "a<b\n"
.text:000000EC 99 18 00 EA          B        __2printf
.text:000000EC          ; End of function f_signed

```

A lot of instructions in ARM mode can be executed only when specific flags are set. This is often used while numbers comparing, for example.

For instance, ADD instruction is ADDAL internally in fact, where AL meaning *Always*, i.e., execute always. Predicates are encoded in 4 high bits of 32-bit ARM instructions (*condition field*). B instruction of unconditional jump is in fact conditional and encoded just like any other conditional jumps, but has AL in *condition field*, and what it means, executing always, ignoring flags.

ADRG T instructions works just like ADR, but will execute only in the case when previous CMP instruction, while comparing two numbers, found one number greater than another (*Greater Than*).

The next BLGT instruction behaves exactly as BL and will be triggered only if result of comparison was the same (*Greater Than*). ADRGT writes a pointer to the string “a>b\n”, into R0 and BLGT calls printf(). Consequently, these instructions with -GT suffix, will be executed only in the case when value in R0 (*a* is there) was bigger than value in R4 (*b* is there).

Then we see ADREQ and BLEQ instructions. They behave just like ADR and BL, but will be executed only in the case when operands were equal to each other while comparison. Another CMP is before them (because printf() call may tamper state of flags).

Then we see LDMGEFD, this instruction works just like LDMFD⁴⁵, but will be triggered only in the case when one value was greater or equal to another while comparison (*Greater or Equal*).

The sense of “LDMGEFD SP!, {R4-R6,PC}” instruction is that is like function epilogue, but it will be triggered only if $a \geq b$, only then function execution will be finished. But if it's not true, i.e., $a < b$, then control flow come to next “LDMFD SP!, {R4-R6,LR}” instruction, this is one more function epilogue, this instruction restores R4-R6 registers state, but also LR instead of PC, thus, it doesn't returns from function. Last two instructions calls printf() with the string «a<b\n» as sole argument. Unconditional jump to printf() instead of function return, is what we already examined in «printf() with several arguments» section, here 1.3.2.

⁴⁵Load Multiple Full Descending

f_unsigned is similar, but ADRHI, BLHI, and LDMCSFD instructions are used there, these predicates (*HI = Unsigned higher, CS = Carry Set (greater than or equal)*) are analogical to those examined before, but serving for unsigned values.

There are not much new in main() for us:

Listing 1.21: main()

```
.text:00000128          EXPORT main
.text:00000128          main
.text:00000128 10 40 2D E9      STMFD    SP!, {R4,LR}
.text:0000012C 02 10 A0 E3      MOV     R1, #2
.text:00000130 01 00 A0 E3      MOV     R0, #1
.text:00000134 DF FF FF EB      BL      f_signed
.text:00000138 02 10 A0 E3      MOV     R1, #2
.text:0000013C 01 00 A0 E3      MOV     R0, #1
.text:00000140 EA FF FF EB      BL      f_unsigned
.text:00000144 00 00 A0 E3      MOV     R0, #0
.text:00000148 10 80 BD E8      LDMFD    SP!, {R4,PC}
.text:00000148          ; End of function main
```

That's how to get rid of conditional jumps in ARM mode.

Why it's so good? Because, ARM is RISC-processor with pipeline for instructions executing. In short, pipelined processor is not very good on jumps at all, so that's why branch predictors are critical here. It's very good if the program has as few jumps as possible, conditional and unconditional, so that's why, predicated instructions can help in reducing conditional jumps count.

There are no such feature in x86, if not to count CMOVcc instruction, it's the same as MOV, but it's triggered only when specific flags are set, usually, set while comparing by CMP.

Optimizing Keil + thumb mode

Listing 1.22: Optimizing Keil + thumb mode

```
.text:00000072          f_signed          ; CODE XREF: main+6
.text:00000072 70 B5          PUSH     {R4-R6,LR}
.text:00000074 0C 00          MOVS    R4, R1
.text:00000076 05 00          MOVS    R5, R0
.text:00000078 A0 42          CMP     R0, R4
.text:0000007A 02 DD          BLE     loc_82
.text:0000007C A4 A0          ADR     R0, aAB          ; "a>b\n"
.text:0000007E 06 F0 B7 F8      BL      __2printf
.text:00000082          loc_82          ; CODE XREF: f_signed+8
.text:00000082 A5 42          CMP     R5, R4
.text:00000084 02 D1          BNE     loc_8C
.text:00000086 A4 A0          ADR     R0, aAB_0        ; "a==b\n"
.text:00000088 06 F0 B2 F8      BL      __2printf
.text:0000008C          loc_8C          ; CODE XREF: f_signed+12
.text:0000008C A5 42          CMP     R5, R4
.text:0000008E 02 DA          BGE     locret_96
.text:00000090 A3 A0          ADR     R0, aAB_1        ; "a<b\n"
.text:00000092 06 F0 AD F8      BL      __2printf
.text:00000096          locret_96        ; CODE XREF: f_signed+1C
.text:00000096 70 BD          POP     {R4-R6,PC}
.text:00000096          ; End of function f_signed
```

Only B instructions in thumb mode may be supplemented by *condition codes*, so the thumb code looks more ordinary.

BLE is usual conditional jump *Less than or Equal*, BNE — *Not Equal*, BGE — *Greater than or Equal*.

f_unsigned function is just the same, but other instructions are used while working with unsigned values: BLS (*Unsigned lower or same*) and BCS (*Carry Set (Greater than or equal)*).

1.9 switch()/case/default

1.9.1 Few number of cases

```
void f (int a)
{
    switch (a)
    {
        case 0: printf ("zero\n"); break;
        case 1: printf ("one\n"); break;
        case 2: printf ("two\n"); break;
        default: printf ("something unknown\n"); break;
    };
};
```

x86

Result (MSVC 2010):

Listing 1.23: MSVC 2010

```
tv64 = -4                ; size = 4
_a$ = 8                  ; size = 4
_f    PROC
    push    ebp
    mov     ebp, esp
    push    ecx
    mov     eax, DWORD PTR _a$[ebp]
    mov     DWORD PTR tv64[ebp], eax
    cmp     DWORD PTR tv64[ebp], 0
    je      SHORT $LN4@f
    cmp     DWORD PTR tv64[ebp], 1
    je      SHORT $LN3@f
    cmp     DWORD PTR tv64[ebp], 2
    je      SHORT $LN2@f
    jmp     SHORT $LN1@f
$LN4@f:
    push    OFFSET $SG739 ; 'zero', 0aH, 00H
    call    _printf
    add     esp, 4
    jmp     SHORT $LN7@f
$LN3@f:
    push    OFFSET $SG741 ; 'one', 0aH, 00H
    call    _printf
    add     esp, 4
    jmp     SHORT $LN7@f
$LN2@f:
    push    OFFSET $SG743 ; 'two', 0aH, 00H
    call    _printf
    add     esp, 4
    jmp     SHORT $LN7@f
$LN1@f:
    push    OFFSET $SG745 ; 'something unknown', 0aH, 00H
    call    _printf
    add     esp, 4
$LN7@f:
    mov     esp, ebp
    pop     ebp
    ret     0
_f    ENDP
```

Out function with a few cases in switch(), in fact, is analogous to this construction:

```
void f (int a)
{
```

```

if (a==0)
    printf ("zero\n");
else if (a==1)
    printf ("one\n");
else if (a==2)
    printf ("two\n");
else
    printf ("something unknown\n");
};

```

When few cases in `switch()`, and we see such code, it's impossible to say with certainty, was it `switch()` in source code, or just pack of `if()`. This mean, `switch()` is syntactic sugar for large number of nested checks constructed using `if()`.

Nothing specially new to us in generated code, with the exception that compiler moving input variable `a` to temporary local variable `tv64`.

If to compile the same in GCC 4.4.1, we'll get almost the same, even with maximal optimization turned on (`-O3` option).

Now let's turn on optimization in MSVC (`/Ox`): `cl 1.c /Fa1.asm /Ox`

Listing 1.24: MSVC

```

_a$ = 8 ; size = 4
_f PROC
    mov     eax, DWORD PTR _a$[esp-4]
    sub     eax, 0
    je      SHORT $LN4@f
    sub     eax, 1
    je      SHORT $LN3@f
    sub     eax, 1
    je      SHORT $LN2@f
    mov     DWORD PTR _a$[esp-4], OFFSET $SG791 ; 'something unknown', 0aH, 00H
    jmp     _printf
$LN2@f:
    mov     DWORD PTR _a$[esp-4], OFFSET $SG789 ; 'two', 0aH, 00H
    jmp     _printf
$LN3@f:
    mov     DWORD PTR _a$[esp-4], OFFSET $SG787 ; 'one', 0aH, 00H
    jmp     _printf
$LN4@f:
    mov     DWORD PTR _a$[esp-4], OFFSET $SG785 ; 'zero', 0aH, 00H
    jmp     _printf
_f ENDP

```

Here we can see even dirty hacks.

First: `a` is placed into `EAX` and `0` subtracted from it. Sounds absurdly, but it may need to check if `0` was in `EAX` before? If yes, flag `ZF` will be set (this also mean that subtracting from zero is zero) and first conditional jump `JE` (*Jump if Equal* or synonym `JZ` — *Jump if Zero*) will be triggered and control flow passed to `$LN4@f` label, where 'zero' message is begin printed. If first jump was not triggered, `1` subtracted from input value and if at some stage `0` will be resulted, corresponding jump will be triggered.

And if no jump triggered at all, control flow passed to `printf()` with argument 'something unknown'.

Second: we see unusual thing for us: string pointer is placed into a variable, and then `printf()` is called not via `CALL`, but via `JMP`. This could be explained simply. Caller pushing to stack some value and via `CALL` calling our function. `CALL` itself pushing returning address to stack and do unconditional jump to our function address. Our function at any place of its execution (since it do not contain any instruction moving stack pointer) has the following stack layout:

- `ESP` — pointing to return address
- `ESP+4` — pointing to a variable

On the other side, when we need to call `printf()` here, we need exactly the same stack layout, except of first `printf()` argument pointing to string. And that is what our code does.

It replaces function's first argument to different and jumping to `printf()`, as if not our function `f()` was called firstly, but immediately `printf()`. `printf()` printing some string to `stdout` and then execute `RET` instruction, which POPping return address from stack and control flow is returned not to `f()`, but to `f()`'s callee, escaping `f()`.

All it's possible because `printf()` is called right at the end of `f()` in any case. In some way, it's all similar to `longjmp()`⁴⁶. And of course, it's all done for the sake of speed.

Similar case with ARM compiler described in "printf() with several arguments", section, here 1.3.2.

ARM: Optimizing Keil + ARM mode

```
.text:0000014C          f1
.text:0000014C 00 00 50 E3      CMP     R0, #0
.text:00000150 13 0E 8F 02      ADREQ   R0, aZero      ; "zero\n"
.text:00000154 05 00 00 0A      BEQ     loc_170
.text:00000158 01 00 50 E3      CMP     R0, #1
.text:0000015C 4B 0F 8F 02      ADREQ   R0, aOne      ; "one\n"
.text:00000160 02 00 00 0A      BEQ     loc_170
.text:00000164 02 00 50 E3      CMP     R0, #2
.text:00000168 4A 0F 8F 12      ADNE    R0, aSomethingUnkno ; "something unknown\n"
.text:0000016C 4E 0F 8F 02      ADREQ   R0, aTwo      ; "two\n"
.text:00000170
.text:00000170          loc_170          ; CODE XREF: f1+8
.text:00000170          ; f1+14
.text:00000170 78 18 00 EA      B       __2printf
```

Again, by investigating this code, we cannot say, was it `switch()` in the original source code, or pack of `if()` statements.

Anyway, we see here predicated instructions again, like `ADREQ` (*Equal*), which will be triggered only in $R0 = 0$ case, and the, address of «zero\n» string will be loaded into `R0`. The next instruction `BEQ` (*Branch Equal*) will redirect control flow to `loc_170`, if $R0 = 0$. By the way, observant reader may ask, will `BEQ` triggered right, because `ADREQ` before it is already filled `R0` register with some other value. Yes, it will, because `BEQ` checking flags set by `CMP` instruction, and `ADREQ` not modifying flags at all.

By the way, there are `-S` suffix for some instructions in ARM, indicative that instruction will not modify flags. For example `ADDS` Will add two number, but flags will not be touched. Such instructions are convenient to use between `CMP` where flags are set and, for example, conditional jumps, where flags are used.

Other instructions are already familiar to us. There are only one call to `printf()`, at the end, and we already examined this trick here 1.3.2. There are three paths to `printf()` at the end.

Also pay attention to what is going on if $a = 2$ and if a is not in range of constants it's comparing against. "CMP `R0, #2`" instruction is needed here to know, if $a = 2$ or not. If it's not true, then `ADNE` will load pointer to the string «something unknown\n» into `R0`, because, a was already checked before to be equal to 0 or 1, so we can be assured that a is not equal to these numbers at this point. And if $R0 = 2$, a pointer to string «two\n» will be loaded by `ADREQ` into `R0`.

ARM: Optimizing Keil + thumb mode

```
.text:000000D4          f1
.text:000000D4 10 B5      PUSH    {R4,LR}
.text:000000D6 00 28      CMP     R0, #0
.text:000000D8 05 D0      BEQ     zero_case
.text:000000DA 01 28      CMP     R0, #1
.text:000000DC 05 D0      BEQ     one_case
.text:000000DE 02 28      CMP     R0, #2
.text:000000E0 05 D0      BEQ     two_case
.text:000000E2 91 A0      ADR     R0, aSomethingUnkno ; "something unknown\n"
.text:000000E4 04 E0      B       default_case
.text:000000E6          ; -----
.text:000000E6          zero_case          ; CODE XREF: f1+4
```

⁴⁶<http://en.wikipedia.org/wiki/Setjmp.h>

```

.text:000000E6 95 A0          ADR     R0, aZero      ; "zero\n"
.text:000000E8 02 E0          B       default_case
.text:000000EA          ; -----
.text:000000EA          one_case          ; CODE XREF: f1+8
.text:000000EA 96 A0          ADR     R0, aOne      ; "one\n"
.text:000000EC 00 E0          B       default_case
.text:000000EE          ; -----
.text:000000EE          two_case          ; CODE XREF: f1+C
.text:000000EE 97 A0          ADR     R0, aTwo      ; "two\n"
.text:000000F0          default_case      ; CODE XREF: f1+10
.text:000000F0          ; f1+14
.text:000000F0 06 F0 7E F8      BL      __2printf
.text:000000F4 10 BD          POP     {R4,PC}
.text:000000F4          ; End of function f1

```

As I already mentioned, there are no feature of *connecting* predicates to majority of functions in thumb mode, so the thumb-code here is somewhat like the x86 code, easily understandable.

1.9.2 A lot of cases

If `switch()` statement contain a lot of case's, it is not very convenient for compiler to emit too large code with a lot JE/JNE instructions.

```

void f (int a)
{
    switch (a)
    {
        case 0: printf ("zero\n"); break;
        case 1: printf ("one\n"); break;
        case 2: printf ("two\n"); break;
        case 3: printf ("three\n"); break;
        case 4: printf ("four\n"); break;
        default: printf ("something unknown\n"); break;
    };
};

```

x86

We got (MSVC 2010):

Listing 1.25: MSVC 2010

```

tv64 = -4          ; size = 4
_a$ = 8           ; size = 4
_f PROC
    push     ebp
    mov     ebp, esp
    push     ecx
    mov     eax, DWORD PTR _a$[ebp]
    mov     DWORD PTR tv64[ebp], eax
    cmp     DWORD PTR tv64[ebp], 4
    ja      SHORT $LN10f
    mov     ecx, DWORD PTR tv64[ebp]
    jmp     DWORD PTR $LN110f[ecx*4]
$LN60f:
    push     OFFSET $SG739 ; 'zero', 0aH, 00H
    call     _printf
    add     esp, 4
    jmp     SHORT $LN90f
$LN50f:
    push     OFFSET $SG741 ; 'one', 0aH, 00H
    call     _printf
    add     esp, 4
    jmp     SHORT $LN90f

```

```

$LN4@f:
    push    OFFSET $SG743 ; 'two', 0aH, 00H
    call    _printf
    add     esp, 4
    jmp     SHORT $LN9@f
$LN3@f:
    push    OFFSET $SG745 ; 'three', 0aH, 00H
    call    _printf
    add     esp, 4
    jmp     SHORT $LN9@f
$LN2@f:
    push    OFFSET $SG747 ; 'four', 0aH, 00H
    call    _printf
    add     esp, 4
    jmp     SHORT $LN9@f
$LN1@f:
    push    OFFSET $SG749 ; 'something unknown', 0aH, 00H
    call    _printf
    add     esp, 4
$LN9@f:
    mov     esp, ebp
    pop     ebp
    ret     0
    npad    2
$LN11@f:
    DD      $LN6@f ; 0
    DD      $LN5@f ; 1
    DD      $LN4@f ; 2
    DD      $LN3@f ; 3
    DD      $LN2@f ; 4
_f      ENDP

```

OK, what we see here is: there are a set of `printf()` calls with various arguments. All them has not only addresses in process memory, but also internal symbolic labels assigned by compiler. Besides, all these labels are also places into `$LN11@f` internal table.

At the function beginning, if `a` is greater than 4, control flow is passed to label `$LN1@f`, where `printf()` with argument 'something unknown' is called.

And if `a` value is less or equals to 4, let's multiply it by 4 and add `$LN1@f` table address. That is how address inside of table is constructed, pointing exactly to the element we need. For example, let's say `a` is equal to 2. $2 * 4 = 8$ (all table elements are addresses within 32-bit process, that is why all elements contain 4 bytes). Address of `$LN11@f` table + 8 = it will be table element where `$LN4@f` label is stored. JMP fetch `$LN4@f` address from the table and jump to it.

This table called sometimes *jumptable*.

Then corresponding `printf()` is called with argument 'two'. Literally, `jmp DWORD PTR $LN11@f[ecx*4]` instruction mean *jump to DWORD, which is stored at address `$LN11@f + ecx * 4`*.

`npad 2.3` is assembly language macro, aligning next label so that it will be stored at address aligned by 4 bytes (or 16). This is very suitable for processor, because it can fetch 32-bit values from memory through memory bus, cache memory, etc, in much effective way if it is aligned.

Let's see what GCC 4.4.1 generate:

Listing 1.26: GCC 4.4.1

```

f      public f
      proc near          ; CODE XREF: main+10

var_18 = dword ptr -18h
arg_0  = dword ptr  8

      push    ebp
      mov     ebp, esp
      sub     esp, 18h    ; char *
      cmp     [ebp+arg_0], 4
      ja      short loc_8048444

```

```

mov     eax, [ebp+arg_0]
shl     eax, 2
mov     eax, ds:off_804855C[eax]
jmp     eax

loc_80483FE:                                ; DATA XREF: .rodata:off_804855C
mov     [esp+18h+var_18], offset aZero ; "zero"
call    _puts
jmp     short locret_8048450

loc_804840C:                                ; DATA XREF: .rodata:08048560
mov     [esp+18h+var_18], offset aOne ; "one"
call    _puts
jmp     short locret_8048450

loc_804841A:                                ; DATA XREF: .rodata:08048564
mov     [esp+18h+var_18], offset aTwo ; "two"
call    _puts
jmp     short locret_8048450

loc_8048428:                                ; DATA XREF: .rodata:08048568
mov     [esp+18h+var_18], offset aThree ; "three"
call    _puts
jmp     short locret_8048450

loc_8048436:                                ; DATA XREF: .rodata:0804856C
mov     [esp+18h+var_18], offset aFour ; "four"
call    _puts
jmp     short locret_8048450

loc_8048444:                                ; CODE XREF: f+A
mov     [esp+18h+var_18], offset aSomethingUnkno ; "something unknown"
call    _puts

locret_8048450:                                ; CODE XREF: f+26
                                                ; f+34...
leave
retn
f
endp

off_804855C  dd offset loc_80483FE  ; DATA XREF: f+12
              dd offset loc_804840C
              dd offset loc_804841A
              dd offset loc_8048428
              dd offset loc_8048436

```

It is almost the same, except little nuance: argument `arg_0` is multiplied by 4 with by shifting it to left by 2 bits (it is almost the same as multiplication by 4) [1.15.3](#). Then label address is taken from `off_804855C` array, address calculated and stored into `EAX`, then “`JMP EAX`” do actual jump.

ARM: Optimizing Keil + ARM mode

```

00000174          f2
00000174 05 00 50 E3          CMP     R0, #5          ; switch 5 cases
00000178 00 F1 8F 30          ADDCC   PC, PC, R0,LSL#2 ; switch jump
0000017C 0E 00 00 EA          B       default_case ; jumptable 00000178 default case
00000180          ; -----
00000180
00000180          loc_180          ; CODE XREF: f2+4
00000180 03 00 00 EA          B       zero_case ; jumptable 00000178 case 0
00000184          ; -----
00000184
00000184          loc_184          ; CODE XREF: f2+4
00000184 04 00 00 EA          B       one_case ; jumptable 00000178 case 1
00000188          ; -----

```



```

00000188
00000188      loc_188                      ; CODE XREF: f2+4
00000188 05 00 00 EA      B      two_case      ; jumptable 00000178 case 2
0000018C      ; -----
0000018C
0000018C      loc_18C                      ; CODE XREF: f2+4
0000018C 06 00 00 EA      B      three_case     ; jumptable 00000178 case 3
00000190      ; -----
00000190
00000190      loc_190                      ; CODE XREF: f2+4
00000190 07 00 00 EA      B      four_case     ; jumptable 00000178 case 4
00000194      ; -----
00000194
00000194      zero_case                      ; CODE XREF: f2+4
00000194      ; f2:loc_180
00000194 EC 00 8F E2      ADR      R0, aZero      ; jumptable 00000178 case 0
00000198 06 00 00 EA      B      loc_1B8
0000019C      ; -----
0000019C
0000019C      one_case                       ; CODE XREF: f2+4
0000019C      ; f2:loc_184
0000019C EC 00 8F E2      ADR      R0, aOne      ; jumptable 00000178 case 1
000001A0 04 00 00 EA      B      loc_1B8
000001A4      ; -----
000001A4
000001A4      two_case                       ; CODE XREF: f2+4
000001A4      ; f2:loc_188
000001A4 01 0C 8F E2      ADR      R0, aTwo      ; jumptable 00000178 case 2
000001A8 02 00 00 EA      B      loc_1B8
000001AC      ; -----
000001AC
000001AC      three_case                     ; CODE XREF: f2+4
000001AC      ; f2:loc_18C
000001AC 01 0C 8F E2      ADR      R0, aThree     ; jumptable 00000178 case 3
000001B0 00 00 00 EA      B      loc_1B8
000001B4      ; -----
000001B4
000001B4      four_case                      ; CODE XREF: f2+4
000001B4      ; f2:loc_190
000001B4 01 0C 8F E2      ADR      R0, aFour      ; jumptable 00000178 case 4
000001B8      loc_1B8                        ; CODE XREF: f2+24
000001B8      ; f2+2C
000001B8 66 18 00 EA      B      __2printf
000001BC      ; -----
000001BC
000001BC      default_case                   ; CODE XREF: f2+4
000001BC      ; f2+8
000001BC D4 00 8F E2      ADR      R0, aSomethingUnkno ; jumptable 00000178 default case
000001C0 FC FF FF EA      B      loc_1B8
000001C0      ; End of function f2

```

This code makes use that ARM mode feature, all instructions in this mode has size 4 bytes.

Let's keep in mind that maximum value for *a* is 4 and any greater value should cause «*something unknown*\n» string printing.

The very first “CMP R0, #5” instruction compares *a* input value with 5.

The next “ADDCC PC, PC, R0, LSL#2”⁴⁷ instruction will execute only if $R0 < 5$ (CC=Carry clear / Less than). Consequently, if ADDCC will not trigger (it's a $R0 \geq 5$ case), a jump to *default_caselabel* will be occurred.

But if $R0 < 5$ and ADDCC will trigger, following events will happen:

Value in R0 is multiplied by 4. In fact, LSL#2 at the instruction's ending mean “shift left by 2 bits”. But as we will see later 1.15.3 in “Shifts” section, shift left by 2 bits is just equivalently to multiplying by 4.

⁴⁷ ADD — addition

Then, $R0 * 4$ value we got, is added to current PC value, thus jumping to one of B (*Branch*) instructions located below.

At the moment of ADDCC execution, PC value is 8 bytes ahead ($0x180$) than address at which ADDCC instruction is located ($0x178$), or, in other words, 2 instructions ahead.

This is how ARM processor pipeline works: when ADDCC instruction is executed, the processor at that moment is beginning to process instruction after the next one, so that's why PC pointing there.

If $a = 0$, then nothing will be added to PC, and actual PC value will be written into PC (which is 8 bytes ahead) and jump to label *loc_180* will be happen, this is 8 bytes ahead of the place where ADDCC instruction is.

In case of $a = 1$, then $PC + 8 + a * 4 = PC + 8 + 1 * 4 = PC + 16 = 0x184$ will be written to PC, this is address of *loc_184* label.

With every 1 added to a , resulting PC increasing by 4. 4 is also instruction length in ARM mode and also, length of each B instruction length, there are 5 of them in row.

Each of these five B instructions passing control further, where something is going on, what was programmed in *switch()*. Pointer loading to corresponding string occurring there, etc.

ARM: Optimizing Keil + thumb mode

```

000000F6          EXPORT f2
000000F6          f2
000000F6 10 B5      PUSH    {R4,LR}
000000F8 03 00      MOVS    R3, R0
000000FA 06 F0 69 F8 BL      __ARM_common_switch8_thumb ; switch 6 cases
000000FA          ; -----
000000FE 05          DCB 5
000000FF 04 06 08 0A 0C 10 DCB 4, 6, 8, 0xA, 0xC, 0x10 ; jump table for switch statement
00000105 00          ALIGN 2
00000106
00000106          zero_case          ; CODE XREF: f2+4
00000106 8D A0          ADR      R0, aZero          ; jumtable 000000FA case 0
00000108 06 E0          B       loc_118
0000010A          ; -----
0000010A
0000010A          one_case           ; CODE XREF: f2+4
0000010A 8E A0          ADR      R0, aOne           ; jumtable 000000FA case 1
0000010C 04 E0          B       loc_118
0000010E          ; -----
0000010E
0000010E          two_case           ; CODE XREF: f2+4
0000010E 8F A0          ADR      R0, aTwo           ; jumtable 000000FA case 2
00000110 02 E0          B       loc_118
00000112          ; -----
00000112
00000112          three_case          ; CODE XREF: f2+4
00000112 90 A0          ADR      R0, aThree          ; jumtable 000000FA case 3
00000114 00 E0          B       loc_118
00000116          ; -----
00000116
00000116          four_case            ; CODE XREF: f2+4
00000116 91 A0          ADR      R0, aFour           ; jumtable 000000FA case 4
00000118
00000118          loc_118                ; CODE XREF: f2+12
00000118                                     ; f2+16
00000118 06 F0 6A F8      BL      __2printf
0000011C 10 BD      POP      {R4,PC}
0000011E          ; -----
0000011E
0000011E          default_case          ; CODE XREF: f2+4
0000011E 82 A0          ADR      R0, aSomethingUnkno ; jumtable 000000FA default case
00000120 FA E7          B       loc_118

000061D0          EXPORT __ARM_common_switch8_thumb

```

```

000061D0      __ARM_common_switch8_thumb      ; CODE XREF: example6_f2+4
000061D0 78 47      BX      PC
000061D0      ; -----
000061D2 00 00      ALIGN 4
000061D2      ; End of function __ARM_common_switch8_thumb
000061D2
000061D4      CODE32
000061D4
000061D4      ; ===== S U B R O U T I N E =====
000061D4
000061D4      __32__ARM_common_switch8_thumb      ; CODE XREF: __ARM_common_switch8_thumb
000061D4 01 C0 5E E5      LDRB      R12, [LR,#-1]
000061D8 0C 00 53 E1      CMP      R3, R12
000061DC 0C 30 DE 27      LDRCB      R3, [LR,R12]
000061E0 03 30 DE 37      LDRCCB     R3, [LR,R3]
000061E4 83 C0 8E E0      ADD      R12, LR, R3,LSL#1
000061E8 1C FF 2F E1      BX      R12
000061E8      ; End of function __32__ARM_common_switch8_thumb

```

One cannot be sure all instructions in thumb and thumb-2 modes will have same size. It's even can be said that in these modes instructions has variable length, just like in x86.

So there are a special table added, containing information about how much cases are there, not including default-case, and offset, for each, each encoding a label, to which control should be passed in corresponding case.

A special function here present in order to work with the table and pass control, named `__ARM_common_switch8_thumb`. It is beginning with 'BX PC' instruction, which function is to switch processor to ARM-mode. Then you may see the function for table processing. It's too complex for describing it here now, so I'll omit elaborations.

But it's interesting to note that the function uses LR register as a pointer to the table. Indeed, after this function calling, LR will contain address after

'BL __ARM_common_switch8_thumb' instruction, and the table is beginning right there.

It's also worth noting that the code is generated as a separate function in order to reuse it, in similar places, in similar cases, for `switch()` processing, so compiler will not generate same code at each place.

IDA 5 successfully perceived it as a service function and table, automatically, and added commentaries to labels like `jumptable 000000FA case 0`.

1.10 Loops

1.10.1 x86

There is a special LOOP instruction in x86 instruction set, it is checking ECX register value and if it is not zero, do ECX decrement⁴⁸ and pass control flow to label mentioned in LOOP operand. Probably, this instruction is not very convenient, so, I didn't ever see any modern compiler emitting it automatically. So, if you see the instruction somewhere in code, it is most likely this is manually written piece of assembly code.

By the way, as home exercise, you could try to explain, why this instruction is not very convenient.

In C/C++ loops are constructed using for(), while(), do/while() statements.

Let's start with for().

This statement defines loop initialization (set loop counter to initial value), loop condition (is counter is bigger than a limit?), what is done at each iteration (increment/decrement) and of course loop body.

```
for (initialization; condition; at each iteration)
{
    loop_body;
}
```

So, generated code will be consisted of four parts too.

Let's start with simple example:

```
int main()
{
    int i;

    for (i=2; i<10; i++)
        f(i);

    return 0;
};
```

Result (MSVC 2010):

Listing 1.27: MSVC 2010

```
_i$ = -4
_main PROC
    push    ebp
    mov     ebp, esp
    push    ecx
    mov     DWORD PTR _i$[ebp], 2    ; loop initialization
    jmp     SHORT $LN3@main
$LN2@main:
    mov     eax, DWORD PTR _i$[ebp] ; here is what we do after each iteration:
    add     eax, 1                  ; add 1 to i value
    mov     DWORD PTR _i$[ebp], eax
$LN3@main:
    cmp     DWORD PTR _i$[ebp], 10  ; this condition is checked *before* each iteration
    jge     SHORT $LN1@main         ; if i is biggest or equals to 10, let's finish loop
    mov     ecx, DWORD PTR _i$[ebp] ; loop body: call f(i)
    push    ecx
    call    _f
    add     esp, 4
    jmp     SHORT $LN2@main         ; jump to loop begin
$LN1@main:
    ; loop end
    xor     eax, eax
    mov     esp, ebp
    pop     ebp
    ret     0
_main ENDP
```

Nothing very special, as we see.

GCC 4.4.1 emitting almost the same code, with small difference:

⁴⁸subtracting 1 from it

Listing 1.28: GCC 4.4.1

```

main          proc near          ; DATA XREF: _start+17

var_20        = dword ptr -20h
var_4         = dword ptr -4

                push    ebp
                mov     ebp, esp
                and     esp, 0FFFFFFF0h
                sub     esp, 20h
                mov     [esp+20h+var_4], 2 ; i initializing
                jmp     short loc_8048476

loc_8048465:
                mov     eax, [esp+20h+var_4]
                mov     [esp+20h+var_20], eax
                call    f
                add     [esp+20h+var_4], 1 ; i increment

loc_8048476:
                cmp     [esp+20h+var_4], 9
                jle     short loc_8048465 ; if i<=9, continue loop
                mov     eax, 0
                leave
                retn

main          endp

```

Now let's see what we will get if optimization is turned on (/Ox):

Listing 1.29: Optimizing MSVC

```

_main  PROC
        push    esi
        mov     esi, 2
$LL3@main:
        push    esi
        call    _f
        inc     esi
        add     esp, 4
        cmp     esi, 10 ; 0000000aH
        jl      SHORT $LL3@main
        xor     eax, eax
        pop     esi
        ret     0
_main  ENDP

```

What is going on here is: place for *i* variable is not allocated in local stack anymore, but even individual register: ESI. This is possible in such small functions where not so many local variables are present.

One very important property is that *f()* function should not change ESI value. Our compiler is sure here. And if compiler decided to use ESI in *f()* too, it would be saved then at *f()* function prologue and restored at *f()* epilogue. Almost like in our listing: please note PUSH ESI/POP ESI at the function begin and end.

Let's try GCC 4.4.1 with maximal optimization turned on (-O3 option):

Listing 1.30: Optimizing GCC 4.4.1

```

main          proc near          ; DATA XREF: _start+17

var_10        = dword ptr -10h

                push    ebp
                mov     ebp, esp
                and     esp, 0FFFFFFF0h
                sub     esp, 10h
                mov     [esp+10h+var_10], 2
                call    f
                mov     [esp+10h+var_10], 3

```

```

        call    f
        mov     [esp+10h+var_10], 4
        call    f
        mov     [esp+10h+var_10], 5
        call    f
        mov     [esp+10h+var_10], 6
        call    f
        mov     [esp+10h+var_10], 7
        call    f
        mov     [esp+10h+var_10], 8
        call    f
        mov     [esp+10h+var_10], 9
        call    f
        xor     eax, eax
        leave
        retn
main     endp

```

Huh, GCC just unwind our loop.

Loop unwinding has advantage in these cases when there are not so much iterations and we could economy some execution speed by removing all loop supporting instructions. On the other side, resulting code is obviously larger.

OK, let's increase maximal value of *i* to 100 and try again. GCC resulting:

Listing 1.31: GCC

```

main     public main
         proc near
var_20    = dword ptr -20h

         push    ebp
         mov     ebp, esp
         and     esp, 0FFFFFFF0h
         push    ebx
         mov     ebx, 2      ; i=2
         sub     esp, 1Ch
         nop     ; aligning label loc_80484D0 (loop body begin) by 16-byte border

loc_80484D0:
         mov     [esp+20h+var_20], ebx ; pass i as first argument to f()
         add     ebx, 1      ; i++
         call    f
         cmp     ebx, 64h    ; i==100?
         jnz     short loc_80484D0 ; if not, continue
         add     esp, 1Ch
         xor     eax, eax    ; return 0
         pop     ebx
         mov     esp, ebp
         pop     ebp
         retn
main     endp

```

It's quite similar to what MSVC 2010 with optimization (/Ox) produce. With the exception that EBX register will be fixed to *i* variable. GCC is sure this register will not be modified inside of *f()* function, and if it will, it will be saved at the function prologue and restored at epilogue, just like here in *main()*.

1.10.2 ARM

Non-optimizing Keil + ARM mode

```

main
        STMFD   SP!, {R4,LR}
        MOV     R4, #2

```

```

B      loc_368
; -----
loc_35C                                ; CODE XREF: main+1C
MOV     R0, R4
BL      f
ADD     R4, R4, #1

loc_368                                ; CODE XREF: main+8
CMP     R4, #0xA
BLT     loc_35C
MOV     R0, #0
LDMFD   SP!, {R4,PC}

```

Iteration counter *i* will be stored in R4 register.

“MOV R4, #2” instruction just initializing *i*.

“MOV R0, R4” and “BL f” instructions are compose loop body, the first instruction preparing argument for *f()* function and the second is calling it.

“ADD R4, R4, #1” instruction is just adding 1 to *i* during each iteration.

“CMP R4, #0xA” comparing *i* with 0xA (10). Next instruction BLT (*Branch Less Than*) will jump if *i* is less than 10.

Otherwise, 0 will be written into R0 (because our function returning 0) and function execution ended.

Optimizing Keil + thumb mode

```

_main
PUSH    {R4,LR}
MOVS    R4, #2

loc_132                                ; CODE XREF: _main+E
MOVS    R0, R4
BL      example7_f
ADDS    R4, R4, #1
CMP     R4, #0xA
BLT     loc_132
MOVS    R0, #0
POP     {R4,PC}

```

Practically, the same.

Optimizing Xcode (LLVM) + thumb-2 mode

```

_main
PUSH    {R4,R7,LR}
MOVW    R4, #0x1124 ; "%d\n"
MOVS    R1, #2
MOVT.W  R4, #0
ADD     R7, SP, #4
ADD     R4, PC
MOV     R0, R4
BLX     _printf
MOV     R0, R4
MOVS    R1, #3
BLX     _printf
MOV     R0, R4
MOVS    R1, #4
BLX     _printf
MOV     R0, R4
MOVS    R1, #5
BLX     _printf
MOV     R0, R4
MOVS    R1, #6

```

```

BLX      _printf
MOV      R0, R4
MOVS     R1, #7
BLX      _printf
MOV      R0, R4
MOVS     R1, #8
BLX      _printf
MOV      R0, R4
MOVS     R1, #9
BLX      _printf
MOVS     R0, #0
POP      {R4,R7,PC}

```

In fact, this was in my `f()` function:

```

void f(int i)
{
    // do something here
    printf ("%d\n", i);
};

```

So, LLVM not just *unrolled* the loop, but also represented my very simple function `f()` as *inlined*, and inserted its body 8 times instead of loop. This is possible when function is so primitive, like mine, and when it's called not many times, like here.

1.10.3 One more thing

On the code generated we can see: after *i* initialization, loop body will not be executed, but *i* condition checked first, and only after loop body is to be executed. And that's correct. Because, if loop condition is not met at the beginning, loop body shouldn't be executed. For example, this is possible in the following case:

```

for (i; i<total_entries_to_process; i++)
    loop_body;

```

If *total_entries_to_process* equals to zero, loop body shouldn't be executed whatsoever. So that's why condition checked before loop body execution.

However, optimizing compiler may swap condition check and loop body, if it sure that the situation described here is not possible, like in case of our very simple example and Keil, Xcode (LLVM), MSVC in optimization mode.

1.11 strlen()

Now let's talk about loops one more time. Often, `strlen()` function⁴⁹ is implemented using `while()` statement. Here is how it's done in MSVC standard libraries:

```
int strlen (const char * str)
{
    const char *eos = str;

    while( *eos++ );

    return( eos - str - 1 );
}
```

1.11.1 x86

Let's compile:

```
_eos$ = -4 ; size = 4
_str$ = 8 ; size = 4
_strlen PROC
    push    ebp
    mov     ebp, esp
    push    ecx
    mov     eax, DWORD PTR _str$[ebp] ; place pointer to string from str
    mov     DWORD PTR _eos$[ebp], eax ; place it to local variable eos
$LN2@strlen_:
    mov     ecx, DWORD PTR _eos$[ebp] ; ECX=eos

    ; take 8-bit byte from address in ECX and place it as 32-bit value to EDX with sign extension

    movsx   edx, BYTE PTR [ecx]
    mov     eax, DWORD PTR _eos$[ebp] ; EAX=eos
    add     eax, 1 ; increment EAX
    mov     DWORD PTR _eos$[ebp], eax ; place EAX back to eos
    test    edx, edx ; EDX is zero?
    je      SHORT $LN1@strlen_ ; yes, then finish loop
    jmp     SHORT $LN2@strlen_ ; continue loop
$LN1@strlen_:

    ; here we calculate the difference between two pointers

    mov     eax, DWORD PTR _eos$[ebp]
    sub     eax, DWORD PTR _str$[ebp]
    sub     eax, 1 ; subtract 1 and return result
    mov     esp, ebp
    pop     ebp
    ret     0
_strlen_ ENDP
```

Two new instructions here: `MOVSX` 1.11.1 and `TEST`.

About first: `MOVSX` 1.11.1 is intended to take byte from some place and store value in 32-bit register. `MOVSX` 1.11.1 meaning *MOV with Sign-Extent*. Rest bits starting at 8th till 31th `MOVSX` 1.11.1 will set to 1 if source byte in memory has *minus* sign or to 0 if *plus*.

And here is why all this.

C/C++ standard defines *char* type as signed. If we have two values, one is *char* and another is *int*, (*int* is signed too), and if first value contain -2 (it is coded as $0xFE$) and we just copying this byte into *int* container, there will be $0x000000FE$, and this, from the point of signed *int* view is 254, but not -2 . In signed *int*, -2 is coded as $0xFFFFFFF$. So if we need to transfer $0xFE$ value from variable of *char* type to *int*, we need to identify its sign and extend it. That is what `MOVSX` 1.11.1 does.

See also in section “Signed number representations” 2.4.

⁴⁹ counting characters in string in C language

I'm not sure if compiler need to store *char* variable in EDX, it could take 8-bit register part (let's say DL). Apparently, compiler's register allocator⁵⁰ works like that.

Then we see TEST EDX, EDX. About TEST instruction, read more in section about bit fields 1.15. But here, this instruction just checking EDX value, if it is equals to 0.

Let's try GCC 4.4.1:

```
strlen      public strlen
strlen      proc near

eos         = dword ptr -4
arg_0       = dword ptr  8

            push    ebp
            mov     ebp, esp
            sub     esp, 10h
            mov     eax, [ebp+arg_0]
            mov     [ebp+eos], eax

loc_80483F0:
            mov     eax, [ebp+eos]
            movzx   eax, byte ptr [eax]
            test    al, al
            setnz   al
            add     [ebp+eos], 1
            test    al, al
            jnz     short loc_80483F0
            mov     edx, [ebp+eos]
            mov     eax, [ebp+arg_0]
            mov     ecx, edx
            sub     ecx, eax
            mov     eax, ecx
            sub     eax, 1
            leave
            retn

strlen      endp
```

The result almost the same as MSVC did, but here we see MOVZX instead of MOVSB 1.11.1. MOVZX mean *MOV with Zero-Extent*. This instruction place 8-bit or 16-bit value into 32-bit register and set the rest bits to zero. In fact, this instruction is convenient only because it enable us to replace two instructions at once: xor eax, eax / mov al, [...].

On other hand, it is obvious to us that compiler could produce that code: mov al, byte ptr [eax] / test al, al — it is almost the same, however, highest EAX register bits will contain random noise. But let's think it is compiler's drawback — it can't produce more understandable code. Strictly speaking, compiler is not obliged to emit understandable (to humans) code at all.

Next new instruction for us is SETNZ. Here, if AL contain not zero, test al, al will set zero to ZF flag, but SETNZ, if ZF==0 (NZ mean *not zero*) will set 1 to AL. Speaking in natural language, *if AL is not zero, let's jump to loc_80483F0*. Compiler emitted slightly redundant code, but let's not forget that optimization is turned off.

Now let's compile all this in MSVC 2010, with optimization turned on (/Ox):

```
_str$ = 8 ; size = 4
_strlen PROC
    mov     ecx, DWORD PTR _str$[esp-4] ; ECX -> pointer to the string
    mov     eax, ecx ; move to EAX
$LL2@strlen_:
    mov     dl, BYTE PTR [eax] ; DL = *EAX
    inc     eax ; EAX++
    test    dl, dl ; DL==0?
    jne     SHORT $LL2@strlen_ ; no, continue loop
    sub     eax, ecx ; calculate pointers difference
    dec     eax ; decrement EAX
    ret     0
_strlen_ ENDP
```

⁵⁰ compiler's function assigning local variables to CPU registers

Now it's all simpler. But it is needless to say that compiler could use registers such efficiently only in small functions with small number of local variables.

INC/DEC — are increment/decrement instruction, in other words: add 1 to variable or subtract.

Let's check GCC 4.4.1 with optimization turned on (-O3 key):

```
strlen      public strlen
            proc near

arg_0       = dword ptr  8

            push    ebp
            mov     ebp, esp
            mov     ecx, [ebp+arg_0]
            mov     eax, ecx

loc_8048418:
            movzx   edx, byte ptr [eax]
            add     eax, 1
            test    dl, dl
            jnz     short loc_8048418
            not     ecx
            add     eax, ecx
            pop     ebp
            retn

strlen      endp
```

Here GCC is almost the same as MSVC, except of MOVZX presence.

However, MOVZX could be replaced here to `mov dl, byte ptr [eax]`.

Probably, it is simpler for GCC compiler's code generator to *remember* that whole register is allocated for *char* variable and it can be sure that highest bits will not contain noise at any point.

After, we also see new instruction NOT. This instruction inverts all bits in operand. It can be said, it is synonym to `XOR ECX, 0xffffffffh` instruction. NOT and following ADD calculating pointer difference and subtracting 1. At the beginning ECX, where pointer to str is stored, inverted and 1 is subtracted from it.

See also: “Signed number representations” [2.4](#).

In other words, at the end of function, just after loop body, these operations are executed:

```
ecx=str;
eax=eos;
ecx=(-ecx)-1;
eax=eax+ecx
return eax
```

...and this is equivalent to:

```
ecx=str;
eax=eos;
eax=eax-ecx;
eax=eax-1;
return eax
```

Why GCC decided it would be better? I cannot be sure. But I'm assure that both variants are equivalent in efficiency sense.

1.11.2 ARM

Non-optimizing Xcode (LLVM) + ARM mode

Listing 1.32: Non-optimizing Xcode (LLVM) + ARM mode

```
_strlen

eos      = -8
str      = -4
```

```

        SUB     SP, SP, #8 ; allocate 8 bytes for local variables
        STR     R0, [SP,#8+str]
        LDR     R0, [SP,#8+str]
        STR     R0, [SP,#8+eos]

loc_2CB8                                ; CODE XREF: _strlen+28
        LDR     R0, [SP,#8+eos]
        ADD     R1, R0, #1
        STR     R1, [SP,#8+eos]
        LDRSB   R0, [R0]
        CMP     R0, #0
        BEQ     loc_2CD4
        B       loc_2CB8

; -----

loc_2CD4                                ; CODE XREF: _strlen+24
        LDR     R0, [SP,#8+eos]
        LDR     R1, [SP,#8+str]
        SUB     R0, R0, R1 ; R0=eos-str
        SUB     R0, R0, #1 ; R0=R0-1
        ADD     SP, SP, #8 ; deallocate 8 bytes for local variables
        BX     LR

```

Non-optimizing LLVM generates too much code, however, here we can see how function works with local variables in stack. There are only two local variables in our function, *eos* and *str*.

In this listing, generated by IDA 5, I renamed *var_8* and *var_4* into *eos* and *str* manually.

So, first instructions are just saves input value in *str* and *eos*.

Loop body is beginning at *loc_2CB8* label.

First three instruction in loop body (LDR, ADD, STR) loads *eos* value into R0, then value is incremented and it's saving back into *eos* local variable located in stack.

The next ‘LDRSB R0, [R0]’ (*Load Register Signed Byte*) instruction loading byte from memory at R0 address and sign-extends it to 32-bit. This is similar to MOVSB instruction in x86. The compiler treating this byte as signed because *char* type in C standard is signed. I already wrote about it 1.11.1 in this section, but related to x86.

It's should be noted, there are two ways to use 8-bit part or 16-bit part of 32-bit register in ARM, as it's possible in x86. Apparently, it's because x86 has a huge history of compatibility with its ancestors like 16-bit 8086 and even 8-bit 8080, but ARM was developed from scratch as 32-bit RISC-processor. Consequently, in order to process separate bytes in ARM, one has to use 32-bit registers anyway.

So, LDRSB loads symbol from string into R0, one by one. Next CMP and BEQ instructions check, if loaded symbol is zero. If not zero, control passing to loop body begin. And if zero, loop is finishing.

At the end of function, a difference between *eos* and *str* is calculated, 1 is also subtracting, and resulting value is returned via R0.

By the way, please note, registers weren't saved in this function. That's because by ARM calling convention, R0-R3 registers are “scratch registers”, they are intended for arguments passing, its values may not be restored upon function exit, because calling function will not use them anymore. Consequently, they may be used for anything we want. Other registers are not used here, so that's why we have nothing to save in stack. Thus, control may be returned back to calling function by simple jump (BX), to address in LR register.

Optimizing Xcode (LLVM) + thumb mode

Listing 1.33: Optimizing Xcode (LLVM) + thumb mode

```

_strlen
        MOV     R1, R0

loc_2DF6                                ; CODE XREF: _strlen+8
        LDRB.W  R2, [R1], #1
        CMP     R2, #0
        BNE     loc_2DF6
        MVNS    R0, R0
        ADD     R0, R1

```

As optimizing LLVM concludes, place in stack for *eos* and *str* may not be allocated, and these variables may always be stored right in registers. Before loop body beginning, *str* will always be in R0, and *eos* — in R1.

‘LDRB.W R2, [R1], #1’ instruction loads byte from memory at the address R1 into R2, sign-extending it to 32-bit value, but not only that. #1 at the instruction’s end calling “Post-indexed addressing”, this mean, 1 is to be added to R1 after byte load. That’s convenient when accessing arrays.

There are no such addressing mode in x86, but it’s present in some other processors, even on PDP-11. There is a legend that pre-increment, post-increment, pre-decrement and post-decrement modes in PDP-11, were “guilty” in appearance such C language (which developed on PDP-11) constructs as **ptr++*, *++ptr*, **ptr--*, *--ptr*. By the way, this is one of hard to memorize C feature. This is how it is:

C term	ARM term	C statement	how it works
Post-increment	post-indexed addressing	<i>*ptr++</i>	use <i>*ptr</i> value, then increment <i>ptr</i> pointer
Post-decrement	post-indexed addressing	<i>*ptr--</i>	use <i>*ptr</i> value, then decrement <i>ptr</i> pointer
Pre-increment	pre-indexed addressing	<i>++ptr</i>	increment <i>ptr</i> pointer, then use <i>*ptr</i> value
Pre-decrement	post-indexed addressing	<i>--ptr</i>	decrement <i>ptr</i> pointer, then use <i>*ptr</i> value

Dennis Ritchie (one of C language creators) mentioned that it’s, probably, was invented by Ken Thompson (another C creator) because this processor feature was present in PDP-7 [Rit86] [Rit93]. Thus, C language compilers may use it, if it’s present in target processor.

Then one may spot CMP and BNE in loop body, these instructions continue operation until 0 will be met in string.

MVNS⁵¹ (inverting all bits, NOT in x86 analogue) instructions and ADD computes *eos* — *str* — 1. In fact, these two instructions computes $R0 = str + eos$, which is equivalent to what was in source code, and why it’s so, I already described here 1.11.1.

Apparently, LLVM, just like GCC, concludes this code will be shorter, or faster.

Optimizing Keil + ARM mode

Listing 1.34: Optimizing Keil + ARM mode

```

_strlen
    MOV    R1, R0

loc_2C8
    ; CODE XREF: _strlen+14
    LDRB   R2, [R1], #1
    CMP    R2, #0
    SUBEQ  R0, R1, R0
    SUBEQ  R0, R0, #1
    BNE    loc_2C8
    BX     LR

```

Almost the same what we saw before, with the exception that *str* — *eos* — 1 expression may be computed not at the function’s end, but right in loop body. —EQsuffix, as we may recall, mean that instruction will be executed only if operands in executed before CMP were equal to each other. Thus, if 0 will be in R0, both SUBEQ instructions are to be executed and result is leaving in R0.

⁵¹ MoVe Not

1.12 Division by 9

Very simple function:

```
int f(int a)
{
    return a/9;
};
```

Is compiled in a very predictable way:

Listing 1.35: MSVC

```
_a$ = 8          ; size = 4
_f  PROC
    push    ebp
    mov     ebp, esp
    mov     eax, DWORD PTR _a$[ebp]
    cdq     ; sign extend EAX to EDX:EAX
    mov     ecx, 9
    idiv    ecx
    pop     ebp
    ret     0
_f  ENDP
```

IDIV divides 64-bit number stored in EDX:EAX register pair by value in ECX register. As a result, EAX will contain quotient⁵², and EDX — remainder. Result is returning from f() function in EAX register, so, that value is not moved anymore after division operation, it is in right place already. Because IDIV require value in EDX:EAX register pair, CDQ instruction (before IDIV) extending EAX value to 64-bit value taking value sign into account, just as MOVSBX 1.11.1 does. If we turn optimization on (/Ox), we got:

Listing 1.36: Optimizing MSVC

```
_a$ = 8          ; size = 4
_f  PROC

    mov     ecx, DWORD PTR _a$[esp-4]
    mov     eax, 954437177    ; 38e38e39H
    imul    ecx
    sar     edx, 1
    mov     eax, edx
    shr     eax, 31          ; 0000001fH
    add     eax, edx
    ret     0
_f  ENDP
```

This is — division using multiplication. Multiplication operation working much faster. And it is possible to use that trick⁵³ to produce a code which is equivalent and faster. GCC 4.4.1 even without optimization turned on, generate almost the same code as MSVC with optimization turned on:

Listing 1.37: Non-optimizing GCC 4.4.1

```
public f
f      proc near

arg_0  = dword ptr 8

    push    ebp
    mov     ebp, esp
    mov     ecx, [ebp+arg_0]
    mov     edx, 954437177
    mov     eax, ecx
    imul    edx
```

⁵²result of division

⁵³Read more about division by multiplication in [War02, chapter 10] and: MSDN: Integer division by constants, <http://www.nynaeve.net/?p=115>

```

sar    edx, 1
mov    eax, ecx
sar    eax, 1Fh
mov    ecx, edx
sub    ecx, eax
mov    eax, ecx
pop    ebp
retn
f      endp

```

1.12.1 ARM

ARM processor, just like in any other "pure" RISC-processors, lacks division instruction. It lacks also a single instruction for multiplication by 32-bit constant. By using one clever trick, it's possible to do division using only three instructions: addition, subtraction and bit shifts [1.15](#).

Here is an example of 32-bit number division by 10 from [[Ltd94](#), 3.3 Division by a Constant]. Quotient and remainder on output.

```

; takes argument in a1
; returns quotient in a1, remainder in a2
; cycles could be saved if only divide or remainder is required
SUB    a2, a1, #10          ; keep (x-10) for later
SUB    a1, a1, a1, lsr #2
ADD    a1, a1, a1, lsr #4
ADD    a1, a1, a1, lsr #8
ADD    a1, a1, a1, lsr #16
MOV    a1, a1, lsr #3
ADD    a3, a1, a1, asl #2
SUBS   a2, a2, a3, asl #1    ; calc (x-10) - (x/10)*10
ADDPL  a1, a1, #1           ; fix-up quotient
ADDMI  a2, a2, #10          ; fix-up remainder
MOV    pc, lr

```

Optimizing Xcode (LLVM) + ARM mode

__text:00002C58 39 1E 08 E3 E3 18 43 E3	MOV	R1, 0x38E38E39
__text:00002C60 10 F1 50 E7	SMMUL	R0, R0, R1
__text:00002C64 C0 10 A0 E1	MOV	R1, R0, ASR#1
__text:00002C68 A0 0F 81 E0	ADD	R0, R1, R0, LSR#31
__text:00002C6C 1E FF 2F E1	BX	LR

This code is mostly the same to what was generated by optimizing MSVC and GCC. Apparently, LLVM use the same algorithm for constants generating.

Observant reader may ask, how MOV writes 32-bit value in register, while this isn't possible in ARM mode. It's not possible indeed, but, as we see, there are 8 bytes per instruction instead of standard 4, in fact, there are two instructions. First instruction loading `0x8E39` value into low 16 bit of register and second instruction is in fact MOVN, it loading `0x383E` into high 16-bit of register. IDA [5](#) is aware of such sequences, and for the sake of compactness, reduced it to one single "pseudo-instruction".

SMMUL (*Signed Most Significant Word Multiply*) instruction multiplies numbers treating them as signed numbers, and leaving high 32-bit part of result in R0, dropping low 32-bit part of result.

"MOV R1, R0, ASR#1" instruction is arithmetic shift right by one bit.

"ADD R0, R1, R0, LSR#31" is $R0 = R1 + R0 >> 31$

As a matter of fact, there are no separate shifting instruction in ARM mode. Instead, some instructions like (MOV, ADD, SUB, RSB)^{[54](#)} May be supplied by option, is the second operand should be shifted, if yes, by what value and how. ASR meaning *Arithmetic Shift Right*, LSR — *Logican Shift Right*.

⁵⁴These instructions are also called "data processing instructions"

Optimizing Xcode (LLVM) + thumb-2 mode

MOV	R1, 0x38E38E39
SMMUL.W	R0, R0, R1
ASRS	R1, R0, #1
ADD.W	R0, R1, R0,LSR#31
BX	LR

There are separate instructions for shifting in thumb mode, and one of them is used here — ASRS (arithmetic shift right).

Non-optimizing Xcode (LLVM) and Keil

Non-optimizing LLVM do not generate code we saw before in this section, but inserting a call to library function `__divsi3` instead.

What about Keil: it inserting call to library function `__aeabi_idivmod` in all cases.

1.13 Work with FPU

FPU (*Floating-point unit*) — is a device within main CPU specially designed to work with floating point numbers. It was called coprocessor in past. It looks like programmable calculator in some way and stay aside of main processor.

It is worth to study stack machines⁵⁵ before FPU studying, or learn Forth language basics⁵⁶.

It is interesting to know that in past (before 80486 CPU) coprocessor was a separate chip and it was not always settled on motherboard. It was possible to buy it separately and install.

Starting at 80486 CPU, FPU is always present in it.

FPU has a stack capable to hold 8 80-bit registers, each register can hold a number in IEEE 754⁵⁷ format.

C/C++ language offer at least two floating number types, *float* (*single-precision*⁵⁸, 32 bits)⁵⁹ and *double* (*double-precision*⁶⁰, 64 bits).

GCC also supports *long double* type (*extended precision*⁶¹, 80 bit) but MSVC is not.

float type require the same number of bits as *int* type in 32-bit environment, but number representation is completely different.

Number consisting of sign, significand (also called *fraction*) and exponent.

Function having *float* or *double* among argument list is getting that value via stack. If function is returning *float* or *double* value, it leave that value in ST(0) register — at top of FPU stack.

1.13.1 Simple example

Let's consider simple example:

```
double f (double a, double b)
{
    return a/3.14 + b*4.1;
};
```

x86

Compile it in MSVC 2010:

Listing 1.38: MSVC 2010

```
CONST    SEGMENT
__real@4010666666666666 DQ 0401066666666666r    ; 4.1
CONST    ENDS
CONST    SEGMENT
__real@40091eb851eb851f DQ 040091eb851eb851fr    ; 3.14
CONST    ENDS
_TEXT    SEGMENT
_a$ = 8          ; size = 8
_b$ = 16         ; size = 8
_f PROC
    push    ebp
    mov     ebp, esp
    fld     QWORD PTR _a$[ebp]

; current stack state: ST(0) = _a

    fdiv    QWORD PTR __real@40091eb851eb851f

; current stack state: ST(0) = result of _a divided by 3.13
```

⁵⁵http://en.wikipedia.org/wiki/Stack_machine

⁵⁶[http://en.wikipedia.org/wiki/Forth_\(programming_language\)](http://en.wikipedia.org/wiki/Forth_(programming_language))

⁵⁷http://en.wikipedia.org/wiki/IEEE_754-2008

⁵⁸http://en.wikipedia.org/wiki/Single-precision_floating-point_format

⁵⁹single precision float numbers format is also addressed in *Working with the float type as with a structure* 1.16.6 section

⁶⁰http://en.wikipedia.org/wiki/Double-precision_floating-point_format

⁶¹http://en.wikipedia.org/wiki/Extended_precision

```

fld    QWORD PTR _b$[ebp]
; current stack state: ST(0) = _b; ST(1) = result of _a divided by 3.13

fmul   QWORD PTR __real@4010666666666666
; current stack state: ST(0) = result of _b * 4.1; ST(1) = result of _a divided by 3.13

faddp  ST(1), ST(0)
; current stack state: ST(0) = result of addition

pop    ebp
ret    0
_f     ENDP

```

FLD takes 8 bytes from stack and load the number into ST(0) register, automatically converting it into internal 80-bit format *extended precision*).

FDIV divide value in register ST(0) by number storing at address `__real@40091eb851eb851f` — 3.14 value is coded there. Assembler syntax missing floating point numbers, so, what we see here is hexadecimal representation of 3.14 number in 64-bit IEEE 754 encoded.

After FDIV execution, ST(0) will hold quotient⁶².

By the way, there are also FDIVP instruction, which divide ST(1) by ST(0), popping both these values from stack and then pushing result. If you know Forth language⁶³, you will quickly understand that this is stack machine⁶⁴.

Next FLD instruction pushing *b* value into stack.

After that, quotient is placed to ST(1), and ST(0) will hold *b* value.

Next FMUL instruction do multiplication: *b* from ST(0) register by value at `__real@4010666666666666` (4.1 number is there) and leaves result in ST(0).

Very last FADDP instruction adds two values at top of stack, placing result at ST(1) register and then popping value at ST(1), hereby leaving result at top of stack in ST(0).

The function must return result in ST(0) register, so, after FADDP there are no any other code except of function epilogue.

GCC 4.4.1 (with -O3 option) emitting the same code, however, slightly different:

Listing 1.39: Optimizing GCC 4.4.1

```

f      public f
      proc near

arg_0   = qword ptr 8
arg_8   = qword ptr 10h

      push    ebp
      fld     ds:dbl_8048608 ; 3.14

; stack state now: ST(0) = 3.13

      mov     ebp, esp
      fdivr   [ebp+arg_0]

; stack state now: ST(0) = result of division

      fld     ds:dbl_8048610 ; 4.1

; stack state now: ST(0) = 4.1, ST(1) = result of division

      fmul    [ebp+arg_8]

```

⁶²result of division

⁶³[http://en.wikipedia.org/wiki/Forth_\(programming_language\)](http://en.wikipedia.org/wiki/Forth_(programming_language))

⁶⁴http://en.wikipedia.org/wiki/Stack_machine

```

; stack state now: ST(0) = result of multiplication, ST(1) = result of division

        pop     ebp
        faddp   st(1), st

; stack state now: ST(0) = result of addition

        retn
f        endp

```

The difference is that, first of all, 3.14 is pushing to stack (into ST(0)), and then value in `arg_0` is dividing by what is in ST(0) register.

FDIVR meaning *Reverse Divide* — to divide with divisor and dividend swapped with each other. There are no such instruction for multiplication, because multiplication is commutative operation, so we have just FMUL without its -R counterpart.

FADDP adding two values but also popping one value from stack. After that operation, ST(0) will contain sum.

This fragment of disassembled code was produced using IDA 5 which named ST(0) register as ST for short.

ARM: Optimizing Xcode (LLVM) + ARM mode

Until ARM has floating standardized point support, several processor manufacturers may add their own instructions extensions. Then, VFP (*Vector Floating Point*) was standardized.

One important difference from x86, there you working with FPU-stack, but here, in ARM, there are no any stack, you work just with registers.

```

f
    VLDR        D16, =3.14
    VMOV        D17, R0, R1 ; load a
    VMOV        D18, R2, R3 ; load b
    VDIV.F64    D16, D17, D16 ; a/3.14
    VLDR        D17, =4.1
    VMUL.F64    D17, D18, D17 ; b*4.1
    VADD.F64    D16, D17, D16 ; +
    VMOV        R0, R1, D16
    BX          LR

dbl_2C98      DCFD 3.14          ; DATA XREF: f
dbl_2CA0      DCFD 4.1          ; DATA XREF: f+10

```

So, we see here new registers used, with D prefix. These are 64-bit registers, there are 32 of them, and these can be used both for floating-point numbers (double), but also for SIMD (it's called NEON here in ARM).

There are also 32 32-bit S-registers, they are intended to be used for single precision floating pointer numbers (float).

It's easy to remember: D-registers are intended for double precision numbers, while S-registers — for single precision numbers.

Both (3.14 and 4.1) constants are stored in memory in IEEE 754 form.

VLDR and VMOV instructions, as it can be easily deduced, are analogous to LDR and MOV, but they works with D-registers. It should be noted that these instructions, just like D-registers, are intended not only for floating point numbers, but can be also used for SIMD (NEON) operations and this will also be revealed soon.

Arguments are passed to function by usual way, via R-registers, however, each number having double precision has size 64-bits, so, for passing each, two R-registers are needed.

“VMOV D17, R0, R1” at the very beginning, composing two 32-bit values from R0 and R1 into one 64-bit value and saves it to D17.

“VMOV R0, R1, D16” is inverse operation, what was in D16 leaving in two R0 and R1 registers, because, double-precision number, needing 64 bits for storage, returning in R0 and R1 registers pair.

VDIV, VMUL and VADD, are instruction for floating point numbers processing, computing, quotient⁶⁵, prod-

⁶⁵result of division

uct⁶⁶ and sum⁶⁷, respectively.
The code for thumb-2 is same.

ARM: Optimizing Keil + thumb mode

```
f
    PUSH    {R3-R7,LR}
    MOVS    R7, R2
    MOVS    R4, R3
    MOVS    R5, R0
    MOVS    R6, R1
    LDR     R2, =0x66666666
    LDR     R3, =0x40106666
    MOVS    R0, R7
    MOVS    R1, R4
    BL      __aeabi_dmul
    MOVS    R7, R0
    MOVS    R4, R1
    LDR     R2, =0x51EB851F
    LDR     R3, =0x40091EB8
    MOVS    R0, R5
    MOVS    R1, R6
    BL      __aeabi_ddiv
    MOVS    R2, R7
    MOVS    R3, R4
    BL      __aeabi_dadd
    POP     {R3-R7,PC}

dword_364      DCD 0x66666666          ; DATA XREF: f+A
dword_368      DCD 0x40106666          ; DATA XREF: f+C
dword_36C      DCD 0x51EB851F          ; DATA XREF: f+1A
dword_370      DCD 0x40091EB8          ; DATA XREF: f+1C
```

Keil generates for processors not supporting FPU or NEON. So, double-precision floating numbers are passed via usual R-registers, and instead of FPU-instructions, service library functions are called, like `__aeabi_dmul`, `__aeabi_ddiv`, `__aeabi_dadd`, which emulating multiplication, division and addition floating-point numbers. Of course, that's slower than FPU-coprocessor, but it's better than nothing.

By the way, similar FPU-emulating libraries were very popular in x86 world when coprocessors were rare and expensive, and were installed only on expensive computers.

FPU-coprocessor emulating calling *soft float* or *armel* in ARM world, while coprocessor's FPU-instructions usage calling *hard float* or *armhf*.

For example, Linux kernel for Raspberry Pi is compiled in two variants. In *soft float* case, arguments will be passed via R-registers, and in *hard float* case — via D-registers.

And that's what don't let you use, for example, *armhf*-libraries from *armel*-code or vice versa, so that's why all code in Linux distribution should be compiled according to chosen calling convention.

1.13.2 Passing floating point number via arguments

```
int main ()
{
    printf ("32.01 ^ 1.54 = %lf\n", pow (32.01,1.54));

    return 0;
}
```

⁶⁶result of multiplication

⁶⁷result of addition

x86

Let's see what we got in (MSVC 2010):

Listing 1.40: MSVC 2010

```
CONST    SEGMENT
__real@40400147ae147ae1 DQ 040400147ae147ae1r    ; 32.01
__real@3ff8a3d70a3d70a4 DQ 03ff8a3d70a3d70a4r    ; 1.54
CONST    ENDS

_main    PROC
    push    ebp
    mov     ebp, esp
    sub     esp, 8 ; allocate place for the first variable
    fld     QWORD PTR __real@3ff8a3d70a3d70a4
    fstp    QWORD PTR [esp]
    sub     esp, 8 ; allocate place for the second variable
    fld     QWORD PTR __real@40400147ae147ae1
    fstp    QWORD PTR [esp]
    call    _pow
    add     esp, 8 ; "return back" place of one variable.

; in local stack here 8 bytes still reserved for us.
; result now in ST(0)

    fstp    QWORD PTR [esp] ; move result from ST(0) to local stack for printf()
    push    OFFSET $SG2651
    call    _printf
    add     esp, 12
    xor     eax, eax
    pop     ebp
    ret     0
_main    ENDP
```

FLD and FSTP are moving variables from/to data segment to FPU stack. `pow()`⁶⁸ taking both values from FPU-stack and returning result in ST(0). `printf()` takes 8 bytes from local stack and interpret them as *double* type variable.

ARM + Non-optimizing Xcode (LLVM) + thumb-2 mode

```
_main
var_C      = -0xC

        PUSH    {R7,LR}
        MOV     R7, SP
        SUB     SP, SP, #4
        VLDR    D16, =32.01
        VMOV    R0, R1, D16
        VLDR    D16, =1.54
        VMOV    R2, R3, D16
        BLX     _pow
        VMOV    D16, R0, R1
        MOV     R0, 0xFC1 ; "32.01 ^ 1.54 = %lf\n"
        ADD     R0, PC
        VMOV    R1, R2, D16
        BLX     _printf
        MOVS    R1, 0
        STR     R0, [SP,#0xC+var_C]
        MOV     R0, R1
        ADD     SP, SP, #4
        POP     {R7,PC}
```

⁶⁸standard C function, raises a number to the given power

dbl_2F90	DCFD 32.01	; DATA XREF: _main+6
dbl_2F98	DCFD 1.54	; DATA XREF: _main+E

As I wrote before, 64-bit floating pointer numbers passing in R-registers pairs. This code is redundant for a little (certainty because optimization is turned off), because, it's actually possible to load values into R-registers straightforwardly without touching D-registers.

So, as we see, `_pow` function receiving first argument in R0 and R1, and the second one in R2 and R3. Function leave result in R0 and R1. Result of `_pow` is moved into D16, then in R1 and R2 pair, from where `printf()` will take this number.

ARM + Non-optimizing Keil + ARM mode

```

_main
    STMFD    SP!, {R4-R6,LR}
    LDR      R2, =0xA3D70A4 ; y
    LDR      R3, =0x3FF8A3D7
    LDR      R0, =0xAE147AE1 ; x
    LDR      R1, =0x40400147
    BL       pow
    MOV      R4, R0
    MOV      R2, R4
    MOV      R3, R1
    ADR      R0, a32_011_54Lf ; "32.01 ^ 1.54 = %lf\n"
    BL       __2printf
    MOV      R0, #0
    LDMFD    SP!, {R4-R6,PC}

y          DCD 0xA3D70A4          ; DATA XREF: _main+4
dword_520  DCD 0x3FF8A3D7        ; DATA XREF: _main+8
; double x
x          DCD 0xAE147AE1        ; DATA XREF: _main+C
dword_528  DCD 0x40400147        ; DATA XREF: _main+10
a32_011_54Lf DCB "32.01 ^ 1.54 = %lf",0xA,0
                                     ; DATA XREF: _main+24

```

D-registers are not used here, only R-register pairs are used.

1.13.3 Comparison example

Let's try this:

```

double d_max (double a, double b)
{
    if (a>b)
        return a;

    return b;
};

```

x86

Despite simplicity of that function, it will be harder to understand how it works.

MSVC 2010 generated:

Listing 1.41: MSVC 2010

```

PUBLIC     _d_max
_TEXT     SEGMENT
_a$ = 8           ; size = 8
_b$ = 16          ; size = 8
_d_max    PROC

```

```

push    ebp
mov     ebp, esp
fld     QWORD PTR _b$[ebp]

; current stack state: ST(0) = _b
; compare _b (ST(0)) and _a, and pop register

fcomp   QWORD PTR _a$[ebp]

; stack is empty here

fnstsw  ax
test    ah, 5
jp      SHORT $LN1@d_max

; we are here only if a>b

fld     QWORD PTR _a$[ebp]
jmp     SHORT $LN2@d_max
$LN1@d_max:
fld     QWORD PTR _b$[ebp]
$LN2@d_max:
pop     ebp
ret     0
_d_max  ENDP

```

So, FLD loading `_b` into ST(0) register.

FCOMP compares ST(0) register state with what is in `_a` value and set C3/C2/C0 bits in FPU status word register. This is 16-bit register reflecting current state of FPU.

For now C3/C2/C0 bits are set, but unfortunately, CPU before Intel P6 ⁶⁹ hasn't any conditional jumps instructions which are checking these bits. Probably, it is a matter of history (remember: FPU was separate chip in past). Modern CPU starting at Intel P6 has FCOMI/FCOMIP/FUCOMI/FUCOMIP instructions — which does that same, but modifies CPU flags ZF/PF/CF.

After bits are set, the FCOMP instruction popping one variable from stack. This is what distinguish it from FCOM, which is just comparing values, leaving the stack at the same state.

FNSTSW copies FPU status word register to AX. Bits C3/C2/C0 are placed at positions 14/10/8, they will be at the same positions in AX registers and all them are placed in high part of AX — AH.

- If `b>a` in our example, then C3/C2/C0 bits will be set as following: 0, 0, 0.
- If `a>b`, then bits will be set: 0, 0, 1.
- If `a=b`, then bits will be set: 1, 0, 0.

After `test ah, 5` execution, bits C3 and C1 will be set to 0, but at positions 0 and 2 (in AH registers) C0 and C2 bits will be leaved.

Now let's talk about parity flag. Another notable epoch rudiment:

One common reason to test the parity flag actually has nothing to do with parity. The FPU has four condition flags (C0 to C3), but they can not be tested directly, and must instead be first copied to the flags register. When this happens, C0 is placed in the carry flag, C2 in the parity flag and C3 in the zero flag. The C2 flag is set when e.g. incomparable floating point values (NaN or unsupported format) are compared with the FUCOM instructions. ⁷⁰

This flag is to be set to 1 if ones number is even. And to zero if odd.

⁶⁹ Intel P6 is Pentium Pro, Pentium II, etc

Thus, PF flag will be set to 1 if both C0 and C2 are set to zero or both are ones. And then following JP (*jump if PF==1*) will be triggered. If we recall values of C3/C2/C0 for different cases, we will see that conditional jump JP will be triggered in two cases: if b>a or a==b (C3 bit is already not considering here, because it was cleared while execution of test ah, 5 instruction).

It is all simple thereafter. If conditional jump was triggered, FLD will load _b value to ST(0), and if it's not triggered, _a will be loaded.

But it is not over yet!

Now let's compile it with MSVC 2010 with optimization option /Ox

Listing 1.42: Optimizing MSVC 2010

```
_a$ = 8          ; size = 8
_b$ = 16         ; size = 8
_d_max PROC
    fld     QWORD PTR _b$[esp-4]
    fld     QWORD PTR _a$[esp-4]

; current stack state: ST(0) = _a, ST(1) = _b

    fcom    ST(1) ; compare _a and ST(1) = (_b)
    fnstsw  ax
    test    ah, 65          ; 00000041H
    jne     SHORT $LN5@d_max
    fstp    ST(1) ; copy ST(0) to ST(1) and pop register, leave (_a) on top

; current stack state: ST(0) = _a

    ret     0
$LN5@d_max:
    fstp    ST(0) ; copy ST(0) to ST(0) and pop register, leave (_b) on top

; current stack state: ST(0) = _b

    ret     0
_d_max ENDP
```

FCOM is distinguished from FCOMP is that sense that it just comparing values and leave FPU stack in the same state. Unlike previous example, operands here in reversed order. And that is why result of comparison in C3/C2/C0 will be different:

- If a>b in our example, then C3/C2/C0 bits will be set as: 0, 0, 0.
- If b>a, then bits will be set as: 0, 0, 1.
- If a=b, then bits will be set as: 1, 0, 0.

It can be said, test ah, 65 instruction just leave two bits — C3 and C0. Both will be zeroes if a>b: in that case JNE jump will not be triggered. Then FSTP ST(1) is following — this instruction copies ST(0) value into operand and popping one value from FPU stack. In other words, that instruction copies ST(0) (where _a value now) into ST(1). After that, two values of _a are at the top of stack now. After that, one value is popping. After that, ST(0) will contain _a and function is finished.

Conditional jump JNE is triggered in two cases: of b>a or a==b. ST(0) into ST(0) will be copied, it is just like idle (NOP) operation, then one value is popping from stack and top of stack (ST(0)) will contain what was in ST(1) before (that is _b). Then function finishes. That instruction used here probably because FPU has no instruction to pop value from stack and not to store it anywhere.

Well, but it is still not over.

Listing 1.43: GCC 4.4.1

```

d_max proc near

b          = qword ptr -10h
a          = qword ptr -8
a_first_half = dword ptr 8
a_second_half = dword ptr 0Ch
b_first_half = dword ptr 10h
b_second_half = dword ptr 14h

    push    ebp
    mov     ebp, esp
    sub     esp, 10h

; put a and b to local stack:

    mov     eax, [ebp+a_first_half]
    mov     dword ptr [ebp+a], eax
    mov     eax, [ebp+a_second_half]
    mov     dword ptr [ebp+a+4], eax
    mov     eax, [ebp+b_first_half]
    mov     dword ptr [ebp+b], eax
    mov     eax, [ebp+b_second_half]
    mov     dword ptr [ebp+b+4], eax

; load a and b to FPU stack:

    fld     [ebp+a]
    fld     [ebp+b]

; current stack state: ST(0) - b; ST(1) - a

    fxch    st(1) ; this instruction swapping ST(1) and ST(0)

; current stack state: ST(0) - a; ST(1) - b

    fucompp ; compare a and b and pop two values from stack, i.e., a and b
    fnstsw ax ; store FPU status to AX
    sahf     ; load SF, ZF, AF, PF, and CF flags state from AH
    setnbe al ; store 1 to AL if CF=0 and ZF=0
    test    al, al ; AL==0 ?
    jz      short loc_8048453 ; yes
    fld     [ebp+a]
    jmp     short locret_8048456

loc_8048453:
    fld     [ebp+b]

locret_8048456:
    leave
    retn
d_max endp

```

FUCOMPP — is almost like FCOM, but popping both values from stack and handling “not-a-numbers” differently. More about *not-a-numbers*:

FPU is able to work with special values which are *not-a-numbers* or NaNs⁷¹. These are infinity, result of dividing by zero, etc. Not-a-numbers can be “quiet” and “signalling”. It is possible to continue to work with “quiet” NaNs, but if one try to do some operation with “signalling” NaNs — an exception will be raised.

FCOM will raise exception if any operand — NaN. FUCOM will raise exception only if any operand — signalling NaN (SNaN).

⁷¹<http://en.wikipedia.org/wiki/NaN>

The following instruction is SAHF — this is rare instruction in the code which is not use FPU. 8 bits from AH is movinto into lower 8 bits of CPU flags in the following order: SF:ZF:-:AF:-:PF:-:CF <- AH.

Let's remember that FNSTSW is moving interesting for us bits C3/C2/C0 into AH and they will be in positions 6, 2, 0 in AH register.

In other words, fnstsw ax / sahf instruction pair is moving C3/C2/C0 into ZF, PF, CF CPU flags.

Now let's also recall, what values of C3/C2/C0 bits will be set:

- If a is greater than b in our example, then C3/C2/C0 bits will be set as: 0, 0, 0.
- if a is less than b, then bits will be set as: 0, 0, 1.
- If a=b, then bits will be set: 1, 0, 0.

In other words, after FUCOMPP/FNSTSW/SAHF instructions, we will have these CPU flags states:

- If a>b, CPU flags will be set as: ZF=0, PF=0, CF=0.
- If a<b, then CPU flags will be set as: ZF=0, PF=0, CF=1.
- If a=b, then CPU flags will be set as: ZF=1, PF=0, CF=0.

How SETNBE instruction will store 1 or 0 to AL: it is depends of CPU flags. It is almost JNBE instruction counterpart, with exception that SETcc⁷² is storing 1 or 0 to AL, but Jcc do actual jump or not. SETNBE store 1 only if CF=0 and ZF=0. If it is not true, zero will be stored into AL.

Both CF is 0 and ZF is 0 simultaneously only in one case: if a>b.

Then one will be stored to AL and the following JZ will not be triggered and function will return _a. On all other cases, _b will be returned.

But it is still not over.

GCC 4.4.1 with -O3 optimization turned on

Listing 1.44: Optimizing GCC 4.4.1

```
d_max      public d_max
d_max      proc near

arg_0      = qword ptr 8
arg_8      = qword ptr 10h

        push    ebp
        mov     ebp, esp
        fld     [ebp+arg_0] ; _a
        fld     [ebp+arg_8] ; _b

; stack state now: ST(0) = _b, ST(1) = _a
        fxch    st(1)

; stack state now: ST(0) = _a, ST(1) = _b
        fucom   st(1) ; compare _a and _b
        fnstsw  ax
        sahf
        ja      short loc_8048448

; store ST(0) to ST(0) (idle operation), pop value at top of stack, leave _b at top
        fstp    st
        jmp     short loc_804844A

loc_8048448:
; store _a to ST(0), pop value at top of stack, leave _a at top
        fstp    st(1)
```

⁷²cc is condition code

```
loc_804844A:
    pop     ebp
    retn
d_max     endp
```

It is almost the same except one: JA usage instead of SAHF. Actually, conditional jump instructions checking “larger”, “lesser” or “equal” for unsigned number comparison (JA, JAE, JBE, JE/JZ, JNA, JNAE, JNB, JNBE, JNE/JNZ) are checking only CF and ZF flags. And C3/C2/C0 bits after comparison are moving into these flags exactly in the same fashion so conditional jumps will work here. JA will work if both CF are ZF zero.

Thereby, conditional jumps instructions listed here can be used after FNSTSW/SAHF instructions pair.

It seems, FPU C3/C2/C0 status bits was placed there deliberately so to map them to base CPU flags without additional permutations.

ARM + Optimizing Xcode (LLVM) + ARM mode

Listing 1.45: Optimizing Xcode (LLVM) + ARM mode

VMOV	D16, R2, R3 ; b
VMOV	D17, R0, R1 ; a
VCMPE.F64	D17, D16
VMRS	APSR_nzcv, FPSCR
VMOVGT.F64	D16, D17 ; copy b to D16
VMOV	R0, R1, D16
BX	LR

A very simple case. Input values are placed into D17 and D16 and then compared with the help of VCMPE instruction. Just like in x86 coprocessor, ARM coprocessor has its own status and flags register, (FPSCR), because there is a need to store coprocessor-specific flags.

And just like in x86, there are no conditional jump instruction in ARM, checking bits in coprocessor status register, so there is VMRS instruction, copying 4 bits (N, Z, C, V) from the coprocessor status word into bits of *general* status (APSR register).

VMOVGT is analogue of MOVGT, instruction, to be executed if one operand is greater than other while comparing (*GT* — *Greater Than*).

If it will be executed, *b* value will be written into D16, stored at that moment in D17.

And if it will not be triggered, then *a* value will stay in D16.

Penultimate instruction VMOV will prepare value in D16 for returning via R0 and R1 registers pair.

ARM + Optimizing Xcode (LLVM) + thumb-2 mode

Listing 1.46: Optimizing Xcode (LLVM) + thumb-2 mode

VMOV	D16, R2, R3 ; b
VMOV	D17, R0, R1 ; a
VCMPE.F64	D17, D16
VMRS	APSR_nzcv, FPSCR
IT GT	
VMOVGT.F64	D16, D17
VMOV	R0, R1, D16
BX	LR

Almost the same as in previous example, however slightly different. As a matter of fact, many instructions in ARM mode can be supplied by condition predicate, and the instruction is to be executed if condition is true.

But there are no such thing in thumb mode. There are no place in 16-bit instructions for spare 4 bits where condition can be encoded.

However, thumb-2 was extended to make possible to specify predicates to old thumb instructions.

Here, is the IDA 5-generated listing, we see VMOVGT instruction, the same as in previous example.

But in fact, usual VMOV is encoded there, but IDA 5 added -GT suffix to it, because there are “IT GT” instruction placed right before.

IT instruction defines so-called *if-then block*. After that instruction, it's possible to place up to 4 instructions, to which predicate suffix will be added. In our example, “IT GT” meaning, the next instruction will be executed, if *GT (Greater Than)* condition is true.

Now more complex code fragment, by the way, from “Angry Birds” (for iOS):

Listing 1.47: Angry Birds Classic

```
ITE NE
VMOVNE      R2, R3, D16
VMOVEQ      R2, R3, D17
```

ITE meaning *if-then-else* and it encode suffixes for two next instructions. First instruction will execute if condition encoded in ITE (*NE, not equal*) will be true at that moment, and the second — if that condition will not be true. (Inverse condition of NE is EQ (*equal*)).

Slightly harder, and this fragment from “Angry Birds” as well:

Listing 1.48: Angry Birds Classic

```
ITTTT EQ
MOVEQ       R0, R4
ADDEQ       SP, SP, #0x20
POPEQ.W     {R8,R10}
POPEQ       {R4-R7,PC}
```

4 “T” symbols in instruction mnemonic mean that 4 next instructions will be executed if condition is true. That's why IDA 5 added -EQ suffix to each 4 instructions.

And if there will be, for example, ITEEE EQ (*if-then-else-else-else*), then suffixes will be set as follows:

```
-EQ
-NE
-NE
-NE
```

Another fragment from “Angry Birds”:

Listing 1.49: Angry Birds Classic

```
CMP.W       R0, #0xFFFFFFFF
ITTE LE
SUBLE.W     R10, R0, #1
NEGLE      R0, R0
MOVGT      R10, R0
```

ITTE (*if-then-then-else*) mean that 1st and 2nd instructions will be executed, if LE (*Less or Equal*) condition is true, and 3rd — if inverse condition (GT — *Greater Than*) is true.

Compilers usually are not generating all possible combinations. For example, it mentioned “Angry Birds” game (*classic* version for iOS) only these cases of IT instruction are used: IT, ITE, ITT, ITTE, ITTT, ITTTT. How I learnt this? In IDA 5 it's possible to produce listing files, so I did it, but I also set in options to show 4 bytes of each opcodes. Then, knowing that high part of 16-bit opcode IT is *0xBF*, I did this with `grep`:

```
cat AngryBirdsClassic.lst | grep " BF" | grep "IT" > results.lst
```

By the way, if to program in ARM assembly language manually for thumb-2 mode, with adding conditional suffixes, assembler will add IT instructions automatically, with respectable flags, where it's need.

ARM + Non-optimizing Xcode (LLVM) + ARM mode

Listing 1.50: Non-optimizing Xcode (LLVM) + ARM mode

```
b          = -0x20
a          = -0x18
val_to_return = -0x10
saved_R7    = -4
```

	STR	R7, [SP,#saved_R7]!
	MOV	R7, SP
	SUB	SP, SP, #0x1C
	BIC	SP, SP, #7
	VMOV	D16, R2, R3
	VMOV	D17, R0, R1
	VSTR	D17, [SP,#0x20+a]
	VSTR	D16, [SP,#0x20+b]
	VLDR	D16, [SP,#0x20+a]
	VLDR	D17, [SP,#0x20+b]
	VCMPE.F64	D16, D17
	VMRS	APSR_nzcv, FPSCR
	BLE	loc_2E08
	VLDR	D16, [SP,#0x20+a]
	VSTR	D16, [SP,#0x20+val_to_return]
	B	loc_2E10
loc_2E08		
	VLDR	D16, [SP,#0x20+b]
	VSTR	D16, [SP,#0x20+val_to_return]
loc_2E10		
	VLDR	D16, [SP,#0x20+val_to_return]
	VMOV	R0, R1, D16
	MOV	SP, R7
	LDR	R7, [SP+0x20+b],#4
	BX	LR

Almost the same we already saw, but too much redundant code because of *a* and *b* variables storage in local stack, as well as returning value.

ARM + Optimizing Keil + thumb mode

Listing 1.51: Optimizing Keil + thumb mode

	PUSH	{R3-R7,LR}
	MOVS	R4, R2
	MOVS	R5, R3
	MOVS	R6, R0
	MOVS	R7, R1
	BL	__aeabi_cdrcmple
	BCS	loc_1C0
	MOVS	R0, R6
	MOVS	R1, R7
	POP	{R3-R7,PC}
loc_1C0		
	MOVS	R0, R4
	MOVS	R1, R5
	POP	{R3-R7,PC}

Keil not generate special instruction for float numbers comparing, because, it cannot rely it will be supported on the target CPU, and it cannot be done by simple bitwise comparing. So there is called external library function for comparing: `__aeabi_cdrcmple`. Please note, comparision result is to be leaved in flags, so the following BCS (*Carry set - Greater than or equal*) instruction may work without any additional code.

1.14 Arrays

Array is just a set of variables in memory, always lying next to each other, always has same type.

1.14.1 Simple example

```
#include <stdio.h>

int main()
{
    int a[20];
    int i;

    for (i=0; i<20; i++)
        a[i]=i*2;

    for (i=0; i<20; i++)
        printf ("a[%d]=%d\n", i, a[i]);

    return 0;
};
```

x86

Let's compile:

Listing 1.52: MSVC

```
_TEXT    SEGMENT
_i$ = -84                ; size = 4
_a$ = -80                ; size = 80
_main    PROC
    push    ebp
    mov     ebp, esp
    sub     esp, 84      ; 00000054H
    mov     DWORD PTR _i$[ebp], 0
    jmp     SHORT $LN6@main
$LN5@main:
    mov     eax, DWORD PTR _i$[ebp]
    add     eax, 1
    mov     DWORD PTR _i$[ebp], eax
$LN6@main:
    cmp     DWORD PTR _i$[ebp], 20    ; 00000014H
    jge     SHORT $LN4@main
    mov     ecx, DWORD PTR _i$[ebp]
    shl     ecx, 1
    mov     edx, DWORD PTR _i$[ebp]
    mov     DWORD PTR _a$[ebp+edx*4], ecx
    jmp     SHORT $LN5@main
$LN4@main:
    mov     DWORD PTR _i$[ebp], 0
    jmp     SHORT $LN3@main
$LN2@main:
    mov     eax, DWORD PTR _i$[ebp]
    add     eax, 1
    mov     DWORD PTR _i$[ebp], eax
$LN3@main:
    cmp     DWORD PTR _i$[ebp], 20    ; 00000014H
    jge     SHORT $LN1@main
    mov     ecx, DWORD PTR _i$[ebp]
    mov     edx, DWORD PTR _a$[ebp+ecx*4]
    push    edx
    mov     eax, DWORD PTR _i$[ebp]
    push    eax
```

```

push    OFFSET $SG2463
call    _printf
add     esp, 12          ; 0000000cH
jmp     SHORT $LN2@main
$LN1@main:
xor     eax, eax
mov     esp, ebp
pop     ebp
ret     0
_main   ENDP

```

Nothing very special, just two loops: first is filling loop and second is printing loop. `shl ecx, 1` instruction is used for ECX value multiplication by 2, more about below [1.15.3](#).

80 bytes are allocated in stack for array, that's 20 elements of 4 bytes.

Here is what GCC 4.4.1 does:

Listing 1.53: GCC 4.4.1

```

main      public main
          proc near          ; DATA XREF: _start+17

var_70    = dword ptr -70h
var_6C    = dword ptr -6Ch
var_68    = dword ptr -68h
i_2       = dword ptr -54h
i         = dword ptr -4

          push    ebp
          mov     ebp, esp
          and     esp, 0FFFFFFF0h
          sub     esp, 70h
          mov     [esp+70h+i], 0          ; i=0
          jmp     short loc_804840A

loc_80483F7:
          mov     eax, [esp+70h+i]
          mov     edx, [esp+70h+i]
          add     edx, edx              ; edx=i*2
          mov     [esp+eax*4+70h+i_2], edx
          add     [esp+70h+i], 1        ; i++

loc_804840A:
          cmp     [esp+70h+i], 13h
          jle     short loc_80483F7
          mov     [esp+70h+i], 0
          jmp     short loc_8048441

loc_804841B:
          mov     eax, [esp+70h+i]
          mov     edx, [esp+eax*4+70h+i_2]
          mov     eax, offset aADD ; "a[%d]=%d\n"
          mov     [esp+70h+var_68], edx
          mov     edx, [esp+70h+i]
          mov     [esp+70h+var_6C], edx
          mov     [esp+70h+var_70], eax
          call    _printf
          add     [esp+70h+i], 1

loc_8048441:
          cmp     [esp+70h+i], 13h
          jle     short loc_804841B
          mov     eax, 0
          leave
          retn

main      endp

```

By the way, *a* variable has *int** type (that is pointer to *int*) — you can try to pass a pointer to array to another function, but it much correctly to say that pointer to the first array element is passed (addresses of another element's places are calculated in obvious way). If to index this pointer as *a[idx]*, *idx* just to be added to the pointer and the element placed there (to which calculated pointer is pointing) returned.

An interesting example: string of characters like “string” is array of characters and it has *const char** type. Index can be applied to this pointer. And that's why it's possible to write like “string”[*i*] — this is correct C/C++ expression!

ARM + Non-optimizing Keil + ARM mode

```

EXPORT _main
_main
    STMFD    SP!, {R4,LR}
    SUB      SP, SP, #0x50      ; allocate place for 20 int variables

; first loop

    MOV      R4, #0             ; i
    B        loc_4A0

loc_494
    MOV      R0, R4,LSL#1       ; R0=R4*2
    STR      R0, [SP,R4,LSL#2]  ; store R0 to SP+R4<<2 (same as SP+R4*4)
    ADD      R4, R4, #1         ; i=i+1

loc_4A0
    CMP      R4, #20            ; i<20?
    BLT      loc_494           ; yes, run loop body again

; second loop

    MOV      R4, #0             ; i
    B        loc_4C4

loc_4B0
    LDR      R2, [SP,R4,LSL#2]  ; (second printf argument) R2=*(SP+R4<<4) (same as *(SP+R4*4))
    MOV      R1, R4             ; (first printf argument) R1=i
    ADR      R0, aADD           ; "a[%d]=%d\n"
    BL       __2printf
    ADD      R4, R4, #1         ; i=i+1

loc_4C4
    CMP      R4, #20            ; i<20?
    BLT      loc_4B0           ; yes, run loop body again
    MOV      R0, #0             ; value to return
    ADD      SP, SP, #0x50      ; deallocate place for 20 int variables
    LDMFD    SP!, {R4,PC}

```

int type require 32 bits for storage, or 4 bytes, so for storage of 20 *int* variables, 80 (0x50) bytes are needed, so that's why “SUB SP, SP, #0x50” instruction in function epilogue allocates exactly this ammount of space in local stack.

In both first and second loops, *i* loop iterator will be placed in R4 register.

A number to be written into array, is calculating as $i * 2$ which is equivalent to shifting left by one bit, so “MOV R0, R4, LSL#1” instruction do this.

“STR R0, [SP, R4, LSL#2]” writes R0 contents into array. Here is how a pointer to array element is to be calculated: SP pointing to array begin, R4 is *i*. So shift *i* left by 2 bits, that's equivalent to multiplication by 4 (because each array element has size of 4 bytes) and add it to address of array begin.

The second loop has inverse “LDR R2, [SP, R4, LSL#2]”, instruction, it loads from array value we need, and the pointer to it is calculated in exactly the same way.

ARM + Optimizing Keil + thumb mode


```

_main
    PUSH    {R4,R5,LR}
    SUB     SP, SP, #0x54    ; allocate place for 20 int variables + one more variable

; first loop

    MOV     R0, #0           ; i
    MOV     R5, SP           ; pointer to first array element

loc_1CE
    LSLS    R1, R0, #1       ; R1=i<<1 (same as i*2)
    LSLS    R2, R0, #2       ; R2=i<<2 (same as i*4)
    ADDS    R0, R0, #1       ; i=i+1
    CMP     R0, #20          ; i<20?
    STR     R1, [R5,R2]      ; store R1 to *(R5+R2) (same R5+i*4)
    BLT     loc_1CE          ; yes, i<20, run loop body again

; second loop

    MOV     R4, #0           ; i=0

loc_1DC
    LSLS    R0, R4, #2       ; R0=i<<2 (same as i*4)
    LDR     R2, [R5,R0]      ; load from *(R5+R0) (same as R5+i*4)
    MOV     R1, R4
    ADR     R0, aADD         ; "a[%d]=%d\n"
    BL      __2printf
    ADDS    R4, R4, #1       ; i=i+1
    CMP     R4, #20          ; i<20?
    BLT     loc_1DC          ; yes, i<20, run loop body again
    MOV     R0, #0           ; value to return
    ADD     SP, SP, #0x54    ; deallocate place for 20 int variables + one more variable
    POP     {R4,R5,PC}

```

Thumb code is very similar. Thumb mode has special instructions for bit shifting (like LSLS), calculating value to be written into array and address of each element in array as well.

Compiler allocates slightly more space in local stack, however, last 4 bytes are not used.

1.14.2 Buffer overflow

So, array indexing is just `array[index]`. If you study generated code closely, you'll probably note missing index bounds checking, which could check index, *if it is less than 20*. What if index will be greater than 20? That's the one C/C++ feature it's often blamed for.

Here is a code successfully compiling and working:

```

#include <stdio.h>

int main()
{
    int a[20];
    int i;

    for (i=0; i<20; i++)
        a[i]=i*2;

    printf ("a[100]=%d\n", a[100]);

    return 0;
};

```

Compilation results (MSVC 2010):

```

_TEXT    SEGMENT
_i$ = -84                ; size = 4
_a$ = -80                ; size = 80

```

```

_main      PROC
    push    ebp
    mov     ebp, esp
    sub     esp, 84          ; 00000054H
    mov     DWORD PTR _i$[ebp], 0
    jmp     SHORT $LN3@main
$LN2@main:
    mov     eax, DWORD PTR _i$[ebp]
    add     eax, 1
    mov     DWORD PTR _i$[ebp], eax
$LN3@main:
    cmp     DWORD PTR _i$[ebp], 20 ; 00000014H
    jge     SHORT $LN1@main
    mov     ecx, DWORD PTR _i$[ebp]
    shl     ecx, 1
    mov     edx, DWORD PTR _i$[ebp]
    mov     DWORD PTR _a$[ebp+edx*4], ecx
    jmp     SHORT $LN2@main
$LN1@main:
    mov     eax, DWORD PTR _a$[ebp+400]
    push    eax
    push    OFFSET $SG2460
    call    _printf
    add     esp, 8
    xor     eax, eax
    mov     esp, ebp
    pop     ebp
    ret     0
_main      ENDP

```

I'm running it, and I got:

```
a[100]=760826203
```

It is just *something*, occasionally lying in the stack near to array, 400 bytes from its first element.

Indeed, how it could be done differently? Compiler may incorporate some code, checking index value to be always in array's bound, like in higher-level programming languages⁷³, but this makes running code slower.

OK, we read some values in stack *illegally*, but what if we could write something to it?

Here is what we will write:

```

#include <stdio.h>

int main()
{
    int a[20];
    int i;

    for (i=0; i<30; i++)
        a[i]=i;

    return 0;
};

```

And what we've got:

```

_TEXT      SEGMENT
_i$ = -84          ; size = 4
_a$ = -80          ; size = 80
_main      PROC
    push    ebp
    mov     ebp, esp
    sub     esp, 84          ; 00000054H
    mov     DWORD PTR _i$[ebp], 0
    jmp     SHORT $LN3@main

```

⁷³ Java, Python, etc

```

$LN2@main:
mov     eax, DWORD PTR _i$[ebp]
add     eax, 1
mov     DWORD PTR _i$[ebp], eax
$LN3@main:
cmp     DWORD PTR _i$[ebp], 30      ; 0000001eH
jge     SHORT $LN1@main
mov     ecx, DWORD PTR _i$[ebp]
mov     edx, DWORD PTR _i$[ebp]    ; that insruction is obviously redundant
mov     DWORD PTR _a$[ebp+ecx*4], edx ; ECX could be used as second operand here instead
jmp     SHORT $LN2@main
$LN1@main:
xor     eax, eax
mov     esp, ebp
pop     ebp
ret     0
_main   ENDP

```

Run compiled program and its crashing. No wonder. Let's see, where exactly it's crashing.

I'm not using debugger anymore, because I tired to run it each time, move mouse, etc, when I need just to spot some register's state at specific point. That's why I wrote very minimalistic tool for myself, *tracer* [5.0.1](#), which is enough for my tasks.

I can also use it just to see, where debuggee is crashed. So let's see:

```

generic tracer 0.4 (WIN32), http://conus.info/gt

New process: C:\PRJ\...\1.exe, PID=7988
EXCEPTION_ACCESS_VIOLATION: 0x15 (<symbol (0x15) is in unknown module>), ExceptionInformation[0]=8
EAX=0x00000000 EBX=0x7EFDE000 ECX=0x0000001D EDX=0x0000001D
ESI=0x00000000 EDI=0x00000000 EBP=0x00000014 ESP=0x0018FF48
EIP=0x00000015
FLAGS=PF ZF IF RF
PID=7988|Process exit, return code -1073740791

```

Now please keep your eyes on registers.

Exception occurred at address 0x15. It's not legal address for code — at least for win32 code! We trapped there somehow against our will. It's also interesting fact that EBP register contain 0x14, ECX and EDX — 0x1D.

Let's study stack layout more.

After control flow was passed into `main()`, EBP register value was saved into stack. Then, 84 bytes was allocated for array and *i* variable. That's $(20+1)*sizeof(int)$. ESP pointing now to `_i` variable in local stack and after execution of next `PUSH something`, *something* will be appeared next to `_i`.

That's stack layout while control is inside `main()`:

ESP	4 bytes for <i>i</i>
ESP+4	80 bytes for a[20] array
ESP+84	saved EBP value
ESP+88	returning address

Instruction `a[19]=something` writes last *int* in array bounds (in bounds yet!)

Instruction `a[20]=something` writes *something* to the place where EBP value is saved.

Please take a look at registers state at the crash moment. In our case, number 20 was written to 20th element. By the function ending, function epilogue restore EBP value. (20 in decimal system is 0x14 in hexadecimal). Then, `RET` instruction was executed, which is equivalent to `POP EIP` instruction.

`RET` instruction taking returning address from stack (that's address in some CRT⁷⁴-function, which was called `main()`), and 21 was stored there (0x15 in hexadecimal). The CPU trapped at the address 0x15, but there are no executable code, so exception was raised.

Welcome! It's called *buffer overflow*⁷⁵.

⁷⁴C Run-Time

⁷⁵http://en.wikipedia.org/wiki/Stack_buffer_overflow

Replace *int* array by string (*char* array), create a long string deliberately, pass it to the program, to the function which is not checking string length and copies it to short buffer, and you'll be able to point to a program an address to which it should jump. Not that simple in reality, but that's how it was emerged⁷⁶.

Let's try the same code in GCC 4.4.1. We got:

```
main      public main
          proc near

a          = dword ptr -54h
i          = dword ptr -4

          push    ebp
          mov     ebp, esp
          sub     esp, 60h
          mov     [ebp+i], 0
          jmp     short loc_80483D1

loc_80483C3:
          mov     eax, [ebp+i]
          mov     edx, [ebp+i]
          mov     [ebp+eax*4+a], edx
          add     [ebp+i], 1

loc_80483D1:
          cmp     [ebp+i], 1Dh
          jle     short loc_80483C3
          mov     eax, 0
          leave
          retn

main      endp
```

Running this in Linux will produce: Segmentation fault.

If we run this in GDB debugger, we get this:

```
(gdb) r
Starting program: /home/dennis/RE/1

Program received signal SIGSEGV, Segmentation fault.
0x00000016 in ?? ()
(gdb) info registers
eax          0x0          0
ecx          0xd2f96388    -755407992
edx          0x1d         29
ebx          0x26eff4      2551796
esp          0xbffff4b0    0xbffff4b0
ebp          0x15         0x15
esi          0x0          0
edi          0x0          0
eip          0x16         0x16
eflags      0x10202      [ IF RF ]
cs           0x73         115
ss           0x7b         123
ds           0x7b         123
es           0x7b         123
fs           0x0          0
gs           0x33         51
(gdb)
```

Register values are slightly different than in win32 example, that's because stack layout is slightly different too.

1.14.3 Buffer overflow protection methods

There are several methods to protect against it, regardless of C/C++ programmers' negligence. MSVC has options like⁷⁷:

⁷⁶Classic article about it: [Smashing The Stack For Fun And Profit](#)

⁷⁷[Wikipedia: compiler-side buffer overflow protection methods](#)

```
/RTCs Stack Frame runtime checking
/GZ Enable stack checks (/RTCs)
```

One of the methods is to write random value among local variables to stack at function prologue and to check it in function epilogue before function exiting. And if value is not the same, do not execute last instruction RET, but halt (or hang). Process will hang, but that's much better then remote attack to your host.

This random value is called “canary” sometimes, it's related to miner's canary⁷⁸, they were used by miners in these days, in order to detect poisonous gases quickly. Canaries are very sensitive to mine gases, they become very agitated in case of danger, or even dead.

If to compile our very simple array example 1.14.1 in MSVC with RTC1 and RTCs option, you will see call to @_RTC_CheckStackVars@8 function at the function end, checking “canary” correctness.

Optimizing Xcode (LLVM) + thumb-2 mode

Let's back to our simple array example 1.14.1, again, now we can see how LLVM will check “canary” correctness:

```
_main
var_64      = -0x64
var_60      = -0x60
var_5C      = -0x5C
var_58      = -0x58
var_54      = -0x54
var_50      = -0x50
var_4C      = -0x4C
var_48      = -0x48
var_44      = -0x44
var_40      = -0x40
var_3C      = -0x3C
var_38      = -0x38
var_34      = -0x34
var_30      = -0x30
var_2C      = -0x2C
var_28      = -0x28
var_24      = -0x24
var_20      = -0x20
var_1C      = -0x1C
var_18      = -0x18
canary      = -0x14
var_10      = -0x10

        PUSH        {R4-R7,LR}
        ADD         R7, SP, #0xC
        STR.W       R8, [SP,#0xC+var_10]!
        SUB         SP, SP, #0x54
        MOVW        R0, #a0bjc_methtype ; "objc_methtype"
        MOVS        R2, #0
        MOVT.W      R0, #0
        MOVS        R5, #0
        ADD         R0, PC
        LDR.W       R8, [R0]
        LDR.W       R0, [R8]
        STR         R0, [SP,#0x64+canary]
        MOVS        R0, #2
        STR         R2, [SP,#0x64+var_64]
        STR         R0, [SP,#0x64+var_60]
        MOVS        R0, #4
        STR         R0, [SP,#0x64+var_5C]
        MOVS        R0, #6
        STR         R0, [SP,#0x64+var_58]
        MOVS        R0, #8
        STR         R0, [SP,#0x64+var_54]
```

⁷⁸ [Wikipedia: Miner's canary](#)

```

        MOVS      R0, #0xA
        STR       R0, [SP,#0x64+var_50]
        MOVS      R0, #0xC
        STR       R0, [SP,#0x64+var_4C]
        MOVS      R0, #0xE
        STR       R0, [SP,#0x64+var_48]
        MOVS      R0, #0x10
        STR       R0, [SP,#0x64+var_44]
        MOVS      R0, #0x12
        STR       R0, [SP,#0x64+var_40]
        MOVS      R0, #0x14
        STR       R0, [SP,#0x64+var_3C]
        MOVS      R0, #0x16
        STR       R0, [SP,#0x64+var_38]
        MOVS      R0, #0x18
        STR       R0, [SP,#0x64+var_34]
        MOVS      R0, #0x1A
        STR       R0, [SP,#0x64+var_30]
        MOVS      R0, #0x1C
        STR       R0, [SP,#0x64+var_2C]
        MOVS      R0, #0x1E
        STR       R0, [SP,#0x64+var_28]
        MOVS      R0, #0x20
        STR       R0, [SP,#0x64+var_24]
        MOVS      R0, #0x22
        STR       R0, [SP,#0x64+var_20]
        MOVS      R0, #0x24
        STR       R0, [SP,#0x64+var_1C]
        MOVS      R0, #0x26
        STR       R0, [SP,#0x64+var_18]
        MOV       R4, 0xFDA ; "a[%d]=%d\n"
        MOV       R0, SP
        ADDS      R6, R0, #4
        ADD       R4, PC
        B         loc_2F1C

; second loop begin

loc_2F14
        ADDS      R0, R5, #1
        LDR.W     R2, [R6,R5,LSL#2]
        MOV       R5, R0

loc_2F1C
        MOV       R0, R4
        MOV       R1, R5
        BLX       _printf
        CMP       R5, #0x13
        BNE       loc_2F14
        LDR.W     R0, [R8]
        LDR       R1, [SP,#0x64+canary]
        CMP       R0, R1
        ITTTT EQ      ; canary still correct?
        MOVEQ     R0, #0
        ADDEQ     SP, SP, #0x54
        LDREQ.W   R8, [SP+0x64+var_64],#4
        POPEQ     {R4-R7,PC}
        BLX       ___stack_chk_fail

```

First of all, as we see, LLVM made loop “unrolled” and all values are written into array one-by-one, already calculated, because LLVM concluded it will be faster. By the way, ARM mode instructions may help to do this even faster, and finding this way could be your homework.

At the function end we see “canaries” comparison — that laying in local stack and correct one, to which R8 register pointing. If they are equal to each other, 4-instruction block is triggered by “ITTTT EQ”, it is writing 0 into R0, function epilogue and exit. If “canaries” are not equal, block will not be triggered, and jump to

___stack_chk_fail function will be occurred, which, as I suppose, will halt execution.

1.14.4 One more word about arrays

Now we understand, why it's not possible to write something like that in C/C++ code ⁷⁹:

```
void f(int size)
{
    int a[size];
    ...
};
```

That's just because compiler should know exact array size to allocate place for it in local stack layout or in data segment (in case of global variable) on compiling stage.

If you need array of arbitrary size, allocate it by `malloc()`, then access allocated memory block as array of variables of type you need. Or use C99 standard feature [ISO07, 6.7.5.2], but it will be looks like `alloca()` 1.2.3 internally.

1.14.5 Multidimensional arrays

Internally, multidimensional array is essentially the same thing as linear array.

Because computer memory in linear, it's one-dimensional array. But this one-dimensional array can be easily represented as multidimensional for convenience.

For example, that's how `a[3][4]` array elements will be placed in one-dimensional array of 12 cells:

0	1	2	3
4	5	6	7
8	9	10	11

So, in order to address elements we need, first multiply first index by 4 (matrix width) and then add second index. That's called *row-major order*, and this method of arrays and matrices representation is used in at least in C/C++, Python. *row-major order* term in plain English language mean: "first, write elements of first row, then second row ... and finally elements of last row".

Another method of representation called *column-major order* (array indices used in reverse order) and it's used at least in FORTRAN, MATLAB, R. *column-major order* term in plain English language mean: "first, write elements of first column, then second column ... and finally elements of last column".

Same thing about multidimensional arrays.

Let's see:

Listing 1.54: simple example

```
#include <stdio.h>

int a[10][20][30];

void insert(int x, int y, int z, int value)
{
    a[x][y][z]=value;
};
```

x86

We got (MSVC 2010):

Listing 1.55: MSVC 2010

```
_DATA    SEGMENT
COMM     _a:DWORD:01770H
```

⁷⁹However, it's possible in C99 standard [ISO07, 6.7.5.2]: GCC is actually do this by allocating array dynamically in stack (like `alloca()` 1.2.3)

```

_DATA    ENDS
PUBLIC  _insert
_TEXT   SEGMENT
_x$ = 8          ; size = 4
_y$ = 12         ; size = 4
_z$ = 16         ; size = 4
_value$ = 20     ; size = 4
_insert  PROC
    push    ebp
    mov     ebp, esp
    mov     eax, DWORD PTR _x$[ebp]
    imul    eax, 2400          ; eax=600*4*x
    mov     ecx, DWORD PTR _y$[ebp]
    imul    ecx, 120          ; ecx=30*4*y
    lea     edx, DWORD PTR _a[eax+ecx] ; edx=a + 600*4*x + 30*4*y
    mov     eax, DWORD PTR _z$[ebp]
    mov     ecx, DWORD PTR _value$[ebp]
    mov     DWORD PTR [edx+eax*4], ecx ; *(edx+z*4)=value
    pop     ebp
    ret     0
_insert  ENDP
_TEXT   ENDS

```

Nothing special. For index calculation, three input arguments are multiplying by formula $address = 600 \cdot 4 \cdot x + 30 \cdot 4 \cdot y + 4z$ to represent array as multidimensional. Do not forget that *int* type is 32-bit (4 bytes), so all coefficients should be multiplied by 4.

Listing 1.56: GCC 4.4.1

```

insert      public insert
            proc near

x            = dword ptr 8
y            = dword ptr 0Ch
z            = dword ptr 10h
value       = dword ptr 14h

            push    ebp
            mov     ebp, esp
            push    ebx
            mov     ebx, [ebp+x]
            mov     eax, [ebp+y]
            mov     ecx, [ebp+z]
            lea     edx, [eax+eax]          ; edx=y*2
            mov     eax, edx                ; eax=y*2
            shl     eax, 4                  ; eax=(y*2)<<4 = y*2*16 = y*32
            sub     eax, edx                ; eax=y*32 - y*2=y*30
            imul    edx, ebx, 600           ; edx=x*600
            add     eax, edx                ; eax=eax+edx=y*30 + x*600
            lea     edx, [eax+ecx]         ; edx=y*30 + x*600 + z
            mov     eax, [ebp+value]
            mov     dword ptr ds:a[edx*4], eax ; *(a+edx*4)=value
            pop     ebx
            pop     ebp
            retn
insert      endp

```

GCC compiler does it differently. For one of operations calculating $(30y)$, GCC produced a code without multiplication instruction. This is how it done: $(y+y) \ll 4 - (y+y) = (2y) \ll 4 - 2y = 2 \cdot 16 \cdot y - 2y = 32y - 2y = 30y$. Thus, for $30y$ calculation, only one addition operation used, one bitwise shift operation and one subtraction operation. That works faster.

ARM + Non-optimizing Xcode (LLVM) + thumb mode

Listing 1.57: Non-optimizing Xcode (LLVM) + thumb mode

```

_insert
value      = -0x10
z          = -0xC
y          = -8
x          = -4

        SUB        SP, SP, #0x10 ; allocate place in local stack for 4 int values
        MOV        R9, 0xFC2 ; a
        ADD        R9, PC
        LDR.W      R9, [R9]
        STR        R0, [SP,#0x10+x]
        STR        R1, [SP,#0x10+y]
        STR        R2, [SP,#0x10+z]
        STR        R3, [SP,#0x10+value]
        LDR        R0, [SP,#0x10+value]
        LDR        R1, [SP,#0x10+z]
        LDR        R2, [SP,#0x10+y]
        LDR        R3, [SP,#0x10+x]
        MOV        R12, 2400
        MUL.W      R3, R3, R12
        ADD        R3, R9
        MOV        R9, 120
        MUL.W      R2, R2, R9
        ADD        R2, R3
        LSLS       R1, R1, #2 ; R1=R1<<2
        ADD        R1, R2
        STR        R0, [R1] ; R1 - address of array element
        ADD        SP, SP, #0x10 ; deallocate place in local stack for 4 int values
        BX        LR

```

Non-optimizing LLVM saves all variables in local stack, however, it's redundant. Address of array element is calculated by formula we already figured out.

ARM + Optimizing Xcode (LLVM) + thumb mode

Listing 1.58: Optimizing Xcode (LLVM) + thumb mode

```

_insert
MOVW     R9, #0x10FC
MOV.W    R12, #2400
MOVT.W   R9, #0
RSB.W    R1, R1, R1, LSL#4 ; R1 = y. R1=y<<4 - y = y*16 - y = y*15
ADD      R9, PC           ; R9 = pointer to a array
LDR.W    R9, [R9]
MLA.W    R0, R0, R12, R9 ; R0 = x, R12 = 2400, R9 = pointer to a. R0=x*2400 + ptr to a
ADD.W    R0, R0, R1, LSL#3 ; R0 = R0+R1<<3 = R0+R1*8 = x*2400 + ptr to a + y*15*8 =
                        ; ptr to a + y*30*4 + x*600*4
STR.W    R3, [R0,R2,LSL#2] ; R2 = z, R3 = value. address=R0+z*4 =
                        ; ptr to a + y*30*4 + x*600*4 + z*4
BX       LR

```

Here is used tricks for replacing multiplication by shift, addition and subtraction we already considered.

Here we also see new instruction for us: RSB (*Reverse Subtract*). It works just as SUB, but swapping operands with each other. Why? SUB, RSB, are those instructions, to the second operand of which shift coefficient may be applied: (LSL#4). But this coefficient may be applied only to second operand. That's fine for commutative operations like addition or multiplication, operands may be swapped there without result affecting. But subtraction is non-commutative operation, so, for these cases, RSB exist.

“LDR.W R9, [R9]” works like LEA 2.1 in x86, but here it does nothing, it's redundant. Apparently, compiler not optimized it.

1.15 Bit fields

A lot of functions defining input flags in arguments using bit fields. Of course, it could be substituted by *bool*-typed variables set, but it's not frugally.

1.15.1 Specific bit checking

x86

Win32 API example:

```
HANDLE fh;

fh=CreateFile ("file", GENERIC_WRITE | GENERIC_READ, FILE_SHARE_READ, NULL, OPEN_ALWAYS,
FILE_ATTRIBUTE_NORMAL, NULL);
```

We got (MSVC 2010):

Listing 1.59: MSVC 2010

```
push    0
push    128                                ; 00000080H
push    4
push    0
push    1
push    -1073741824                       ; c0000000H
push    OFFSET $SG78813
call    DWORD PTR __imp__CreateFileA@28
mov     DWORD PTR _fh$[ebp], eax
```

Let's take a look into WinNT.h:

Listing 1.60: WinNT.h

```
#define GENERIC_READ      (0x80000000L)
#define GENERIC_WRITE     (0x40000000L)
#define GENERIC_EXECUTE   (0x20000000L)
#define GENERIC_ALL       (0x10000000L)
```

Everything is clear, `GENERIC_READ | GENERIC_WRITE = 0x80000000 | 0x40000000 = 0xC0000000`, and that's value is used as second argument for `CreateFile()`⁸⁰ function.

How `CreateFile()` will check flags?

Let's take a look into `KERNEL32.DLL` in Windows XP SP3 x86 and we'll find this fragment of code in the function `CreateFileW`:

Listing 1.61: `KERNEL32.DLL` (Windows XP SP3 x86)

```
.text:7C83D429      test    byte ptr [ebp+dwDesiredAccess+3], 40h
.text:7C83D42D      mov     [ebp+var_8], 1
.text:7C83D434      jz      short loc_7C83D417
.text:7C83D436      jmp     loc_7C810817
```

Here we see `TEST` instruction, it takes, however, not the whole second argument, but only most significant byte (`ebp+dwDesiredAccess+3`) and checks it for `0x40` flag (meaning `GENERIC_WRITE` flag here)

`TEST` is merely the same instruction as `AND`, but without result saving (recall the fact `CMP` instruction is merely the same as `SUB`, but without result saving 1.4.5).

This fragment of code logic is as follows:

```
if ((dwDesiredAccess&0x40000000) == 0) goto loc_7C83D417
```

If `AND` instruction leaving this bit, `ZF` flag will be cleared and `JZ` conditional jump will not be triggered. Conditional jump will be triggered only if `0x40000000` bit is absent in `dwDesiredAccess` variable — then `AND` result will be 0, `ZF` flag will be set and conditional jump is to be triggered.

Let's try `GCC 4.4.1` and `Linux`:

⁸⁰ [MSDN: CreateFile function](#)

```
#include <stdio.h>
#include <fcntl.h>

void main()
{
    int handle;

    handle=open ("file", O_RDWR | O_CREAT);
};
```

We got:

```
main      public main
          proc near

var_20     = dword ptr -20h
var_1C     = dword ptr -1Ch
var_4      = dword ptr -4

          push    ebp
          mov     ebp, esp
          and     esp, 0FFFFFFF0h
          sub     esp, 20h
          mov     [esp+20h+var_1C], 42h
          mov     [esp+20h+var_20], offset aFile ; "file"
          call    _open
          mov     [esp+20h+var_4], eax
          leave
          retn
main      endp
```

[caption=GCC 4.4.1]

Let's take a look into `open()` function in `libc.so.6` library, but there is only `syscall` calling:

Listing 1.62: `open()` (`libc.so.6`)

```
.text:000BE69B      mov     edx, [esp+4+mode] ; mode
.text:000BE69F      mov     ecx, [esp+4+flags] ; flags
.text:000BE6A3      mov     ebx, [esp+4+filename] ; filename
.text:000BE6A7      mov     eax, 5
.text:000BE6AC      int     80h                ; LINUX - sys_open
```

So, `open()` bit fields apparently checked somewhere in Linux kernel.

Of course, it is easily to download both Glibc and Linux kernel source code, but we are interesting to understand the matter without it.

So, as of Linux 2.6, when `sys_open` syscall is called, control eventually passed into `do_sys_open` kernel function. From there — to `do_filp_open()` function (this function located in kernel source tree in the file `fs/namei.c`).

Important note. Aside from usual passing arguments via stack, there are also method to pass some of them via registers. This is also called `fastcall` 2.5.3. This works faster, because CPU not needed to access a stack in memory to read argument values. GCC has option `regparm`⁸¹, and it's possible to set a number of arguments which might be passed via registers.

Linux 2.6 kernel compiled with `-mregparm=3` option^{82 83}.

What it means to us, the first 3 arguments will be passed via `EAX`, `EDX` and `ECX` registers, the other ones via stack. Of course, if arguments number is less than 3, only part of registers will be used.

So, let's download Linux Kernel 2.6.31, compile it in Ubuntu: `make vmlinux`, open it in IDA 5, find the `do_filp_open()` function. At the beginning, we will see (comments are mine):

```
do_filp_open  proc near
...
          push    ebp
```

⁸¹<http://ohse.de/uwe/articles/gcc-attributes.html#func-regparm>

⁸²http://kernelnewbies.org/Linux_2_6_20#head-042c62f290834eb1fe0a1942bbf5bb9a4accbc8f

⁸³See also `arch\i386\include\asm\calling.h` file in kernel tree

```

mov     ebp, esp
push    edi
push    esi
push    ebx
mov     ebx, ecx
add     ebx, 1
sub     esp, 98h
mov     esi, [ebp+arg_4] ; acc_mode (5th arg)
test    bl, 3
mov     [ebp+var_80], eax ; dfd (1th arg)
mov     [ebp+var_7C], edx ; pathname (2th arg)
mov     [ebp+var_78], ecx ; open_flag (3th arg)
jnz     short loc_C01EF684
mov     ebx, ecx          ; ebx <- open_flag

```

[caption=do_filp_open() (linux kernel 2.6.31)]

GCC saves first 3 arguments values in local stack. Otherwise, if compiler would not touch these registers, it would be too tight environment for compiler's register allocator.

Let's find this fragment of code:

```

loc_C01EF6B4:                ; CODE XREF: do_filp_open+4F
                                ; O_CREAT
test     bl, 40h
jnz     loc_C01EF810
mov     edi, ebx
shr     edi, 11h
xor     edi, 1
and     edi, 1
test     ebx, 10000h
jz      short loc_C01EF6D3
or      edi, 2

```

[caption=do_filp_open() (linux kernel 2.6.31)]

0x40 — is what O_CREAT macro equals to. open_flag checked for 0x40 bit presence, and if this bit is 1, next JNZ instruction is triggered.

ARM

O_CREAT bit is checked differently in Linux kernel 3.8.0.

Listing 1.63: linux kernel 3.8.0

```

struct file *do_filp_open(int dfd, struct filename *pathname,
                          const struct open_flags *op)
{
    ...
    filp = path_openat(dfd, pathname, &nd, op, flags | LOOKUP_RCU);
    ...
}

static struct file *path_openat(int dfd, struct filename *pathname,
                                struct nameidata *nd, const struct open_flags *op, int flags)
{
    ...
    error = do_last(nd, &path, file, op, &opened, pathname);
    ...
}

static int do_last(struct nameidata *nd, struct path *path,
                  struct file *file, const struct open_flags *op,
                  int *opened, struct filename *name)
{
    ...
    if (!(open_flag & O_CREAT)) {
        ...
    }
}

```

```

        error = lookup_fast(nd, path, &inode);
    ...
    } else {
    ...
        error = complete_walk(nd);
    }
    ...
}

```

Here is how kernel compiled for ARM mode looks like in IDA 5:

Listing 1.64: do_last() (vmlinux)

```

...
.text:C0169EA8      MOV            R9, R3 ; R3 - (4th argument) open_flag
...
.text:C0169ED4      LDR            R6, [R9] ; R6 - open_flag
...
.text:C0169F68      TST            R6, #0x40 ; jumtable C0169F00 default case
.text:C0169F6C      BNE            loc_C016A128
.text:C0169F70      LDR            R2, [R4,#0x10]
.text:C0169F74      ADD            R12, R4, #8
.text:C0169F78      LDR            R3, [R4,#0xC]
.text:C0169F7C      MOV            R0, R4
.text:C0169F80      STR            R12, [R11,#var_50]
.text:C0169F84      LDRB           R3, [R2,R3]
.text:C0169F88      MOV            R2, R8
.text:C0169F8C      CMP            R3, #0
.text:C0169F90      ORRNE          R1, R1, #3
.text:C0169F94      STRNE          R1, [R4,#0x24]
.text:C0169F98      ANDS           R3, R6, #0x200000
.text:C0169F9C      MOV            R1, R12
.text:C0169FA0      LDRNE          R3, [R4,#0x24]
.text:C0169FA4      ANDNE          R3, R3, #1
.text:C0169FA8      EORNE          R3, R3, #1
.text:C0169FAC      STR            R3, [R11,#var_54]
.text:C0169FB0      SUB            R3, R11, #-var_38
.text:C0169FB4      BL             lookup_fast
...
.text:C016A128      loc_C016A128                                ; CODE XREF: do_last.isra.14+DC
.text:C016A128      MOV            R0, R4
.text:C016A12C      BL             complete_walk
...

```

TST is analogical to TEST instruction in x86.

We can “spot” visually this code fragment by the fact that `lookup_fast()` will be executed in one case and `complete_walk()` in another case. This is corresponding to `do_last()` function source code.

`O_CREAT` macro is equals to `0x40` here too.

1.15.2 Specific bit setting/clearing

For example:

```

#define IS_SET(flag, bit)    ((flag) & (bit))
#define SET_BIT(var, bit)    ((var) |= (bit))
#define REMOVE_BIT(var, bit) ((var) &= ~(bit))

int f(int a)
{
    int rt=a;

    SET_BIT (rt, 0x4000);
    REMOVE_BIT (rt, 0x200);

    return rt;
};

```

x86

We got (MSVC 2010):

Listing 1.65: MSVC 2010

```
_rt$ = -4          ; size = 4
_a$ = 8           ; size = 4
_f PROC
    push    ebp
    mov     ebp, esp
    push    ecx
    mov     eax, DWORD PTR _a$[ebp]
    mov     DWORD PTR _rt$[ebp], eax
    mov     ecx, DWORD PTR _rt$[ebp]
    or      ecx, 16384          ; 00004000H
    mov     DWORD PTR _rt$[ebp], ecx
    mov     edx, DWORD PTR _rt$[ebp]
    and     edx, -513          ; ffffffffH
    mov     DWORD PTR _rt$[ebp], edx
    mov     eax, DWORD PTR _rt$[ebp]
    mov     esp, ebp
    pop     ebp
    ret     0
_f ENDP
```

OR instruction adding one more bit to value, ignoring others.

AND resetting one bit. It can be said, AND just copies all bits except one. Indeed, in the second AND operand only those bits are set, which are needed to be saved, except one bit we wouldn't like to copy (which is 0 in bitmask). It's easier way to memorize the logic.

If we compile it in MSVC with optimization turned on (/Ox), the code will be even shorter:

Listing 1.66: Optimizing MSVC

```
_a$ = 8           ; size = 4
_f PROC
    mov     eax, DWORD PTR _a$[esp-4]
    and     eax, -513          ; ffffffffH
    or      eax, 16384          ; 00004000H
    ret     0
_f ENDP
```

Let's try GCC 4.4.1 without optimization:

Listing 1.67: Non-optimizing GCC

```
f          public f
           proc near

var_4      = dword ptr -4
arg_0      = dword ptr 8

           push    ebp
           mov     ebp, esp
           sub     esp, 10h
           mov     eax, [ebp+arg_0]
           mov     [ebp+var_4], eax
           or      [ebp+var_4], 4000h
           and     [ebp+var_4], 0FFFFFFDFFh
           mov     eax, [ebp+var_4]
           leave
           retn
f          endp
```

There are some redundant code present, however, it's shorter than MSVC version without optimization. Now let's try GCC with optimization turned on -O3:

Listing 1.68: Optimizing GCC

```

f      public f
      proc near
arg_0      = dword ptr 8

      push    ebp
      mov     ebp, esp
      mov     eax, [ebp+arg_0]
      pop     ebp
      or      ah, 40h
      and     ah, 0FDh
      retn
f      endp

```

That's shorter. It is important to note that compiler works with EAX register part via AH register — that's EAX register part from 8th to 15th bits inclusive.

Important note: 16-bit CPU 8086 accumulator was named AX and consisted of two 8-bit halves — AL (lower byte) and AH (higher byte). In 80386 almost all registers were extended to 32-bit, accumulator was named EAX, but for the sake of compatibility, its *older parts* may be still accessed as AX/AH/AL registers.

Because all x86 CPUs are 16-bit 8086 CPU successors, these *older* 16-bit opcodes are shorter than newer 32-bit opcodes. That's why “or ah, 40h” instruction occupying only 3 bytes. It would be more logical way to emit here “or eax, 04000h”, but that's 5 bytes, or even 6 (if register in first operand is not EAX).

It would be even shorter if to turn on -O3 optimization flag and also set regparm=3.

Listing 1.69: Optimizing GCC

```

f      public f
      proc near
      push    ebp
      or      ah, 40h
      mov     ebp, esp
      and     ah, 0FDh
      pop     ebp
      retn
f      endp

```

Indeed — first argument is already loaded into EAX, so it's possible to work with it in-place. It's worth noting that both function prologue (“push ebp / mov ebp, esp”) and epilogue (“pop ebp”) can easily be omitted here, but GCC probably isn't good enough for such code size optimizations. However, such short functions are better to be *inlined functions*⁸⁴.

ARM + Optimizing Keil + ARM mode

Listing 1.70: Optimizing Keil + ARM mode

```

02 0C C0 E3      BIC      R0, R0, #0x200
01 09 80 E3      ORR      R0, R0, #0x4000
1E FF 2F E1      BX       LR

```

BIC is “logical and”, analogical to AND in x86. ORR is “logical or”, analogical to OR in x86.

So far, so easy.

ARM + Optimizing Keil + thumb mode

Listing 1.71: Optimizing Keil + thumb mode

```

01 21 89 03      MOVS     R1, 0x4000
08 43            ORRS     R0, R1
49 11            ASRS     R1, R1, #5    ; generate 0x200 and place to R1

```

⁸⁴http://en.wikipedia.org/wiki/Inline_function

88 43	BICS	R0, R1
70 47	BX	LR

Apparently, Keil concludes that the code in thumb mode, making $0x200$ from $0x4000$, will be more compact than code, writting $0x200$ to arbitrary register.

So that's why, with the help of ASRS (arithmetic shift right), this value is calculating as $0x4000 \gg 5$.

ARM + Optimizing Xcode (LLVM) + ARM mode

Listing 1.72: Optimizing Xcode (LLVM) + ARM mode

42 0C C0 E3	BIC	R0, R0, #0x4200
01 09 80 E3	ORR	R0, R0, #0x4000
1E FF 2F E1	BX	LR

The code was generated by LLVM, in source code form, in fact, could be looks like:

```
REMOVE_BIT (rt, 0x4200);
SET_BIT (rt, 0x4000);
```

And it does exactly the same we need. But why $0x4200$? Perhaps, that's LLVM optimizer's artifact⁸⁵. Probably, compiler's optimizer error, but generated code works correct anyway.

More about compiler's anomalies, read here [8](#).

For thumb mode, Optimizing Xcode (LLVM) generates exactly the same code.

1.15.3 Shifts

Bit shifts in C/C++ are implemented via \ll and \gg operators.

Here is a simple example of function, calculating number of 1 bits in input variable:

```
#define IS_SET(flag, bit)      ((flag) & (bit))

int f(unsigned int a)
{
    int i;
    int rt=0;

    for (i=0; i<32; i++)
        if (IS_SET (a, 1<<i))
            rt++;

    return rt;
};
```

In this loop, iteration count value i counting from 0 to 31, $1 \ll i$ statement will be counting from 1 to $0x80000000$. Describing this operation in natural language, we would say *shift 1 by n bits left*. In other words, $1 \ll i$ statement will consequently produce all possible bit positions in 32-bit number. By the way, freed bit at right is always cleared. IS_SET macro is checking bit presence in a.




Figure 1.1: How SHL instruction works⁸⁶

⁸⁵It was LLVM build 2410.2.00 bundled with Apple Xcode 4.6.3

⁸⁶illustration taken from wikipedia

The IS_SET macro is in fact logical and operation (*AND*) and it returns zero if specific bit is absent there, or bit mask, if the bit is present. *if()* operator triggered in C/C++ if expression in it isn't zero, it might be even 123456, that's why it always working correctly.

x86

Let's compile (MSVC 2010):

Listing 1.73: MSVC 2010

```
_rt$ = -8          ; size = 4
_i$ = -4          ; size = 4
_a$ = 8           ; size = 4
_f PROC
    push    ebp
    mov     ebp, esp
    sub     esp, 8
    mov     DWORD PTR _rt$[ebp], 0
    mov     DWORD PTR _i$[ebp], 0
    jmp     SHORT $LN4@f
$LN3@f:
    mov     eax, DWORD PTR _i$[ebp]    ; increment of 1
    add     eax, 1
    mov     DWORD PTR _i$[ebp], eax
$LN4@f:
    cmp     DWORD PTR _i$[ebp], 32     ; 00000020H
    jge     SHORT $LN2@f               ; loop finished?
    mov     edx, 1
    mov     ecx, DWORD PTR _i$[ebp]
    shl     edx, cl                    ; EDX=EDX<<CL
    and     edx, DWORD PTR _a$[ebp]
    je      SHORT $LN1@f               ; result of AND instruction was 0?
                                         ; then skip next instructions
    mov     eax, DWORD PTR _rt$[ebp]   ; no, not zero
    add     eax, 1                     ; increment rt
    mov     DWORD PTR _rt$[ebp], eax
$LN1@f:
    jmp     SHORT $LN3@f
$LN2@f:
    mov     eax, DWORD PTR _rt$[ebp]
    mov     esp, ebp
    pop     ebp
    ret     0
_f ENDP
```

That's how SHL (*Shift Left*) working.

Let's compile it in GCC 4.4.1:

Listing 1.74: GCC 4.4.1

```
public f
f      proc near

rt     = dword ptr -0Ch
i      = dword ptr -8
arg_0  = dword ptr 8

        push    ebp
        mov     ebp, esp
        push    ebx
        sub     esp, 10h
        mov     [ebp+rt], 0
        mov     [ebp+i], 0
        jmp     short loc_80483EF
loc_80483D0:
        mov     eax, [ebp+i]
```

```

mov     edx, 1
mov     ebx, edx
mov     ecx, eax
shl     ebx, cl
mov     eax, ebx
and     eax, [ebp+arg_0]
test    eax, eax
jz      short loc_80483EB
add     [ebp+rt], 1
loc_80483EB:
add     [ebp+i], 1
loc_80483EF:
cmp     [ebp+i], 1Fh
jle     short loc_80483D0
mov     eax, [ebp+rt]
add     esp, 10h
pop     ebx
pop     ebp
retn
f      endp

```

Shift instructions are often used in division and multiplications by power of two numbers (1, 2, 4, 8, etc).
For example:

```

unsigned int f(unsigned int a)
{
    return a/4;
};

```

We got (MSVC 2010):

Listing 1.75: MSVC 2010

```

_a$ = 8 ; size = 4
_f      PROC
    mov     eax, DWORD PTR _a$[esp-4]
    shr     eax, 2
    ret     0
_f      ENDP

```

SHR (*SH*ift *R*ight) instruction in this example is shifting a number by 2 bits right. Two freed bits at left (e.g., two most significant bits) are set to zero. Two least significant bits are dropped. In fact, these two dropped bits — division operation remainder.

SHR instruction works just like as SHL but in other direction.




Figure 1.2: How SHR instruction works⁸⁷

It can be easily understood if to imagine decimal numeral system and number 23. 23 can be easily divided by 10 just by dropping last digit (3 — is division remainder). 2 is leaving after operation as a quotient ⁸⁸.

The same story about multiplication. Multiplication by 4 is just shifting the number to the left by 2 bits, inserting 2 zero bits at right (as the last two bits). It's just like to multiply 3 by 100 — we need just to add two zeroes at the right.

⁸⁷ illustration taken from wikipedia

⁸⁸ division result

ARM + Optimizing Xcode (LLVM) + ARM mode

Listing 1.76: Optimizing Xcode (LLVM) + ARM mode

```

MOV      R1, R0
MOV      R0, #0
MOV      R2, #1
MOV      R3, R0
loc_2E54
TST      R1, R2,LSL R3 ; set flags according to R1 & (R2<<R3)
ADD      R3, R3, #1    ; R3++
ADDNE    R0, R0, #1    ; if ZF flag is cleared by TST, R0++
CMP      R3, #32
BNE      loc_2E54
BX       LR
```

TST is the same things as TEST in x86.

As I mentioned before [1.12.1](#), there are no separate shifting instructions in ARM mode. However, there are modifiers LSL (*Logical Shift Left*), LSR (*Logical Shift Right*), ASR (*Arithmetic Shift Right*), ROR (*Rotate Right*) and RRX (*Rotate Right with Extend*), which may be added to such instructions as MOV, TST, CMP, ADD, SUB, RSB⁸⁹.

These modifiers are defines, how to shift second operand and by how many bits.

Thus “TST R1, R2,LSL R3” instruction works here as $R1 \wedge (R2 \ll R3)$.

ARM + Optimizing Xcode (LLVM) + thumb-2 mode

Almost the same, but here are two LSL.W/TST instructions are used instead of single TST, because, in thumb mode, it's not possible to define LSL modifier right in TST.

```

MOV      R1, R0
MOVS     R0, #0
MOV.W    R9, #1
MOVS     R3, #0
loc_2F7A
LSL.W    R2, R9, R3
TST      R2, R1
ADD.W    R3, R3, #1
IT NE
ADDNE    R0, #1
CMP      R3, #32
BNE      loc_2F7A
BX       LR
```

1.15.4 CRC32 calculation example

This is very popular table-based CRC32 hash calculation method⁹⁰.

```
/* By Bob Jenkins, (c) 2006, Public Domain */

#include <stdio.h>
#include <stddef.h>
#include <string.h>

typedef unsigned long ub4;
typedef unsigned char ub1;

static const ub4 crctab[256] = {
    0x00000000, 0x77073096, 0xee0e612c, 0x990951ba, 0x076dc419, 0x706af48f,
    0xe963a535, 0x9e6495a3, 0x0edb8832, 0x79dcb8a4, 0xe0d5e91e, 0x97d2d988,
    0x09b64c2b, 0x7eb17cbd, 0xe7b82d07, 0x90bf1d91, 0x1db71064, 0x6ab020f2,
    0xf3b97148, 0x84be41de, 0x1adad47d, 0x6ddde4eb, 0xf4d4b551, 0x83d385c7,
```

⁸⁹These instructions are also called “data processing instructions”

⁹⁰Source code was taken here: <http://burtleburtle.net/bob/c/crc.c>

```

0x136c9856, 0x646ba8c0, 0xfd62f97a, 0x8a65c9ec, 0x14015c4f, 0x63066cd9,
0xfa0f3d63, 0x8d080df5, 0x3b6e20c8, 0x4c69105e, 0xd56041e4, 0xa2677172,
0x3c03e4d1, 0x4b04d447, 0xd20d85fd, 0xa50ab56b, 0x35b5a8fa, 0x42b2986c,
0xdbbbc9d6, 0xacbcf940, 0x32d86ce3, 0x45df5c75, 0xdc6d0dcf, 0xabd13d59,
0x26d930ac, 0x51de003a, 0xc8d75180, 0xbfd06116, 0x21b4f4b5, 0x56b3c423,
0xcfba9599, 0xb8bda50f, 0x2802b89e, 0x5f058808, 0xc60cd9b2, 0xb10be924,
0x2f6f7c87, 0x58684c11, 0xc1611dab, 0xb6662d3d, 0x76dc4190, 0x01db7106,
0x98d220bc, 0xefd5102a, 0x71b18589, 0x06b6b51f, 0x9fbfe4a5, 0xe8b8d433,
0x7807c9a2, 0x0f00f934, 0x9609a88e, 0xe10e9818, 0x7f6a0dbb, 0x086d3d2d,
0x91646c97, 0xe6635c01, 0x6b6b51f4, 0x1c6c6162, 0x856530d8, 0xf262004e,
0x6c0695ed, 0x1b01a57b, 0x8208f4c1, 0xf50fc457, 0x65b0d9c6, 0x12b7e950,
0x8bbeb8ea, 0xfcb9887c, 0x62dd1ddf, 0x15da2d49, 0x8cd37cf3, 0xfbd44c65,
0x4db26158, 0x3ab551ce, 0xa3bc0074, 0xd4bb30e2, 0x4adfa541, 0x3dd895d7,
0xa4d1c46d, 0xd3d6f4fb, 0x4369e96a, 0x346ed9fc, 0xad678846, 0xda60b8d0,
0x44042d73, 0x33031de5, 0xaa0a4c5f, 0xdd0d7cc9, 0x5005713c, 0x270241aa,
0xbe0b1010, 0xc90c2086, 0x5768b525, 0x206f85b3, 0xb966d409, 0xce61e49f,
0x5edef90e, 0x29d9c998, 0xb0d09822, 0xc7d7a8b4, 0x59b33d17, 0x2eb40d81,
0xb7bd5c3b, 0xc0ba6cad, 0xedb88320, 0x9abfb3b6, 0x03b6e20c, 0x74b1d29a,
0xead54739, 0x9dd277af, 0x04db2615, 0x73dc1683, 0xe3630b12, 0x94643b84,
0x0d6d6a3e, 0x7a6a5aa8, 0xe40ecf0b, 0x9309ff9d, 0x0a00ae27, 0x7d079eb1,
0xf00f9344, 0x8708a3d2, 0x1e01f268, 0x6906c2fe, 0xf762575d, 0x806567cb,
0x196c3671, 0x6e6b06e7, 0xfed41b76, 0x89d32be0, 0x10da7a5a, 0x67dd4acc,
0xf9b9df6f, 0x8ebeeef9, 0x17b7be43, 0x60b08ed5, 0xd6d6a3e8, 0xa1d1937e,
0x38d8c2c4, 0x4fdff252, 0xd1bb67f1, 0xa6bc5767, 0x3fb506dd, 0x48b2364b,
0xd80d2bda, 0xaf0a1b4c, 0x36034faf, 0x41047a60, 0xdf60efc3, 0xa867df55,
0x316e8eef, 0x4669be79, 0xcb61b38c, 0xbc66831a, 0x256fd2a0, 0x5268e236,
0xcc0c7795, 0xbb0b4703, 0x220216b9, 0x5505262f, 0xc5ba3bbe, 0xb2bd0b28,
0x2bb45a92, 0x5cb36a04, 0xc2d7ffa7, 0xb5d0cf31, 0x2cd99e8b, 0x5bdeae1d,
0x9b64c2b0, 0xec63f226, 0x756aa39c, 0x026d930a, 0x9c0906a9, 0xeb0e363f,
0x72076785, 0x05005713, 0x95bf4a82, 0xe2b87a14, 0x7bb12bae, 0x0cb61b38,
0x92d28e9b, 0xe5d5be0d, 0x7cdcefb7, 0x0bdbdf21, 0x86d3d2d4, 0xf1d4e242,
0x68ddb3f8, 0x1fda836e, 0x81be16cd, 0xf6b9265b, 0xf6b077e1, 0x18b74777,
0x88085ae6, 0xff0f6a70, 0x66063bca, 0x11010b5c, 0x8f659eff, 0xf862ae69,
0x616bffd3, 0x166ccc45, 0xa00ae278, 0xd70dd2ee, 0x4e048354, 0x3903b3c2,
0xa7672661, 0xd06016f7, 0x4969474d, 0x3e6e77db, 0xaed16a4a, 0xd9d65adc,
0x40df0b66, 0x37d83bf0, 0xa9bcae53, 0xdeb9ec5, 0x47b2cf7f, 0x30b5ffe9,
0xbdbdf21c, 0xcabac28a, 0x53b39330, 0x24b4a3a6, 0xbad03605, 0xcdd70693,
0x54de5729, 0x23d967bf, 0xb3667a2e, 0xc4614ab8, 0x5d681b02, 0x2a6f2b94,
0xb40bbe37, 0xc30c8ea1, 0x5a05df1b, 0x2d02ef8d,

```

```
};
```

```
/* how to derive the values in crctab[] from polynomial 0xedb88320 */
```

```
void build_table()
```

```

{
    ub4 i, j;
    for (i=0; i<256; ++i) {
        j = i;
        j = (j>>1) ^ ((j&1) ? 0xedb88320 : 0);
        j = (j>>1) ^ ((j&1) ? 0xedb88320 : 0);
        j = (j>>1) ^ ((j&1) ? 0xedb88320 : 0);
        j = (j>>1) ^ ((j&1) ? 0xedb88320 : 0);
        j = (j>>1) ^ ((j&1) ? 0xedb88320 : 0);
        j = (j>>1) ^ ((j&1) ? 0xedb88320 : 0);
        j = (j>>1) ^ ((j&1) ? 0xedb88320 : 0);
        j = (j>>1) ^ ((j&1) ? 0xedb88320 : 0);
        printf("0x%.8lx, ", j);
        if (i%6 == 5) printf("\n");
    }
}

```

```
/* the hash function */
```

```
ub4 crc(const void *key, ub4 len, ub4 hash)
```

```

{
    ub4 i;
    const ub1 *k = key;

```

```

    for (hash=len, i=0; i<len; ++i)
        hash = (hash >> 8) ^ crctab[(hash & 0xff) ^ k[i]];
    return hash;
}

/* To use, try "gcc -O crc.c -o crc; crc < crc.c" */
int main()
{
    char s[1000];
    while (gets(s)) printf("%.8lx\n", crc(s, strlen(s), 0));
    return 0;
}

```

We are interesting in `crc()` function only. By the way, pay attention to two loop initializers in `for()` statement: `hash=len, i=0`. C/C++ standard allows this, of course. Emited code will contain two operations in loop initialization part instead of usual one.

Let's compile it in MSVC with optimization (`/Ox`). For the sake of brevity, only `crc()` function is listed here, with my comments.

```

_key$ = 8           ; size = 4
_len$ = 12          ; size = 4
_hash$ = 16         ; size = 4
_crc    PROC
    mov     edx, DWORD PTR _len$[esp-4]
    xor     ecx, ecx ; i will be stored in ECX
    mov     eax, edx
    test    edx, edx
    jbe     SHORT $LN1@crc
    push    ebx
    push    esi
    mov     esi, DWORD PTR _key$[esp+4] ; ESI = key
    push    edi
$LL3@crc:
; work with bytes using only 32-bit registers. byte from address key+i we store into EDI

    movzx   edi, BYTE PTR [ecx+esi]
    mov     ebx, eax ; EBX = (hash = len)
    and     ebx, 255 ; EBX = hash & 0xff

; XOR EDI, EBX (EDI=EDI^EBX) - this operation uses all 32 bits of each register
; but other bits (8-31) are cleared all time, so it's OK
; these are cleared because, as for EDI, it was done by MOVZX instruction above
; high bits of EBX was cleared by AND EBX, 255 instruction above (255 = 0xff)

    xor     edi, ebx

; EAX=EAX>>8; bits 24-31 taken "from nowhere" will be cleared
    shr     eax, 8

; EAX=EAX^crctab[EDI*4] - choose EDI-th element from crctab[] table
    xor     eax, DWORD PTR _crctab[edi*4]
    inc     ecx           ; i++
    cmp     ecx, edx      ; i<len ?
    jb      SHORT $LL3@crc ; yes
    pop     edi
    pop     esi
    pop     ebx
$LN1@crc:
    ret     0
_crc    ENDP

```

Let's try the same in GCC 4.4.1 with `-O3` option:

```

crc          public crc
crc          proc near

```

```

key      = dword ptr 8
hash     = dword ptr 0Ch

        push    ebp
        xor     edx, edx
        mov     ebp, esp
        push    esi
        mov     esi, [ebp+key]
        push    ebx
        mov     ebx, [ebp+hash]
        test    ebx, ebx
        mov     eax, ebx
        jz      short loc_80484D3
        nop
        lea     esi, [esi+0] ; padding; ESI doesn't changing here

loc_80484B8:
        mov     ecx, eax ; save previous state of hash to ECX
        xor     al, [esi+edx] ; AL=*(key+i)
        add     edx, 1 ; i++
        shr     ecx, 8 ; ECX=hash>>8
        movzx   eax, al ; EAX=*(key+i)
        mov     eax, dword ptr ds:crctab[eax*4] ; EAX=crctab[EAX]
        xor     eax, ecx ; hash=EAX^ECX
        cmp     ebx, edx
        ja      short loc_80484B8

loc_80484D3:
        pop     ebx
        pop     esi
        pop     ebp
        retn

crc
\
endp

```

GCC aligned loop start by 8-byte border by adding NOP and `lea esi, [esi+0]` (that's *idle operation* too). Read more about it in npad section [2.3](#).

1.16 Structures

It can be defined that C/C++ structure, with some assumptions, just a set of variables, always stored in memory together, not necessary of the same type.

1.16.1 SYSTEMTIME example

Let's take SYSTEMTIME⁹¹ win32 structure describing time.

That's how it's defined:

Listing 1.77: WinBase.h

```
typedef struct _SYSTEMTIME {
    WORD wYear;
    WORD wMonth;
    WORD wDayOfWeek;
    WORD wDay;
    WORD wHour;
    WORD wMinute;
    WORD wSecond;
    WORD wMilliseconds;
} SYSTEMTIME, *PSYSTEMTIME;
```

Let's write a C function to get current time:

```
#include <windows.h>
#include <stdio.h>

void main()
{
    SYSTEMTIME t;
    GetSystemTime (&t);

    printf ("%04d-%02d-%02d %02d:%02d:%02d\n",
        t.wYear, t.wMonth, t.wDay,
        t.wHour, t.wMinute, t.wSecond);

    return;
};
```

We got (MSVC 2010):

Listing 1.78: MSVC 2010

```
_t$ = -16          ; size = 16
_main             PROC
    push         ebp
    mov         ebp, esp
    sub         esp, 16          ; 00000010H
    lea         eax, DWORD PTR _t$[ebp]
    push        eax
    call        DWORD PTR __imp__GetSystemTime@4
    movzx       ecx, WORD PTR _t$[ebp+12] ; wSecond
    push        ecx
    movzx       edx, WORD PTR _t$[ebp+10] ; wMinute
    push        edx
    movzx       eax, WORD PTR _t$[ebp+8] ; wHour
    push        eax
    movzx       ecx, WORD PTR _t$[ebp+6] ; wDay
    push        ecx
    movzx       edx, WORD PTR _t$[ebp+2] ; wMonth
    push        edx
    movzx       eax, WORD PTR _t$[ebp] ; wYear
    push        eax
```

⁹¹ [MSDN: SYSTEMTIME structure](#)

```

push    OFFSET $SG78811 ; '%04d-%02d-%02d %02d:%02d:%02d', 0aH, 00H
call    _printf
add     esp, 28          ; 0000001cH
xor     eax, eax
mov     esp, ebp
pop     ebp
ret     0
_main   ENDP

```

16 bytes are allocated for this structure in local stack — that's exactly `sizeof(WORD)*8` (there are 8 WORD variables in the structure).

Pay attention to the fact the structure beginning with `wYear` field. It can be said, an pointer to `SYSTEMTIME` structure is passed to `GetSystemTime()`⁹², but it's also can be said, pointer to `wYear` field is passed, and that's the same! `GetSystemTime()` writting current year to the WORD pointer pointing to, then shifting 2 bytes ahead, then writting current month, etc, etc.

The fact that structure fields are just variables located side-by-side, I can illustrate by the following method. Keeping in ming `SYSTEMTIME` structure description, I can rewrite this simple example like this:

```

#include <windows.h>
#include <stdio.h>

void main()
{
    WORD array[8];
    GetSystemTime (array);

    printf ("%04d-%02d-%02d %02d:%02d:%02d\n",
        array[0] /* wYear */, array[1] /* wMonth */, array[3] /* wDay */,
        array[4] /* wHour */, array[5] /* wMinute */, array[6] /* wSecond */);

    return;
};

```

Compiler will grumble for a little:

```
systemtime2.c(7) : warning C4133: 'function' : incompatible types - from 'WORD [8]' to 'LPSYSTEMTIME'
```

But nevertheless, it will produce this code:

Listing 1.79: MSVC 2010

```

$SG78573 DB      '%04d-%02d-%02d %02d:%02d:%02d', 0aH, 00H

_array$ = -16                                         ; size = 16
_main  PROC
    push    ebp
    mov     ebp, esp
    sub     esp, 16                                  ; 00000010H
    lea     eax, DWORD PTR _array$[ebp]
    push    eax
    call    DWORD PTR __imp__GetSystemTime@4
    movzx   ecx, WORD PTR _array$[ebp+12] ; wSecond
    push    ecx
    movzx   edx, WORD PTR _array$[ebp+10] ; wMinute
    push    edx
    movzx   eax, WORD PTR _array$[ebp+8] ; wHour
    push    eax
    movzx   ecx, WORD PTR _array$[ebp+6] ; wDay
    push    ecx
    movzx   edx, WORD PTR _array$[ebp+2] ; wMonth
    push    edx
    movzx   eax, WORD PTR _array$[ebp] ; wYear
    push    eax
    push    OFFSET $SG78573

```

⁹²[MSDN: SYSTEMTIME structure](#)


```

        call    _printf
        add     esp, 28                      ; 0000001cH
        xor     eax, eax
        mov     esp, ebp
        pop     ebp
        ret     0
_main    ENDP

```

And it works just as the same!

It's very interesting fact that result in assembly form cannot be distinguished from the result of previous compilation. So by looking at this code, one cannot say for sure, was there structure declared, or just pack of variables.

Nevertheless, no one will do it in sane state of mind. Because it's not convenient. Also, structure fields may be changed by developers, swapped, etc.

1.16.2 Let's allocate place for structure using malloc()

However, sometimes it's simpler to place structures not in local stack, but in heap:

```

#include <windows.h>
#include <stdio.h>

void main()
{
    SYSTEMTIME *t;

    t=(SYSTEMTIME *)malloc (sizeof (SYSTEMTIME));

    GetSystemTime (t);

    printf ("%04d-%02d-%02d %02d:%02d:%02d\n",
        t->wYear, t->wMonth, t->wDay,
        t->wHour, t->wMinute, t->wSecond);

    free (t);

    return;
};

```

Let's compile it now with optimization (/Ox) so to easily see what we need.

Listing 1.80: Optimizing MSVC

```

_main    PROC
        push    esi
        push    16                      ; 00000010H
        call    _malloc
        add     esp, 4
        mov     esi, eax
        push    esi
        call    DWORD PTR __imp__GetSystemTime@4
        movzx   eax, WORD PTR [esi+12] ; wSecond
        movzx   ecx, WORD PTR [esi+10] ; wMinute
        movzx   edx, WORD PTR [esi+8]  ; wHour
        push    eax
        movzx   eax, WORD PTR [esi+6] ; wDay
        push    ecx
        movzx   ecx, WORD PTR [esi+2] ; wMonth
        push    edx
        movzx   edx, WORD PTR [esi]   ; wYear
        push    eax
        push    ecx
        push    edx
        push    OFFSET $SG78833
        call    _printf
        push    esi

```

```

call    _free
add     esp, 32                ; 00000020H
xor     eax, eax
pop     esi
ret     0
_main   ENDP

```

So, `sizeof(SYSTEMTIME) = 16`, that's exact number of bytes to be allocated by `malloc()`. It return the pointer to freshly allocated memory block in `EAX`, which is then moved into `ESI`. `GetSystemTime()` win32 function undertake to save `ESI` value, and that's why it is not saved here and continue to be used after `GetSystemTime()` call.

New instruction — `MOVZX` (*Move with Zero extent*). It may be used almost in those cases as `MOVSX` 1.11.1, but, it clearing other bits to 0. That's because `printf()` require 32-bit *int*, but we got `WORD` in structure — that's 16-bit unsigned type. That's why by copying value from `WORD` into *int*, bits from 16 to 31 should be cleared, because there will be random noise otherwise, leaved from previous operations on registers.

In this example, I can represent structure as array of `WORD`-s:

```

#include <windows.h>
#include <stdio.h>

void main()
{
    WORD *t;

    t=(WORD *)malloc (16);

    GetSystemTime (t);

    printf ("%04d-%02d-%02d %02d:%02d:%02d\n",
        t[0] /* wYear */, t[1] /* wMonth */, t[3] /* wDay */,
        t[4] /* wHour */, t[5] /* wMinute */, t[6] /* wSecond */);

    free (t);

    return;
};

```

We got:

Listing 1.81: Optimizing MSVC

```

$SG78594 DB      '%04d-%02d-%02d %02d:%02d:%02d', 0aH, 00H

_main  PROC
    push    esi
    push    16                ; 00000010H
    call    _malloc
    add     esp, 4
    mov     esi, eax
    push    esi
    call    DWORD PTR __imp__GetSystemTime@4
    movzx   eax, WORD PTR [esi+12]
    movzx   ecx, WORD PTR [esi+10]
    movzx   edx, WORD PTR [esi+8]
    push    eax
    movzx   eax, WORD PTR [esi+6]
    push    ecx
    movzx   ecx, WORD PTR [esi+2]
    push    edx
    movzx   edx, WORD PTR [esi]
    push    eax
    push    ecx
    push    edx
    push    OFFSET $SG78594
    call    _printf
    push    esi
    call    _free

```

```

        add     esp, 32                                ; 00000020H
        xor     eax, eax
        pop     esi
        ret     0
_main    ENDP

```

Again, we got a code that cannot be distinguished from previous. And again I should note, one shouldn't do this in practice.

1.16.3 struct tm

Linux

As of Linux, let's take tm structure from time.h for example:

```

#include <stdio.h>
#include <time.h>

void main()
{
    struct tm t;
    time_t unix_time;

    unix_time=time(NULL);

    localtime_r (&unix_time, &t);

    printf ("Year: %d\n", t.tm_year+1900);
    printf ("Month: %d\n", t.tm_mon);
    printf ("Day: %d\n", t.tm_mday);
    printf ("Hour: %d\n", t.tm_hour);
    printf ("Minutes: %d\n", t.tm_min);
    printf ("Seconds: %d\n", t.tm_sec);
};

```

Let's compile it in GCC 4.4.1:

```

main      proc near
          push    ebp
          mov     ebp, esp
          and     esp, 0FFFFFFF0h
          sub     esp, 40h
          mov     dword ptr [esp], 0 ; first argument for time()
          call    time
          mov     [esp+3Ch], eax
          lea     eax, [esp+3Ch] ; take pointer to what time() returned
          lea     edx, [esp+10h] ; at ESP+10h struct tm will begin
          mov     [esp+4], edx ; pass pointer to the structure begin
          mov     [esp], eax ; pass pointer to result of time()
          call    localtime_r
          mov     eax, [esp+24h] ; tm_year
          lea     edx, [eax+76Ch] ; edx=eax+1900
          mov     eax, offset format ; "Year: %d\n"
          mov     [esp+4], edx
          mov     [esp], eax
          call    printf
          mov     edx, [esp+20h] ; tm_mon
          mov     eax, offset aMonthD ; "Month: %d\n"
          mov     [esp+4], edx
          mov     [esp], eax
          call    printf
          mov     edx, [esp+1Ch] ; tm_mday
          mov     eax, offset aDayD ; "Day: %d\n"
          mov     [esp+4], edx
          mov     [esp], eax
          call    printf

```

```

        mov     edx, [esp+18h]      ; tm_hour
        mov     eax, offset aHourD ; "Hour: %d\n"
        mov     [esp+4], edx
        mov     [esp], eax
        call    printf
        mov     edx, [esp+14h]     ; tm_min
        mov     eax, offset aMinutesD ; "Minutes: %d\n"
        mov     [esp+4], edx
        mov     [esp], eax
        call    printf
        mov     edx, [esp+10h]
        mov     eax, offset aSecondsD ; "Seconds: %d\n"
        mov     [esp+4], edx      ; tm_sec
        mov     [esp], eax
        call    printf
        leave
        retn
main     endp

```

Somehow, IDA 5 didn't create local variable names in local stack. But since we already experienced reverse engineers :-), we may do it without this information in this simple example.

Please also pay attention to `lea edx, [eax+76Ch]` — this instruction just adds `0x76C` to EAX, but not modify any flags. See also relevant section about LEA [2.1](#).

In order to illustrate that structure is just variables laying side-by-side in one place, let's rework example, while looking at the file *time.h*:

Listing 1.82: time.h

```

struct tm
{
    int    tm_sec;
    int    tm_min;
    int    tm_hour;
    int    tm_mday;
    int    tm_mon;
    int    tm_year;
    int    tm_wday;
    int    tm_yday;
    int    tm_isdst;
};

```

```

#include <stdio.h>
#include <time.h>

void main()
{
    int tm_sec, tm_min, tm_hour, tm_mday, tm_mon, tm_year, tm_wday, tm_yday, tm_isdst;
    time_t unix_time;

    unix_time=time(NULL);

    localtime_r (&unix_time, &tm_sec);

    printf ("Year: %d\n", tm_year+1900);
    printf ("Month: %d\n", tm_mon);
    printf ("Day: %d\n", tm_mday);
    printf ("Hour: %d\n", tm_hour);
    printf ("Minutes: %d\n", tm_min);
    printf ("Seconds: %d\n", tm_sec);
};

```

Please note that pointer to exactly `tm_sec` is passed into `localtime_r`, i.e., to the first “structure” element. Compiler will warn us:

Listing 1.83: GCC 4.7.3

```

GCC_tm2.c: In function 'main':
GCC_tm2.c:11:5: warning: passing argument 2 of 'localtime_r' from incompatible pointer type [enabled by
    default]
In file included from GCC_tm2.c:2:0:
/usr/include/time.h:59:12: note: expected 'struct tm *' but argument is of type 'int *'

```

But nevertheless, will generate this:

Listing 1.84: GCC 4.7.3

```

main          proc near

var_30        = dword ptr -30h
var_2C        = dword ptr -2Ch
unix_time     = dword ptr -1Ch
tm_sec        = dword ptr -18h
tm_min        = dword ptr -14h
tm_hour       = dword ptr -10h
tm_mday       = dword ptr -0Ch
tm_mon        = dword ptr -8
tm_year       = dword ptr -4

                push    ebp
                mov     ebp, esp
                and     esp, 0FFFFFF0h
                sub     esp, 30h
                call    __main
                mov     [esp+30h+var_30], 0 ; arg 0
                call    time
                mov     [esp+30h+unix_time], eax
                lea     eax, [esp+30h+tm_sec]
                mov     [esp+30h+var_2C], eax
                lea     eax, [esp+30h+unix_time]
                mov     [esp+30h+var_30], eax
                call    localtime_r
                mov     eax, [esp+30h+tm_year]
                add     eax, 1900
                mov     [esp+30h+var_2C], eax
                mov     [esp+30h+var_30], offset aYearD ; "Year: %d\n"
                call    printf
                mov     eax, [esp+30h+tm_mon]
                mov     [esp+30h+var_2C], eax
                mov     [esp+30h+var_30], offset aMonthD ; "Month: %d\n"
                call    printf
                mov     eax, [esp+30h+tm_mday]
                mov     [esp+30h+var_2C], eax
                mov     [esp+30h+var_30], offset aDayD ; "Day: %d\n"
                call    printf
                mov     eax, [esp+30h+tm_hour]
                mov     [esp+30h+var_2C], eax
                mov     [esp+30h+var_30], offset aHourD ; "Hour: %d\n"
                call    printf
                mov     eax, [esp+30h+tm_min]
                mov     [esp+30h+var_2C], eax
                mov     [esp+30h+var_30], offset aMinutesD ; "Minutes: %d\n"
                call    printf
                mov     eax, [esp+30h+tm_sec]
                mov     [esp+30h+var_2C], eax
                mov     [esp+30h+var_30], offset aSecondsD ; "Seconds: %d\n"
                call    printf
                leave
                retn
main          endp

```

This code is identical to what we saw previously and it's not possible to say, was it structure in original source code or just pack of variables.

And this works. However, it's not recommended to do this in practice. Usually, compiler allocate variables in local stack in the same order as they were declared in function. Nevertheless, there are no any guarantee.

By the way, some other compiler may warn that `tm_year`, `tm_mon`, `tm_mday`, `tm_hour`, `tm_min` variables, but not `tm_sec` are used without being initialized. Indeed, compiler don't know that these will be filled when calling to `localtime_r()`.

I chose exactly this example for illustration, because all structure fields has *int* type, and `SYSTEMTIME` structure fields — 16-bit `WORD`, and if to declare them as a local variables, they will be aligned by 32-bit border, and nothing will work (because `GetSystemTime()` will fill them incorrectly). Read more about it in next section: "Fields packing in structure".

So, structure is just variables pack laying on one place, side-by-side. I could say that structure is a syntactic sugar, directing compiler to hold them in one place. However, I'm not programming languages expert, so, most likely, I'm wrong with this term. By the way, there were a times, in very early C versions (before 1972), in which there were no structures at all [Rit93].

ARM + Optimizing Keil + thumb mode

Same example:

Listing 1.85: Optimizing Keil + thumb mode

```
var_38      = -0x38
var_34      = -0x34
var_30      = -0x30
var_2C      = -0x2C
var_28      = -0x28
var_24      = -0x24
timer       = -0xC

        PUSH    {LR}
        MOVS    R0, #0          ; timer
        SUB     SP, SP, #0x34
        BL      time
        STR     R0, [SP,#0x38+timer]
        MOV     R1, SP          ; tp
        ADD     R0, SP, #0x38+timer ; timer
        BL      localtime_r
        LDR     R1, =0x76C
        LDR     R0, [SP,#0x38+var_24]
        ADDS    R1, R0, R1
        ADR     R0, aYearD      ; "Year: %d\n"
        BL      __2printf
        LDR     R1, [SP,#0x38+var_28]
        ADR     R0, aMonthD     ; "Month: %d\n"
        BL      __2printf
        LDR     R1, [SP,#0x38+var_2C]
        ADR     R0, aDayD       ; "Day: %d\n"
        BL      __2printf
        LDR     R1, [SP,#0x38+var_30]
        ADR     R0, aHourD      ; "Hour: %d\n"
        BL      __2printf
        LDR     R1, [SP,#0x38+var_34]
        ADR     R0, aMinutesD   ; "Minutes: %d\n"
        BL      __2printf
        LDR     R1, [SP,#0x38+var_38]
        ADR     R0, aSecondsD   ; "Seconds: %d\n"
        BL      __2printf
        ADD     SP, SP, #0x34
        POP     {PC}
```

ARM + Optimizing Xcode (LLVM) + thumb-2 mode

IDA 5 “get to know” tm structure (because IDA 5 “knows” argument types of library functions like `localtime_r()`), so it shows here structure elements accesses and also names are assigned to them.

Listing 1.86: Optimizing Xcode (LLVM) + thumb-2 mode

```
var_38      = -0x38
var_34      = -0x34

        PUSH      {R7,LR}
        MOV       R7, SP
        SUB       SP, SP, #0x30
        MOVS      R0, #0 ; time_t *
        BLX       _time
        ADD       R1, SP, #0x38+var_34 ; struct tm *
        STR       R0, [SP,#0x38+var_38]
        MOV       R0, SP ; time_t *
        BLX       _localtime_r
        LDR       R1, [SP,#0x38+var_34.tm_year]
        MOV       R0, 0xF44 ; "Year: %d\n"
        ADD       R0, PC ; char *
        ADDW      R1, R1, #0x76C
        BLX       _printf
        LDR       R1, [SP,#0x38+var_34.tm_mon]
        MOV       R0, 0xF3A ; "Month: %d\n"
        ADD       R0, PC ; char *
        BLX       _printf
        LDR       R1, [SP,#0x38+var_34.tm_mday]
        MOV       R0, 0xF35 ; "Day: %d\n"
        ADD       R0, PC ; char *
        BLX       _printf
        LDR       R1, [SP,#0x38+var_34.tm_hour]
        MOV       R0, 0xF2E ; "Hour: %d\n"
        ADD       R0, PC ; char *
        BLX       _printf
        LDR       R1, [SP,#0x38+var_34.tm_min]
        MOV       R0, 0xF28 ; "Minutes: %d\n"
        ADD       R0, PC ; char *
        BLX       _printf
        LDR       R1, [SP,#0x38+var_34]
        MOV       R0, 0xF25 ; "Seconds: %d\n"
        ADD       R0, PC ; char *
        BLX       _printf
        ADD       SP, SP, #0x30
        POP       {R7,PC}
```

...

```
00000000 tm          struc ; (sizeof=0x2C, standard type)
00000000 tm_sec      DCD ?
00000004 tm_min      DCD ?
00000008 tm_hour     DCD ?
0000000C tm_mday     DCD ?
00000010 tm_mon      DCD ?
00000014 tm_year     DCD ?
00000018 tm_wday     DCD ?
0000001C tm_yday     DCD ?
00000020 tm_isdst    DCD ?
00000024 tm_gmtimeoff DCD ?
00000028 tm_zone     DCD ? ; offset
0000002C tm          ends
```

1.16.4 Fields packing in structure

One important thing is fields packing in structures⁹³.

Let's take a simple example:

```
#include <stdio.h>

struct s
{
    char a;
    int b;
    char c;
    int d;
};

void f(struct s s)
{
    printf ("a=%d; b=%d; c=%d; d=%d\n", s.a, s.b, s.c, s.d);
};
```

As we see, we have two *char* fields (each is exactly one byte) and two more — *int* (each - 4 bytes).

x86

That's all compiling into:

```
_s$ = 8          ; size = 16
?f@@YAXUs@@@Z PROC    ; f
    push     ebp
    mov     ebp, esp
    mov     eax, DWORD PTR _s$[ebp+12]
    push     eax
    movsx    ecx, BYTE PTR _s$[ebp+8]
    push     ecx
    mov     edx, DWORD PTR _s$[ebp+4]
    push     edx
    movsx    eax, BYTE PTR _s$[ebp]
    push     eax
    push     OFFSET $SG3842
    call     _printf
    add     esp, 20      ; 00000014H
    pop     ebp
    ret     0
?f@@YAXUs@@@Z ENDP    ; f
_TEXT      ENDS
```

As we can see, each field's address is aligned by 4-bytes border. That's why each *char* using 4 bytes here, like *int*. Why? Thus it's easier for CPU to access memory at aligned addresses and to cache data from it.

However, it's not very economical in size sense.

Let's try to compile it with option (/Zp1) (/Zp[n] pack structs on n-byte boundary).

Listing 1.87: MSVC /Zp1

```
_TEXT      SEGMENT
_s$ = 8          ; size = 10
?f@@YAXUs@@@Z PROC    ; f
    push     ebp
    mov     ebp, esp
    mov     eax, DWORD PTR _s$[ebp+6]
    push     eax
    movsx    ecx, BYTE PTR _s$[ebp+5]
    push     ecx
    mov     edx, DWORD PTR _s$[ebp+1]
    push     edx
```

⁹³See also: [Wikipedia: Data structure alignment](#)


```

movsx  eax, BYTE PTR _s$[ebp]
push   eax
push   OFFSET $SG3842
call   _printf
add     esp, 20      ; 00000014H
pop     ebp
ret     0
?f@@YAXUs@@@Z ENDP      ; f

```

Now the structure takes only 10 bytes and each *char* value takes 1 byte. What it give to us? Size economy. And as drawback — CPU will access these fields without maximal performance it can.

As it can be easily guessed, if the structure is used in many source and object files, all these should be compiled with the same convention about structures packing.

Aside from MSVC /Zp option which set how to align each structure field, here is also `#pragma pack` compiler option, it can be defined right in source code. It's available in both MSVC⁹⁴ and GCC⁹⁵.

Let's back to SYSTEMTIME structure consisting in 16-bit fields. How our compiler know to pack them on 1-byte alignment method?

WinNT.h file has this:

Listing 1.88: WinNT.h

```
#include "pshpack1.h"
```

And this:

Listing 1.89: WinNT.h

```
#include "pshpack4.h"          // 4 byte packing is the default
```

The file PshPack1.h looks like:

Listing 1.90: PshPack1.h

```

#if ! (defined(lint) || defined(RC_INVOKED))
#if ( _MSC_VER >= 800 && !defined(_M_I86)) || defined(_PUSHPOP_SUPPORTED)
#pragma warning(disable:4103)
#if !(defined( MIDL_PASS )) || defined( __midl )
#pragma pack(push,1)
#else
#pragma pack(1)
#endif
#else
#pragma pack(1)
#endif
#endif /* ! (defined(lint) || defined(RC_INVOKED)) */

```

That's how compiler will pack structures defined after `#pragma pack`.

ARM + Optimizing Keil + thumb mode

Listing 1.91: Optimizing Keil + thumb mode

```

.text:0000003E          exit                                ; CODE XREF: f+16
.text:0000003E 05 B0          ADD      SP, SP, #0x14
.text:00000040 00 BD          POP      {PC}

.text:00000280          f
.text:00000280
.text:00000280          var_18      = -0x18
.text:00000280          a          = -0x14
.text:00000280          b          = -0x10

```

⁹⁴MSDN: Working with Packing Structures

⁹⁵Structure-Packing Pragma

```

.text:00000280      c      = -0xC
.text:00000280      d      = -8
.text:00000280
.text:00000280 0F B5      PUSH    {R0-R3,LR}
.text:00000282 81 B0      SUB     SP, SP, #4
.text:00000284 04 98      LDR     R0, [SP,#16]    ; d
.text:00000286 02 9A      LDR     R2, [SP,#8]     ; b
.text:00000288 00 90      STR     R0, [SP]
.text:0000028A 68 46      MOV     R0, SP
.text:0000028C 03 7B      LDRB    R3, [R0,#12]    ; c
.text:0000028E 01 79      LDRB    R1, [R0,#4]     ; a
.text:00000290 59 A0      ADR     R0, aADBDCDDD   ; "a=%d; b=%d; c=%d; d=%d\n"
.text:00000292 05 F0 AD FF      BL      __2printf
.text:00000296 D2 E6      B       exit

```

As we may recall, here a structure passed instead of pointer to structure, and because first 4 function arguments in ARM are passed via registers, so then structure fields are passed via R0–R3.

LDRB loads one byte from memory and extending it to 32-bit, taking into account its sign. This is similar to MOVSBX 1.11.1 instruction in x86. Here it's used for loading fields *a* and *c* from structure.

One more thing we spot easily, instead of function epilogue, here is jump to another function's epilogue! Indeed, that was another function, not related in any way to our function, however, it has exactly the same epilogue (probably because, it hold 5 local variables too ($5 * 4 = 0x14$)). Also, it's located nearby (take a look on addresses). Indeed, there are no difference, which epilogue to execute, if it works just as we need. Apparently, Keil decides to do this because of economy. Epilogue takes 4 bytes while jump — only 2.

ARM + Optimizing Xcode (LLVM) + thumb-2 mode

Listing 1.92: Optimizing Xcode (LLVM) + thumb-2 mode

```

var_C      = -0xC

      PUSH    {R7,LR}
      MOV     R7, SP
      SUB     SP, SP, #4
      MOV     R9, R1 ; b
      MOV     R1, R0 ; a
      MOVW    R0, #0xF10 ; "a=%d; b=%d; c=%d; d=%d\n"
      SXTB    R1, R1 ; prepare a
      MOVT.W  R0, #0
      STR     R3, [SP,#0xC+var_C] ; place d to stack for printf()
      ADD     R0, PC ; format-string
      SXTB    R3, R2 ; prepare c
      MOV     R2, R9 ; b
      BLX     _printf
      ADD     SP, SP, #4
      POP     {R7,PC}

```

SXTB (*Signed Extend Byte*) is analogous to MOVSBX 1.11.1 in x86 as well, but works not with memory, but with register. All the rest — just the same.

1.16.5 Nested structures

Now what about situations when one structure define another structure inside?

```

#include <stdio.h>

struct inner_struct
{
    int a;
    int b;
};

struct outer_struct

```

```

{
    char a;
    int b;
    struct inner_struct c;
    char d;
    int e;
};

void f(struct outer_struct s)
{
    printf ("a=%d; b=%d; c.a=%d; c.b=%d; d=%d; e=%d\n",
           s.a, s.b, s.c.a, s.c.b, s.d, s.e);
};

```

...in this case, both `inner_struct` fields will be placed between `a,b` and `d,e` fields of `outer_struct`.
Let's compile (MSVC 2010):

Listing 1.93: MSVC 2010

```

_s$ = 8          ; size = 24
_f  PROC
    push    ebp
    mov     ebp, esp
    mov     eax, DWORD PTR _s$[ebp+20] ; e
    push    eax
    movsx   ecx, BYTE PTR _s$[ebp+16] ; d
    push    ecx
    mov     edx, DWORD PTR _s$[ebp+12] ; c.b
    push    edx
    mov     eax, DWORD PTR _s$[ebp+8] ; c.a
    push    eax
    mov     ecx, DWORD PTR _s$[ebp+4] ; b
    push    ecx
    movsx   edx, BYTE PTR _s$[ebp] ; a
    push    edx
    push    OFFSET $SG2466
    call    _printf
    add     esp, 28 ; 0000001cH
    pop     ebp
    ret     0
_f  ENDP

```

One curious point here is that by looking onto this assembly code, we do not even see that another structure was used inside of it! Thus, we would say, nested structures are finally unfolds into *linear* or *one-dimensional* structure.

Of course, if to replace `struct inner_struct c;` declaration to `struct inner_struct *c;` (thus making a pointer here) situation will be significantly different.

1.16.6 Bit fields in structure

CPUID example

C/C++ language allow to define exact number of bits for each structure fields. It's very useful if one need to save memory space. For example, one bit is enough for variable of *bool* type. But of course, it's not rational if speed is important.

Let's consider CPUID⁹⁶ instruction example. This instruction return information about current CPU and its features.

If EAX is set to 1 before instruction execution, CPUID will return this information packed into EAX register:

⁹⁶<http://en.wikipedia.org/wiki/CPUID>

3:0	Stepping
7:4	Model
11:8	Family
13:12	Processor Type
19:16	Extended Model
27:20	Extended Family

MSVC 2010 has CPUID macro, but GCC 4.4.1 — hasn't. So let's make this function by yourself for GCC, using its built-in assembler⁹⁷.

```
#include <stdio.h>

#ifdef __GNUC__
static inline void cpuid(int code, int *a, int *b, int *c, int *d) {
    asm volatile("cpuid":"=a"(*a), "=b"(*b), "=c"(*c), "=d"(*d):"a"(code));
}
#endif

#ifdef _MSC_VER
#include <intrin.h>
#endif

struct CPUID_1_EAX
{
    unsigned int stepping:4;
    unsigned int model:4;
    unsigned int family_id:4;
    unsigned int processor_type:2;
    unsigned int reserved1:2;
    unsigned int extended_model_id:4;
    unsigned int extended_family_id:8;
    unsigned int reserved2:4;
};

int main()
{
    struct CPUID_1_EAX *tmp;
    int b[4];

#ifdef _MSC_VER
    __cpuid(b,1);
#endif

#ifdef __GNUC__
    cpuid (1, &b[0], &b[1], &b[2], &b[3]);
#endif

    tmp=(struct CPUID_1_EAX *)&b[0];

    printf ("stepping=%d\n", tmp->stepping);
    printf ("model=%d\n", tmp->model);
    printf ("family_id=%d\n", tmp->family_id);
    printf ("processor_type=%d\n", tmp->processor_type);
    printf ("extended_model_id=%d\n", tmp->extended_model_id);
    printf ("extended_family_id=%d\n", tmp->extended_family_id);

    return 0;
};
```

After CPUID will fill EAX/EBX/ECX/EDX, these registers will be reflected in b[] array. Then, we have a pointer to CPUID_1_EAX structure and we point it to EAX value from b[] array.

In other words, we treat 32-bit *int* value as a structure.

Then we read from the stucture.

⁹⁷ [More about internal GCC assembler](#)

Let's compile it in MSVC 2008 with /Ox option:

Listing 1.94: Optimizing MSVC 2008

```
_b$ = -16          ; size = 16
_main PROC
    sub     esp, 16          ; 00000010H
    push    ebx

    xor     ecx, ecx
    mov     eax, 1
    cpuid
    push    esi
    lea     esi, DWORD PTR _b$[esp+24]
    mov     DWORD PTR [esi], eax
    mov     DWORD PTR [esi+4], ebx
    mov     DWORD PTR [esi+8], ecx
    mov     DWORD PTR [esi+12], edx

    mov     esi, DWORD PTR _b$[esp+24]
    mov     eax, esi
    and     eax, 15          ; 0000000fH
    push    eax
    push    OFFSET $SG15435 ; 'stepping=%d', 0aH, 00H
    call    _printf

    mov     ecx, esi
    shr     ecx, 4
    and     ecx, 15          ; 0000000fH
    push    ecx
    push    OFFSET $SG15436 ; 'model=%d', 0aH, 00H
    call    _printf

    mov     edx, esi
    shr     edx, 8
    and     edx, 15          ; 0000000fH
    push    edx
    push    OFFSET $SG15437 ; 'family_id=%d', 0aH, 00H
    call    _printf

    mov     eax, esi
    shr     eax, 12          ; 0000000cH
    and     eax, 3
    push    eax
    push    OFFSET $SG15438 ; 'processor_type=%d', 0aH, 00H
    call    _printf

    mov     ecx, esi
    shr     ecx, 16          ; 00000010H
    and     ecx, 15          ; 0000000fH
    push    ecx
    push    OFFSET $SG15439 ; 'extended_model_id=%d', 0aH, 00H
    call    _printf

    shr     esi, 20          ; 00000014H
    and     esi, 255         ; 000000ffH
    push    esi
    push    OFFSET $SG15440 ; 'extended_family_id=%d', 0aH, 00H
    call    _printf
    add     esp, 48          ; 00000030H
    pop     esi

    xor     eax, eax
    pop     ebx

    add     esp, 16          ; 00000010H
    ret     0
```

SHR instruction shifting value in EAX by number of bits should be *skipped*, e.g., we ignore some bits *at right*. AND instruction clearing not needed bits *at left*, or, in other words, leave only those bits in EAX we need now. Let's try GCC 4.4.1 with -O3 option.

Listing 1.95: Optimizing GCC 4.4.1

```
main          proc near          ; DATA XREF: _start+17
    push      ebp
    mov       ebp, esp
    and       esp, 0FFFFFF0h
    push      esi
    mov       esi, 1
    push      ebx
    mov       eax, esi
    sub       esp, 18h
    cpuid
    mov       esi, eax
    and       eax, 0Fh
    mov       [esp+8], eax
    mov       dword ptr [esp+4], offset aSteppingD ; "stepping=%d\n"
    mov       dword ptr [esp], 1
    call      ___printf_chk
    mov       eax, esi
    shr       eax, 4
    and       eax, 0Fh
    mov       [esp+8], eax
    mov       dword ptr [esp+4], offset aModelD ; "model=%d\n"
    mov       dword ptr [esp], 1
    call      ___printf_chk
    mov       eax, esi
    shr       eax, 8
    and       eax, 0Fh
    mov       [esp+8], eax
    mov       dword ptr [esp+4], offset aFamily_idD ; "family_id=%d\n"
    mov       dword ptr [esp], 1
    call      ___printf_chk
    mov       eax, esi
    shr       eax, 0Ch
    and       eax, 3
    mov       [esp+8], eax
    mov       dword ptr [esp+4], offset aProcessor_type ; "processor_type=%d\n"
    mov       dword ptr [esp], 1
    call      ___printf_chk
    mov       eax, esi
    shr       eax, 10h
    shr       esi, 14h
    and       eax, 0Fh
    and       esi, 0FFh
    mov       [esp+8], eax
    mov       dword ptr [esp+4], offset aExtended_model ; "extended_model_id=%d\n"
    mov       dword ptr [esp], 1
    call      ___printf_chk
    mov       [esp+8], esi
    mov       dword ptr [esp+4], offset unk_80486D0
    mov       dword ptr [esp], 1
    call      ___printf_chk
    add       esp, 18h
    xor       eax, eax
    pop       ebx
    pop       esi
    mov       esp, ebp
    pop       ebp
    retn
main          endp
```

Almost the same. The only thing to note is that GCC somehow united calculation of `extended_model_id` and `extended_family_id` into one block, instead of calculating them separately, before corresponding each `printf()` call.

Working with the float type as with a structure

As it was already noted in section about FPU 1.13, both *float* and *double* types consisted of sign, significand (or fraction) and exponent. But will we able to work with these fields directly? Let's try with *float*.




Figure 1.3: float value format⁹⁸

```
#include <stdio.h>
#include <assert.h>
#include <stdlib.h>
#include <memory.h>

struct float_as_struct
{
    unsigned int fraction : 23; // fractional part
    unsigned int exponent : 8; // exponent + 0x3FF
    unsigned int sign : 1; // sign bit
};

float f(float _in)
{
    float f=_in;
    struct float_as_struct t;

    assert (sizeof (struct float_as_struct) == sizeof (float));

    memcpy (&t, &f, sizeof (float));

    t.sign=1; // set negative sign
    t.exponent=t.exponent+2; // multiple d by 2^n (n here is 2)

    memcpy (&f, &t, sizeof (float));

    return f;
};

int main()
{
    printf ("%f\n", f(1.234));
};
```

`float_as_struct` structure occupies as much space in memory as *float*, e.g., 4 bytes or 32 bits.

Now we setting negative sign in input value and also by adding 2 to exponent we thereby multiplying the whole number by 2^2 , e.g., by 4.

Let's compile in MSVC 2008 without optimization:

Listing 1.96: Non-optimizing MSVC 2008

```
_t$ = -8          ; size = 4
_f$ = -4          ; size = 4
__in$ = 8         ; size = 4
?f@@YAMM@Z PROC  ; f
    push    ebp
```

⁹⁸illustration taken from wikipedia

```

mov     ebp, esp
sub     esp, 8

fld     DWORD PTR __in$[ebp]
fstp    DWORD PTR _f$[ebp]

push    4
lea     eax, DWORD PTR _f$[ebp]
push    eax
lea     ecx, DWORD PTR _t$[ebp]
push    ecx
call    _memcpy
add     esp, 12          ; 0000000cH

mov     edx, DWORD PTR _t$[ebp]
or      edx, -2147483648 ; 80000000H - set minus sign
mov     DWORD PTR _t$[ebp], edx

mov     eax, DWORD PTR _t$[ebp]
shr     eax, 23          ; 00000017H - drop significand
and     eax, 255         ; 000000ffH - leave here only exponent
add     eax, 2           ; add 2 to it
and     eax, 255         ; 000000ffH
shl     eax, 23          ; 00000017H - shift result to place of bits 30:23
mov     ecx, DWORD PTR _t$[ebp]
and     ecx, -2139095041 ; 807fffffH - drop exponent
or      ecx, eax         ; add original value without exponent with new calculated exponent
mov     DWORD PTR _t$[ebp], ecx

push    4
lea     edx, DWORD PTR _t$[ebp]
push    edx
lea     eax, DWORD PTR _f$[ebp]
push    eax
call    _memcpy
add     esp, 12          ; 0000000cH

fld     DWORD PTR _f$[ebp]

mov     esp, ebp
pop     ebp
ret     0
?f@@YAMM@Z ENDP          ; f

```

Redundant for a bit. If it compiled with /Ox flag there are no memcpy() call, f variable is used directly. But it's easier to understand it all considering unoptimized version.

What GCC 4.4.1 with -O3 will do?

Listing 1.97: Optimizing GCC 4.4.1

```

; f(float)
public _Ziff
_Ziff proc near

var_4 = dword ptr -4
arg_0 = dword ptr 8

push    ebp
mov     ebp, esp
sub     esp, 4
mov     eax, [ebp+arg_0]
or      eax, 80000000h ; set minus sign
mov     edx, eax
and     eax, 807FFFFh ; leave only significand and exponent in EAX
shr     edx, 23        ; prepare exponent
add     edx, 2         ; add 2
movzx   edx, dl        ; clear all bits except 7:0 in EAX

```



```

    shl     edx, 23          ; shift new calculated exponent to its place
    or      eax, edx         ; add new exponent and original value without exponent
    mov     [ebp+var_4], eax
    fld     [ebp+var_4]
    leave
    retn
_Z1ff endp

    public main
main proc near
    push    ebp
    mov     ebp, esp
    and     esp, 0FFFFFF0h
    sub     esp, 10h
    fld     ds:dword_8048614 ; -4.936
    fstp    qword ptr [esp+8]
    mov     dword ptr [esp+4], offset asc_8048610 ; "%f\n"
    mov     dword ptr [esp], 1
    call    ___printf_chk
    xor     eax, eax
    leave
    retn
main endp

```

The `f()` function is almost understandable. However, what is interesting, GCC was able to calculate `f(1.234)` result during compilation stage despite all this hodge-podge with structure fields and prepared this argument to `printf()` as precalculated!

1.17 C++ classes

1.17.1 Simple example

I placed a C++ classes description here intentionally after structures description, because internally, C++ classes representation is almost the same as structures representation.

Let's try an example with two variables, two constructors and one method:

```
#include <stdio.h>

class c
{
private:
    int v1;
    int v2;
public:
    c() // default ctor
    {
        v1=667;
        v2=999;
    };

    c(int a, int b) // ctor
    {
        v1=a;
        v2=b;
    };

    void dump()
    {
        printf ("%d; %d\n", v1, v2);
    };
};

int main()
{
    class c c1;
    class c c2(5,6);

    c1.dump();
    c2.dump();

    return 0;
};
```

Here is how main() function looks like translated into assembly language:

```
_c2$ = -16      ; size = 8
_c1$ = -8       ; size = 8
_main PROC
    push     ebp
    mov     ebp, esp
    sub     esp, 16      ; 00000010H
    lea     ecx, DWORD PTR _c1$[ebp]
    call    ??0c@@QAE@XZ ; c::c
    push    6
    push    5
    lea     ecx, DWORD PTR _c2$[ebp]
    call    ??0c@@QAE@HH@Z ; c::c
    lea     ecx, DWORD PTR _c1$[ebp]
    call    ?dump@c@@QAE@XZ ; c::dump
    lea     ecx, DWORD PTR _c2$[ebp]
    call    ?dump@c@@QAE@XZ ; c::dump
    xor     eax, eax
    mov     esp, ebp
    pop     ebp
```

```

    ret     0
_main     ENDP

```

So what's going on. For each object (instance of class *c*) 8 bytes allocated, that's exactly size of 2 variables storage.

For *c1* a default argumentless constructor `??0c@@QAE@XZ` is called. For *c2* another constructor `??0c@@QAE@HH@Z` is called and two numbers are passed as arguments.

A pointer to object (*this* in C++ terminology) is passed in ECX register. This is called *thiscall* 2.5.4 — a pointer to object passing method.

Msvc doing it using ECX register. Needless to say, it's not a standardized method, other compilers could do it differently, e.g., via first function argument (like GCC).

Why these functions has so odd names? That's *name mangling*⁹⁹.

C++ class may contain several methods sharing the same name but having different arguments — that's polymorphism. And of course, different classes may own methods sharing the same name.

Name mangling enable us to encode class name + method name + all method argument types in one ASCII-string, which will be used as internal function name. That's all because neither linker, nor DLL operation system loader (mangled names may be among DLL exports as well) knows nothing about C++ or OOP.

`dump()` function called two times after.

Now let's see constructors' code:

```

_this$ = -4          ; size = 4
??0c@@QAE@XZ PROC   ; c::c, COMDAT
; _this$ = ecx
    push     ebp
    mov      ebp, esp
    push     ecx
    mov      DWORD PTR _this$[ebp], ecx
    mov      eax, DWORD PTR _this$[ebp]
    mov      DWORD PTR [eax], 667      ; 0000029bH
    mov      ecx, DWORD PTR _this$[ebp]
    mov      DWORD PTR [ecx+4], 999     ; 000003e7H
    mov      eax, DWORD PTR _this$[ebp]
    mov      esp, ebp
    pop      ebp
    ret      0
??0c@@QAE@XZ ENDP    ; c::c

_this$ = -4          ; size = 4
_a$ = 8              ; size = 4
_b$ = 12             ; size = 4
??0c@@QAE@HH@Z PROC ; c::c, COMDAT
; _this$ = ecx
    push     ebp
    mov      ebp, esp
    push     ecx
    mov      DWORD PTR _this$[ebp], ecx
    mov      eax, DWORD PTR _this$[ebp]
    mov      ecx, DWORD PTR _a$[ebp]
    mov      DWORD PTR [eax], ecx
    mov      edx, DWORD PTR _this$[ebp]
    mov      eax, DWORD PTR _b$[ebp]
    mov      DWORD PTR [edx+4], eax
    mov      eax, DWORD PTR _this$[ebp]
    mov      esp, ebp
    pop      ebp
    ret      8
??0c@@QAE@HH@Z ENDP ; c::c

```

Constructors are just functions, they use pointer to structure in ECX, moving the pointer into own local variable, however, it's not necessary.

Now `dump()` method:

⁹⁹ [Wikipedia: Name mangling](#)

```

_this$ = -4          ; size = 4
?dump@C@QAE@XZ PROC ; c::dump, COMDAT
; _this$ = ecx
    push    ebp
    mov     ebp, esp
    push    ecx
    mov     DWORD PTR _this$[ebp], ecx
    mov     eax, DWORD PTR _this$[ebp]
    mov     ecx, DWORD PTR [eax+4]
    push    ecx
    mov     edx, DWORD PTR _this$[ebp]
    mov     eax, DWORD PTR [edx]
    push    eax
    push    OFFSET ??_C@_07NJBDCIEC@?$CFd?$DL?5?$CFd?6?$AA@
    call    _printf
    add     esp, 12      ; 0000000cH
    mov     esp, ebp
    pop     ebp
    ret     0
?dump@C@QAE@XZ ENDP ; c::dump

```

Simple enough: `dump()` taking pointer to the structure containing two *int*'s in `ECX`, takes two values from it and passing it into `printf()`.

The code is much shorter if compiled with optimization (`/Ox`):

```

??0C@QAE@XZ PROC ; c::c, COMDAT
; _this$ = ecx
    mov     eax, ecx
    mov     DWORD PTR [eax], 667 ; 0000029bH
    mov     DWORD PTR [eax+4], 999 ; 000003e7H
    ret     0
??0C@QAE@XZ ENDP ; c::c

_a$ = 8 ; size = 4
_b$ = 12 ; size = 4
??0C@QAE@HH@Z PROC ; c::c, COMDAT
; _this$ = ecx
    mov     edx, DWORD PTR _b$[esp-4]
    mov     eax, ecx
    mov     ecx, DWORD PTR _a$[esp-4]
    mov     DWORD PTR [eax], ecx
    mov     DWORD PTR [eax+4], edx
    ret     8
??0C@QAE@HH@Z ENDP ; c::c

?dump@C@QAE@XZ PROC ; c::dump, COMDAT
; _this$ = ecx
    mov     eax, DWORD PTR [ecx+4]
    mov     ecx, DWORD PTR [ecx]
    push    eax
    push    ecx
    push    OFFSET ??_C@_07NJBDCIEC@?$CFd?$DL?5?$CFd?6?$AA@
    call    _printf
    add     esp, 12      ; 0000000cH
    ret     0
?dump@C@QAE@XZ ENDP ; c::dump

```

That's all. One more thing to say is that stack pointer wasn't corrected with `add esp, X` after constructor called. Withal, constructor has `ret 8` instead of `RET` at the end.

This is all because here used [thiscall 2.5.4](#) calling convention, the method of passing values through the stack, which is, together with [stdcall 2.5.2](#) method, offers to correct stack to callee rather than to caller. `ret x` instruction adding `X` to `ESP`, then passes control to caller function.

See also section about calling conventions [2.5](#).

It's also should be noted that compiler deciding when to call constructor and destructor — but that's we already know from C++ language basics.

GCC

It's almost the same situation in GCC 4.4.1, with few exceptions.

```
main      public main
          proc near          ; DATA XREF: _start+17

var_20    = dword ptr -20h
var_1C    = dword ptr -1Ch
var_18    = dword ptr -18h
var_10    = dword ptr -10h
var_8     = dword ptr -8

          push    ebp
          mov     ebp, esp
          and     esp, 0FFFFFFF0h
          sub     esp, 20h
          lea     eax, [esp+20h+var_8]
          mov     [esp+20h+var_20], eax
          call    _ZN1cC1Ev
          mov     [esp+20h+var_18], 6
          mov     [esp+20h+var_1C], 5
          lea     eax, [esp+20h+var_10]
          mov     [esp+20h+var_20], eax
          call    _ZN1cC1Eii
          lea     eax, [esp+20h+var_8]
          mov     [esp+20h+var_20], eax
          call    _ZN1c4dumpEv
          lea     eax, [esp+20h+var_10]
          mov     [esp+20h+var_20], eax
          call    _ZN1c4dumpEv
          mov     eax, 0
          leave
          retn
main      endp
```

Here we see another *name mangling* style, specific to GNU¹⁰⁰ standards. It's also can be noted that pointer to object is passed as first function argument — hiddenly from programmer, of course.

First constructor:

```
_ZN1cC1Ev  public _ZN1cC1Ev ; weak
          proc near          ; CODE XREF: main+10

arg_0     = dword ptr 8

          push    ebp
          mov     ebp, esp
          mov     eax, [ebp+arg_0]
          mov     dword ptr [eax], 667
          mov     eax, [ebp+arg_0]
          mov     dword ptr [eax+4], 999
          pop     ebp
          retn
_ZN1cC1Ev  endp
```

What it does is just writes two numbers using pointer passed in first (and single) argument.

Second constructor:

```
_ZN1cC1Eii public _ZN1cC1Eii
          proc near
```

¹⁰⁰One more document about different compilers name mangling types: http://www.agner.org/optimize/calling_conventions.pdf

```

arg_0      = dword ptr  8
arg_4      = dword ptr  0Ch
arg_8      = dword ptr  10h

        push    ebp
        mov     ebp, esp
        mov     eax, [ebp+arg_0]
        mov     edx, [ebp+arg_4]
        mov     [eax], edx
        mov     eax, [ebp+arg_0]
        mov     edx, [ebp+arg_8]
        mov     [eax+4], edx
        pop     ebp
        retn
_ZN1cC1Eii endp

```

This is a function, analog of which could be looks like:

```

void ZN1cC1Eii (int *obj, int a, int b)
{
    *obj=a;
    *(obj+1)=b;
};

```

...and that's completely predictable.

Now dump() function:

```

_ZN1c4dumpEv public _ZN1c4dumpEv
               proc near

var_18        = dword ptr -18h
var_14        = dword ptr -14h
var_10        = dword ptr -10h
arg_0         = dword ptr  8

        push    ebp
        mov     ebp, esp
        sub     esp, 18h
        mov     eax, [ebp+arg_0]
        mov     edx, [eax+4]
        mov     eax, [ebp+arg_0]
        mov     eax, [eax]
        mov     [esp+18h+var_10], edx
        mov     [esp+18h+var_14], eax
        mov     [esp+18h+var_18], offset aDD ; "%d; %d\n"
        call    _printf
        leave
        retn
_ZN1c4dumpEv endp

```

This function in its *internal representation* has sole argument, used as pointer to the object (*this*).

Thus, if to base our judgment on these simple examples, the difference between MSVC and GCC is style of function names encoding (*name mangling*) and passing pointer to object (via ECX register or via first argument).

1.17.2 Class inheritance in C++

It can be said about inherited classes that it's simple structure we already considered, but extending in inherited classes.

Let's take simple example:

```
#include <stdio.h>

class object
{
    public:
        int color;
        object() { };
        object (int color) { this->color=color; };
        void print_color() { printf ("color=%d\n", color); };
};

class box : public object
{
    private:
        int width, height, depth;
    public:
        box(int color, int width, int height, int depth)
        {
            this->color=color;
            this->width=width;
            this->height=height;
            this->depth=depth;
        };
        void dump()
        {
            printf ("this is box. color=%d, width=%d, height=%d, depth=%d\n", color, width, height, depth);
        };
};

class sphere : public object
{
    private:
        int radius;
    public:
        sphere(int color, int radius)
        {
            this->color=color;
            this->radius=radius;
        };
        void dump()
        {
            printf ("this is sphere. color=%d, radius=%d\n", color, radius);
        };
};

int main()
{
    box b(1, 10, 20, 30);
    sphere s(2, 40);

    b.print_color();
    s.print_color();

    b.dump();
    s.dump();

    return 0;
};
```

Let's investigate generated code of dump() functions/methods and also object::print_color(), let's see

memory layout for structures-objects (as of 32-bit code).

So, dump() methods for several classes, generated by MSVC 2008 with /Ox and /Ob0 options ¹⁰¹

Listing 1.98: Optimizing MSVC 2008 /Ob0

```
??_C@_09GCEDOLPA@color?$DN?$CFd?6?$AA@ DB 'color=%d', 0aH, 00H ; 'string'
?print_color@object@@QAEXXZ PROC ; object::print_color, COMDAT
; _this$ = ecx
    mov     eax, DWORD PTR [ecx]
    push    eax

; 'color=%d', 0aH, 00H
    push    OFFSET ??_C@_09GCEDOLPA@color?$DN?$CFd?6?$AA@
    call    _printf
    add     esp, 8
    ret     0
?print_color@object@@QAEXXZ ENDP ; object::print_color
```

Listing 1.99: Optimizing MSVC 2008 /Ob0

```
?dump@box@@QAEXXZ PROC ; box::dump, COMDAT
; _this$ = ecx
    mov     eax, DWORD PTR [ecx+12]
    mov     edx, DWORD PTR [ecx+8]
    push    eax
    mov     eax, DWORD PTR [ecx+4]
    mov     ecx, DWORD PTR [ecx]
    push    edx
    push    eax
    push    ecx

; 'this is box. color=%d, width=%d, height=%d, depth=%d', 0aH, 00H ; 'string'
    push    OFFSET ??_C@_0DG@NCNGAADL@this?5is?5box?4?5color?$DN?$CFd?0?5width?$DN?$CFd?0@
    call    _printf
    add     esp, 20 ; 00000014H
    ret     0
?dump@box@@QAEXXZ ENDP ; box::dump
```

Listing 1.100: Optimizing MSVC 2008 /Ob0

```
?dump@sphere@@QAEXXZ PROC ; sphere::dump, COMDAT
; _this$ = ecx
    mov     eax, DWORD PTR [ecx+4]
    mov     ecx, DWORD PTR [ecx]
    push    eax
    push    ecx

; 'this is sphere. color=%d, radius=%d', 0aH, 00H
    push    OFFSET ??_C@_0CF@EFEDJLDC@this?5is?5sphere?4?5color?$DN?$CFd?0?5radius@
    call    _printf
    add     esp, 12 ; 0000000cH
    ret     0
?dump@sphere@@QAEXXZ ENDP ; sphere::dump
```

So, here is memory layout:
(base class *object*)

offset	description
+0x0	int color

(inherited classes)

box:

¹⁰¹ /Ob0 options mean inline expansion disabling, because, function inlining right into the code where the function is called will make our experiment harder

offset	description
+0x0	int color
+0x4	int width
+0x8	int height
+0xC	int depth

sphere:

offset	description
+0x0	int color
+0x4	int radius

Let's see `main()` function body:

Listing 1.101: Optimizing MSVC 2008 /Ob0

```

PUBLIC  _main
_TEXT   SEGMENT
_s$ = -24                ; size = 8
_b$ = -16                ; size = 16
_main   PROC
    sub     esp, 24                ; 00000018H
    push    30                    ; 0000001eH
    push    20                    ; 00000014H
    push    10                    ; 0000000aH
    push    1
    lea     ecx, DWORD PTR _b$[esp+40]
    call    ??0box@@QAE@HHHH@Z    ; box::box
    push    40                    ; 00000028H
    push    2
    lea     ecx, DWORD PTR _s$[esp+32]
    call    ??0sphere@@QAE@HH@Z   ; sphere::sphere
    lea     ecx, DWORD PTR _b$[esp+24]
    call    ?print_color@object@@QAEXXZ ; object::print_color
    lea     ecx, DWORD PTR _s$[esp+24]
    call    ?print_color@object@@QAEXXZ ; object::print_color
    lea     ecx, DWORD PTR _b$[esp+24]
    call    ?dump@box@@QAEXXZ      ; box::dump
    lea     ecx, DWORD PTR _s$[esp+24]
    call    ?dump@sphere@@QAEXXZ   ; sphere::dump
    xor     eax, eax
    add     esp, 24                ; 00000018H
    ret     0
_main   ENDP

```

Inherited classes should always add their fields after base classes' fields, so to make possible for base class methods to work with their fields.

When `object::print_color()` method is called, a pointers to both *box* object and *sphere* object are passed as *this*, it can work with these objects easily because *color* field in these objects is always at the pinned address (at +0x0 offset).

It can be said, `object::print_color()` method is agnostic in relation to input object type as long as fields will be *pinned* at the same addresses, and this condition is always true.

And if you create inherited class of *box* class, for example, compiler will add new fields after *depth* field, leaving *box* class fields at the pinned addresses.

Thus, `box::dump()` method will work fine accessing *color/width/height/depts* fields always pinned on known addresses.

GCC-generated code is almost the same, with the sole exception of *this* pointer passing (as it was described above, it passing as first argument instead of ECX registers).

1.17.3 Encapsulation in C++

Encapsulation is data hiding in *private* sections of class, for example, to allow access to them only from this class methods, but no more than.

However, are there any marks in code about the fact that some field is private and some other — not?

No, there are no such marks.

Let's try simple example:

```
#include <stdio.h>

class box
{
    private:
        int color, width, height, depth;
    public:
        box(int color, int width, int height, int depth)
        {
            this->color=color;
            this->width=width;
            this->height=height;
            this->depth=depth;
        };
        void dump()
        {
            printf ("this is box. color=%d, width=%d, height=%d, depth=%d\n", color, width, height, depth);
        };
};
```

Let's compile it again in MSVC 2008 with /Ox and /Ob0 options and let's see box::dump() method code:

```
?dump@box@@QAEXXZ PROC                ; box::dump, COMDAT
; _this$ = ecx
        mov     eax, DWORD PTR [ecx+12]
        mov     edx, DWORD PTR [ecx+8]
        push    eax
        mov     eax, DWORD PTR [ecx+4]
        mov     ecx, DWORD PTR [ecx]
        push    edx
        push    eax
        push    ecx
; 'this is box. color=%d, width=%d, height=%d, depth=%d', 0aH, 00H
        push    OFFSET ??_C@_ODG@NCNGAADL@this?5is?5box?4?5color?$DN?$CFd?0?5width?$DN?$CFd?0@
        call    _printf
        add     esp, 20                    ; 00000014H
        ret     0
?dump@box@@QAEXXZ ENDP                ; box::dump
```

Here is a memory layout of the class:

offset	description
+0x0	int color
+0x4	int width
+0x8	int height
+0xC	int depth

All fields are private and not allowed to access from any other functions, but, knowing this layout, can we create a code modifying these fields?

So I added hack_oop_encapsulation() function, which, if has the body as follows, won't compile:

```
void hack_oop_encapsulation(class box * o)
{
    o->width=1; // that code can't be compiled: "error C2248: 'box::width' : cannot access private member
    declared in class 'box'"
};
```

Nevertheless, if to cast *box* type to *pointer to int array*, and if to modify array of *int*-s we got, then we have success.

```
void hack_oop_encapsulation(class box * o)
{
    unsigned int *ptr_to_object=reinterpret_cast<unsigned int*>(o);
    ptr_to_object[1]=123;
};
```

This functions' code is very simple — it can be said, the function taking pointer to array of *int*-s on input and writing 123 to the second *int*:

```
?hack_oop_encapsulation@@YAXPAVbox@@@Z PROC                ; hack_oop_encapsulation
    mov     eax, DWORD PTR _o$[esp-4]
    mov     DWORD PTR [eax+4], 123                          ; 0000007bH
    ret     0
?hack_oop_encapsulation@@YAXPAVbox@@@Z ENDP                ; hack_oop_encapsulation
```

Let's check, how it works:

```
int main()
{
    box b(1, 10, 20, 30);

    b.dump();

    hack_oop_encapsulation(&b);

    b.dump();

    return 0;
};
```

Let's run:

```
this is box. color=1, width=10, height=20, depth=30
this is box. color=1, width=123, height=20, depth=30
```

We see, encapsulation is just class fields protection only on compiling stage. C++ compiler won't allow to generate a code modifying protected fields straightforwardly, nevertheless, it's possible using *dirty hacks*.

1.17.4 Multiple inheritance in C++

Multiple inheritance is a class creation which inherits fields and methods from two or more classes.

Let's write simple example again:

```
#include <stdio.h>

class box
{
public:
    int width, height, depth;
    box() { };
    box(int width, int height, int depth)
    {
        this->width=width;
        this->height=height;
        this->depth=depth;
    };
    void dump()
    {
        printf ("this is box. width=%d, height=%d, depth=%d\n", width, height, depth);
    };
    int get_volume()
    {
        return width * height * depth;
    };
};

class solid_object
{
public:
    int density;
    solid_object() { };
    solid_object(int density)
    {
        this->density=density;
    };
    int get_density()
    {
        return density;
    };
    void dump()
    {
        printf ("this is solid_object. density=%d\n", density);
    };
};

class solid_box: box, solid_object
{
public:
    solid_box (int width, int height, int depth, int density)
    {
        this->width=width;
        this->height=height;
        this->depth=depth;
        this->density=density;
    };
    void dump()
    {
        printf ("this is solid_box. width=%d, height=%d, depth=%d, density=%d\n", width, height, depth,
density);
    };
    int get_weight() { return get_volume() * get_density(); };
};

int main()
```

```

{
    box b(10, 20, 30);
    solid_object so(100);
    solid_box sb(10, 20, 30, 3);

    b.dump();
    so.dump();
    sb.dump();
    printf ("%d\n", sb.get_weight());

    return 0;
};

```

Let's compile it in MSVC 2008 with /Ox and /Ob0 options and let's see `box::dump()`, `solid_object::dump()` and `solid_box::dump()` methods code:

Listing 1.102: Optimizing MSVC 2008 /Ob0

```

?dump@box@@QAEXXZ PROC                                ; box::dump, COMDAT
; _this$ = ecx
    mov     eax, DWORD PTR [ecx+8]
    mov     edx, DWORD PTR [ecx+4]
    push    eax
    mov     eax, DWORD PTR [ecx]
    push    edx
    push    eax
; 'this is box. width=%d, height=%d, depth=%d', 0aH, 00H
    push    OFFSET ??_C@_OCM@DIKPHDFI@this?5is?5box?4?5width?$DN?$CFd?0?5height?$DN?$CFd@
    call    _printf
    add     esp, 16                                ; 00000010H
    ret     0
?dump@box@@QAEXXZ ENDP                                ; box::dump

```

Listing 1.103: Optimizing MSVC 2008 /Ob0

```

?dump@solid_object@@QAEXXZ PROC                        ; solid_object::dump, COMDAT
; _this$ = ecx
    mov     eax, DWORD PTR [ecx]
    push    eax
; 'this is solid_object. density=%d', 0aH
    push    OFFSET ??_C@_OCC@KICFJINL@this?5is?5solid_object?4?5density?$DN?$CFd@
    call    _printf
    add     esp, 8
    ret     0
?dump@solid_object@@QAEXXZ ENDP                        ; solid_object::dump

```

Listing 1.104: Optimizing MSVC 2008 /Ob0

```

?dump@solid_box@@QAEXXZ PROC                          ; solid_box::dump, COMDAT
; _this$ = ecx
    mov     eax, DWORD PTR [ecx+12]
    mov     edx, DWORD PTR [ecx+8]
    push    eax
    mov     eax, DWORD PTR [ecx+4]
    mov     ecx, DWORD PTR [ecx]
    push    edx
    push    eax
    push    ecx
; 'this is solid_box. width=%d, height=%d, depth=%d, density=%d', 0aH
    push    OFFSET ??_C@_ODO@HNCNIHNN@this?5is?5solid_box?4?5width?$DN?$CFd?0?5hei@
    call    _printf
    add     esp, 20                                ; 00000014H
    ret     0
?dump@solid_box@@QAEXXZ ENDP                          ; solid_box::dump

```

So, the memory layout for all three classes is:
box class:

offset	description
+0x0	width
+0x4	height
+0x8	depth

solid_object class:

offset	description
+0x0	density

It can be said, *solid_box* class memory layout will be *united*:

solid_box class:

offset	description
+0x0	width
+0x4	height
+0x8	depth
+0xC	density

The code of *box::get_volume()* and *solid_object::get_density()* methods is trivial:

Listing 1.105: Optimizing MSVC 2008 /Ob0

```
?get_volume@box@@QAEHXZ PROC                                ; box::get_volume, COMDAT
; _this$ = ecx
    mov     eax, DWORD PTR [ecx+8]
    imul    eax, DWORD PTR [ecx+4]
    imul    eax, DWORD PTR [ecx]
    ret     0
?get_volume@box@@QAEHXZ ENDP                                ; box::get_volume
```

Listing 1.106: Optimizing MSVC 2008 /Ob0

```
?get_density@solid_object@@QAEHXZ PROC                      ; solid_object::get_density, COMDAT
; _this$ = ecx
    mov     eax, DWORD PTR [ecx]
    ret     0
?get_density@solid_object@@QAEHXZ ENDP                      ; solid_object::get_density
```

But the code of *solid_box::get_weight()* method is much more interesting:

Listing 1.107: Optimizing MSVC 2008 /Ob0

```
?get_weight@solid_box@@QAEHXZ PROC                          ; solid_box::get_weight, COMDAT
; _this$ = ecx
    push    esi
    mov     esi, ecx
    push    edi
    lea     ecx, DWORD PTR [esi+12]
    call    ?get_density@solid_object@@QAEHXZ              ; solid_object::get_density
    mov     ecx, esi
    mov     edi, eax
    call    ?get_volume@box@@QAEHXZ                        ; box::get_volume
    imul    eax, edi
    pop     edi
    pop     esi
    ret     0
?get_weight@solid_box@@QAEHXZ ENDP                          ; solid_box::get_weight
```

get_weight() just calling two methods, but for *get_volume()* it just passing pointer to this, and for *get_density()* it passing pointer to this shifted by 12 (or 0xC) bytes, and there, in *solid_box* class memory layout, fields of *solid_object* class are beginning.

Thus, *solid_object::get_density()* method will believe it is working with usual *solid_object* class, and *box::get_volume()* method will work with its three fields, believing this is usual object of *box* class.

Thus, we can say, an object of some class, inheriting from several classes, representing in memory *united* class, containing all inherited fields. And each inherited method called with a pointer to corresponding structure's part passed.

1.17.5 C++ virtual methods

Yet another simple example:

```
#include <stdio.h>

class object
{
public:
    int color;
    object() { };
    object (int color) { this->color=color; };
    virtual void dump()
    {
        printf ("color=%d\n", color);
    };
};

class box : public object
{
private:
    int width, height, depth;
public:
    box(int color, int width, int height, int depth)
    {
        this->color=color;
        this->width=width;
        this->height=height;
        this->depth=depth;
    };
    void dump()
    {
        printf ("this is box. color=%d, width=%d, height=%d, depth=%d\n", color, width, height, depth);
    };
};

class sphere : public object
{
private:
    int radius;
public:
    sphere(int color, int radius)
    {
        this->color=color;
        this->radius=radius;
    };
    void dump()
    {
        printf ("this is sphere. color=%d, radius=%d\n", color, radius);
    };
};

int main()
{
    box b(1, 10, 20, 30);
    sphere s(2, 40);

    object *o1=&b;
    object *o2=&s;

    o1->dump();
    o2->dump();
    return 0;
};
```

Class *object* has virtual method `dump()`, being replaced in *box* and *sphere* class-inheritors.

If in some environment, where it's not known what type has object, as in `main()` function in example, a virtual method `dump()` is called, somewhere, the information about its type should be stored, so to call relevant virtual method.

Let's compile it in MSVC 2008 with `/Ox` and `/Ob0` options and let's see `main()` function code:

```
_s$ = -32 ; size = 12
_b$ = -20 ; size = 20
_main PROC
    sub     esp, 32 ; 00000020H
    push    30 ; 0000001eH
    push    20 ; 00000014H
    push    10 ; 0000000aH
    push    1
    lea     ecx, DWORD PTR _b$[esp+48]
    call    ??0box@@QAE@HHH@Z ; box::box
    push    40 ; 00000028H
    push    2
    lea     ecx, DWORD PTR _s$[esp+40]
    call    ??0sphere@@QAE@HH@Z ; sphere::sphere
    mov     eax, DWORD PTR _b$[esp+32]
    mov     edx, DWORD PTR [eax]
    lea     ecx, DWORD PTR _b$[esp+32]
    call    edx
    mov     eax, DWORD PTR _s$[esp+32]
    mov     edx, DWORD PTR [eax]
    lea     ecx, DWORD PTR _s$[esp+32]
    call    edx
    xor     eax, eax
    add     esp, 32 ; 00000020H
    ret     0
_main ENDP
```

Pointer to `dump()` function is taken somewhere from object. Where the address of new method would be written there? Only somewhere in constructors: there are no other place, because, nothing more is called in `main()` function. ¹⁰²

Let's see `box` class constructor's code:

```
??_R0?AVbox@@@8 DD FLAT:??_7type_info@@6B@ ; box 'RTTI Type Descriptor'
                DD 00H
                DB '.?AVbox@@', 00H

??_R1A?0A@EA@box@@@8 DD FLAT:??_R0?AVbox@@@8 ; box::'RTTI Base Class Descriptor at (0,-1,0,64)'
                DD 01H
                DD 00H
                DD 0fffffffH
                DD 00H
                DD 040H
                DD FLAT:??_R3box@@@8

??_R2box@@@8 DD FLAT:??_R1A?0A@EA@box@@@8 ; box::'RTTI Base Class Array'
                DD FLAT:??_R1A?0A@EA@object@@@8

??_R3box@@@8 DD 00H ; box::'RTTI Class Hierarchy Descriptor'
                DD 00H
                DD 02H
                DD FLAT:??_R2box@@@8

??_R4box@@6B@ DD 00H ; box::'RTTI Complete Object Locator'
                DD 00H
                DD 00H
                DD FLAT:??_R0?AVbox@@@8
                DD FLAT:??_R3box@@@8

??_7box@@6B@ DD FLAT:??_R4box@@6B@ ; box::'vtable'
```

¹⁰² About pointers to functions, read more in relevant section: [1.19](#)

```

        DD      FLAT:??dump@box@@UAEXXZ

_color$ = 8                                ; size = 4
_width$ = 12                              ; size = 4
_height$ = 16                             ; size = 4
_depth$ = 20                              ; size = 4
??0box@@QAE@HHHH@Z PROC                  ; box::box, COMDAT
; _this$ = ecx
    push     esi
    mov      esi, ecx
    call     ??0object@@QAE@XZ            ; object::object
    mov      eax, DWORD PTR _color$[esp]
    mov      ecx, DWORD PTR _width$[esp]
    mov      edx, DWORD PTR _height$[esp]
    mov      DWORD PTR [esi+4], eax
    mov      eax, DWORD PTR _depth$[esp]
    mov      DWORD PTR [esi+16], eax
    mov      DWORD PTR [esi], OFFSET ??_7box@@6B@
    mov      DWORD PTR [esi+8], ecx
    mov      DWORD PTR [esi+12], edx
    mov      eax, esi
    pop      esi
    ret      16                          ; 00000010H
??0box@@QAE@HHHH@Z ENDP                  ; box::box

```

Here we see slightly different memory layout: the first field is a pointer to some table `box::'vftable'` (name was set by MSVC compiler).

In this table we see a link to the table named `box::'RTTI Complete Object Locator'` and also a link to `box::dump()` method. So this is named virtual methods table and RTTI¹⁰³. Table of virtual methods contain addresses of methods and RTTI table contain information about types. By the way, RTTI-tables are the tables enumerated while calling to *dynamic_cast* and *typeid* in C++. You can also see here class name as plain text string. Thus, some method of base *object* class may call virtual method *object::dump()*, which in turn, will call a method of inherited class, because that information is present right in the object's structure.

Some CPU time needed for enumerating these tables and finding right virtual method address, thus virtual methods are widely considered as slightly slower than usual methods.

In GCC-generated code RTTI-tables constructed slightly differently.

¹⁰³Run-time type information

1.18 Unions

1.18.1 Pseudo-random number generator example

If we need float random numbers from 0 to 1, the most simplest thing is to use random numbers generator like Mersenne twister producing random 32-bit values in DWORD form, transform this value to *float* and then divide it by `RAND_MAX` (`0xffffffff` in our case) — value we got will be in 0..1 interval.

But as we know, division operation is almost always very slow. Will it be possible to get rid of it, as in case of division by multiplication? [1.12](#)

Let's recall what float number consisted of: sign bit, significand bits and exponent bits. We need just to store random bits to all significand bits for getting random float number!

Exponent cannot be zero (number will be denormalized in this case), so we will store `01111111` to exponent — this mean exponent will be 1. Then fill significand with random bits, set sign bit to 0 (which mean positive number) and voilà. Generated numbers will be in 1 to 2 interval, so we also should subtract 1 from it.

Very simple linear congruential random numbers generator is used in my example¹⁰⁴, producing 32-bit numbers. The PRNG initializing by current time in UNIX-style.

Then, *float* type represented as *union* — that is the C/C++ construction enabling us to interpret piece of memory differently typed. In our case, we are able to create a variable of `union` type and then access to it as it's *float* or as it's *uint32_t*. It can be said, it's just a hack. A dirty one.

```
#include <stdio.h>
#include <stdint.h>
#include <time.h>

union uint32_t_float
{
    uint32_t i;
    float f;
};

// from the Numerical Recipes book
const uint32_t RNG_a=1664525;
const uint32_t RNG_c=1013904223;

int main()
{
    uint32_t_float tmp;

    uint32_t RNG_state=time(NULL); // initial seed
    for (int i=0; i<100; i++)
    {
        RNG_state=RNG_state*RNG_a+RNG_c;
        tmp.i=RNG_state & 0x007fffff | 0x3f800000;
        float x=tmp.f-1;
        printf ("%f\n", x);
    };
    return 0;
};
```

MSVC 2010 (/Ox):

```
$SG4232    DB    '%f', 0aH, 00H

__real@3ff0000000000000 DQ 03ff000000000000r    ; 1

tv140 = -4                                ; size = 4
_tmp$ = -4                                ; size = 4
_main    PROC
    push    ebp
    mov     ebp, esp
```

¹⁰⁴idea was taken from: <http://xor0110.wordpress.com/2010/09/24/how-to-generate-floating-point-random-numbers-efficiently>

```

    and     esp, -64                ; ffffffff0H
    sub     esp, 56                ; 00000038H
    push    esi
    push    edi
    push    0
    call    __time64
    add     esp, 4
    mov     esi, eax
    mov     edi, 100              ; 00000064H
$LN3@main:

; let's generate random 32-bit number

    imul    esi, 1664525          ; 0019660dH
    add     esi, 1013904223       ; 3c6ef35fH
    mov     eax, esi

; leave bits for significand only

    and     eax, 8388607          ; 007ffffffH

; set exponent to 1

    or      eax, 1065353216       ; 3f800000H

; store this value as int

    mov     DWORD PTR _tmp$[esp+64], eax
    sub     esp, 8

; load this value as float

    fld     DWORD PTR _tmp$[esp+72]

; subtract one from it

    fsub    QWORD PTR __real@3ff0000000000000
    fstp    DWORD PTR tv140[esp+72]
    fld     DWORD PTR tv140[esp+72]
    fstp    QWORD PTR [esp]
    push    OFFSET $SG4232
    call    _printf
    add     esp, 12              ; 0000000cH
    dec     edi
    jne     SHORT $LN3@main
    pop     edi
    xor     eax, eax
    pop     esi
    mov     esp, ebp
    pop     ebp
    ret     0
_main     ENDP
_TEXT     ENDS
END

```

GCC producing very same code.

1.19 Pointers to functions

Pointer to function, as any other pointer, is just an address of function beginning in its code segment.

It is often used in callbacks ¹⁰⁵.

Well-known examples are:

- `qsort()` ¹⁰⁶, `atexit()` ¹⁰⁷ from the standard C library;
- signals in *NIX OS ¹⁰⁸;
- thread starting: `CreateThread()` (win32), `pthread_create()` (POSIX);
- a lot of win32 functions, e.g. `EnumChildWindows()` ¹⁰⁹.

So, `qsort()` function is a C/C++ standard library quicksort implementation. The function is able to sort anything, any types of data, if you have a function for two elements comparison and `qsort()` is able to call it.

The comparison function can be defined as:

```
int (*compare)(const void *, const void *)
```

Let's use slightly modified example I found [here](#):

```
/* ex3 Sorting ints with qsort */

#include <stdio.h>
#include <stdlib.h>

int comp(const void * _a, const void * _b)
{
    const int *a=(const int *)_a;
    const int *b=(const int *)_b;

    if (*a==*b)
        return 0;
    else
        if (*a < *b)
            return -1;
        else
            return 1;
}

int main(int argc, char* argv[])
{
    int numbers[10]={1892,45,200,-98,4087,5,-12345,1087,88,-100000};
    int i;

    /* Sort the array */
    qsort(numbers,10,sizeof(int),comp) ;
    for (i=0;i<9;i++)
        printf("Number = %d\n",numbers[ i ] ) ;
    return 0;
}
```

Let's compile it in MSVC 2010 (I omitted some parts for the sake of brevity) with `/Ox` option:

Listing 1.108: Optimizing MSVC 2010

```
__a$ = 8          ; size = 4
__b$ = 12         ; size = 4
_comp          PROC
```

¹⁰⁵[http://en.wikipedia.org/wiki/Callback_\(computer_science\)](http://en.wikipedia.org/wiki/Callback_(computer_science))

¹⁰⁶[http://en.wikipedia.org/wiki/Qsort_\(C_standard_library\)](http://en.wikipedia.org/wiki/Qsort_(C_standard_library))

¹⁰⁷<http://www.opengroup.org/onlinepubs/009695399/functions/atexit.html>

¹⁰⁸<http://en.wikipedia.org/wiki/Signal.h>

¹⁰⁹[http://msdn.microsoft.com/en-us/library/ms633494\(VS.85\).aspx](http://msdn.microsoft.com/en-us/library/ms633494(VS.85).aspx)

```

mov     eax, DWORD PTR __a$[esp-4]
mov     ecx, DWORD PTR __b$[esp-4]
mov     eax, DWORD PTR [eax]
mov     ecx, DWORD PTR [ecx]
cmp     eax, ecx
jne     SHORT $LN4@comp
xor     eax, eax
ret     0
$LN4@comp:
xor     edx, edx
cmp     eax, ecx
setge   dl
lea     eax, DWORD PTR [edx+edx-1]
ret     0
_comp   ENDP

...

_numbers$ = -44      ; size = 40
_i$ = -4             ; size = 4
_argc$ = 8           ; size = 4
_argv$ = 12          ; size = 4
_main      PROC
push     ebp
mov     ebp, esp
sub     esp, 44      ; 0000002cH
mov     DWORD PTR _numbers$[ebp], 1892      ; 00000764H
mov     DWORD PTR _numbers$[ebp+4], 45      ; 0000002dH
mov     DWORD PTR _numbers$[ebp+8], 200     ; 000000c8H
mov     DWORD PTR _numbers$[ebp+12], -98    ; ffffffff9eH
mov     DWORD PTR _numbers$[ebp+16], 4087   ; 00000ff7H
mov     DWORD PTR _numbers$[ebp+20], 5      ;
mov     DWORD PTR _numbers$[ebp+24], -12345 ; fffffcfc7H
mov     DWORD PTR _numbers$[ebp+28], 1087   ; 0000043fH
mov     DWORD PTR _numbers$[ebp+32], 88     ; 00000058H
mov     DWORD PTR _numbers$[ebp+36], -100000 ; fffe7960H
push     OFFSET _comp
push     4
push     10          ; 0000000aH
lea     eax, DWORD PTR _numbers$[ebp]
push     eax
call    _qsort
add     esp, 16      ; 00000010H

...

```

Nothing surprising so far. As a fourth argument, an address of label `_comp` is passed, that's just a place where function `comp()` located.

How `qsort()` calling it?

Let's take a look into this function located in `MSVCR80.DLL` (a `MSVC` DLL module with C standard library functions):

Listing 1.109: `MSVCR80.DLL`

```

.text:7816CBF0 ; void __cdecl qsort(void *, unsigned int, unsigned int, int (__cdecl *)(const void *, const
void *))
.text:7816CBF0      public _qsort
.text:7816CBF0 _qsort      proc near
.text:7816CBF0
.text:7816CBF0 lo          = dword ptr -104h
.text:7816CBF0 hi          = dword ptr -100h
.text:7816CBF0 var_FC      = dword ptr -0FCh
.text:7816CBF0 stkptr      = dword ptr -0F8h
.text:7816CBF0 lostk       = dword ptr -0F4h
.text:7816CBF0 histk       = dword ptr -7Ch
.text:7816CBF0 base        = dword ptr 4

```

```

.text:7816CBF0 num          = dword ptr 8
.text:7816CBF0 width       = dword ptr 0Ch
.text:7816CBF0 comp        = dword ptr 10h
.text:7816CBF0
                sub      esp, 100h

....

.text:7816CCE0 loc_7816CCE0:                                ; CODE XREF: _qsort+B1
.text:7816CCE0                shr      eax, 1
.text:7816CCE2                imul    eax, ebp
.text:7816CCE5                add     eax, ebx
.text:7816CCE7                mov     edi, eax
.text:7816CCE9                push   edi
.text:7816CCEA                push   ebx
.text:7816CCEB                call   [esp+118h+comp]
.text:7816CCF2                add     esp, 8
.text:7816CCF5                test    eax, eax
.text:7816CCF7                jle     short loc_7816CD04

```

comp — is fourth function argument. Here the control is just passed to the address in comp. Before it, two arguments prepared for comp(). Its result is checked after its execution.

That's why it's dangerous to use pointers to functions. First of all, if you call qsort() with incorrect pointer to function, qsort() may pass control to incorrect place, process may crash and this bug will be hard to find.

Second reason is that callback function types should comply strictly, calling wrong function with wrong arguments of wrong types may lead to serious problems, however, process crashing is not a big problem — big problem is to determine a reason of crashing — because compiler may be silent about potential trouble while compiling.

1.19.1 GCC

Not a big difference:

Listing 1.110: GCC

```

lea     eax, [esp+40h+var_28]
mov     [esp+40h+var_40], eax
mov     [esp+40h+var_28], 764h
mov     [esp+40h+var_24], 2Dh
mov     [esp+40h+var_20], 0C8h
mov     [esp+40h+var_1C], 0FFFFFF9Eh
mov     [esp+40h+var_18], 0FF7h
mov     [esp+40h+var_14], 5
mov     [esp+40h+var_10], 0FFFCFC7h
mov     [esp+40h+var_C], 43Fh
mov     [esp+40h+var_8], 58h
mov     [esp+40h+var_4], 0FFFE7960h
mov     [esp+40h+var_34], offset comp
mov     [esp+40h+var_38], 4
mov     [esp+40h+var_3C], 0Ah
call    _qsort

```

comp() function:

```

comp      public comp
comp      proc near

arg_0     = dword ptr 8
arg_4     = dword ptr 0Ch

                push   ebp
                mov     ebp, esp
                mov     eax, [ebp+arg_4]
                mov     ecx, [ebp+arg_0]
                mov     edx, [eax]
                xor     eax, eax

```

```

        cmp     [ecx], edx
        jnz     short loc_8048458
        pop     ebp
        retn

loc_8048458:
        setnl   al
        movzx   eax, al
        lea     eax, [eax+eax-1]
        pop     ebp
        retn

comp    endp

```

qsort() implementation is located in libc.so.6 and it is in fact just a wrapper for qsort_r(). It will call then quicksort(), where our defined function will be called via passed pointer:

Listing 1.111: (File libc.so.6, glibc version — 2.10.1)

```

.text:0002DDF6      mov     edx, [ebp+arg_10]
.text:0002DDF9      mov     [esp+4], esi
.text:0002DDFD      mov     [esp], edi
.text:0002DE00      mov     [esp+8], edx
.text:0002DE04      call    [ebp+arg_C]
...

```


1.20 SIMD

SIMD is just acronym: *Single Instruction, Multiple Data*.

As it's said, it's multiple data processing using only one instruction.

Just as FPU, that CPU subsystem looks like separate processor inside x86.

SIMD began as MMX in x86. 8 new 64-bit registers appeared: MM0-MM7.

Each MMX register may hold 2 32-bit values, 4 16-bit values or 8 bytes. For example, it is possible to add 8 8-bit values (bytes) simultaneously by adding two values in MMX-registers.

One simple example is graphics editor, representing image as a two dimensional array. When user change image brightness, the editor should add some coefficient to each pixel value, or to subtract. For the sake of brevity, our image may be grayscale and each pixel defined by one 8-bit byte, then it's possible to change brightness of 8 pixels simultaneously.

When MMX appeared, these registers was actually located in FPU registers. It was possible to use either FPU or MMX at the same time. One might think, Intel saved on transistors, but in fact, the reason of such symbiosis is simpler — older operation system may not aware of additional CPU registers wouldn't save them at the context switching, but will save FPU registers. Thus, MMX-enabled CPU + old operation system + process using MMX = that all will work together.

SSE — is extension of SIMD registers up to 128 bits, now separately from FPU.

AVX — another extension to 256 bits.

Now about practical usage.

Of course, memory copying (`memcpy`), memory comparing (`memcmp`) and so on.

One more example: we got DES encryption algorithm, it takes 64-bit block, 56-bit key, encrypt block and produce 64-bit result. DES algorithm may be considered as a very large electronic circuit, with wires and AND/OR/NOT gates.

Bitslice DES¹¹⁰ — is an idea of processing group of blocks and keys simultaneously. Let's say, variable of type *unsigned int* on x86 may hold up to 32 bits, so, it's possible to store there intermediate results for 32 blocks-keys pairs simultaneously, using 64+56 variables of *unsigned int* type.

I wrote an utility to brute-force Oracle RDBMS passwords/hashes (ones based on DES), slightly modified bitslice DES algorithm for SSE2 and AVX — now it's possible to encrypt 128 or 256 block-keys pairs simultaneously.

http://conus.info/utils/ops_SIMD/

1.20.1 Vectorization

Vectorization¹¹¹, for example, is when you have a loop taking couple of arrays at input and producing one array. Loop body takes values from input arrays, do something and put result into output array. It's important that there is only one single operation applied to each element. Vectorization — is to process several elements simultaneously.

For example:

```
for (i = 0; i < 1024; i++)
{
    C[i] = A[i]*B[i];
}
```

This fragment of code takes elements from A and B, multiplies them and save result into C.

If each array element we have is 32-bit *int*, then it's possible to load 4 elements from A into 128-bit XMM-register, from B to another XMM-registers, and by executing *PMULLD* (*Multiply Packed Signed Dword Integers and Store Low Result*) and *PMULHW* (*Multiply Packed Signed Integers and Store High Result*), it's possible to get 4 64-bit products¹¹² at once.

Thus, loop body count is 1024/4 instead of 1024, that's 4 times less and, of course, faster.

Some compilers can do vectorization automatically in some simple cases, e.g., Intel C++¹¹³.

I wrote tiny function:

¹¹⁰<http://www.darkside.com.au/bitslice/>

¹¹¹[Wikipedia: vectorization](#)

¹¹²multiplication result

¹¹³More about Intel C++ automatic vectorization: [Excerpt: Effective Automatic Vectorization](#)

```
int f (int sz, int *ar1, int *ar2, int *ar3)
{
    for (int i=0; i<sz; i++)
        ar3[i]=ar1[i]+ar2[i];

    return 0;
};
```

Intel C++

Let's compile it with Intel C++ 11.1.051 win32:

```
icl intel.cpp /QaxSSE2 /Faintel.asm /Ox
```

We got (in IDA 5):

```
; int __cdecl f(int, int *, int *, int *)
        public ?f@@YAHHPAH00@Z
?f@@YAHHPAH00@Z proc near

var_10      = dword ptr -10h
sz          = dword ptr  4
ar1         = dword ptr  8
ar2         = dword ptr 0Ch
ar3         = dword ptr 10h

        push    edi
        push    esi
        push    ebx
        push    esi
        mov     edx, [esp+10h+sz]
        test    edx, edx
        jle     loc_15B
        mov     eax, [esp+10h+ar3]
        cmp     edx, 6
        jle     loc_143
        cmp     eax, [esp+10h+ar2]
        jbe     short loc_36
        mov     esi, [esp+10h+ar2]
        sub     esi, eax
        lea     ecx, ds:0[edx*4]
        neg     esi
        cmp     ecx, esi
        jbe     short loc_55

loc_36:                                ; CODE XREF: f(int,int *,int *,int *)+21
        cmp     eax, [esp+10h+ar2]
        jnb     loc_143
        mov     esi, [esp+10h+ar2]
        sub     esi, eax
        lea     ecx, ds:0[edx*4]
        cmp     esi, ecx
        jb      loc_143

loc_55:                                ; CODE XREF: f(int,int *,int *,int *)+34
        cmp     eax, [esp+10h+ar1]
        jbe     short loc_67
        mov     esi, [esp+10h+ar1]
        sub     esi, eax
        neg     esi
        cmp     ecx, esi
        jbe     short loc_7F

loc_67:                                ; CODE XREF: f(int,int *,int *,int *)+59
```

```

        cmp     eax, [esp+10h+ar1]
        jnb     loc_143
        mov     esi, [esp+10h+ar1]
        sub     esi, eax
        cmp     esi, ecx
        jb      loc_143

loc_7F:
        mov     edi, eax          ; CODE XREF: f(int,int *,int *,int *)+65
        and     edi, 0Fh         ; edi = ar1
        jz      short loc_9A     ; is ar1 16-byte aligned?
        test    edi, 3
        jnz     loc_162
        neg     edi
        add     edi, 10h
        shr     edi, 2

loc_9A:
        lea     ecx, [edi+4]      ; CODE XREF: f(int,int *,int *,int *)+84
        cmp     edx, ecx
        jl      loc_162
        mov     ecx, edx
        sub     ecx, edi
        and     ecx, 3
        neg     ecx
        add     ecx, edx
        test    edi, edi
        jbe     short loc_D6
        mov     ebx, [esp+10h+ar2]
        mov     [esp+10h+var_10], ecx
        mov     ecx, [esp+10h+ar1]
        xor     esi, esi

loc_C1:
        mov     edx, [ecx+esi*4]  ; CODE XREF: f(int,int *,int *,int *)+CD
        add     edx, [ebx+esi*4]
        mov     [eax+esi*4], edx
        inc     esi
        cmp     esi, edi
        jb      short loc_C1
        mov     ecx, [esp+10h+var_10]
        mov     edx, [esp+10h+sz]

loc_D6:
        mov     esi, [esp+10h+ar2] ; CODE XREF: f(int,int *,int *,int *)+B2
        lea     esi, [esi+edi*4] ; is ar2+i*4 16-byte aligned?
        test    esi, 0Fh
        jz      short loc_109    ; yes!
        mov     ebx, [esp+10h+ar1]
        mov     esi, [esp+10h+ar2]

loc_ED:
        movdqu  xmm1, xmmword ptr [ebx+edi*4]
        movdqu  xmm0, xmmword ptr [esi+edi*4] ; ar2+i*4 is not 16-byte aligned, so load it to xmm0
        paddb   xmm1, xmm0
        movdqa  xmmword ptr [eax+edi*4], xmm1
        add     edi, 4
        cmp     edi, ecx
        jb      short loc_ED
        jmp     short loc_127

; -----

loc_109:
        mov     ebx, [esp+10h+ar1] ; CODE XREF: f(int,int *,int *,int *)+E3
        mov     esi, [esp+10h+ar2]

```

```

loc_111:                                ; CODE XREF: f(int,int *,int *,int *)+125
        movdqu  xmm0, xmmword ptr [ebx+edi*4]
        paddb  xmm0, xmmword ptr [esi+edi*4]
        movdqa  xmmword ptr [eax+edi*4], xmm0
        add     edi, 4
        cmp     edi, ecx
        jnb     short loc_111

loc_127:                                ; CODE XREF: f(int,int *,int *,int *)+107
                                           ; f(int,int *,int *,int *)+164
        cmp     ecx, edx
        jnb     short loc_15B
        mov     esi, [esp+10h+ar1]
        mov     edi, [esp+10h+ar2]

loc_133:                                ; CODE XREF: f(int,int *,int *,int *)+13F
        mov     ebx, [esi+ecx*4]
        add     ebx, [edi+ecx*4]
        mov     [eax+ecx*4], ebx
        inc     ecx
        cmp     ecx, edx
        jnb     short loc_133
        jmp     short loc_15B

; -----

loc_143:                                ; CODE XREF: f(int,int *,int *,int *)+17
                                           ; f(int,int *,int *,int *)+3A ...
        mov     esi, [esp+10h+ar1]
        mov     edi, [esp+10h+ar2]
        xor     ecx, ecx

loc_14D:                                ; CODE XREF: f(int,int *,int *,int *)+159
        mov     ebx, [esi+ecx*4]
        add     ebx, [edi+ecx*4]
        mov     [eax+ecx*4], ebx
        inc     ecx
        cmp     ecx, edx
        jnb     short loc_14D

loc_15B:                                ; CODE XREF: f(int,int *,int *,int *)+A
                                           ; f(int,int *,int *,int *)+129 ...
        xor     eax, eax
        pop     ecx
        pop     ebx
        pop     esi
        pop     edi
        ret     4

; -----

loc_162:                                ; CODE XREF: f(int,int *,int *,int *)+8C
                                           ; f(int,int *,int *,int *)+9F
        xor     ecx, ecx
        jmp     short loc_127
?f@@YAHHPAH00@Z endp

```

SSE2-related instructions are:

- **MOVDQU** (*Move Unaligned Double Quadword*) — it just load 16 bytes from memory into XMM-register.
- **PADDQ** (*Add Packed Integers*) — adding 4 pairs of 32-bit numbers and leaving result in first operand. By the way, no exception raised in case of overflow and no flags will be set, just low 32-bit of result will be stored. If one of PADDQ operands — address of value in memory, address should be aligned by 16-byte border. If it's not aligned, exception will be raised ¹¹⁴.

¹¹⁴More about data aligning: [Wikipedia: Data structure alignment](#)

- **MOVDQA (Move Aligned Double Quadword)** — the same as **MOVDQU**, but requires address of value in memory to be aligned by 16-bit border. If it's not aligned, exception will be raised. **MOVDQA** works faster than **MOVDQU**, but requires aforesaid.

So, these SSE2-instructions will be executed only in case if there are more 4 pairs to work on plus pointer **ar3** is aligned on 16-byte border.

More than that, if **ar2** is aligned on 16-byte border too, this fragment of code will be executed:

```
movdqu xmm0, xmmword ptr [ebx+edi*4] ; ar1+i*4
padd    xmm0, xmmword ptr [esi+edi*4] ; ar2+i*4
movdqa  xmmword ptr [eax+edi*4], xmm0 ; ar3+i*4
```

Otherwise, value from **ar2** will be loaded to **XMM0** using **MOVDQU**, it doesn't require aligned pointer, but may work slower:

```
movdqu  xmm1, xmmword ptr [ebx+edi*4] ; ar1+i*4
movdqu  xmm0, xmmword ptr [esi+edi*4] ; ar2+i*4 is not 16-byte aligned, so load it to xmm0
padd    xmm1, xmm0
movdqa  xmmword ptr [eax+edi*4], xmm1 ; ar3+i*4
```

In all other cases, non-SSE2 code will be executed.

GCC

GCC may also vectorize in some simple cases¹¹⁵, if to use **-O3** option and to turn on SSE2 support: **-msse2**.

What we got (GCC 4.4.1):

```
; f(int, int *, int *, int *)
public _Z1fiPiS_S_
_Z1fiPiS_S_
proc near

var_18      = dword ptr -18h
var_14      = dword ptr -14h
var_10      = dword ptr -10h
arg_0       = dword ptr  8
arg_4       = dword ptr  0Ch
arg_8       = dword ptr  10h
arg_C       = dword ptr  14h

        push    ebp
        mov     ebp, esp
        push    edi
        push    esi
        push    ebx
        sub     esp, 0Ch
        mov     ecx, [ebp+arg_0]
        mov     esi, [ebp+arg_4]
        mov     edi, [ebp+arg_8]
        mov     ebx, [ebp+arg_C]
        test    ecx, ecx
        jle     short loc_80484D8
        cmp     ecx, 6
        lea     eax, [ebx+10h]
        ja      short loc_80484E8

loc_80484C1:                                ; CODE XREF: f(int,int *,int *,int *)+4B
                                                ; f(int,int *,int *,int *)+61 ...
        xor     eax, eax
        nop
        lea     esi, [esi+0]

loc_80484C8:                                ; CODE XREF: f(int,int *,int *,int *)+36
        mov     edx, [edi+eax*4]
```

¹¹⁵ More about GCC vectorization support: <http://gcc.gnu.org/projects/tree-ssa/vectorization.html>

```

        add     edx, [esi+eax*4]
        mov     [ebx+eax*4], edx
        add     eax, 1
        cmp     eax, ecx
        jnz     short loc_80484C8

loc_80484D8:                                ; CODE XREF: f(int,int *,int *,int *)+17
                                           ; f(int,int *,int *,int *)+A5
        add     esp, 0Ch
        xor     eax, eax
        pop     ebx
        pop     esi
        pop     edi
        pop     ebp
        retn

; -----
        align 8

loc_80484E8:                                ; CODE XREF: f(int,int *,int *,int *)+1F
        test    bl, 0Fh
        jnz     short loc_80484C1
        lea     edx, [esi+10h]
        cmp     ebx, edx
        jbe     loc_8048578

loc_80484F8:                                ; CODE XREF: f(int,int *,int *,int *)+E0
        lea     edx, [edi+10h]
        cmp     ebx, edx
        ja      short loc_8048503
        cmp     edi, eax
        jbe     short loc_80484C1

loc_8048503:                                ; CODE XREF: f(int,int *,int *,int *)+5D
        mov     eax, ecx
        shr     eax, 2
        mov     [ebp+var_14], eax
        shl     eax, 2
        test    eax, eax
        mov     [ebp+var_10], eax
        jz      short loc_8048547
        mov     [ebp+var_18], ecx
        mov     ecx, [ebp+var_14]
        xor     eax, eax
        xor     edx, edx
        nop

loc_8048520:                                ; CODE XREF: f(int,int *,int *,int *)+9B
        movdqu  xmm1, xmmword ptr [edi+eax]
        movdqu  xmm0, xmmword ptr [esi+eax]
        add     edx, 1
        padd    xmm0, xmm1
        movdqa  xmmword ptr [ebx+eax], xmm0
        add     eax, 10h
        cmp     edx, ecx
        jb      short loc_8048520
        mov     ecx, [ebp+var_18]
        mov     eax, [ebp+var_10]
        cmp     ecx, eax
        jz      short loc_80484D8

loc_8048547:                                ; CODE XREF: f(int,int *,int *,int *)+73
        lea     edx, ds:0[eax*4]
        add     esi, edx
        add     edi, edx
        add     ebx, edx
        lea     esi, [esi+0]

```

```

loc_8048558:                                ; CODE XREF: f(int,int *,int *,int *)+CC
        mov     edx, [edi]
        add     eax, 1
        add     edi, 4
        add     edx, [esi]
        add     esi, 4
        mov     [ebx], edx
        add     ebx, 4
        cmp     ecx, eax
        jg      short loc_8048558
        add     esp, 0Ch
        xor     eax, eax
        pop     ebx
        pop     esi
        pop     edi
        pop     ebp
        retn

; -----
loc_8048578:                                ; CODE XREF: f(int,int *,int *,int *)+52
        cmp     eax, esi
        jnb     loc_80484C1
        jmp     loc_80484F8
_Z1fiPiS_S_    endp

```

Almost the same, however, not as meticulously as Intel C++ doing it.

1.20.2 SIMD strlen() implementation

It should be noted that SIMD-instructions may be inserted into C/C++ code via special macros¹¹⁶. As of MSVC, some of them are located in `intrin.h` file.

It is possible to implement `strlen()` function¹¹⁷ using SIMD-instructions, working 2-2.5 times faster than usual implementation. This function will load 16 characters into XMM-register and check each against zero.

```

size_t strlen_sse2(const char *str)
{
    register size_t len = 0;
    const char *s=str;
    bool str_is_aligned=(((unsigned int)str)&0xFFFFFFF0) == (unsigned int)str;

    if (str_is_aligned==false)
        return strlen (str);

    __m128i xmm0 = _mm_setzero_si128();
    __m128i xmm1;
    int mask = 0;

    for (;;)
    {
        xmm1 = _mm_load_si128((__m128i *)s);
        xmm1 = _mm_cmpeq_epi8(xmm1, xmm0);
        if ((mask = _mm_movemask_epi8(xmm1)) != 0)
        {
            unsigned long pos;
            _BitScanForward(&pos, mask);
            len += (size_t)pos;

            break;
        }
        s += sizeof(__m128i);
        len += sizeof(__m128i);
    }
}

```

¹¹⁶MSDN: MMX, SSE, and SSE2 Intrinsics

¹¹⁷strlen() — standard C library function for calculating string length

```
};

return len;
}
```

(the example is based on source code from [there](#)).

Let's compile in MSVC 2010 with /Ox option:

```
_pos$75552 = -4          ; size = 4
_str$ = 8                ; size = 4
?strlen_sse2@@YAIPBD@Z PROC ; strlen_sse2

    push    ebp
    mov     ebp, esp
    and     esp, -16      ; ffffffff0H
    mov     eax, DWORD PTR _str$[ebp]
    sub     esp, 12       ; 0000000cH
    push    esi
    mov     esi, eax
    and     esi, -16      ; ffffffff0H
    xor     edx, edx
    mov     ecx, eax
    cmp     esi, eax
    je      SHORT $LN4@strlen_sse
    lea     edx, DWORD PTR [eax+1]
    npad    3
$LL11@strlen_sse:
    mov     cl, BYTE PTR [eax]
    inc     eax
    test    cl, cl
    jne     SHORT $LL11@strlen_sse
    sub     eax, edx
    pop     esi
    mov     esp, ebp
    pop     ebp
    ret     0
$LN4@strlen_sse:
    movdqa  xmm1, XMMWORD PTR [eax]
    pxor    xmm0, xmm0
    pcmpeqb xmm1, xmm0
    pmovmskb eax, xmm1
    test    eax, eax
    jne     SHORT $LN9@strlen_sse
$LL3@strlen_sse:
    movdqa  xmm1, XMMWORD PTR [ecx+16]
    add     ecx, 16        ; 00000010H
    pcmpeqb xmm1, xmm0
    add     edx, 16        ; 00000010H
    pmovmskb eax, xmm1
    test    eax, eax
    je      SHORT $LL3@strlen_sse
$LN9@strlen_sse:
    bsf     eax, eax
    mov     ecx, eax
    mov     DWORD PTR _pos$75552[esp+16], eax
    lea     eax, DWORD PTR [ecx+edx]
    pop     esi
    mov     esp, ebp
    pop     ebp
    ret     0
?strlen_sse2@@YAIPBD@Z ENDP          ; strlen_sse2
```

First of all, we check `str` pointer, if it's aligned by 16-byte border. If not, let's call usual `strlen()` implementation.

Then, load next 16 bytes into XMM1 register using `MOVDQA` instruction.

Observant reader might ask, why `MOVDQU` cannot be used here, because it can load data from the memory regardless the fact if the pointer aligned or not.

Yes, it might be done in this way: if pointer is aligned, load data using `MOVDQA`, if not — use slower `MOVDQU`.

But here we are may stick into hard to notice caveat:

In Windows NT line of operation systems¹¹⁸, but not limited to it, memory allocated by pages of 4 KiB (4096 bytes). Each win32-process has ostensibly 4 GiB, but in fact, only some parts of address space are connected to real physical memory. If the process accessing to the absent memory block, exception will be raised. That's how virtual memory works¹¹⁹.

So, some function loading 16 bytes at once, may step over a border of allocated memory block. Let's consider, OS allocated 8192 (0x2000) bytes at the address 0x008c0000. Thus, the block is the bytes starting from address 0x008c0000 to 0x008c1fff inclusive.

After that block, that is, starting from address 0x008c2008 there are nothing at all, e.g., OS not allocated any memory there. Attempt to access a memory starting from that address will raise exception.

And let's consider, the program holding some string containing 5 characters almost at the end of block, and that's not a crime.

0x008c1ff8	'h'
0x008c1ff9	'e'
0x008c1ffa	'l'
0x008c1ffb	'l'
0x008c1ffc	'o'
0x008c1ffd	'\x00'
0x008c1ffe	random noise
0x008c1fff	random noise

So, in usual conditions the program calling `strlen()` passing it a pointer to string 'hello' lying in memory at address 0x008c1ff8. `strlen()` will read one byte at a time until 0x008c1ffd, where zero-byte, and so here it will stop working.

Now if we implement own `strlen()` reading 16 byte at once, starting at any address, will it be aligned or not, `MOVDQU` may attempt to load 16 bytes at once at address 0x008c1ff8 up to 0x008c2008, and then exception will be raised. That's the situation to be avoided, of course.

So then we'll work only with the addresses aligned by 16 byte border, what in combination with a knowledge of operation system page size is usually aligned by 16 byte too, give us some warranty our function will not read from unallocated memory.

Let's back to our function.

`_mm_setzero_si128()` — is a macro generating `pxor xmm0, xmm0` — instruction just clear /XMMZERO register

`_mm_load_si128()` — is a macro for `MOVDQA`, it just loading 16 bytes from the address in XMM1.

`_mm_cmpeq_epi8()` — is a macro for `PCMPEQB`, is an instruction comparing two XMM-registers bitwise.

And if some byte was equals to other, there will be 0xff at this place in result or 0 if otherwise.

For example.

XMM1: 11223344556677880000000000000000

XMM0: 11ab3444007877881111111111111111

After `pcmpeqb xmm1, xmm0` execution, XMM1 register will contain:

XMM1: ff0000ff0000ffff0000000000000000

In our case, this instruction comparing each 16-byte block with the block of 16 zero-bytes, was set in XMM0 by `pxor xmm0, xmm0`.

The next macro is `_mm_movemask_epi8()` — that is `PMOVBMSKB` instruction.

It is very useful if to use it with `PCMPEQB`.

¹¹⁸Windows NT, 2000, XP, Vista, 7, 8

¹¹⁹[http://en.wikipedia.org/wiki/Page_\(computer_memory\)](http://en.wikipedia.org/wiki/Page_(computer_memory))

```
pmovmskb eax, xmm1
```

This instruction will set first EAX bit into 1 if most significant bit of the first byte in XMM1 is 1. In other words, if first byte of XMM1 register is 0xff, first EAX bit will be set to 1 too.

If second byte in XMM1 register is 0xff, then second EAX bit will be set to 1 too. In other words, the instruction is answer to the question *which bytes in XMM1 are 0xff?* And will prepare 16 bits in EAX. Other EAX bits will be cleared.

By the way, do not forget about this feature of our algorithm:

There might be 16 bytes on input like `hello\x00garbage\x00ab`

It's a 'hello' string, terminating zero after and some random noise in memory.

If we load these 16 bytes into XMM1 and compare them with zeroed XMM0, we will get something like (I use here order from MSB¹²⁰ to LSB¹²¹):

```
XMM1: 0000ff00000000000000ff0000000000
```

This mean, the instruction found two zero bytes, and that's not surprising.

PMOVMSKB in our case will prepare EAX like (in binary representation): `0010000000100000b`.

Obviously, our function should consider only first zero and ignore others.

The next instruction — BSF (*Bit Scan Forward*). This instruction find first bit set to 1 and stores its position into first operand.

```
EAX=0010000000100000b
```

After `bsf eax, eax` instruction execution, EAX will contain 5, this mean, 1 found at 5th bit position (starting from zero).

MSVC has a macro for this instruction: `_BitScanForward`.

Now it's simple. If zero byte found, its position added to what we already counted and now we have ready to return result.

Almost all.

By the way, it's also should be noted, MSVC compiler emitted two loop bodies side by side, for optimization.

By the way, SSE 4.2 (appeared in Intel Core i7) offers more instructions where these string manipulations might be even easier: http://www.strchr.com/strcmp_and_strlen_using_sse_4.2

¹²⁰ most significant bit

¹²¹ least significant bit

1.21 64 bits

1.21.1 x86-64

It's a 64-bit extension to x86-architecture.

From the reverse engineer's perspective, most important differences are:

- Almost all registers (except FPU and SIMD) are extended to 64 bits and got *r-* prefix. 8 additional registers added. Now general purpose registers are: *rax, rbx, rcx, rdx, rbp, rsp, rsi, rdi, r8, r9, r10, r11, r12, r13, r14, r15*.

It's still possible to access to *older* register parts as usual. For example, it's possible to access lower 32-bit part of *RAX* using *EAX*.

New *r8-r15* registers also has its *lower parts*: *r8d-r15d* (lower 32-bit parts), *r8w-r15w* (lower 16-bit parts), *r8b-r15b* (lower 8-bit parts).

SIMD-registers number are doubled: from 8 to 16: *XMM0-XMM15*.

- In Win64, function calling convention is slightly different, somewhat resembling *fastcall* [2.5.3](#). First 4 arguments stored in *RCX, RDX, R8, R9* registers, others — in stack. Caller function should also allocate 32 bytes so the callee may save there 4 first arguments and use these registers for own needs. Short functions may use arguments just from registers, but larger may save their values into stack.

See also section about calling conventions [2.5](#).

- C *int* type is still 32-bit for compatibility.
- All pointers are 64-bit now.

Since now registers number are doubled, compilers has more space now for maneuvering calling *register allocation* ¹²². What it meanings for us, emitted code will contain less local variables.

For example, function calculating first S-box of DES encryption algorithm, it processing 32/64/128/256 values at once (depending on *DES_type* type (*uint32, uint64, SSE2* or *AVX*)) using *bitslice* DES method (read more about this method here [1.20](#)):

```
/*
 * Generated S-box files.
 *
 * This software may be modified, redistributed, and used for any purpose,
 * so long as its origin is acknowledged.
 *
 * Produced by Matthew Kwan - March 1998
 */

#ifdef _WIN64
#define DES_type unsigned __int64
#else
#define DES_type unsigned int
#endif

void
s1 (
    DES_type    a1,
    DES_type    a2,
    DES_type    a3,
    DES_type    a4,
    DES_type    a5,
    DES_type    a6,
    DES_type    *out1,
    DES_type    *out2,
    DES_type    *out3,
```

¹²² assigning variables to registers

```

DES_type    *out4
) {
    DES_type    x1, x2, x3, x4, x5, x6, x7, x8;
    DES_type    x9, x10, x11, x12, x13, x14, x15, x16;
    DES_type    x17, x18, x19, x20, x21, x22, x23, x24;
    DES_type    x25, x26, x27, x28, x29, x30, x31, x32;
    DES_type    x33, x34, x35, x36, x37, x38, x39, x40;
    DES_type    x41, x42, x43, x44, x45, x46, x47, x48;
    DES_type    x49, x50, x51, x52, x53, x54, x55, x56;

    x1 = a3 & ~a5;
    x2 = x1 ^ a4;
    x3 = a3 & ~a4;
    x4 = x3 | a5;
    x5 = a6 & x4;
    x6 = x2 ^ x5;
    x7 = a4 & ~a5;
    x8 = a3 ^ a4;
    x9 = a6 & ~x8;
    x10 = x7 ^ x9;
    x11 = a2 | x10;
    x12 = x6 ^ x11;
    x13 = a5 ^ x5;
    x14 = x13 & x8;
    x15 = a5 & ~a4;
    x16 = x3 ^ x14;
    x17 = a6 | x16;
    x18 = x15 ^ x17;
    x19 = a2 | x18;
    x20 = x14 ^ x19;
    x21 = a1 & x20;
    x22 = x12 ^ ~x21;
    *out2 ^= x22;
    x23 = x1 | x5;
    x24 = x23 ^ x8;
    x25 = x18 & ~x2;
    x26 = a2 & ~x25;
    x27 = x24 ^ x26;
    x28 = x6 | x7;
    x29 = x28 ^ x25;
    x30 = x9 ^ x24;
    x31 = x18 & ~x30;
    x32 = a2 & x31;
    x33 = x29 ^ x32;
    x34 = a1 & x33;
    x35 = x27 ^ x34;
    *out4 ^= x35;
    x36 = a3 & x28;
    x37 = x18 & ~x36;
    x38 = a2 | x3;
    x39 = x37 ^ x38;
    x40 = a3 | x31;
    x41 = x24 & ~x37;
    x42 = x41 | x3;
    x43 = x42 & ~a2;
    x44 = x40 ^ x43;
    x45 = a1 & ~x44;
    x46 = x39 ^ ~x45;
    *out1 ^= x46;
    x47 = x33 & ~x9;
    x48 = x47 ^ x39;
    x49 = x4 ^ x36;
    x50 = x49 & ~x5;
    x51 = x42 | x18;
    x52 = x51 ^ a5;
    x53 = a2 & ~x52;

```

```

x54 = x50 ^ x53;
x55 = a1 | x54;
x56 = x48 ^ ~x55;
*out3 ^= x56;
}

```

There is a lot of local variables. Of course, not them all will be in local stack. Let's compile it with MSVC 2008 with /Ox option:

Listing 1.112: Optimizing MSVC 2008

```

PUBLIC      _s1
; Function compile flags: /Ogtpy
_TEXT      SEGMENT
_x6$ = -20      ; size = 4
_x3$ = -16      ; size = 4
_x1$ = -12      ; size = 4
_x8$ = -8       ; size = 4
_x4$ = -4       ; size = 4
_a1$ = 8        ; size = 4
_a2$ = 12       ; size = 4
_a3$ = 16       ; size = 4
_x33$ = 20      ; size = 4
_x7$ = 20       ; size = 4
_a4$ = 20       ; size = 4
_a5$ = 24       ; size = 4
tv326 = 28      ; size = 4
_x36$ = 28      ; size = 4
_x28$ = 28      ; size = 4
_a6$ = 28       ; size = 4
_out1$ = 32     ; size = 4
_x24$ = 36     ; size = 4
_out2$ = 36     ; size = 4
_out3$ = 40     ; size = 4
_out4$ = 44     ; size = 4
_s1        PROC
    sub     esp, 20                ; 00000014H
    mov     edx, DWORD PTR _a5$[esp+16]
    push    ebx
    mov     ebx, DWORD PTR _a4$[esp+20]
    push    ebp
    push    esi
    mov     esi, DWORD PTR _a3$[esp+28]
    push    edi
    mov     edi, ebx
    not     edi
    mov     ebp, edi
    and     edi, DWORD PTR _a5$[esp+32]
    mov     ecx, edx
    not     ecx
    and     ebp, esi
    mov     eax, ecx
    and     eax, esi
    and     ecx, ebx
    mov     DWORD PTR _x1$[esp+36], eax
    xor     eax, ebx
    mov     esi, ebp
    or      esi, edx
    mov     DWORD PTR _x4$[esp+36], esi
    and     esi, DWORD PTR _a6$[esp+32]
    mov     DWORD PTR _x7$[esp+32], ecx
    mov     edx, esi
    xor     edx, eax
    mov     DWORD PTR _x6$[esp+36], edx
    mov     edx, DWORD PTR _a3$[esp+32]
    xor     edx, ebx
    mov     ebx, esi

```

```

xor     ebx, DWORD PTR _a5$[esp+32]
mov     DWORD PTR _x8$[esp+36], edx
and     ebx, edx
mov     ecx, edx
mov     edx, ebx
xor     edx, ebp
or      edx, DWORD PTR _a6$[esp+32]
not     ecx
and     ecx, DWORD PTR _a6$[esp+32]
xor     edx, edi
mov     edi, edx
or      edi, DWORD PTR _a2$[esp+32]
mov     DWORD PTR _x3$[esp+36], ebp
mov     ebp, DWORD PTR _a2$[esp+32]
xor     edi, ebx
and     edi, DWORD PTR _a1$[esp+32]
mov     ebx, ecx
xor     ebx, DWORD PTR _x7$[esp+32]
not     edi
or      ebx, ebp
xor     edi, ebx
mov     ebx, edi
mov     edi, DWORD PTR _out2$[esp+32]
xor     ebx, DWORD PTR [edi]
not     eax
xor     ebx, DWORD PTR _x6$[esp+36]
and     eax, edx
mov     DWORD PTR [edi], ebx
mov     ebx, DWORD PTR _x7$[esp+32]
or      ebx, DWORD PTR _x6$[esp+36]
mov     edi, esi
or      edi, DWORD PTR _x1$[esp+36]
mov     DWORD PTR _x28$[esp+32], ebx
xor     edi, DWORD PTR _x8$[esp+36]
mov     DWORD PTR _x24$[esp+32], edi
xor     edi, ecx
not     edi
and     edi, edx
mov     ebx, edi
and     ebx, ebp
xor     ebx, DWORD PTR _x28$[esp+32]
xor     ebx, eax
not     eax
mov     DWORD PTR _x33$[esp+32], ebx
and     ebx, DWORD PTR _a1$[esp+32]
and     eax, ebp
xor     eax, ebx
mov     ebx, DWORD PTR _out4$[esp+32]
xor     eax, DWORD PTR [ebx]
xor     eax, DWORD PTR _x24$[esp+32]
mov     DWORD PTR [ebx], eax
mov     eax, DWORD PTR _x28$[esp+32]
and     eax, DWORD PTR _a3$[esp+32]
mov     ebx, DWORD PTR _x3$[esp+36]
or      edi, DWORD PTR _a3$[esp+32]
mov     DWORD PTR _x36$[esp+32], eax
not     eax
and     eax, edx
or      ebx, ebp
xor     ebx, eax
not     eax
and     eax, DWORD PTR _x24$[esp+32]
not     ebp
or      eax, DWORD PTR _x3$[esp+36]
not     esi
and     ebp, eax

```

```

or     eax, edx
xor     eax, DWORD PTR _a5$[esp+32]
mov     edx, DWORD PTR _x36$[esp+32]
xor     edx, DWORD PTR _x4$[esp+36]
xor     ebp, edi
mov     edi, DWORD PTR _out1$[esp+32]
not     eax
and     eax, DWORD PTR _a2$[esp+32]
not     ebp
and     ebp, DWORD PTR _a1$[esp+32]
and     edx, esi
xor     eax, edx
or     eax, DWORD PTR _a1$[esp+32]
not     ebp
xor     ebp, DWORD PTR [edi]
not     ecx
and     ecx, DWORD PTR _x33$[esp+32]
xor     ebp, ebx
not     eax
mov     DWORD PTR [edi], ebp
xor     eax, ecx
mov     ecx, DWORD PTR _out3$[esp+32]
xor     eax, DWORD PTR [ecx]
pop     edi
pop     esi
xor     eax, ebx
pop     ebp
mov     DWORD PTR [ecx], eax
pop     ebx
add     esp, 20                ; 00000014H
ret     0
_s1     ENDP

```

5 variables was allocated in local stack by compiler.

Now let's try the same thing in 64-bit version of MSVC 2008:

Listing 1.113: Optimizing MSVC 2008

```

a1$ = 56
a2$ = 64
a3$ = 72
a4$ = 80
x36$1$ = 88
a5$ = 88
a6$ = 96
out1$ = 104
out2$ = 112
out3$ = 120
out4$ = 128
s1     PROC
$LN3:
mov     QWORD PTR [rsp+24], rbx
mov     QWORD PTR [rsp+32], rbp
mov     QWORD PTR [rsp+16], rdx
mov     QWORD PTR [rsp+8], rcx
push     rsi
push     rdi
push     r12
push     r13
push     r14
push     r15
mov     r15, QWORD PTR a5$[rsp]
mov     rcx, QWORD PTR a6$[rsp]
mov     rbp, r8
mov     r10, r9
mov     rax, r15
mov     rdx, rbp

```

```

not    rax
xor    rdx, r9
not    r10
mov    r11, rax
and    rax, r9
mov    rsi, r10
mov    QWORD PTR x36$1$[rsp], rax
and    r11, r8
and    rsi, r8
and    r10, r15
mov    r13, rdx
mov    rbx, r11
xor    rbx, r9
mov    r9, QWORD PTR a2$[rsp]
mov    r12, rsi
or     r12, r15
not    r13
and    r13, rcx
mov    r14, r12
and    r14, rcx
mov    rax, r14
mov    r8, r14
xor    r8, rbx
xor    rax, r15
not    rbx
and    rax, rdx
mov    rdi, rax
xor    rdi, rsi
or     rdi, rcx
xor    rdi, r10
and    rbx, rdi
mov    rcx, rdi
or     rcx, r9
xor    rcx, rax
mov    rax, r13
xor    rax, QWORD PTR x36$1$[rsp]
and    rcx, QWORD PTR a1$[rsp]
or     rax, r9
not    rcx
xor    rcx, rax
mov    rax, QWORD PTR out2$[rsp]
xor    rcx, QWORD PTR [rax]
xor    rcx, r8
mov    QWORD PTR [rax], rcx
mov    rax, QWORD PTR x36$1$[rsp]
mov    rcx, r14
or     rax, r8
or     rcx, r11
mov    r11, r9
xor    rcx, rdx
mov    QWORD PTR x36$1$[rsp], rax
mov    r8, rsi
mov    rdx, rcx
xor    rdx, r13
not    rdx
and    rdx, rdi
mov    r10, rdx
and    r10, r9
xor    r10, rax
xor    r10, rbx
not    rbx
and    rbx, r9
mov    rax, r10
and    rax, QWORD PTR a1$[rsp]
xor    rbx, rax
mov    rax, QWORD PTR out4$[rsp]

```



```

xor    rbx, QWORD PTR [rax]
xor    rbx, rcx
mov    QWORD PTR [rax], rbx
mov    rbx, QWORD PTR x36$1$[rsp]
and    rbx, rbp
mov    r9, rbx
not    r9
and    r9, rdi
or     r8, r11
mov    rax, QWORD PTR out1$[rsp]
xor    r8, r9
not    r9
and    r9, rcx
or     rdx, rbp
mov    rbp, QWORD PTR [rsp+80]
or     r9, rsi
xor    rbx, r12
mov    rcx, r11
not    rcx
not    r14
not    r13
and    rcx, r9
or     r9, rdi
and    rbx, r14
xor    r9, r15
xor    rcx, rdx
mov    rdx, QWORD PTR a1$[rsp]
not    r9
not    rcx
and    r13, r10
and    r9, r11
and    rcx, rdx
xor    r9, rbx
mov    rbx, QWORD PTR [rsp+72]
not    rcx
xor    rcx, QWORD PTR [rax]
or     r9, rdx
not    r9
xor    rcx, r8
mov    QWORD PTR [rax], rcx
mov    rax, QWORD PTR out3$[rsp]
xor    r9, r13
xor    r9, QWORD PTR [rax]
xor    r9, r8
mov    QWORD PTR [rax], r9
pop    r15
pop    r14
pop    r13
pop    r12
pop    rdi
pop    rsi
ret    0
s1     ENDP

```

Nothing allocated in local stack by compiler, x36 is synonym for a5.

By the way, we can see here, the function saved RCX and RDX registers in allocated by caller space, but R8 and R9 are not saved but used from the beginning.

By the way, there are CPUs with much more general purpose registers, Itanium, for example — 128 registers.

1.21.2 ARM

In ARM, 64-bit instructions are appeared in ARMv8.

1.22 C99 restrict

Here is a reason why FORTRAN programs, in some cases, works faster than C/C++ ones.

```
void f1 (int* x, int* y, int* sum, int* product, int* sum_product, int* update_me, size_t s)
{
    for (int i=0; i<s; i++)
    {
        sum[i]=x[i]+y[i];
        product[i]=x[i]*y[i];
        update_me[i]=i*123; // some dummy value
        sum_product[i]=sum[i]+product[i];
    };
};
```

That's very simple example with one specific thing in it: pointer to `update_me` array could be a pointer to `sum` array, `product` array, or even `sum_product` array — because, there are nothing criminal in it, right?

Compiler is fully aware about it, so it generates a code with four stages in loop body:

- calculate next `sum[i]`
- calculate next `product[i]`
- calculate next `update_me[i]`
- calculate next `sum_product[i]` — on this stage, we need to load from memory already calculated `sum[i]` and `product[i]`

Is it possible to optimize the last stage? Because already calculated `sum[i]` and `product[i]` are not necessary to load from memory again, because we already calculated them. Yes, but compiler isn't sure that nothing was overwritten on 3rd stage! This is called “pointer aliasing”, a situation, when compiler can't be sure that a memory to which some pointer is pointing, wasn't changed.

restrict in C99 standard [ISO07, 6.7.3.1] is a promise, given by programmer to compiler that function arguments marked by this keyword will always be pointing to different memory locations and never be crossed.

If to be more precise and describe this formally, *restrict* shows that only this pointer will be used to access an object, with which we are working via this pointer, and no other pointer will be used for it. It can be even said that some object will be accessed only via one single pointer, if it's marked as *restrict*.

Let's add this keyword to each argument-pointer:

```
void f2 (int* restrict x, int* restrict y, int* restrict sum, int* restrict product, int* restrict
sum_product,
        int* restrict update_me, size_t s)
{
    for (int i=0; i<s; i++)
    {
        sum[i]=x[i]+y[i];
        product[i]=x[i]*y[i];
        update_me[i]=i*123; // some dummy value
        sum_product[i]=sum[i]+product[i];
    };
};
```

Let's see results:

Listing 1.114: GCC x64: f1()

```
f1:
    push    r15 r14 r13 r12 rbp rdi rsi rbx
    mov     r13, QWORD PTR 120[rsi]
    mov     rbp, QWORD PTR 104[rsi]
    mov     r12, QWORD PTR 112[rsi]
    test    r13, r13
    je      .L1
    add     r13, 1
```

```

        xor     ebx, ebx
        mov     edi, 1
        xor     r11d, r11d
        jmp     .L4
.L6:
        mov     r11, rdi
        mov     rdi, rax
.L4:
        lea     rax, 0[0+r11*4]
        lea     r10, [rcx+rax]
        lea     r14, [rdx+rax]
        lea     rsi, [r8+rax]
        add     rax, r9
        mov     r15d, DWORD PTR [r10]
        add     r15d, DWORD PTR [r14]
        mov     DWORD PTR [rsi], r15d          ; store to sum[]
        mov     r10d, DWORD PTR [r10]
        imul    r10d, DWORD PTR [r14]
        mov     DWORD PTR [rax], r10d          ; store to product[]
        mov     DWORD PTR [r12+r11*4], ebx     ; store to update_me[]
        add     ebx, 123
        mov     r10d, DWORD PTR [rsi]          ; reload sum[i]
        add     r10d, DWORD PTR [rax]          ; reload product[i]
        lea     rax, 1[rdi]
        cmp     rax, r13
        mov     DWORD PTR 0[rbp+r11*4], r10d    ; store to sum_product[]
        jne     .L6
.L1:
        pop     rbx rsi rdi rbp r12 r13 r14 r15
        ret

```

Listing 1.115: GCC x64: f2()

```

f2:
        push    r13 r12 rbp rdi rsi rbx
        mov     r13, QWORD PTR 104[rsp]
        mov     rbp, QWORD PTR 88[rsp]
        mov     r12, QWORD PTR 96[rsp]
        test    r13, r13
        je      .L7
        add     r13, 1
        xor     r10d, r10d
        mov     edi, 1
        xor     eax, eax
        jmp     .L10
.L11:
        mov     rax, rdi
        mov     rdi, r11
.L10:
        mov     esi, DWORD PTR [rcx+rax*4]
        mov     r11d, DWORD PTR [rdx+rax*4]
        mov     DWORD PTR [r12+rax*4], r10d    ; store to update_me[]
        add     r10d, 123
        lea     ebx, [rsi+r11]
        imul    r11d, esi
        mov     DWORD PTR [r8+rax*4], ebx     ; store to sum[]
        mov     DWORD PTR [r9+rax*4], r11d    ; store to product[]
        add     r11d, ebx
        mov     DWORD PTR 0[rbp+rax*4], r11d    ; store to sum_product[]
        lea     r11, 1[rdi]
        cmp     r11, r13
        jne     .L11
.L7:
        pop     rbx rsi rdi rbp r12 r13
        ret

```

The difference between compiled `f1()` and `f2()` function as follows: in `f1()`, `sum[i]` and `product[i]` are reloaded in the middle of loop, and in `f2()` there are no such thing, already calculated values are used, because we “promised” to compiler, that no one and nothing will change values in `sum[i]` and `product[i]` while execution of loop body, so it is “sure” that value from memory may not be again. Obviously, second example will work faster.

But what if pointers in function arguments will be crossed somehow? This will be on programmer’s conscience, but results will be incorrect.

Let’s back to FORTRAN. Compilers from this programming language treats all pointers as such, so when there are was not possible to set *restrict*, FORTRAN in these cases may generate faster code.

How practical is it? In the cases when function works with several big blocks in memory. There are a lot of such in linear algebra, for example. A lot of linear algebra used on supercomputers/HPC, probably, that’s why, traditionally, FORTRAN is still used there [Loh10].

But when there are not a big number of iterations, certainly, speed boost will not be significant.

Chapter 2

Couple things to add

2.1 LEA instruction

LEA (*Load Effective Address*) is instruction intended not for values summing but for address forming, e.g., for forming address of array element by adding array address, element index, with multiplication of element size¹.

Important property of LEA instruction is that it do not alter processor flags.

```
int f(int a, int b)
{
    return a*8+b;
};
```

MSVC 2010 with /Ox option:

```
_a$ = 8                                ; size = 4
_b$ = 12                              ; size = 4
_f      PROC
    mov     eax, DWORD PTR _b$[esp-4]
    mov     ecx, DWORD PTR _a$[esp-4]
    lea     eax, DWORD PTR [eax+ecx*8]
    ret     0
_f      ENDP
```

¹See also: http://en.wikipedia.org/wiki/Addressing_mode

2.2 Function prologue and epilogue

Function prologue is instructions at function start. It is often something like the following code fragment:

```
push    ebp
mov     ebp, esp
sub     esp, X
```

What these instructions do: save EBP register value, set EBP to ESP and then allocate space in stack for local variables.

EBP value is fixed over a period of function execution and it will be used for local variables and arguments access. One can use ESP, but it changes over time and it is not convenient.

Function epilogue annuls allocated space in stack, returning EBP value to initial state and returning control flow to callee:

```
mov     esp, ebp
pop     ebp
ret     0
```

Epilogue and prologue can make recursion performance worse.

For example, once upon a time I wrote a function to seek right node in binary tree. As a recursive function it would look stylish but because some time was spent at each function call for prologue/epilogue, it was working couple of times slower than the implementation without recursion.

By the way, that is the reason of tail call² existence: when compiler (or interpreter) transforms recursion (with which it's possible: *tail recursion*) into iteration for efficiency.

²http://en.wikipedia.org/wiki/Tail_call

2.3 npad

It's an assembly language macro for label aligning by some specific border.

That's often need for the busy labels to where control flow is often passed, e.g., loop body begin. So the CPU will effectively load data or code from the memory, through memory bus, cache lines, etc.

Taken from listing.inc (MSVC):

By the way, it's curious example of different NOP variations. All these instructions has no effects whatsoever, but has different size.

```
;; LISTING.INC
;;
;; This file contains assembler macros and is included by the files created
;; with the -FA compiler switch to be assembled by MASM (Microsoft Macro
;; Assembler).
;;
;; Copyright (c) 1993-2003, Microsoft Corporation. All rights reserved.

;; non destructive nops
npad macro size
if size eq 1
    nop
else
    if size eq 2
        mov edi, edi
    else
        if size eq 3
            ; lea ecx, [ecx+00]
            DB 8DH, 49H, 00H
        else
            if size eq 4
                ; lea esp, [esp+00]
                DB 8DH, 64H, 24H, 00H
            else
                if size eq 5
                    add eax, DWORD PTR 0
                else
                    if size eq 6
                        ; lea ebx, [ebx+00000000]
                        DB 8DH, 9BH, 00H, 00H, 00H, 00H
                    else
                        if size eq 7
                            ; lea esp, [esp+00000000]
                            DB 8DH, 0A4H, 24H, 00H, 00H, 00H, 00H
                        else
                            if size eq 8
                                ; jmp .+8; .npad 6
                                DB 0EBH, 06H, 8DH, 9BH, 00H, 00H, 00H, 00H
                            else
                                if size eq 9
                                    ; jmp .+9; .npad 7
                                    DB 0EBH, 07H, 8DH, 0A4H, 24H, 00H, 00H, 00H, 00H
                                else
                                    if size eq 10
                                        ; jmp .+A; .npad 7; .npad 1
                                        DB 0EBH, 08H, 8DH, 0A4H, 24H, 00H, 00H, 00H, 00H, 90H
                                    else
                                        if size eq 11
                                            ; jmp .+B; .npad 7; .npad 2
                                            DB 0EBH, 09H, 8DH, 0A4H, 24H, 00H, 00H, 00H, 00H, 8BH, 0FFH
                                        else
                                            if size eq 12
                                                ; jmp .+C; .npad 7; .npad 3
                                                DB 0EBH, 0AH, 8DH, 0A4H, 24H, 00H, 00H, 00H, 00H, 8DH, 49H, 00H
                                            else
                                                if size eq 13
```


2.4 Signed number representations

There are several methods of representing signed numbers³, but in x86 architecture used “two’s complement”.

Difference between signed and unsigned numbers is that if we represent `0xFFFFFFFF` and `0x00000002` as unsigned, then first number (4294967294) is bigger than second (2). If to represent them both as signed, first will be -2 , and it is lesser than second (2). That is the reason why conditional jumps 1.8 are present both for signed (e.g. JG, JL) and unsigned (JA, JBE) operations.

2.4.1 Integer overflow

It is worth noting that incorrect representation of number can lead integer overflow vulnerability.

For example, we have some network service, it receives network packets. In that packets there are also field where subpacket length is coded. It is 32-bit value. After network packet received, service checking that field, and if it is larger than, for example, some `MAX_PACKET_SIZE` (let’s say, 10 kilobytes), packet ignored as incorrect. Comparison is signed. Intruder set this value to `0xFFFFFFFF`. While comparison, this number is considered as signed -1 and it’s lesser than 10 kilobytes. No error here. Service would like to copy that subpacket to another place in memory and call `memcpy (dst, src, 0xFFFFFFFF)` function: this operation, rapidly scratching a lot of inside of process memory.

More about it: <http://www.phrack.org/issues.html?issue=60&id=10>

³http://en.wikipedia.org/wiki/Signed_number_representations

2.5 Arguments passing methods (calling conventions)

2.5.1 cdecl

This is the most popular method for arguments passing to functions in C/C++ languages.

Caller pushing arguments to stack in reverse order: last argument, then penultimate element and finally — first argument. Caller also should return back ESP to its initial state after callee function exit.

Listing 2.1: cdecl

```
push arg3
push arg2
push arg1
call function
add esp, 12 ; return ESP
```

2.5.2 stdcall

Almost the same thing as *cdecl*, with the exception that callee set ESP to initial state executing `RET x` instruction instead of `RET`, where $x = \text{arguments number} * \text{sizeof(int)}$ ⁴. Caller will not adjust stack pointer by `add esp, x` instruction.

Listing 2.2: stdcall

```
push arg3
push arg2
push arg1
call function

function:
... do something ...
ret 12
```

This method is ubiquitous in win32 standard libraries, but not in win64 (see below about win64).

Variable arguments number functions

`printf()`-like functions are, probably, the only case of variable arguments functions in C/C++, but it's easy to illustrate an important difference between *cdecl* and *stdcall* with help of it. Let's start with the idea that compiler knows argument count of each `printf()` function calling. However, called `printf()`, which is already compiled and located in `MSVCRT.DLL` (if to talk about Windows), has not any information about how much arguments were passed, however it can determine it from format string. Thus, if `printf()` would be *stdcall*-function and restored stack pointer to its initial state by counting number of arguments in format string, this could be dangerous situation, when one programmer's typo may provoke sudden program crash. Thus it's not suitable for such functions to use *stdcall*, *cdecl* is better.

2.5.3 fastcall

That's general naming for a method for passing some of arguments via registers and all others — via stack. It worked faster than *cdecl/stdcall* on older CPUs. It's not a standardized way, so, various compilers may do it differently. Well known caveat: if you have two DLLs, one uses another, and they are built by different compilers with different *fastcall* calling conventions.

Both MSVC and GCC passing first and second argument via `ECX` and `EDX` and other arguments via stack. Caller should restore stack pointer into initial state.

Stack pointer should be restored to initial state by callee, like in *stdcall*.

⁴Size of *int* type variable is 4 in x86 systems and 8 in x64 systems

Listing 2.3: fastcall

```
push arg3
mov edx, arg2
mov ecx, arg1
call function

function:
.. do something ..
ret 4
```

GCC regparm

It's *fastcall* evolution⁵ is some sense. With the `-mregparm` option it's possible to set, how many arguments will be passed via registers. 3 at maximum. Thus, EAX, EDX and ECX registers will be used.

Of course, if number of arguments is less then 3, not all registers 3 will be used.

Caller restores stack pointer to its initial state.

2.5.4 thiscall

In C++, it's a *this* pointer to object passing into function-method.

In MSVC, *this* is usually passed in ECX register.

In GCC, *this* pointer is passed as a first function-method argument. Thus it will be seen that internally all function-methods has extra argument for it.

2.5.5 x86-64

win64

The method of arguments passing in Win64 is somewhat resembling to *fastcall*. First 4 arguments are passed via RCX, RDX, R8, R9, others — via stack. Caller also must prepare a place for 32 bytes or 4 64-bit values, so then callee can save there first 4 arguments. Short functions may use argument values just from registers, but larger may save its values for further use.

Caller also should return stack pointer into initial state.

This calling convention is also used in Windows x86-64 system DLLs (instead if *stdcall* in win32).

2.5.6 Returning values of *float* and *double* type

In all conventions except of Win64, values of type *float* or *double* returning via FPU register ST(0).

In Win64, return values of *float* and *double* types are returned in XMM0 register instead of ST(0).

⁵<http://www.ohse.de/uwe/articles/gcc-attributes.html#func-regparm>

2.6 position-independent code

While analyzing Linux shared (.so) libraries, one may frequently spot such code pattern:

Listing 2.4: libc-2.17.so x86

```
.text:0012D5E3 __x86_get_pc_thunk_bx proc near          ; CODE XREF: sub_17350+3
.text:0012D5E3                                     ; sub_173CC+4 ...
.text:0012D5E3             mov     ebx, [esp+0]
.text:0012D5E6             retn
.text:0012D5E6 __x86_get_pc_thunk_bx endp

...

.text:000576C0 sub_576C0      proc near                ; CODE XREF: tmpfile+73
...

.text:000576C0             push    ebp
.text:000576C1             mov     ecx, large gs:0
.text:000576C8             push    edi
.text:000576C9             push    esi
.text:000576CA             push    ebx
.text:000576CB             call   __x86_get_pc_thunk_bx
.text:000576D0             add     ebx, 157930h
.text:000576D6             sub     esp, 9Ch

...

.text:000579F0             lea     eax, (a__gen_tempname - 1AF000h)[ebx] ; "__gen_tempname"
.text:000579F6             mov     [esp+0ACh+var_A0], eax
.text:000579FA             lea     eax, (a__SysdepsPosix - 1AF000h)[ebx] ; "../sysdeps/posix/tempname.c"
.text:00057A00             mov     [esp+0ACh+var_A8], eax
.text:00057A04             lea     eax, (aInvalidKindIn_ - 1AF000h)[ebx] ; "! \"invalid KIND in
__gen_tempname\""
.text:00057A0A             mov     [esp+0ACh+var_A4], 14Ah
.text:00057A12             mov     [esp+0ACh+var_AC], eax
.text:00057A15             call   __assert_fail
```

All pointers to strings are corrected by some constant and by EBX value, which calculated at the beginning of each function. This is so called position-independent code (PIC), it is intended to execute placed in any random place of memory, that's why it can't contain any absolute memory addresses.

PIC was crucial in early computer systems and crucial now in embedded systems without virtual memory support (where processes are all placed in single continuous memory block). It is also still used in *NIX systems for shared libraries, because, shared libraries are shared across many processes while loaded in memory only once. But all these processes may map the same shared library on different addresses, so that's why shared library should be working correctly without fixing on any absolute address.

Let's do some simple experiment:

```
#include <stdio.h>

int global_variable=123;

int f1(int var)
{
    int rt=global_variable+var;
    printf ("returning %d\n", rt);
    return rt;
};
```

Let's compile it in GCC 4.7.3 and see resulting .so file in IDA 5:

```
gcc -fPIC -shared -O3 -o 1.so 1.c
```

Listing 2.5: GCC 4.7.3

```

.text:00000440      public __x86_get_pc_thunk_bx
.text:00000440 __x86_get_pc_thunk_bx proc near          ; CODE XREF: _init_proc+4
.text:00000440                                         ; deregister_tm_clones+4 ...
.text:00000440      mov     ebx, [esp+0]
.text:00000443      retn
.text:00000443 __x86_get_pc_thunk_bx endp

.text:00000570      public f1
.text:00000570 f1      proc near
.text:00000570
.text:00000570 var_1C      = dword ptr -1Ch
.text:00000570 var_18      = dword ptr -18h
.text:00000570 var_14      = dword ptr -14h
.text:00000570 var_8       = dword ptr -8
.text:00000570 var_4       = dword ptr -4
.text:00000570 arg_0      = dword ptr 4
.text:00000570
.text:00000570      sub     esp, 1Ch
.text:00000573      mov     [esp+1Ch+var_8], ebx
.text:00000577      call    __x86_get_pc_thunk_bx
.text:0000057C      add     ebx, 1A84h
.text:00000582      mov     [esp+1Ch+var_4], esi
.text:00000586      mov     eax, ds:(global_variable_ptr - 2000h)[ebx]
.text:0000058C      mov     esi, [eax]
.text:0000058E      lea     eax, (aReturningD - 2000h)[ebx] ; "returning %d\n"
.text:00000594      add     esi, [esp+1Ch+arg_0]
.text:00000598      mov     [esp+1Ch+var_18], eax
.text:0000059C      mov     [esp+1Ch+var_1C], 1
.text:000005A3      mov     [esp+1Ch+var_14], esi
.text:000005A7      call    ___printf_chk
.text:000005AC      mov     eax, esi
.text:000005AE      mov     ebx, [esp+1Ch+var_8]
.text:000005B2      mov     esi, [esp+1Ch+var_4]
.text:000005B6      add     esp, 1Ch
.text:000005B9      retn
.text:000005B9 f1      endp

```

That's it: pointers to «*returning %d\n*» string and *global_variable* are to be corrected at each function execution. The *__x86_get_pc_thunk_bx()* function return address of the point after call to itself (0x57C here) in EBX. That's the simple way to get value of program counter (EIP) at some place. The 0x1A84 constant is related to the difference between this function begin and so called *Global Offset Table Procedure Linkage Table* (GOT PLT), the section right after *Global Offset Table* (GOT), where pointer to *global_variable* is. IDA 5 shows these offset processed, so to understand them easily, but in fact the code is:

```

.text:00000577      call    __x86_get_pc_thunk_bx
.text:0000057C      add     ebx, 1A84h
.text:00000582      mov     [esp+1Ch+var_4], esi
.text:00000586      mov     eax, [ebx-0Ch]
.text:0000058C      mov     esi, [eax]
.text:0000058E      lea     eax, [ebx-1A30h]

```

So, EBX pointing to GOT PLT section and to calculate pointer to *global_variable* which stored in GOT, 0xC should be subtracted. To calculate pointer to the «*returning %d\n*» string, 0x1A30 should be subtracted.

By the way, that is the reason why AMD64 instruction set supports RIP⁶-relative addressing, just to simplify PIC-code.

Let's compile the same C code in the same GCC version, but for x64.

IDA 5 would simplify output code but suppressing RIP-relative addressing details, so I'll run *objdump* instead to see the details:

```

0000000000000720 <f1>:
720:  48 8b 05 b9 08 20 00      mov     rax,QWORD PTR [rip+0x2008b9]      # 200fe0 <_DYNAMIC+0x1d0>

```

⁶program counter in AMD64

```

727: 53          push    rbx
728: 89 fb       mov     ebx,edi
72a: 48 8d 35 20 00 00 00 lea     rsi,[rip+0x20]      # 751 <_fini+0x9>
731: bf 01 00 00 00 mov     edi,0x1
736: 03 18       add     ebx,DWORD PTR [rax]
738: 31 c0       xor     eax,eax
73a: 89 da       mov     edx,ebx
73c: e8 df fe ff ff call    620 <__printf_chk@plt>
741: 89 d8       mov     eax,ebx
743: 5b         pop     rbx
744: c3         ret

```

0x2008b9 is the difference between address of instruction at 0x720 and *global_variable* and 0x20 is the difference between that address of instruction at 0x72A and «*returning %d\n*» string.

The PIC mechanism is not used in Windows DLLs. If Windows loader needs to load DLL to another place, it “patches” DLL in memory (at the *FIXUP* places) in order to correct all addresses. This means, several Windows processes can’t share once loaded DLL on different addresses in different process’ memory blocks — because each loaded into memory DLL instance *fixed* to be work only at these addresses..

Chapter 3

Finding important/interesting stuff in the code

Minimalism it's not a significant feature of modern software.

But not because programmers writing a lot, but because all libraries are usually linked statically to executable files. If all external libraries were shifted into external DLL files, the world would be different. (Another reason for C++ — STL and other template libraries.)

Thus, it's very important to determine origin of some function, if it's from standard library or well-known library (like Boost¹, libpng²), and which one — is related to what we are trying to find in the code.

It's just absurdly to rewrite all code to C/C++ to find what we looking for.

One of the primary reverse engineer's task is to find quickly in the code what is needed.

IDA 5 disassembler allow us search among text strings, byte sequences, constants. It's even possible to export the code into .lst or .asm text file and then use `grep`, `awk`, etc.

When you try to understand what some code is doing, this easily could be some open-source library like libpng. So when you see some constants or text strings looks familiar, it's always worth to google it. And if you find the opensource project where it's used, then it will be enough just to compare the functions. It may solve some part of problem.

For example, once upon a time I tried to understand how SAP 6.0 network packets compression/decompression is working. It's a huge software, but a detailed .PDB with debugging information is present, and that's cosily. I finally came to idea that one of the functions doing decompressing of network packet called `CsDecomprLZC()`. Immediately I tried to google its name and I quickly found that the function named as the same is used in MaxDB (it's open-source SAP project)³.

<http://www.google.com/search?q=CsDecomprLZC>

Astoundingly, MaxDB and SAP 6.0 software shared the same code for network packets compression/decompression.

3.1 Communication with the outer world

First what to look on is which functions from operation system API and standard libraries are used.

If the program is divided into main executable file and a group of DLL-files, sometimes, these function's names may be helpful.

If we are interesting, what exactly may lead to `MessageBox()` call with specific text, first what we can try to do: find this text in data segment, find references to it and find the points from which a control may be passed to `MessageBox()` call we're interesting in.

If we are talking about some game and we're interesting, which events are more or less random in it, we may try to find `rand()` function or its replacement (like Mersenne twister algorithm) and find a places from which this function called and most important: how the results are used.

But if it's not a game, but `rand()` is used, it's also interesting, why. There are cases of unexpected `rand()` usage in data compression algorithm (for encryption imitation): <http://blog.yurichev.com/node/44>.

¹<http://www.boost.org/>

²<http://www.libpng.org/pub/png/libpng.html>

³More about it in releval section 7.2.1

3.2 String

Debugging messages are often very helpful if present. In some sense, debugging messages are reporting about what's going on in program right now. Often these are `printf()`-like functions, which writes to log-files, and sometimes, not writing anything but calls are still present, because this build is not debug build but release one. If local or global variables are dumped in debugging messages, it might be helpful as well because it's possible to get variable names at least. For example, one of such functions in Oracle RDBMS is `ksdwrt()`.

Sometimes `assert()` macro presence is useful too: usually, this macro leave in code source file name, line number and condition.

Meaningful text strings are often helpful. IDA 5 disassembler may show from which function and from which point this specific string is used. Funny cases [sometimes happen](#).

Paradoxically, but error messages may help us as well. In Oracle RDBMS, errors are reporting using group of functions. [More about it](#).

It's possible to find very quickly, which functions reporting about errors and in which conditions. By the way, it's often a reason why copy-protection systems has inarticulate cryptic error messages or just error numbers. No one happy when software cracker quickly understand why copy-protection is triggered just by error message.

3.3 Constants

Some algorithms, especially cryptographical, use distinct constants, which is easy to find in code using IDA 5.

For example, MD5⁴ algorithm initializes its own internal variables like:

```
var int h0 := 0x67452301
var int h1 := 0xEFCDAB89
var int h2 := 0x98BADCFE
var int h3 := 0x10325476
```

If you find these four constants usage in the code in a row — it's very high probability this function is related to MD5.

3.3.1 Magic numbers

A lot of file formats defining a standard file header where *magic number*⁵ is used.

For example, all Win32 and MS-DOS executables are started with two characters “MZ”⁶.

At the MIDI-file beginning “MThd” signature must be present. If we have a program that using MIDI-files for something, very likely, it will check MIDI-files for validity by checking at least first 4 bytes.

This could be done like:

(*buf* pointing to the beginning of loaded file into memory)

```
cmp [buf], 0x6468544D ; "MThd"
jnz _error_not_a_MIDI_file
```

...or by calling function for comparing memory blocks `memcmp()` or any other equivalent code up to `CMPSPB` instruction.

When you find such place you already may say where MIDI-file loading is beginning, also, we could see a location of MIDI-file contents buffer and what is used from that buffer and how.

⁴<http://en.wikipedia.org/wiki/MD5>

⁵[http://en.wikipedia.org/wiki/Magic_number_\(programming\)](http://en.wikipedia.org/wiki/Magic_number_(programming))

⁶http://en.wikipedia.org/wiki/DOS_MZ_executable

DHCP

This applies to network protocols as well. For example, DHCP protocol network packets contains so called *magic cookie*: 0x63538263. Any code generating DHCP protocol packets somewhere and somehow should embed this constant into packet. If we find it in the code we may find where it happen and not only this. *Something* that received DHCP packet should check *magic cookie*, comparing it with the constant.

For example, let's take dhcpcore.dll file from Windows 7 x64 and search for the constant. And we found it, two times: it seems, that constant is used in two functions eloquently named as DhcpExtractOptionsForValidation() and DhcpExtractFullOptions():

Listing 3.1: dhcpcore.dll (Windows 7 x64)

```
.rdata:000007FF6483CBE8 dword_7FF6483CBE8 dd 63538263h ; DATA XREF: DhcpExtractOptionsForValidation+79
.rdata:000007FF6483CBEC dword_7FF6483CBEC dd 63538263h ; DATA XREF: DhcpExtractFullOptions+97
```

And the places where these constants accessed:

Listing 3.2: dhcpcore.dll (Windows 7 x64)

```
.text:000007FF6480875F mov     eax, [rsi]
.text:000007FF64808761 cmp     eax, cs:dword_7FF6483CBE8
.text:000007FF64808767 jnz     loc_7FF64817179
```

And:

Listing 3.3: dhcpcore.dll (Windows 7 x64)

```
.text:000007FF648082C7 mov     eax, [r12]
.text:000007FF648082CB cmp     eax, cs:dword_7FF6483CBEC
.text:000007FF648082D1 jnz     loc_7FF648173AF
```

3.4 Finding the right instructions

If the program is using FPU instructions and there are very few of them in a code, one can try to check each by debugger.

For example, we may be interesting, how Microsoft Excel calculating formulae entered by user. For example, division operation.

If to load excel.exe (from Office 2010) version 14.0.4756.1000 into IDA 5, then make a full listing and to find each FDIV instructions (except ones which use constants as a second operand — obviously, it's not suits us):

```
cat EXCEL.lst | grep fdiv | grep -v dbl_ > EXCEL.fdiv
```

...then we realizing they are just 144.

We can enter string like $(1/3)$ in Excel and check each instruction.

Checking each instruction in debugger or *tracer* 5.0.1 (one may check 4 instruction at a time), it seems, we are lucky here and sought-for instruction is just 14th:

```
.text:3011E919 DC 33 fdiv qword ptr [ebx]
```

```
PID=13944|TID=28744|(0) 0x2f64e919 (Excel.exe!BASE+0x11e919)
EAX=0x02088006 EBX=0x02088018 ECX=0x00000001 EDX=0x00000001
ESI=0x02088000 EDI=0x00544804 EBP=0x0274FA3C ESP=0x0274F9F8
EIP=0x2F64E919
FLAGS=PF IF
FPU ControlWord=IC RC=NEAR PC=64bits PM UM OM ZM DM IM
FPU StatusWord=
FPU ST(0): 1.000000
```

ST(0) holding first argument (1) and second one is in [ebx].

Next instruction after FDIV writes result into memory:

```
.text:3011E91B DD 1E fstp qword ptr [esi]
```

If to set breakpoint on it, we may see result:

```
PID=32852|TID=36488|(0) 0x2f40e91b (Excel.exe!BASE+0x11e91b)
EAX=0x00598006 EBX=0x00598018 ECX=0x00000001 EDX=0x00000001
ESI=0x00598000 EDI=0x00294804 EBP=0x026CF93C ESP=0x026CF8F8
EIP=0x2F40E91B
FLAGS=PF IF
FPU ControlWord=IC RC=NEAR PC=64bits PM UM OM ZM DM IM
FPU StatusWord=C1 P
FPU ST(0): 0.333333
```

Also, as a practical joke, we can modify it on-fly:

```
tracer -l:excel.exe bpx=excel.exe!BASE+0x11E91B,set(st0,666)
```

```
PID=36540|TID=24056|(0) 0x2f40e91b (Excel.exe!BASE+0x11e91b)
EAX=0x00680006 EBX=0x00680018 ECX=0x00000001 EDX=0x00000001
ESI=0x00680000 EDI=0x00395404 EBP=0x0290FD9C ESP=0x0290FD58
EIP=0x2F40E91B
FLAGS=PF IF
FPU ControlWord=IC RC=NEAR PC=64bits PM UM OM ZM DM IM
FPU StatusWord=C1 P
FPU ST(0): 0.333333
Set ST0 register to 666.000000
```

Excel showing 666 in that cell what finally convincing us we find the right place.

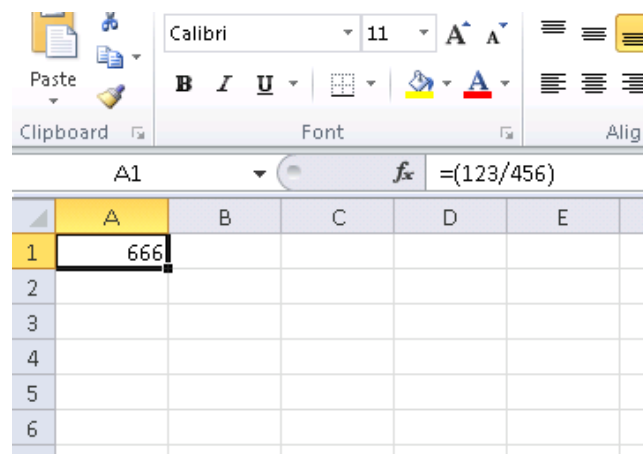


Figure 3.1: Practical joke worked

If to try the same Excel version, but x64, we'll see there are only 12 FDIV instructions, and the one we looking for — third.

```
tracer.exe -l:excel.exe bpx=excel.exe!BASE+0x1B7FCC,set(st0,666)
```

It seems, a lot of division operations of *float* and *double* types, compiler replaced by SSE-instructions like DIVSD (DIVSD present here 268 in total).

3.5 Suspicious code patterns

Modern compilers do not emit LOOP and RCL instructions. On the other hand, these instructions are well-known to coders who like to code in straight assembly language. If you spot these, it can be said, with a high probability, this fragment of code is hand-written. Also, function prologue/epilogue is not usually present in hand-written assembly copy.

3.6 Magic numbers usage while tracing

Often, main goal is to get to know, how some value was read from file, or received via network, being used. Often, manual tracing of some value is very labouring task. One of the simple methods (however, not 100% reliable) is to use your own *magic number*.

This resembling X-ray computed tomography is some sense: radiocontrast agent is often injected into patient's blood, which is used for improving visibility of internal structures in X-rays. For example, it's well known how blood of healthy man/woman percolates in kidneys and if agent is in blood, it will be easily seen on tomography, how good and normal blood was percolating, are there any stones or tumors.

We can take some 32-bit number like *0x0badf00d*, or someone birth date like *0x11101979* and to write this, 4 byte holding number, to some place in file used by the program we investigate.

Then, while tracing this program, with *tracer 5.0.1* in *code coverage* mode, and then, with help if *grep* or just by searching in the text file (of tracing results), we can easily see, where the value was used and how.

Example of *grepable tracer 5.0.1* results in *cc* mode:

0x150bf66 (_kziaia+0x14), e=	1 [MOV EBX, [EBP+8]] [EBP+8]=0xf59c934
0x150bf69 (_kziaia+0x17), e=	1 [MOV EDX, [69AEB08h]] [69AEB08h]=0
0x150bf6f (_kziaia+0x1d), e=	1 [FS: MOV EAX, [2Ch]]
0x150bf75 (_kziaia+0x23), e=	1 [MOV ECX, [EAX+EDX*4]] [EAX+EDX*4]=0xf1ac360
0x150bf78 (_kziaia+0x26), e=	1 [MOV [EBP-4], ECX] ECX=0xf1ac360

This can be used for network packets as well. It's important to be unique for *magic number* and not to be present in the program's code.

Aside of *tracer 5.0.1*, DosBox (MS-DOS emulator) in *heavydebug* mode, can write information about all register's states for each executed instruction of program to plain text file⁷, so this method may be useful for DOS programs as well.

3.7 Old-school methods, nevertheless, interesting to know

3.7.1 Memory “snapshots” comparing

The method of simple two memory snapshots comparing in order to see changes, was often used to hack 8-bit computer games and hacking “high score” files.

For example, if you got some loaded game on 8-bit computer (it's not much memory on these, but game is usually consumes even less memory) and you know that you have now, let's say, 100 bullets, you can do a “snapshot” of all memory and save it to some place. Then shoot somewhere, bullet count now 99, do second “snapshot” and then compare both: somewhere should be a byte which was 100 in the beginning and now it's 99. Considering a fact that these 8-bit games were often written in assembly language and such variables were global, it can be said for sure, which address in memory holding bullets count. If to search all references to that address in disassembled game code, it's not very hard to find a piece of code decrementing bullets count, write NOP instruction⁸ there, or couple of NOP-s, we'll have a game with 100 (for example) bullets forever. Games on these 8-bit computers was usually loaded on the same address, also, there were no much different versions of each game (usually, just one version is popular), enthusiastic gamers knew, which byte should be written (using BASIC instruction *POKE*⁹) to which address in order to hack it. This led to “cheat” lists containing of *POKE* instructions published in magazines related to 8-bit games. See also: http://en.wikipedia.org/wiki/PEEK_and_POKE.

The same story about modifying “high score” files, this may work not only with 8-bit games. Let's notice your score count and save the file somewhere. When “high score” count will be different, just compare two files, it can be even done with DOS-utility FC¹⁰ (“high score” files are often in binary form). There will be some place where couple of bytes will be different and it will be easy to see which ones are holding score number. However, game developers are aware of such tricks and may protect against it.

⁷ See also my blog post about this DosBox feature: <http://blog.yurichev.com/node/55>

⁸ “no operation”, idle operation

⁹ BASIC language instruction writting byte on specific address

¹⁰ MS-DOS utility for binary files comparing

Chapter 4

Tasks

There are two questions almost for every task, if otherwise isn't specified:

- 1) What this function does? Answer in one-sentence form.
- 2) Rewrite this function into C/C++.

Hints and solutions are in the appendix of this book.

4.1 Easy level

4.1.1 Task 1.1

This is standard C library function. Source code taken from OpenWatcom. Compiled in MSVC 2010.

```
_TEXT    SEGMENT
_input$ = 8                      ; size = 1
_f PROC
    push    ebp
    mov     ebp, esp
    movsx   eax, BYTE PTR _input$[ebp]
    cmp     eax, 97                ; 00000061H
    jl      SHORT $LN1@f
    movsx   ecx, BYTE PTR _input$[ebp]
    cmp     ecx, 122                ; 0000007aH
    jg      SHORT $LN1@f
    movsx   edx, BYTE PTR _input$[ebp]
    sub     edx, 32                 ; 00000020H
    mov     BYTE PTR _input$[ebp], dl
$LN1@f:
    mov     al, BYTE PTR _input$[ebp]
    pop     ebp
    ret     0
_f ENDP
_TEXT     ENDS
```

It is the same code compiled by GCC 4.4.1 with -O3 option (maximum optimization):

```
_f          proc near
input       = dword ptr 8

            push    ebp
            mov     ebp, esp
            movzx   eax, byte ptr [ebp+input]
            lea     edx, [eax-61h]
            cmp     dl, 19h
            ja      short loc_80483F2
            sub     eax, 20h

loc_80483F2:
            pop     ebp
```

```

        retn
_f      endp

```

4.1.2 Task 1.2

This is also standard C library function. Source code is taken from OpenWatcom and modified slightly. Compiled in MSVC 2010 with /Ox optimization flag.

This function also use these standard C functions: isspace() and isdigit().

```

EXTRN    _isdigit:PROC
EXTRN    _isspace:PROC
EXTRN    ___ptr_check:PROC
; Function compile flags: /Ogtpy
_TEXT    SEGMENT
_p$ = 8                                     ; size = 4
_f      PROC
    push    ebx
    push    esi
    mov     esi, DWORD PTR _p$[esp+4]
    push    edi
    push    0
    push    esi
    call    ___ptr_check
    mov     eax, DWORD PTR [esi]
    push    eax
    call    _isspace
    add     esp, 12                         ; 0000000cH
    test    eax, eax
    je      SHORT $LN6@f
    npad    2
$LL7@f:
    mov     ecx, DWORD PTR [esi+4]
    add     esi, 4
    push    ecx
    call    _isspace
    add     esp, 4
    test    eax, eax
    jne     SHORT $LL7@f
$LN6@f:
    mov     bl, BYTE PTR [esi]
    cmp     bl, 43                         ; 0000002bH
    je      SHORT $LN4@f
    cmp     bl, 45                         ; 0000002dH
    jne     SHORT $LN5@f
$LN4@f:
    add     esi, 4
$LN5@f:
    mov     edx, DWORD PTR [esi]
    push    edx
    xor     edi, edi
    call    _isdigit
    add     esp, 4
    test    eax, eax
    je      SHORT $LN2@f
$LL3@f:
    mov     ecx, DWORD PTR [esi]
    mov     edx, DWORD PTR [esi+4]
    add     esi, 4
    lea     eax, DWORD PTR [edi+edi*4]
    push    edx
    lea     edi, DWORD PTR [ecx+eax*2-48]
    call    _isdigit
    add     esp, 4
    test    eax, eax
    jne     SHORT $LL3@f

```

```

$LN20f:
    cmp     bl, 45                ; 0000002dH
    jne     SHORT $LN140f
    neg     edi
$LN140f:
    mov     eax, edi
    pop     edi
    pop     esi
    pop     ebx
    ret     0
_f        ENDP
_TEXT     ENDS

```

Same code compiled in GCC 4.4.1. This task is slightly harder because GCC compiled `isspace()` and `isdigit()` functions like inline-functions and inserted their bodies right into code.

```

_f        proc near

var_10    = dword ptr -10h
var_9     = byte ptr -9
input     = dword ptr  8

    push    ebp
    mov     ebp, esp
    sub     esp, 18h
    jmp     short loc_8048410

loc_804840C:
    add     [ebp+input], 4

loc_8048410:
    call    ___ctype_b_loc
    mov     edx, [eax]
    mov     eax, [ebp+input]
    mov     eax, [eax]
    add     eax, eax
    lea     eax, [edx+eax]
    movzx   eax, word ptr [eax]
    movzx   eax, ax
    and     eax, 2000h
    test    eax, eax
    jnz     short loc_804840C
    mov     eax, [ebp+input]
    mov     eax, [eax]
    mov     [ebp+var_9], al
    cmp     [ebp+var_9], '+'
    jz      short loc_8048444
    cmp     [ebp+var_9], '-'
    jnz     short loc_8048448

loc_8048444:
    add     [ebp+input], 4

loc_8048448:
    mov     [ebp+var_10], 0
    jmp     short loc_8048471

loc_8048451:
    mov     edx, [ebp+var_10]
    mov     eax, edx
    shl     eax, 2
    add     eax, edx
    add     eax, eax
    mov     edx, eax
    mov     eax, [ebp+input]
    mov     eax, [eax]
    lea     eax, [edx+eax]

```

```

        sub     eax, 30h
        mov     [ebp+var_10], eax
        add     [ebp+input], 4

loc_8048471:
        call    ___ctype_b_loc
        mov     edx, [eax]
        mov     eax, [ebp+input]
        mov     eax, [eax]
        add     eax, eax
        lea     eax, [edx+eax]
        movzx   eax, word ptr [eax]
        movzx   eax, ax
        and     eax, 800h
        test    eax, eax
        jnz     short loc_8048451
        cmp     [ebp+var_9], 2Dh
        jnz     short loc_804849A
        neg     [ebp+var_10]

loc_804849A:
        mov     eax, [ebp+var_10]
        leave
        retn

_f      endp

```

4.1.3 Task 1.3

This is standard C function too, actually, two functions working in pair. Source code taken from MSVC 2010 and modified slightly.

The matter of modification is that this function can work properly in multi-threaded environment, and I removed its support for simplification (or for confusion).

Compiled in MSVC 2010 with /Ox flag.

```

_BSS     SEGMENT
_v      DD     01H DUP (?)
_BSS     ENDS

_TEXT    SEGMENT
_s$ = 8                                     ; size = 4
f1      PROC
        push    ebp
        mov     ebp, esp
        mov     eax, DWORD PTR _s$[ebp]
        mov     DWORD PTR _v, eax
        pop     ebp
        ret     0
f1      ENDP
_TEXT    ENDS
PUBLIC   f2

_TEXT    SEGMENT
f2      PROC
        push    ebp
        mov     ebp, esp
        mov     eax, DWORD PTR _v
        imul    eax, 214013                ; 000343fdH
        add     eax, 2531011                ; 00269ec3H
        mov     DWORD PTR _v, eax
        mov     eax, DWORD PTR _v
        shr     eax, 16                    ; 00000010H
        and     eax, 32767                 ; 00007fffH
        pop     ebp
        ret     0

```



```
f2      ENDP
_TEXT   ENDS
END
```

Same code compiled in GCC 4.4.1:

```
f1      public f1
proc near

arg_0    = dword ptr 8

        push    ebp
        mov     ebp, esp
        mov     eax, [ebp+arg_0]
        mov     ds:v, eax
        pop     ebp
        retn
f1      endp

f2      public f2
proc near
        push    ebp
        mov     ebp, esp
        mov     eax, ds:v
        imul    eax, 343FDh
        add     eax, 269EC3h
        mov     ds:v, eax
        mov     eax, ds:v
        shr     eax, 10h
        and     eax, 7FFFh
        pop     ebp
        retn
f2      endp

bss      segment dword public 'BSS' use32
        assume cs:_bss
        dd ?
bss      ends
```

4.1.4 Task 1.4

This is standard C library function. Source code taken from MSVC 2010. Compiled in MSVC 2010 with /Ox flag.

```
PUBLIC    _f
_TEXT    SEGMENT
_arg1$ = 8          ; size = 4
_arg2$ = 12         ; size = 4
_f      PROC
        push    esi
        mov     esi, DWORD PTR _arg1$[esp]
        push    edi
        mov     edi, DWORD PTR _arg2$[esp+4]
        cmp     BYTE PTR [edi], 0
        mov     eax, esi
        je      SHORT $LN7@f
        mov     dl, BYTE PTR [esi]
        push    ebx
        test    dl, dl
        je      SHORT $LN4@f
        sub     esi, edi
        npad    6
$LL5@f:
        mov     ecx, edi
        test    dl, dl
        je      SHORT $LN2@f
```

```

$LL30f:
    mov     dl, BYTE PTR [ecx]
    test    dl, dl
    je      SHORT $LN140f
    movsx   ebx, BYTE PTR [esi+ecx]
    movsx   edx, dl
    sub     ebx, edx
    jne     SHORT $LN20f
    inc     ecx
    cmp     BYTE PTR [esi+ecx], bl
    jne     SHORT $LL30f
$LN20f:
    cmp     BYTE PTR [ecx], 0
    je      SHORT $LN140f
    mov     dl, BYTE PTR [eax+1]
    inc     eax
    inc     esi
    test    dl, dl
    jne     SHORT $LL50f
    xor     eax, eax
    pop     ebx
    pop     edi
    pop     esi
    ret     0
_f        ENDP
_TEXT     ENDS
END

```

Same code compiled in GCC 4.4.1:

```

f                public f
f                proc near

var_C            = dword ptr -0Ch
var_8            = dword ptr -8
var_4            = dword ptr -4
arg_0            = dword ptr 8
arg_4            = dword ptr 0Ch

                push    ebp
                mov     ebp, esp
                sub     esp, 10h
                mov     eax, [ebp+arg_0]
                mov     [ebp+var_4], eax
                mov     eax, [ebp+arg_4]
                movzx   eax, byte ptr [eax]
                test    al, al
                jnz     short loc_8048443
                mov     eax, [ebp+arg_0]
                jmp     short locret_8048453

loc_80483F4:
                mov     eax, [ebp+var_4]
                mov     [ebp+var_8], eax
                mov     eax, [ebp+arg_4]
                mov     [ebp+var_C], eax
                jmp     short loc_804840A

loc_8048402:
                add     [ebp+var_8], 1
                add     [ebp+var_C], 1

loc_804840A:
                mov     eax, [ebp+var_8]
                movzx   eax, byte ptr [eax]
                test    al, al
                jz      short loc_804842E

```

```

        mov     eax, [ebp+var_C]
        movzx   eax, byte ptr [eax]
        test    al, al
        jz      short loc_804842E
        mov     eax, [ebp+var_8]
        movzx   edx, byte ptr [eax]
        mov     eax, [ebp+var_C]
        movzx   eax, byte ptr [eax]
        cmp     dl, al
        jz      short loc_8048402

loc_804842E:
        mov     eax, [ebp+var_C]
        movzx   eax, byte ptr [eax]
        test    al, al
        jnz     short loc_804843D
        mov     eax, [ebp+var_4]
        jmp     short locret_8048453

loc_804843D:
        add     [ebp+var_4], 1
        jmp     short loc_8048444

loc_8048443:
        nop

loc_8048444:
        mov     eax, [ebp+var_4]
        movzx   eax, byte ptr [eax]
        test    al, al
        jnz     short loc_80483F4
        mov     eax, 0

locret_8048453:
        leave
        retn
f      endp

```

4.1.5 Task 1.5

This task is rather on knowledge than on reading code.

The function is taken from OpenWatcom. Compiled in MSVC 2010 with /Ox flag.

```

_DATA      SEGMENT
COMM       __v:DWORD
_DATA      ENDS
PUBLIC     __real@3e45798ee2308c3a
PUBLIC     __real@4147ffff80000000
PUBLIC     __real@4150017ec0000000
PUBLIC     _f
EXTRN      __fltused:DWORD
CONST      SEGMENT
__real@3e45798ee2308c3a DQ 03e45798ee2308c3ar      ; 1e-008
__real@4147ffff80000000 DQ 04147ffff80000000r      ; 3.14573e+006
__real@4150017ec0000000 DQ 04150017ec0000000r      ; 4.19584e+006
CONST      ENDS
_TEXT      SEGMENT
_v1$ = -16          ; size = 8
_v2$ = -8           ; size = 8
_f        PROC
        sub     esp, 16          ; 00000010H
        fld     QWORD PTR __real@4150017ec0000000
        fstp    QWORD PTR _v1$[esp+16]
        fld     QWORD PTR __real@4147ffff80000000

```

```

fstp    QWORD PTR _v2$[esp+16]
fld     QWORD PTR _v1$[esp+16]
fld     QWORD PTR _v1$[esp+16]
fdiv    QWORD PTR _v2$[esp+16]
fmul    QWORD PTR _v2$[esp+16]
fsubp   ST(1), ST(0)
fcomp   QWORD PTR __real@3e45798ee2308c3a
fnstsw  ax
test    ah, 65          ; 00000041H
jne     SHORT $LN1@f
or      DWORD PTR __v, 1
$LN1@f:
add     esp, 16         ; 00000010H
ret     0
_f      ENDP
_TEXT   ENDS

```

4.1.6 Task 1.6

Compiled in MSVC 2010 with /Ox option.

```

PUBLIC  _f
; Function compile flags: /Ogtpy
_TEXT   SEGMENT
_k0$ = -12          ; size = 4
_k3$ = -8           ; size = 4
_k2$ = -4           ; size = 4
_v$ = 8             ; size = 4
_k1$ = 12           ; size = 4
_k$ = 12            ; size = 4
_f      PROC
    sub     esp, 12    ; 0000000cH
    mov     ecx, DWORD PTR _v$[esp+8]
    mov     eax, DWORD PTR [ecx]
    mov     ecx, DWORD PTR [ecx+4]
    push    ebx
    push    esi
    mov     esi, DWORD PTR _k$[esp+16]
    push    edi
    mov     edi, DWORD PTR [esi]
    mov     DWORD PTR _k0$[esp+24], edi
    mov     edi, DWORD PTR [esi+4]
    mov     DWORD PTR _k1$[esp+20], edi
    mov     edi, DWORD PTR [esi+8]
    mov     esi, DWORD PTR [esi+12]
    xor     edx, edx
    mov     DWORD PTR _k2$[esp+24], edi
    mov     DWORD PTR _k3$[esp+24], esi
    lea     edi, DWORD PTR [edx+32]
$LL8@f:
    mov     esi, ecx
    shr     esi, 5
    add     esi, DWORD PTR _k1$[esp+20]
    mov     ebx, ecx
    shl     ebx, 4
    add     ebx, DWORD PTR _k0$[esp+24]
    sub     edx, 1640531527 ; 61c88647H
    xor     esi, ebx
    lea     ebx, DWORD PTR [edx+ecx]
    xor     esi, ebx
    add     eax, esi
    mov     esi, eax
    shr     esi, 5
    add     esi, DWORD PTR _k3$[esp+24]
    mov     ebx, eax

```

```

shl     ebx, 4
add     ebx, DWORD PTR _k2$[esp+24]
xor     esi, ebx
lea     ebx, DWORD PTR [edx+eax]
xor     esi, ebx
add     ecx, esi
dec     edi
jne     SHORT $LL80f
mov     edx, DWORD PTR _v$[esp+20]
pop     edi
pop     esi
mov     DWORD PTR [edx], eax
mov     DWORD PTR [edx+4], ecx
pop     ebx
add     esp, 12                ; 0000000cH
ret     0
_f      ENDP

```

4.1.7 Task 1.7

This function is taken from Linux 2.6 kernel.

Compiled in MSVC 2010 with /Ox option:

```

_table  db 000h, 080h, 040h, 0c0h, 020h, 0a0h, 060h, 0e0h
         db 010h, 090h, 050h, 0d0h, 030h, 0b0h, 070h, 0f0h
         db 008h, 088h, 048h, 0c8h, 028h, 0a8h, 068h, 0e8h
         db 018h, 098h, 058h, 0d8h, 038h, 0b8h, 078h, 0f8h
         db 004h, 084h, 044h, 0c4h, 024h, 0a4h, 064h, 0e4h
         db 014h, 094h, 054h, 0d4h, 034h, 0b4h, 074h, 0f4h
         db 00ch, 08ch, 04ch, 0cch, 02ch, 0ach, 06ch, 0ech
         db 01ch, 09ch, 05ch, 0dch, 03ch, 0bch, 07ch, 0fch
         db 002h, 082h, 042h, 0c2h, 022h, 0a2h, 062h, 0e2h
         db 012h, 092h, 052h, 0d2h, 032h, 0b2h, 072h, 0f2h
         db 00ah, 08ah, 04ah, 0cah, 02ah, 0aah, 06ah, 0eah
         db 01ah, 09ah, 05ah, 0dah, 03ah, 0bah, 07ah, 0fah
         db 006h, 086h, 046h, 0c6h, 026h, 0a6h, 066h, 0e6h
         db 016h, 096h, 056h, 0d6h, 036h, 0b6h, 076h, 0f6h
         db 00eh, 08eh, 04eh, 0ceh, 02eh, 0aeh, 06eh, 0eeh
         db 01eh, 09eh, 05eh, 0deh, 03eh, 0beh, 07eh, 0feh
         db 001h, 081h, 041h, 0c1h, 021h, 0a1h, 061h, 0e1h
         db 011h, 091h, 051h, 0d1h, 031h, 0b1h, 071h, 0f1h
         db 009h, 089h, 049h, 0c9h, 029h, 0a9h, 069h, 0e9h
         db 019h, 099h, 059h, 0d9h, 039h, 0b9h, 079h, 0f9h
         db 005h, 085h, 045h, 0c5h, 025h, 0a5h, 065h, 0e5h
         db 015h, 095h, 055h, 0d5h, 035h, 0b5h, 075h, 0f5h
         db 00dh, 08dh, 04dh, 0cdh, 02dh, 0adh, 06dh, 0edh
         db 01dh, 09dh, 05dh, 0ddh, 03dh, 0bdh, 07dh, 0fdh
         db 003h, 083h, 043h, 0c3h, 023h, 0a3h, 063h, 0e3h
         db 013h, 093h, 053h, 0d3h, 033h, 0b3h, 073h, 0f3h
         db 00bh, 08bh, 04bh, 0cbh, 02bh, 0abh, 06bh, 0ebh
         db 01bh, 09bh, 05bh, 0dbh, 03bh, 0bbh, 07bh, 0fbh
         db 007h, 087h, 047h, 0c7h, 027h, 0a7h, 067h, 0e7h
         db 017h, 097h, 057h, 0d7h, 037h, 0b7h, 077h, 0f7h
         db 00fh, 08fh, 04fh, 0cfh, 02fh, 0afh, 06fh, 0efh
         db 01fh, 09fh, 05fh, 0dfh, 03fh, 0bfh, 07fh, 0ffh

f        proc near
arg_0    = dword ptr 4

        mov     edx, [esp+arg_0]
        movzx   eax, dl
        movzx   eax, _table[eax]
        mov     ecx, edx
        shr     edx, 8

```

```

        movzx    edx, dl
        movzx    edx, _table[edx]
        shl     ax, 8
        movzx    eax, ax
        or      eax, edx
        shr     ecx, 10h
        movzx    edx, cl
        movzx    edx, _table[edx]
        shr     ecx, 8
        movzx    ecx, cl
        movzx    ecx, _table[ecx]
        shl     dx, 8
        movzx    edx, dx
        shl     eax, 10h
        or      edx, ecx
        or      eax, edx
        retn
f      endp

```

4.1.8 Task 1.8

Compiled in MSVC 2010 with /O1 option¹:

```

_a$ = 8      ; size = 4
_b$ = 12     ; size = 4
_c$ = 16     ; size = 4
?s@@YAXPAN00@Z PROC      ; s, COMDAT
    mov     eax, DWORD PTR _b$[esp-4]
    mov     ecx, DWORD PTR _a$[esp-4]
    mov     edx, DWORD PTR _c$[esp-4]
    push    esi
    push    edi
    sub     ecx, eax
    sub     edx, eax
    mov     edi, 200      ; 000000c8H
$LL6@s:
    push    100          ; 00000064H
    pop     esi
$LL3@s:
    fld     QWORD PTR [ecx+eax]
    fadd    QWORD PTR [eax]
    fstp    QWORD PTR [edx+eax]
    add     eax, 8
    dec     esi
    jne     SHORT $LL3@s
    dec     edi
    jne     SHORT $LL6@s
    pop     edi
    pop     esi
    ret     0
?s@@YAXPAN00@Z ENDP      ; s

```

4.1.9 Task 1.9

Compiled in MSVC 2010 with /O1 option:

```

tv315 = -8      ; size = 4
tv291 = -4      ; size = 4
_a$ = 8      ; size = 4
_b$ = 12     ; size = 4
_c$ = 16     ; size = 4
?m@@YAXPAN00@Z PROC ; m, COMDAT

```

¹/O1: minimize space

```

push    ebp
mov     ebp, esp
push    ecx
push    ecx
mov     edx, DWORD PTR _a$[ebp]
push    ebx
mov     ebx, DWORD PTR _c$[ebp]
push    esi
mov     esi, DWORD PTR _b$[ebp]
sub     edx, esi
push    edi
sub     esi, ebx
mov     DWORD PTR tv315[ebp], 100 ; 00000064H
$LL9@m:
mov     eax, ebx
mov     DWORD PTR tv291[ebp], 300 ; 0000012cH
$LL6@m:
fldz
lea     ecx, DWORD PTR [esi+eax]
fstp    QWORD PTR [eax]
mov     edi, 200 ; 000000c8H
$LL3@m:
dec     edi
fld     QWORD PTR [ecx+edx]
fmul    QWORD PTR [ecx]
fadd    QWORD PTR [eax]
fstp    QWORD PTR [eax]
jne     HORT $LL3@m
add     eax, 8
dec     DWORD PTR tv291[ebp]
jne     SHORT $LL6@m
add     ebx, 800 ; 00000320H
dec     DWORD PTR tv315[ebp]
jne     SHORT $LL9@m
pop     edi
pop     esi
pop     ebx
leave
ret     0
?m@@YAXPAN00@Z ENDP ; m

```

4.1.10 Task 1.10

If to compile this piece of code and run, some number will be printed. Where it came from? Where it came from if to compile it in MSVC with optimization (/Ox)?

```

#include <stdio.h>

int main()
{
    printf ("%d\n");

    return 0;
};

```

4.2 Middle level

4.2.1 Task 2.1

Well-known algorithm, also included in standard C library. Source code was taken from glibc 2.11.1. Compiled in GCC 4.4.1 with -Os option (code size optimization). Listing was done by IDA 4.9 disassembler from ELF-file generated by GCC and linker.

For those who wants use IDA while learning, here you may find .elf and .idb files, .idb can be opened with freeware IDA 4.9:

<http://conus.info/RE-tasks/middle/1/>

```
f                proc near

var_150          = dword ptr -150h
var_14C          = dword ptr -14Ch
var_13C          = dword ptr -13Ch
var_138          = dword ptr -138h
var_134          = dword ptr -134h
var_130          = dword ptr -130h
var_128          = dword ptr -128h
var_124          = dword ptr -124h
var_120          = dword ptr -120h
var_11C          = dword ptr -11Ch
var_118          = dword ptr -118h
var_114          = dword ptr -114h
var_110          = dword ptr -110h
var_C            = dword ptr -0Ch
arg_0            = dword ptr  8
arg_4            = dword ptr  0Ch
arg_8            = dword ptr  10h
arg_C            = dword ptr  14h
arg_10           = dword ptr  18h


                push    ebp
                mov     ebp, esp
                push    edi
                push    esi
                push    ebx
                sub     esp, 14Ch
                mov     ebx, [ebp+arg_8]
                cmp     [ebp+arg_4], 0
                jz      loc_804877D
                cmp     [ebp+arg_4], 4
                lea     eax, ds:0[ebx*4]
                mov     [ebp+var_130], eax
                jbe     loc_804864C
                mov     eax, [ebp+arg_4]
                mov     ecx, ebx
                mov     esi, [ebp+arg_0]
                lea     edx, [ebp+var_110]
                neg     ecx
                mov     [ebp+var_118], 0
                mov     [ebp+var_114], 0
                dec     eax
                imul    eax, ebx
                add     eax, [ebp+arg_0]
                mov     [ebp+var_11C], edx
                mov     [ebp+var_134], ecx
                mov     [ebp+var_124], eax
                lea     eax, [ebp+var_118]
                mov     [ebp+var_14C], eax
                mov     [ebp+var_120], ebx

loc_8048433:                                ; CODE XREF: f+28C
                mov     eax, [ebp+var_124]
                xor     edx, edx
                push    edi
                push    [ebp+arg_10]
                sub     eax, esi
                div     [ebp+var_120]
                push    esi
                shr     eax, 1
                imul    eax, [ebp+var_120]
```



```

    lea     edx, [esi+eax]
    push    edx
    mov     [ebp+var_138], edx
    call    [ebp+arg_C]
    add     esp, 10h
    mov     edx, [ebp+var_138]
    test    eax, eax
    jns     short loc_8048482
    xor     eax, eax

loc_804846D:                                ; CODE XREF: f+CC
    mov     cl, [edx+eax]
    mov     bl, [esi+eax]
    mov     [edx+eax], bl
    mov     [esi+eax], cl
    inc     eax
    cmp     [ebp+var_120], eax
    jnz     short loc_804846D

loc_8048482:                                ; CODE XREF: f+B5
    push    ebx
    push    [ebp+arg_10]
    mov     [ebp+var_138], edx
    push    edx
    push    [ebp+var_124]
    call    [ebp+arg_C]
    mov     edx, [ebp+var_138]
    add     esp, 10h
    test    eax, eax
    jns     short loc_80484F6
    mov     ecx, [ebp+var_124]
    xor     eax, eax

loc_80484AB:                                ; CODE XREF: f+10D
    movzx   edi, byte ptr [edx+eax]
    mov     bl, [ecx+eax]
    mov     [edx+eax], bl
    mov     ebx, edi
    mov     [ecx+eax], bl
    inc     eax
    cmp     [ebp+var_120], eax
    jnz     short loc_80484AB
    push    ecx
    push    [ebp+arg_10]
    mov     [ebp+var_138], edx
    push    esi
    push    edx
    call    [ebp+arg_C]
    add     esp, 10h
    mov     edx, [ebp+var_138]
    test    eax, eax
    jns     short loc_80484F6
    xor     eax, eax

loc_80484E1:                                ; CODE XREF: f+140
    mov     cl, [edx+eax]
    mov     bl, [esi+eax]
    mov     [edx+eax], bl
    mov     [esi+eax], cl
    inc     eax
    cmp     [ebp+var_120], eax
    jnz     short loc_80484E1

loc_80484F6:                                ; CODE XREF: f+ED
                                           ; f+129
    mov     eax, [ebp+var_120]

```

```

        mov     edi, [ebp+var_124]
        add     edi, [ebp+var_134]
        lea     ebx, [esi+eax]
        jmp     short loc_8048513
; -----
loc_804850D:                ; CODE XREF: f+17B
        add     ebx, [ebp+var_120]

loc_8048513:                ; CODE XREF: f+157
                                ; f+1F9
        push    eax
        push    [ebp+arg_10]
        mov     [ebp+var_138], edx
        push    edx
        push    ebx
        call    [ebp+arg_C]
        add     esp, 10h
        mov     edx, [ebp+var_138]
        test    eax, eax
        jns     short loc_8048537
        jmp     short loc_804850D
; -----
loc_8048531:                ; CODE XREF: f+19D
        add     edi, [ebp+var_134]

loc_8048537:                ; CODE XREF: f+179
        push    ecx
        push    [ebp+arg_10]
        mov     [ebp+var_138], edx
        push    edi
        push    edx
        call    [ebp+arg_C]
        add     esp, 10h
        mov     edx, [ebp+var_138]
        test    eax, eax
        js      short loc_8048531
        cmp     ebx, edi
        jnb     short loc_8048596
        xor     eax, eax
        mov     [ebp+var_128], edx

loc_804855F:                ; CODE XREF: f+1BE
        mov     cl, [ebx+eax]
        mov     dl, [edi+eax]
        mov     [ebx+eax], dl
        mov     [edi+eax], cl
        inc     eax
        cmp     [ebp+var_120], eax
        jnz     short loc_804855F
        mov     edx, [ebp+var_128]
        cmp     edx, ebx
        jnz     short loc_8048582
        mov     edx, edi
        jmp     short loc_8048588
; -----
loc_8048582:                ; CODE XREF: f+1C8
        cmp     edx, edi
        jnz     short loc_8048588
        mov     edx, ebx

loc_8048588:                ; CODE XREF: f+1CC
                                ; f+1D0
        add     ebx, [ebp+var_120]

```

```

        add     edi, [ebp+var_134]
        jmp     short loc_80485AB
; -----

loc_8048596:                ; CODE XREF: f+1A1
        jnz     short loc_80485AB
        mov     ecx, [ebp+var_134]
        mov     eax, [ebp+var_120]
        lea     edi, [ebx+ecx]
        add     ebx, eax
        jmp     short loc_80485B3
; -----

loc_80485AB:                ; CODE XREF: f+1E0
                                ; f:loc_8048596
        cmp     ebx, edi
        jbe     loc_8048513

loc_80485B3:                ; CODE XREF: f+1F5
        mov     eax, edi
        sub     eax, esi
        cmp     eax, [ebp+var_130]
        ja      short loc_80485EB
        mov     eax, [ebp+var_124]
        mov     esi, ebx
        sub     eax, ebx
        cmp     eax, [ebp+var_130]
        ja      short loc_8048634
        sub     [ebp+var_11C], 8
        mov     edx, [ebp+var_11C]
        mov     ecx, [edx+4]
        mov     esi, [edx]
        mov     [ebp+var_124], ecx
        jmp     short loc_8048634
; -----

loc_80485EB:                ; CODE XREF: f+209
        mov     edx, [ebp+var_124]
        sub     edx, ebx
        cmp     edx, [ebp+var_130]
        jbe     short loc_804862E
        cmp     eax, edx
        mov     edx, [ebp+var_11C]
        lea     eax, [edx+8]
        jle     short loc_8048617
        mov     [edx], esi
        mov     esi, ebx
        mov     [edx+4], edi
        mov     [ebp+var_11C], eax
        jmp     short loc_8048634
; -----

loc_8048617:                ; CODE XREF: f+252
        mov     ecx, [ebp+var_11C]
        mov     [ebp+var_11C], eax
        mov     [ecx], ebx
        mov     ebx, [ebp+var_124]
        mov     [ecx+4], ebx

loc_804862E:                ; CODE XREF: f+245
        mov     [ebp+var_124], edi

loc_8048634:                ; CODE XREF: f+21B
                                ; f+235 ...
        mov     eax, [ebp+var_14C]
        cmp     [ebp+var_11C], eax

```

```

        ja      loc_8048433
        mov     ebx, [ebp+var_120]

loc_804864C:                                ; CODE XREF: f+2A
        mov     eax, [ebp+arg_4]
        mov     ecx, [ebp+arg_0]
        add     ecx, [ebp+var_130]
        dec     eax
        imul    eax, ebx
        add     eax, [ebp+arg_0]
        cmp     ecx, eax
        mov     [ebp+var_120], eax
        jbe     short loc_804866B
        mov     ecx, eax

loc_804866B:                                ; CODE XREF: f+2B3
        mov     esi, [ebp+arg_0]
        mov     edi, [ebp+arg_0]
        add     esi, ebx
        mov     edx, esi
        jmp     short loc_80486A3
; -----

loc_8048677:                                ; CODE XREF: f+2F1
        push    eax
        push    [ebp+arg_10]
        mov     [ebp+var_138], edx
        mov     [ebp+var_13C], ecx
        push    edi
        push    edx
        call    [ebp+arg_C]
        add     esp, 10h
        mov     edx, [ebp+var_138]
        mov     ecx, [ebp+var_13C]
        test    eax, eax
        jns     short loc_80486A1
        mov     edi, edx

loc_80486A1:                                ; CODE XREF: f+2E9
        add     edx, ebx

loc_80486A3:                                ; CODE XREF: f+2C1
        cmp     edx, ecx
        jbe     short loc_8048677
        cmp     edi, [ebp+arg_0]
        jz      loc_8048762
        xor     eax, eax

loc_80486B2:                                ; CODE XREF: f+313
        mov     ecx, [ebp+arg_0]
        mov     dl, [edi+eax]
        mov     cl, [ecx+eax]
        mov     [edi+eax], cl
        mov     ecx, [ebp+arg_0]
        mov     [ecx+eax], dl
        inc     eax
        cmp     ebx, eax
        jnz     short loc_80486B2
        jmp     loc_8048762
; -----

loc_80486CE:                                ; CODE XREF: f+3C3
        lea     edx, [esi+edi]
        jmp     short loc_80486D5
; -----

```

```

loc_80486D3:                                ; CODE XREF: f+33B
        add     edx, edi

loc_80486D5:                                ; CODE XREF: f+31D
        push    eax
        push    [ebp+arg_10]
        mov     [ebp+var_138], edx
        push    edx
        push    esi
        call    [ebp+arg_C]
        add     esp, 10h
        mov     edx, [ebp+var_138]
        test    eax, eax
        js      short loc_80486D3
        add     edx, ebx
        cmp     edx, esi
        mov     [ebp+var_124], edx
        jz      short loc_804876F
        mov     edx, [ebp+var_134]
        lea     eax, [esi+ebx]
        add     edx, eax
        mov     [ebp+var_11C], edx
        jmp     short loc_804875B
; -----

loc_8048710:                                ; CODE XREF: f+3AA
        mov     cl, [eax]
        mov     edx, [ebp+var_11C]
        mov     [ebp+var_150], eax
        mov     byte ptr [ebp+var_130], cl
        mov     ecx, eax
        jmp     short loc_8048733
; -----

loc_8048728:                                ; CODE XREF: f+391
        mov     al, [edx+ebx]
        mov     [ecx], al
        mov     ecx, [ebp+var_128]

loc_8048733:                                ; CODE XREF: f+372
        mov     [ebp+var_128], edx
        add     edx, edi
        mov     eax, edx
        sub     eax, edi
        cmp     [ebp+var_124], eax
        jbe     short loc_8048728
        mov     dl, byte ptr [ebp+var_130]
        mov     eax, [ebp+var_150]
        mov     [ecx], dl
        dec     [ebp+var_11C]

loc_804875B:                                ; CODE XREF: f+35A
        dec     eax
        cmp     eax, esi
        jnb     short loc_8048710
        jmp     short loc_804876F
; -----

loc_8048762:                                ; CODE XREF: f+2F6
                                                ; f+315
        mov     edi, ebx
        neg     edi
        lea     ecx, [edi-1]
        mov     [ebp+var_134], ecx

loc_804876F:                                ; CODE XREF: f+347

```

```

                                ; f+3AC
    add     esi, ebx
    cmp     esi, [ebp+var_120]
    jbe     loc_80486CE

loc_804877D:                    ; CODE XREF: f+13
    lea     esp, [ebp-0Ch]
    pop     ebx
    pop     esi
    pop     edi
    pop     ebp
    retn
f      endp

```

4.3 crackme / keygenme

Couple of my keygenmes²:

<http://crackmes.de/users/yonkie/>

²program which imitates fictional software protection, for which one need to make a keys/licenses generator

Chapter 5

Tools

- IDA as disassembler. Older freeware version is available for downloading: <http://www.hex-rays.com/idapro/idadownloadfreeware.htm>.
- Microsoft Visual Studio Express¹: Stripped-down Visual Studio version, convenient for simple experiments.
- Hiew² for small modifications of code in binary files.

5.0.1 Debugger

*tracer*³ instead of debugger.

I stopped to use debugger eventually, because all I need from it is to spot some function's arguments while execution, or registers' state at some point. To load debugger each time is too much, so I wrote a small utility *tracer*. It has console-interface, working from command-line, enable us to intercept function execution, set breakpoints at arbitrary places, spot registers' state, modify it, etc.

However, as for learning, it's highly advisable to trace code in debugger manually, watch how register's state changing (for example, classic SoftICE, OllyDbg, WinDbg highlighting changed registers), flags, data, change them manually, watch reaction, etc.

¹<http://www.microsoft.com/express/Downloads/>

²<http://www.hiew.ru/>

³<http://conus.info/gt/>

Chapter 6

Books/blogs worth reading

6.1 Books

6.1.1 Windows

- Windows® Internals (Mark E. Russinovich and David A. Solomon with Alex Ionescu)¹

6.1.2 C/C++

- C++ language standard: ISO/IEC 14882:2003²

6.1.3 x86 / x86-64

- Intel manuals: <http://www.intel.com/products/processor/manuals/>
- AMD manuals: <http://developer.amd.com/documentation/guides/Pages/default.aspx#manuals>

6.1.4 ARM

ARM manuals: <http://infocenter.arm.com/help/index.jsp?topic=/com.arm.doc.subset.architecture.reference/index.html>

6.2 Blogs

6.2.1 Windows

- Microsoft: Raymond Chen
- <http://www.nynaeve.net/>

¹<http://www.microsoft.com/learning/en/us/book.aspx?ID=12069&locale=en-us>

²http://www.iso.org/iso/catalogue_detail.htm?csnumber=38110

Chapter 7

More examples

7.1 “QR9”: Rubik’s cube inspired amateur crypto-algorithm

Sometimes amateur cryptosystems appear to be pretty bizarre.

I was asked to reverse engineer an amateur cryptoalgorithm of some data crypting utility, source code of which was lost¹.

Here is also IDA 5 exported listing from original crypting utility:

```
.text:00541000 set_bit      proc near                ; CODE XREF: rotate1+42
.text:00541000                                     ; rotate2+42 ...
.text:00541000
.text:00541000 arg_0      = dword ptr  4
.text:00541000 arg_4      = dword ptr  8
.text:00541000 arg_8      = dword ptr  0Ch
.text:00541000 arg_C      = byte ptr  10h
.text:00541000
.text:00541000          mov     al, [esp+arg_C]
.text:00541004          mov     ecx, [esp+arg_8]
.text:00541008          push    esi
.text:00541009          mov     esi, [esp+4+arg_0]
.text:0054100D          test    al, al
.text:0054100F          mov     eax, [esp+4+arg_4]
.text:00541013          mov     dl, 1
.text:00541015          jz      short loc_54102B
.text:00541017          shl     dl, cl
.text:00541019          mov     cl, cube64[eax+esi*8]
.text:00541020          or      cl, dl
.text:00541022          mov     cube64[eax+esi*8], cl
.text:00541029          pop     esi
.text:0054102A          retn
.text:0054102B ; -----
.text:0054102B
.text:0054102B loc_54102B:                ; CODE XREF: set_bit+15
.text:0054102B          shl     dl, cl
.text:0054102D          mov     cl, cube64[eax+esi*8]
.text:00541034          not     dl
.text:00541036          and     cl, dl
.text:00541038          mov     cube64[eax+esi*8], cl
.text:0054103F          pop     esi
.text:00541040          retn
.text:00541040 set_bit      endp
.text:00541040 ; -----
.text:00541041          align 10h
.text:00541050
.text:00541050 ; ===== S U B R O U T I N E =====
.text:00541050
.text:00541050
```

¹I also got permit from customer to publish the algorithm details

```

.text:00541050 get_bit      proc near                ; CODE XREF: rotate1+16
.text:00541050                                     ; rotate2+16 ...
.text:00541050
.text:00541050 arg_0      = dword ptr  4
.text:00541050 arg_4      = dword ptr  8
.text:00541050 arg_8      = byte ptr  0Ch
.text:00541050
.text:00541050          mov     eax, [esp+arg_4]
.text:00541054          mov     ecx, [esp+arg_0]
.text:00541058          mov     al, cube64[eax+ecx*8]
.text:0054105F          mov     cl, [esp+arg_8]
.text:00541063          shr     al, cl
.text:00541065          and     al, 1
.text:00541067          retn
.text:00541067 get_bit      endp
.text:00541067 ; -----
.text:00541068          align 10h
.text:00541070
.text:00541070 ; ===== S U B R O U T I N E =====
.text:00541070
.text:00541070 rotate1      proc near                ; CODE XREF: rotate_all_with_password+8E
.text:00541070
.text:00541070 internal_array_64= byte ptr -40h
.text:00541070 arg_0      = dword ptr  4
.text:00541070
.text:00541070          sub     esp, 40h
.text:00541073          push    ebx
.text:00541074          push    ebp
.text:00541075          mov     ebp, [esp+48h+arg_0]
.text:00541079          push    esi
.text:0054107A          push    edi
.text:0054107B          xor     edi, edi        ; EDI is loop1 counter
.text:0054107D          lea     ebx, [esp+50h+internal_array_64]
.text:00541081
.text:00541081 first_loop1_begin:                ; CODE XREF: rotate1+2E
.text:00541081          xor     esi, esi        ; ESI is loop2 counter
.text:00541083
.text:00541083 first_loop2_begin:                ; CODE XREF: rotate1+25
.text:00541083          push    ebp                ; arg_0
.text:00541084          push    esi
.text:00541085          push    edi
.text:00541086          call   get_bit
.text:0054108B          add     esp, 0Ch
.text:0054108E          mov     [ebx+esi], al    ; store to internal array
.text:00541091          inc     esi
.text:00541092          cmp     esi, 8
.text:00541095          jnl     short first_loop2_begin
.text:00541097          inc     edi
.text:00541098          add     ebx, 8
.text:0054109B          cmp     edi, 8
.text:0054109E          jnl     short first_loop1_begin
.text:005410A0          lea     ebx, [esp+50h+internal_array_64]
.text:005410A4          mov     edi, 7        ; EDI is loop1 counter, initial state is 7
.text:005410A9
.text:005410A9 second_loop1_begin:                ; CODE XREF: rotate1+57
.text:005410A9          xor     esi, esi        ; ESI is loop2 counter
.text:005410AB
.text:005410AB second_loop2_begin:                ; CODE XREF: rotate1+4E
.text:005410AB          mov     al, [ebx+esi]    ; value from internal array
.text:005410AE          push    eax
.text:005410AF          push    ebp                ; arg_0
.text:005410B0          push    edi
.text:005410B1          push    esi
.text:005410B2          call   set_bit

```

```

.text:005410B7      add     esp, 10h
.text:005410BA      inc     esi             ; increment loop2 counter
.text:005410BB      cmp     esi, 8
.text:005410BE      jnl    short second_loop2_begin
.text:005410C0      dec     edi             ; decrement loop2 counter
.text:005410C1      add     ebx, 8
.text:005410C4      cmp     edi, 0FFFFFFFh
.text:005410C7      jg     short second_loop1_begin
.text:005410C9      pop     edi
.text:005410CA      pop     esi
.text:005410CB      pop     ebp
.text:005410CC      pop     ebx
.text:005410CD      add     esp, 40h
.text:005410D0      retn
.text:005410D0 rotate1      endp
.text:005410D0
.text:005410D0 ; -----
.text:005410D1      align 10h
.text:005410E0
.text:005410E0 ; ===== S U B R O U T I N E =====
.text:005410E0
.text:005410E0
.text:005410E0 rotate2      proc near             ; CODE XREF: rotate_all_with_password+7A
.text:005410E0
.text:005410E0 internal_array_64= byte ptr -40h
.text:005410E0 arg_0          = dword ptr 4
.text:005410E0
.text:005410E0      sub     esp, 40h
.text:005410E3      push    ebx
.text:005410E4      push    ebp
.text:005410E5      mov     ebp, [esp+48h+arg_0]
.text:005410E9      push    esi
.text:005410EA      push    edi
.text:005410EB      xor     edi, edi             ; loop1 counter
.text:005410ED      lea     ebx, [esp+50h+internal_array_64]
.text:005410F1
.text:005410F1 loc_5410F1:             ; CODE XREF: rotate2+2E
.text:005410F1      xor     esi, esi             ; loop2 counter
.text:005410F3
.text:005410F3 loc_5410F3:             ; CODE XREF: rotate2+25
.text:005410F3      push    esi             ; loop2
.text:005410F4      push    edi             ; loop1
.text:005410F5      push    ebp             ; arg_0
.text:005410F6      call    get_bit
.text:005410FB      add     esp, 0Ch
.text:005410FE      mov     [ebx+esi], al      ; store to internal array
.text:00541101      inc     esi             ; increment loop1 counter
.text:00541102      cmp     esi, 8
.text:00541105      jnl    short loc_5410F3
.text:00541107      inc     edi             ; increment loop2 counter
.text:00541108      add     ebx, 8
.text:0054110B      cmp     edi, 8
.text:0054110E      jnl    short loc_5410F1
.text:00541110      lea     ebx, [esp+50h+internal_array_64]
.text:00541114      mov     edi, 7             ; loop1 counter is initial state 7
.text:00541119
.text:00541119 loc_541119:             ; CODE XREF: rotate2+57
.text:00541119      xor     esi, esi             ; loop2 counter
.text:0054111B
.text:0054111B loc_54111B:             ; CODE XREF: rotate2+4E
.text:0054111B      mov     al, [ebx+esi]      ; get byte from internal array
.text:0054111E      push    eax
.text:0054111F      push    edi             ; loop1 counter
.text:00541120      push    esi             ; loop2 counter
.text:00541121      push    ebp             ; arg_0
.text:00541122      call    set_bit

```

```

.text:00541127      add     esp, 10h
.text:0054112A      inc     esi             ; increment loop2 counter
.text:0054112B      cmp     esi, 8
.text:0054112E      jl      short loc_54111B
.text:00541130      dec     edi             ; decrement loop2 counter
.text:00541131      add     ebx, 8
.text:00541134      cmp     edi, 0FFFFFFFh
.text:00541137      jg      short loc_541119
.text:00541139      pop     edi
.text:0054113A      pop     esi
.text:0054113B      pop     ebp
.text:0054113C      pop     ebx
.text:0054113D      add     esp, 40h
.text:00541140      retn
.text:00541140 rotate2      endp
.text:00541140 ; -----
.text:00541141      align 10h
.text:00541150 ; ===== S U B R O U T I N E =====
.text:00541150 rotate3      proc near             ; CODE XREF: rotate_all_with_password+66
.text:00541150      var_40      = byte ptr -40h
.text:00541150      arg_0       = dword ptr 4
.text:00541150      sub     esp, 40h
.text:00541153      push    ebx
.text:00541154      push    ebp
.text:00541155      mov     ebp, [esp+48h+arg_0]
.text:00541159      push    esi
.text:0054115A      push    edi
.text:0054115B      xor     edi, edi
.text:0054115D      lea     ebx, [esp+50h+var_40]
.text:00541161      loc_541161:             ; CODE XREF: rotate3+2E
.text:00541161      xor     esi, esi
.text:00541163      loc_541163:             ; CODE XREF: rotate3+25
.text:00541163      push    esi
.text:00541164      push    ebp
.text:00541165      push    edi
.text:00541166      call   get_bit
.text:0054116B      add     esp, 0Ch
.text:0054116E      mov     [ebx+esi], al
.text:00541171      inc     esi
.text:00541172      cmp     esi, 8
.text:00541175      jl      short loc_541163
.text:00541177      inc     edi
.text:00541178      add     ebx, 8
.text:0054117B      cmp     edi, 8
.text:0054117E      jl      short loc_541161
.text:00541180      xor     ebx, ebx
.text:00541182      lea     edi, [esp+50h+var_40]
.text:00541186      loc_541186:             ; CODE XREF: rotate3+54
.text:00541186      mov     esi, 7
.text:0054118B      loc_54118B:             ; CODE XREF: rotate3+4E
.text:0054118B      mov     al, [edi]
.text:0054118D      push    eax
.text:0054118E      push    ebx
.text:0054118F      push    ebp
.text:00541190      push    esi
.text:00541191      call   set_bit

```

```

.text:00541196      add     esp, 10h
.text:00541199      inc     edi
.text:0054119A      dec     esi
.text:0054119B      cmp     esi, 0FFFFFFFh
.text:0054119E      jg      short loc_54118B
.text:005411A0      inc     ebx
.text:005411A1      cmp     ebx, 8
.text:005411A4      jl      short loc_541186
.text:005411A6      pop     edi
.text:005411A7      pop     esi
.text:005411A8      pop     ebp
.text:005411A9      pop     ebx
.text:005411AA      add     esp, 40h
.text:005411AD      retn
.text:005411AD rotate3      endp
.text:005411AD
.text:005411AD ; -----
.text:005411AE      align 10h
.text:005411B0
.text:005411B0 ; ===== S U B R O U T I N E =====
.text:005411B0
.text:005411B0 rotate_all_with_password proc near      ; CODE XREF: crypt+1F
.text:005411B0                                     ; decrypt+36
.text:005411B0
.text:005411B0 arg_0      = dword ptr 4
.text:005411B0 arg_4      = dword ptr 8
.text:005411B0
.text:005411B0      mov     eax, [esp+arg_0]
.text:005411B4      push    ebp
.text:005411B5      mov     ebp, eax
.text:005411B7      cmp     byte ptr [eax], 0
.text:005411BA      jz      exit
.text:005411C0      push    ebx
.text:005411C1      mov     ebx, [esp+8+arg_4]
.text:005411C5      push    esi
.text:005411C6      push    edi
.text:005411C7
.text:005411C7 loop_begin:      ; CODE XREF: rotate_all_with_password+9F
.text:005411C7      movsx   eax, byte ptr [ebp+0]
.text:005411CB      push    eax      ; C
.text:005411CC      call    _tolower
.text:005411D1      add     esp, 4
.text:005411D4      cmp     al, 'a'
.text:005411D6      jl      short next_character_in_password
.text:005411D8      cmp     al, 'z'
.text:005411DA      jg      short next_character_in_password
.text:005411DC      movsx   ecx, al
.text:005411DF      sub     ecx, 'a'
.text:005411E2      cmp     ecx, 24
.text:005411E5      jle     short skip_subtracting
.text:005411E7      sub     ecx, 24
.text:005411EA
.text:005411EA skip_subtracting:      ; CODE XREF: rotate_all_with_password+35
.text:005411EA      mov     eax, 55555556h
.text:005411EF      imul    ecx
.text:005411F1      mov     eax, edx
.text:005411F3      shr     eax, 1Fh
.text:005411F6      add     edx, eax
.text:005411F8      mov     eax, ecx
.text:005411FA      mov     esi, edx
.text:005411FC      mov     ecx, 3
.text:00541201      cdq
.text:00541202      idiv    ecx
.text:00541204      sub     edx, 0
.text:00541207      jz      short call_rotate1

```

```

.text:00541209      dec     edx
.text:0054120A      jz      short call_rotate2
.text:0054120C      dec     edx
.text:0054120D      jnz     short next_character_in_password
.text:0054120F      test    ebx, ebx
.text:00541211      jle     short next_character_in_password
.text:00541213      mov     edi, ebx
.text:00541215
.text:00541215 call_rotate3:                                ; CODE XREF: rotate_all_with_password+6F
.text:00541215      push    esi
.text:00541216      call    rotate3
.text:0054121B      add     esp, 4
.text:0054121E      dec     edi
.text:0054121F      jnz     short call_rotate3
.text:00541221      jmp     short next_character_in_password
.text:00541223 ; -----
.text:00541223
.text:00541223 call_rotate2:                                ; CODE XREF: rotate_all_with_password+5A
.text:00541223      test    ebx, ebx
.text:00541225      jle     short next_character_in_password
.text:00541227      mov     edi, ebx
.text:00541229
.text:00541229 loc_541229:                                ; CODE XREF: rotate_all_with_password+83
.text:00541229      push    esi
.text:0054122A      call    rotate2
.text:0054122F      add     esp, 4
.text:00541232      dec     edi
.text:00541233      jnz     short loc_541229
.text:00541235      jmp     short next_character_in_password
.text:00541237 ; -----
.text:00541237
.text:00541237 call_rotate1:                                ; CODE XREF: rotate_all_with_password+57
.text:00541237      test    ebx, ebx
.text:00541239      jle     short next_character_in_password
.text:0054123B      mov     edi, ebx
.text:0054123D
.text:0054123D loc_54123D:                                ; CODE XREF: rotate_all_with_password+97
.text:0054123D      push    esi
.text:0054123E      call    rotate1
.text:00541243      add     esp, 4
.text:00541246      dec     edi
.text:00541247      jnz     short loc_54123D
.text:00541249
.text:00541249 next_character_in_password:                ; CODE XREF: rotate_all_with_password+26
.text:00541249                                         ; rotate_all_with_password+2A ...
.text:00541249      mov     al, [ebp+1]
.text:0054124C      inc     ebp
.text:0054124D      test    al, al
.text:0054124F      jnz     loop_begin
.text:00541255      pop     edi
.text:00541256      pop     esi
.text:00541257      pop     ebx
.text:00541258
.text:00541258 exit:                                ; CODE XREF: rotate_all_with_password+A
.text:00541258      pop     ebp
.text:00541259      retn
.text:00541259 rotate_all_with_password endp
.text:00541259
.text:00541259 ; -----
.text:0054125A      align 10h
.text:00541260
.text:00541260 ; ===== S U B R O U T I N E =====
.text:00541260
.text:00541260
.text:00541260 crypt      proc near                                ; CODE XREF: crypt_file+8A

```

```

.text:00541260 arg_0      = dword ptr 4
.text:00541260 arg_4      = dword ptr 8
.text:00541260 arg_8      = dword ptr 0Ch
.text:00541260
.text:00541260          push    ebx
.text:00541261          mov     ebx, [esp+4+arg_0]
.text:00541265          push    ebp
.text:00541266          push    esi
.text:00541267          push    edi
.text:00541268          xor     ebp, ebp
.text:0054126A
.text:0054126A loc_54126A:                                ; CODE XREF: crypt+41
.text:0054126A          mov     eax, [esp+10h+arg_8]
.text:0054126E          mov     ecx, 10h
.text:00541273          mov     esi, ebx
.text:00541275          mov     edi, offset cube64
.text:0054127A          push    1
.text:0054127C          push    eax
.text:0054127D          rep movsd
.text:0054127F          call   rotate_all_with_password
.text:00541284          mov     eax, [esp+18h+arg_4]
.text:00541288          mov     edi, ebx
.text:0054128A          add     ebp, 40h
.text:0054128D          add     esp, 8
.text:00541290          mov     ecx, 10h
.text:00541295          mov     esi, offset cube64
.text:0054129A          add     ebx, 40h
.text:0054129D          cmp     ebp, eax
.text:0054129F          rep movsd
.text:005412A1          jl      short loc_54126A
.text:005412A3          pop     edi
.text:005412A4          pop     esi
.text:005412A5          pop     ebp
.text:005412A6          pop     ebx
.text:005412A7          retn
.text:005412A7 crypt      endp
.text:005412A7
.text:005412A7 ; -----
.text:005412A8          align 10h
.text:005412B0
.text:005412B0 ; ===== S U B R O U T I N E =====
.text:005412B0
.text:005412B0
.text:005412B0 ; int __cdecl decrypt(int, int, void *Src)
.text:005412B0 decrypt      proc near                                ; CODE XREF: decrypt_file+99
.text:005412B0
.text:005412B0 arg_0      = dword ptr 4
.text:005412B0 arg_4      = dword ptr 8
.text:005412B0 Src        = dword ptr 0Ch
.text:005412B0
.text:005412B0          mov     eax, [esp+Src]
.text:005412B4          push    ebx
.text:005412B5          push    ebp
.text:005412B6          push    esi
.text:005412B7          push    edi
.text:005412B8          push    eax                ; Src
.text:005412B9          call   __strdup
.text:005412BE          push    eax                ; Str
.text:005412BF          mov     [esp+18h+Src], eax
.text:005412C3          call   __strrev
.text:005412C8          mov     ebx, [esp+18h+arg_0]
.text:005412CC          add     esp, 8
.text:005412CF          xor     ebp, ebp
.text:005412D1
.text:005412D1 loc_5412D1:                                ; CODE XREF: decrypt+58
.text:005412D1          mov     ecx, 10h

```

```

.text:005412D6      mov     esi, ebx
.text:005412D8      mov     edi, offset cube64
.text:005412DD      push    3
.text:005412DF      rep movsd
.text:005412E1      mov     ecx, [esp+14h+Src]
.text:005412E5      push    ecx
.text:005412E6      call   rotate_all_with_password
.text:005412EB      mov     eax, [esp+18h+arg_4]
.text:005412EF      mov     edi, ebx
.text:005412F1      add     ebp, 40h
.text:005412F4      add     esp, 8
.text:005412F7      mov     ecx, 10h
.text:005412FC      mov     esi, offset cube64
.text:00541301      add     ebx, 40h
.text:00541304      cmp     ebp, eax
.text:00541306      rep movsd
.text:00541308      jl      short loc_5412D1
.text:0054130A      mov     edx, [esp+10h+Src]
.text:0054130E      push    edx                ; Memory
.text:0054130F      call   _free
.text:00541314      add     esp, 4
.text:00541317      pop     edi
.text:00541318      pop     esi
.text:00541319      pop     ebp
.text:0054131A      pop     ebx
.text:0054131B      retn
.text:0054131B decrypt      endp
.text:0054131B
.text:0054131B ; -----
.text:0054131C      align 10h
.text:00541320
.text:00541320 ; ===== S U B R O U T I N E =====
.text:00541320
.text:00541320 ; int __cdecl crypt_file(int Str, char *Filename, int password)
.text:00541320 crypt_file      proc near                ; CODE XREF: _main+42
.text:00541320
.text:00541320 Str                = dword ptr 4
.text:00541320 Filename          = dword ptr 8
.text:00541320 password          = dword ptr 0Ch
.text:00541320
.text:00541320      mov     eax, [esp+Str]
.text:00541324      push    ebp
.text:00541325      push    offset Mode        ; "rb"
.text:0054132A      push    eax                ; Filename
.text:0054132B      call   _fopen              ; open file
.text:00541330      mov     ebp, eax
.text:00541332      add     esp, 8
.text:00541335      test    ebp, ebp
.text:00541337      jnz     short loc_541348
.text:00541339      push    offset Format        ; "Cannot open input file!\n"
.text:0054133E      call   _printf
.text:00541343      add     esp, 4
.text:00541346      pop     ebp
.text:00541347      retn
.text:00541348 ; -----
.text:00541348
.text:00541348 loc_541348:                ; CODE XREF: crypt_file+17
.text:00541348      push    ebx
.text:00541349      push    esi
.text:0054134A      push    edi
.text:0054134B      push    2                  ; Origin
.text:0054134D      push    0                  ; Offset
.text:0054134F      push    ebp                ; File
.text:00541350      call   _fseek
.text:00541355      push    ebp                ; File

```



```

.text:00541356      call     _ftell          ; get file size
.text:0054135B      push     0                ; Origin
.text:0054135D      push     0                ; Offset
.text:0054135F      push     ebp              ; File
.text:00541360      mov     [esp+2Ch+Str], eax
.text:00541364      call     _fseek          ; rewind to start
.text:00541369      mov     esi, [esp+2Ch+Str]
.text:0054136D      and     esi, 0FFFFFFC0h ; reset all lowest 6 bits
.text:00541370      add     esi, 40h         ; align size to 64-byte border
.text:00541373      push     esi             ; Size
.text:00541374      call     _malloc
.text:00541379      mov     ecx, esi
.text:0054137B      mov     ebx, eax         ; allocated buffer pointer -> to EBX
.text:0054137D      mov     edx, ecx
.text:0054137F      xor     eax, eax
.text:00541381      mov     edi, ebx
.text:00541383      push     ebp             ; File
.text:00541384      shr     ecx, 2
.text:00541387      rep stosd
.text:00541389      mov     ecx, edx
.text:0054138B      push     1              ; Count
.text:0054138D      and     ecx, 3
.text:00541390      rep stosb              ; memset (buffer, 0, aligned_size)
.text:00541392      mov     eax, [esp+38h+Str]
.text:00541396      push     eax             ; ElementSize
.text:00541397      push     ebx             ; DstBuf
.text:00541398      call     _fread          ; read file
.text:0054139D      push     ebp             ; File
.text:0054139E      call     _fclose
.text:005413A3      mov     ecx, [esp+44h+password]
.text:005413A7      push     ecx             ; password
.text:005413A8      push     esi             ; aligned size
.text:005413A9      push     ebx             ; buffer
.text:005413AA      call     crypt           ; do crypt
.text:005413AF      mov     edx, [esp+50h+Filename]
.text:005413B3      add     esp, 40h
.text:005413B6      push     offset aWb       ; "wb"
.text:005413BB      push     edx             ; Filename
.text:005413BC      call     _fopen
.text:005413C1      mov     edi, eax
.text:005413C3      push     edi             ; File
.text:005413C4      push     1              ; Count
.text:005413C6      push     3              ; Size
.text:005413C8      push     offset aQr9      ; "QR9"
.text:005413CD      call     _fwrite         ; write file signature
.text:005413D2      push     edi             ; File
.text:005413D3      push     1              ; Count
.text:005413D5      lea     eax, [esp+30h+Str]
.text:005413D9      push     4              ; Size
.text:005413DB      push     eax             ; Str
.text:005413DC      call     _fwrite         ; write original file size
.text:005413E1      push     edi             ; File
.text:005413E2      push     1              ; Count
.text:005413E4      push     esi             ; Size
.text:005413E5      push     ebx             ; Str
.text:005413E6      call     _fwrite         ; write crypted file
.text:005413EB      push     edi             ; File
.text:005413EC      call     _fclose
.text:005413F1      push     ebx             ; Memory
.text:005413F2      call     _free
.text:005413F7      add     esp, 40h
.text:005413FA      pop     edi
.text:005413FB      pop     esi
.text:005413FC      pop     ebx
.text:005413FD      pop     ebp
.text:005413FE      retn

```

```

.text:005413FE crypt_file      endp
.text:005413FE
.text:005413FE ; -----
.text:005413FF                align 10h
.text:00541400
.text:00541400 ; ===== S U B R O U T I N E =====
.text:00541400
.text:00541400 ; int __cdecl decrypt_file(char *Filename, int, void *Src)
.text:00541400 decrypt_file      proc near                ; CODE XREF: _main+6E
.text:00541400
.text:00541400 Filename          = dword ptr 4
.text:00541400 arg_4              = dword ptr 8
.text:00541400 Src                = dword ptr 0Ch
.text:00541400
.text:00541400                mov     eax, [esp+Filename]
.text:00541404                push    ebx
.text:00541405                push    ebp
.text:00541406                push    esi
.text:00541407                push    edi
.text:00541408                push    offset aRb          ; "rb"
.text:0054140D                push    eax                ; Filename
.text:0054140E                call    _fopen
.text:00541413                mov     esi, eax
.text:00541415                add     esp, 8
.text:00541418                test    esi, esi
.text:0054141A                jnz     short loc_54142E
.text:0054141C                push    offset aCannotOpenIn_0 ; "Cannot open input file!\n"
.text:00541421                call    _printf
.text:00541426                add     esp, 4
.text:00541429                pop     edi
.text:0054142A                pop     esi
.text:0054142B                pop     ebp
.text:0054142C                pop     ebx
.text:0054142D                retn
.text:0054142E ; -----
.text:0054142E
.text:0054142E loc_54142E:                ; CODE XREF: decrypt_file+1A
.text:0054142E                push    2                ; Origin
.text:00541430                push    0                ; Offset
.text:00541432                push    esi                ; File
.text:00541433                call    _fseek
.text:00541438                push    esi                ; File
.text:00541439                call    _ftell
.text:0054143E                push    0                ; Origin
.text:00541440                push    0                ; Offset
.text:00541442                push    esi                ; File
.text:00541443                mov     ebp, eax
.text:00541445                call    _fseek
.text:0054144A                push    ebp                ; Size
.text:0054144B                call    _malloc
.text:00541450                push    esi                ; File
.text:00541451                mov     ebx, eax
.text:00541453                push    1                ; Count
.text:00541455                push    ebp                ; ElementSize
.text:00541456                push    ebx                ; DstBuf
.text:00541457                call    _fread
.text:0054145C                push    esi                ; File
.text:0054145D                call    _fclose
.text:00541462                add     esp, 34h
.text:00541465                mov     ecx, 3
.text:0054146A                mov     edi, offset aQr9_0 ; "QR9"
.text:0054146F                mov     esi, ebx
.text:00541471                xor     edx, edx
.text:00541473                repe cmps
.text:00541475                jz      short loc_541489

```

```

.text:00541477      push    offset aFileIsNotCrypt ; "File is not crypted!\n"
.text:0054147C      call     _printf
.text:00541481      add     esp, 4
.text:00541484      pop     edi
.text:00541485      pop     esi
.text:00541486      pop     ebp
.text:00541487      pop     ebx
.text:00541488      retn
.text:00541489 ; -----
.text:00541489
.text:00541489 loc_541489:                                ; CODE XREF: decrypt_file+75
.text:00541489      mov     eax, [esp+10h+Src]
.text:0054148D      mov     edi, [ebx+3]
.text:00541490      add     ebp, 0FFFFFFF9h
.text:00541493      lea     esi, [ebx+7]
.text:00541496      push    eax                ; Src
.text:00541497      push    ebp                ; int
.text:00541498      push    esi                ; int
.text:00541499      call    decrypt
.text:0054149E      mov     ecx, [esp+1Ch+arg_4]
.text:005414A2      push    offset aWb_0       ; "wb"
.text:005414A7      push    ecx                ; Filename
.text:005414A8      call    _fopen
.text:005414AD      mov     ebp, eax
.text:005414AF      push    ebp                ; File
.text:005414B0      push    1                  ; Count
.text:005414B2      push    edi                ; Size
.text:005414B3      push    esi                ; Str
.text:005414B4      call    _fwrite
.text:005414B9      push    ebp                ; File
.text:005414BA      call    _fclose
.text:005414BF      push    ebx                ; Memory
.text:005414C0      call    _free
.text:005414C5      add     esp, 2Ch
.text:005414C8      pop     edi
.text:005414C9      pop     esi
.text:005414CA      pop     ebp
.text:005414CB      pop     ebx
.text:005414CC      retn
.text:005414CC decrypt_file      endp

```

All function and label names are given by me while analysis.

I started from top. Here is a function taking two file names and password.

```

.text:00541320 ; int __cdecl crypt_file(int Str, char *Filename, int password)
.text:00541320 crypt_file      proc near
.text:00541320
.text:00541320 Str                = dword ptr 4
.text:00541320 Filename          = dword ptr 8
.text:00541320 password          = dword ptr 0Ch
.text:00541320

```

Open file and report error in case of error:

```

.text:00541320      mov     eax, [esp+Str]
.text:00541324      push    ebp
.text:00541325      push    offset Mode       ; "rb"
.text:0054132A      push    eax                ; Filename
.text:0054132B      call    _fopen            ; open file
.text:00541330      mov     ebp, eax
.text:00541332      add     esp, 8
.text:00541335      test    ebp, ebp
.text:00541337      jnz     short loc_541348
.text:00541339      push    offset Format      ; "Cannot open input file!\n"
.text:0054133E      call    _printf
.text:00541343      add     esp, 4
.text:00541346      pop     ebp

```

```
.text:00541347          retn
.text:00541348 ; -----
.text:00541348
.text:00541348 loc_541348:
```

Get file size via fseek()/ftell():

```
.text:00541348 push    ebx
.text:00541349 push    esi
.text:0054134A push    edi
.text:0054134B push    2          ; Origin
.text:0054134D push    0          ; Offset
.text:0054134F push    ebp          ; File

; move current file position to the end
.text:00541350 call     _fseek
.text:00541355 push    ebp          ; File
.text:00541356 call     _ftell      ; get current file position
.text:0054135B push    0          ; Origin
.text:0054135D push    0          ; Offset
.text:0054135F push    ebp          ; File
.text:00541360 mov     [esp+2Ch+Str], eax

; move current file position to the start
.text:00541364 call     _fseek
```

This fragment of code calculates file size aligned to 64-byte border. This is because this cryptoalgorithm works with only 64-byte blocks. Its operation is pretty simple: divide file size by 64, forget about remainder and add 1, then multiple by 64. The following code removes remainder as if value was already divided by 64 and adds 64. It is almost the same.

```
.text:00541369 mov     esi, [esp+2Ch+Str]
.text:0054136D and     esi, 0FFFFFFC0h ; reset all lowest 6 bits
.text:00541370 add     esi, 40h       ; align size to 64-byte border
```

Allocate buffer with aligned size:

```
.text:00541373          push    esi          ; Size
.text:00541374          call     _malloc
```

Call memset(), e.g, clear allocated buffer².

```
.text:00541379 mov     ecx, esi
.text:0054137B mov     ebx, eax          ; allocated buffer pointer -> to EBX
.text:0054137D mov     edx, ecx
.text:0054137F xor     eax, eax
.text:00541381 mov     edi, ebx
.text:00541383 push    ebp          ; File
.text:00541384 shr     ecx, 2
.text:00541387 rep stosd
.text:00541389 mov     ecx, edx
.text:0054138B push    1          ; Count
.text:0054138D and     ecx, 3
.text:00541390 rep stosb          ; memset (buffer, 0, aligned_size)
```

Read file via standard C function fread().

```
.text:00541392          mov     eax, [esp+38h+Str]
.text:00541396          push    eax          ; ElementSize
.text:00541397          push    ebx          ; DstBuf
.text:00541398          call     _fread      ; read file
.text:0054139D          push    ebp          ; File
.text:0054139E          call     _fclose
```

Call crypt(). This function takes buffer, buffer size (aligned) and password string.

²malloc() + memset() could be replaced by calloc()

.text:005413A3	mov	ecx, [esp+44h+password]
.text:005413A7	push	ecx ; password
.text:005413A8	push	esi ; aligned size
.text:005413A9	push	ebx ; buffer
.text:005413AA	call	crypt ; do crypt

Create output file. By the way, developer forgot to check if it is was created correctly! File opening result is being checked though.

.text:005413AF	mov	edx, [esp+50h+Filename]
.text:005413B3	add	esp, 40h
.text:005413B6	push	offset aWb ; "wb"
.text:005413BB	push	edx ; Filename
.text:005413BC	call	_fopen
.text:005413C1	mov	edi, eax

Newly created file handle is in EDI register now. Write signature "QR9".

.text:005413C3	push	edi ; File
.text:005413C4	push	1 ; Count
.text:005413C6	push	3 ; Size
.text:005413C8	push	offset aQr9 ; "QR9"
.text:005413CD	call	_fwrite ; write file signature

Write actual file size (not aligned):

.text:005413D2	push	edi ; File
.text:005413D3	push	1 ; Count
.text:005413D5	lea	eax, [esp+30h+Str]
.text:005413D9	push	4 ; Size
.text:005413DB	push	eax ; Str
.text:005413DC	call	_fwrite ; write original file size

Write crypted buffer:

.text:005413E1	push	edi ; File
.text:005413E2	push	1 ; Count
.text:005413E4	push	esi ; Size
.text:005413E5	push	ebx ; Str
.text:005413E6	call	_fwrite ; write crypted file

Close file and free allocated buffer:

.text:005413EB	push	edi ; File
.text:005413EC	call	_fclose
.text:005413F1	push	ebx ; Memory
.text:005413F2	call	_free
.text:005413F7	add	esp, 40h
.text:005413FA	pop	edi
.text:005413FB	pop	esi
.text:005413FC	pop	ebx
.text:005413FD	pop	ebp
.text:005413FE	retn	
.text:005413FE crypt_file	endp	

Here is reconstructed C-code:

```
void crypt_file(char *fin, char* fout, char *pw)
{
    FILE *f;
    int flen, flen_aligned;
    BYTE *buf;

    f=fopen(fin, "rb");

    if (f==NULL)
    {
```

```

        printf ("Cannot open input file!\n");
        return;
};

fseek (f, 0, SEEK_END);
flen=ftell (f);
fseek (f, 0, SEEK_SET);

flen_aligned=(flen&0xFFFFFC0)+0x40;

buf=(BYTE*)malloc (flen_aligned);
memset (buf, 0, flen_aligned);

fread (buf, flen, 1, f);

fclose (f);

crypt (buf, flen_aligned, pw);

f=fopen(fout, "wb");

fwrite ("QR9", 3, 1, f);
fwrite (&flen, 4, 1, f);
fwrite (buf, flen_aligned, 1, f);

fclose (f);

free (buf);
};

```

Decrypting procedure is almost the same:

```

.text:00541400 ; int __cdecl decrypt_file(char *Filename, int, void *Src)
.text:00541400 decrypt_file      proc near
.text:00541400
.text:00541400 Filename          = dword ptr 4
.text:00541400 arg_4              = dword ptr 8
.text:00541400 Src                = dword ptr 0Ch
.text:00541400
.text:00541400                mov     eax, [esp+Filename]
.text:00541404                push    ebx
.text:00541405                push    ebp
.text:00541406                push    esi
.text:00541407                push    edi
.text:00541408                push    offset aRb      ; "rb"
.text:0054140D                push    eax          ; Filename
.text:0054140E                call    _fopen
.text:00541413                mov     esi, eax
.text:00541415                add     esp, 8
.text:00541418                test    esi, esi
.text:0054141A                jnz     short loc_54142E
.text:0054141C                push    offset aCannotOpenIn_0 ; "Cannot open input file!\n"
.text:00541421                call    _printf
.text:00541426                add     esp, 4
.text:00541429                pop     edi
.text:0054142A                pop     esi
.text:0054142B                pop     ebp
.text:0054142C                pop     ebx
.text:0054142D                retn
.text:0054142E ; -----
.text:0054142E
.text:0054142E loc_54142E:
.text:0054142E                push    2              ; Origin
.text:00541430                push    0              ; Offset
.text:00541432                push    esi          ; File
.text:00541433                call    _fseek
.text:00541438                push    esi          ; File

```

```

.text:00541439      call     _ftell
.text:0054143E      push     0           ; Origin
.text:00541440      push     0           ; Offset
.text:00541442      push     esi          ; File
.text:00541443      mov     ebp, eax
.text:00541445      call     _fseek
.text:0054144A      push     ebp          ; Size
.text:0054144B      call     _malloc
.text:00541450      push     esi          ; File
.text:00541451      mov     ebx, eax
.text:00541453      push     1           ; Count
.text:00541455      push     ebp          ; ElementSize
.text:00541456      push     ebx          ; DstBuf
.text:00541457      call     _fread
.text:0054145C      push     esi          ; File
.text:0054145D      call     _fclose

```

Check signature (first 3 bytes):

```

.text:00541462      add     esp, 34h
.text:00541465      mov     ecx, 3
.text:0054146A      mov     edi, offset aQr9_0 ; "QR9"
.text:0054146F      mov     esi, ebx
.text:00541471      xor     edx, edx
.text:00541473      repe    cmpsb
.text:00541475      jz      short loc_541489

```

Report an error if signature is absent:

```

.text:00541477      push    offset aFileIsNotCrypt ; "File is not crypted!\n"
.text:0054147C      call    _printf
.text:00541481      add     esp, 4
.text:00541484      pop     edi
.text:00541485      pop     esi
.text:00541486      pop     ebp
.text:00541487      pop     ebx
.text:00541488      retn
.text:00541489 ; -----
.text:00541489
.text:00541489 loc_541489:

```

Call decrypt().

```

.text:00541489      mov     eax, [esp+10h+Src]
.text:0054148D      mov     edi, [ebx+3]
.text:00541490      add     ebp, 0FFFFFFF9h
.text:00541493      lea     esi, [ebx+7]
.text:00541496      push    eax           ; Src
.text:00541497      push    ebp           ; int
.text:00541498      push    esi           ; int
.text:00541499      call    decrypt
.text:0054149E      mov     ecx, [esp+1Ch+arg_4]
.text:005414A2      push    offset aWb_0   ; "wb"
.text:005414A7      push    ecx           ; Filename
.text:005414A8      call    _fopen
.text:005414AD      mov     ebp, eax
.text:005414AF      push    ebp           ; File
.text:005414B0      push    1             ; Count
.text:005414B2      push    edi           ; Size
.text:005414B3      push    esi           ; Str
.text:005414B4      call    _fwrite
.text:005414B9      push    ebp           ; File
.text:005414BA      call    _fclose
.text:005414BF      push    ebx           ; Memory
.text:005414C0      call    _free
.text:005414C5      add     esp, 2Ch
.text:005414C8      pop     edi

```

```

.text:005414C9      pop     esi
.text:005414CA      pop     ebp
.text:005414CB      pop     ebx
.text:005414CC      retn
.text:005414CC decrypt_file  endp

```

Here is reconstructed C-code:

```

void decrypt_file(char *fin, char* fout, char *pw)
{
    FILE *f;
    int real_flen, flen;
    BYTE *buf;

    f=fopen(fin, "rb");

    if (f==NULL)
    {
        printf ("Cannot open input file!\n");
        return;
    };

    fseek (f, 0, SEEK_END);
    flen=ftell (f);
    fseek (f, 0, SEEK_SET);

    buf=(BYTE*)malloc (flen);

    fread (buf, flen, 1, f);

    fclose (f);

    if (memcmp (buf, "QR9", 3)!=0)
    {
        printf ("File is not crypted!\n");
        return;
    };

    memcpy (&real_flen, buf+3, 4);

    decrypt (buf+(3+4), flen-(3+4), pw);

    f=fopen(fout, "wb");

    fwrite (buf+(3+4), real_flen, 1, f);

    fclose (f);

    free (buf);
};

```

OK, now let's go deeper.

Function crypt():

```

.text:00541260 crypt      proc near
.text:00541260
.text:00541260 arg_0      = dword ptr  4
.text:00541260 arg_4      = dword ptr  8
.text:00541260 arg_8      = dword ptr  0Ch
.text:00541260
.text:00541260      push     ebx
.text:00541261      mov      ebx, [esp+4+arg_0]
.text:00541265      push     ebp
.text:00541266      push     esi
.text:00541267      push     edi
.text:00541268      xor      ebp, ebp
.text:0054126A

```



```
.text:0054126A loc_54126A:
```

This fragment of code copies part of input buffer to internal array I named later “cube64”. The size is in ECX register. MOVSD means *move 32-bit dword*, so, 16 of 32-bit dwords are exactly 64 bytes.

```
.text:0054126A      mov     eax, [esp+10h+arg_8]
.text:0054126E      mov     ecx, 10h
.text:00541273      mov     esi, ebx    ; EBX is pointer within input buffer
.text:00541275      mov     edi, offset cube64
.text:0054127A      push    1
.text:0054127C      push    eax
.text:0054127D      rep movsd
```

Call rotate_all_with_password():

```
.text:0054127F      call    rotate_all_with_password
```

Copy crypted contents back from “cube64” to buffer:

```
.text:00541284      mov     eax, [esp+18h+arg_4]
.text:00541288      mov     edi, ebx
.text:0054128A      add     ebp, 40h
.text:0054128D      add     esp, 8
.text:00541290      mov     ecx, 10h
.text:00541295      mov     esi, offset cube64
.text:0054129A      add     ebx, 40h    ; add 64 to input buffer pointer
.text:0054129D      cmp     ebp, eax    ; EBP contain ammount of crypted data.
.text:0054129F      rep movsd
```

If EBP is not bigger that input argument size, then continue to next block.

```
.text:005412A1      jl      short loc_54126A
.text:005412A3      pop     edi
.text:005412A4      pop     esi
.text:005412A5      pop     ebp
.text:005412A6      pop     ebx
.text:005412A7      retn
.text:005412A7 crypt      endp
```

Reconstructed crypt() function:

```
void crypt (BYTE *buf, int sz, char *pw)
{
    int i=0;

    do
    {
        memcpy (cube, buf+i, 8*8);
        rotate_all (pw, 1);
        memcpy (buf+i, cube, 8*8);
        i+=64;
    }
    while (i<sz);
};
```

OK, now let’s go deeper into function rotate_all_with_password(). It takes two arguments: password string and number. In crypt(), number 1 is used, and in decrypt() (where rotate_all_with_password() function is called too), number is 3.

```
.text:005411B0 rotate_all_with_password proc near
.text:005411B0
.text:005411B0 arg_0      = dword ptr  4
.text:005411B0 arg_4      = dword ptr  8
.text:005411B0
.text:005411B0      mov     eax, [esp+arg_0]
.text:005411B4      push    ebp
.text:005411B5      mov     ebp, eax
```

Check for character in password. If it is zero, exit:

```
.text:005411B7      cmp     byte ptr [eax], 0
.text:005411BA      jz      exit
.text:005411C0      push   ebx
.text:005411C1      mov     ebx, [esp+8+arg_4]
.text:005411C5      push   esi
.text:005411C6      push   edi
.text:005411C7
.text:005411C7 loop_begin:
```

Call `tolower()`, standard C function.

```
.text:005411C7      movsx   eax, byte ptr [ebp+0]
.text:005411CB      push   eax ; C
.text:005411CC      call   _tolower
.text:005411D1      add     esp, 4
```

Hmm, if password contains non-alphabetical latin character, it is skipped! Indeed, if we run crypting utility and try non-alphabetical latin characters in password, they seem to be ignored.

```
.text:005411D4      cmp     al, 'a'
.text:005411D6      jl      short next_character_in_password
.text:005411D8      cmp     al, 'z'
.text:005411DA      jg      short next_character_in_password
.text:005411DC      movsx   ecx, al
```

Subtract “a” value (97) from character.

```
.text:005411DF      sub     ecx, 'a' ; 97
```

After subtracting, we’ll get 0 for “a” here, 1 for “b”, etc. And 25 for “z”.

```
.text:005411E2      cmp     ecx, 24
.text:005411E5      jle     short skip_subtracting
.text:005411E7      sub     ecx, 24
```

It seems, “y” and “z” are exceptional characters too. After that fragment of code, “y” becomes 0 and “z” — 1. This means, 26 latin alphabet symbols will become values in range 0..23, (24 in total).

```
.text:005411EA
.text:005411EA skip_subtracting: ; CODE XREF: rotate_all_with_password+35
```

This is actually division via multiplication. Read more about it in “Division by 9” section [1.12](#).

The code actually divides password character value by 3.

```
.text:005411EA      mov     eax, 55555556h
.text:005411EF      imul    ecx
.text:005411F1      mov     eax, edx
.text:005411F3      shr     eax, 1Fh
.text:005411F6      add     edx, eax
.text:005411F8      mov     eax, ecx
.text:005411FA      mov     esi, edx
.text:005411FC      mov     ecx, 3
.text:00541201      cdq
.text:00541202      idiv    ecx
```

EDX is the remainder of division.

```
.text:00541204 sub     edx, 0
.text:00541207 jz      short call_rotate1 ; if remainder is zero, go to rotate1
.text:00541209 dec     edx
.text:0054120A jz      short call_rotate2 ; .. it it is 1, go to rotate2
.text:0054120C dec     edx
.text:0054120D jnz     short next_character_in_password
.text:0054120F test    ebx, ebx
.text:00541211 jle     short next_character_in_password
.text:00541213 mov     edi, ebx
```

If remainder is 2, call rotate3(). EDI is a second argument of rotate_all_with_password(). As I already wrote, 1 is for crypting operations and 3 is for decrypting. So, here is a loop. When crypting, rotate1/2/3 will be called the same number of times as given in the first argument.

```
.text:00541215 call_rotate3:
.text:00541215         push    esi
.text:00541216         call    rotate3
.text:0054121B         add     esp, 4
.text:0054121E         dec     edi
.text:0054121F         jnz     short call_rotate3
.text:00541221         jmp     short next_character_in_password
.text:00541223
.text:00541223 call_rotate2:
.text:00541223         test    ebx, ebx
.text:00541225         jle     short next_character_in_password
.text:00541227         mov     edi, ebx
.text:00541229
.text:00541229 loc_541229:
.text:00541229         push    esi
.text:0054122A         call    rotate2
.text:0054122F         add     esp, 4
.text:00541232         dec     edi
.text:00541233         jnz     short loc_541229
.text:00541235         jmp     short next_character_in_password
.text:00541237
.text:00541237 call_rotate1:
.text:00541237         test    ebx, ebx
.text:00541239         jle     short next_character_in_password
.text:0054123B         mov     edi, ebx
.text:0054123D
.text:0054123D loc_54123D:
.text:0054123D         push    esi
.text:0054123E         call    rotate1
.text:00541243         add     esp, 4
.text:00541246         dec     edi
.text:00541247         jnz     short loc_54123D
.text:00541249
```

Fetch next character from password string.

```
.text:00541249 next_character_in_password:
.text:00541249         mov     al, [ebp+1]
```

Increment character pointer within password string:

```
.text:0054124C         inc     ebp
.text:0054124D         test    al, al
.text:0054124F         jnz     loop_begin
.text:00541255         pop     edi
.text:00541256         pop     esi
.text:00541257         pop     ebx
.text:00541258
.text:00541258 exit:
.text:00541258         pop     ebp
.text:00541259         retn
.text:00541259 rotate_all_with_password endp
```

Here is reconstructed C code:

```
void rotate_all (char *pwd, int v)
{
    char *p=pwd;

    while (*p)
    {
        char c=*p;
        int q;
```

```

        c=tolower (c);

        if (c>='a' && c<='z')
        {
            q=c-'a';
            if (q>24)
                q-=24;

            int quotient=q/3;
            int remainder=q % 3;

            switch (remainder)
            {
            case 0: for (int i=0; i<v; i++) rotate1 (quotient); break;
            case 1: for (int i=0; i<v; i++) rotate2 (quotient); break;
            case 2: for (int i=0; i<v; i++) rotate3 (quotient); break;
            };

        };

        p++;
    };
};

```

Now let's go deeper and investigate rotate1/2/3 functions. Each function calls two another functions. I eventually gave them names set_bit() and get_bit().

Let's start with get_bit():

```

.text:00541050 get_bit      proc near
.text:00541050
.text:00541050 arg_0        = dword ptr 4
.text:00541050 arg_4        = dword ptr 8
.text:00541050 arg_8        = byte ptr 0Ch
.text:00541050
.text:00541050          mov     eax, [esp+arg_4]
.text:00541054          mov     ecx, [esp+arg_0]
.text:00541058          mov     al, cube64[eax+ecx*8]
.text:0054105F          mov     cl, [esp+arg_8]
.text:00541063          shr     al, cl
.text:00541065          and     al, 1
.text:00541067          retn
.text:00541067 get_bit      endp

```

...in other words: calculate an index in the array cube64: $arg_4 + arg_0 * 8$. Then shift a byte from an array by arg_8 bits right. Isolate lowest bit and return it.

Let's see another function, set_bit():

```

.text:00541000 set_bit      proc near
.text:00541000
.text:00541000 arg_0        = dword ptr 4
.text:00541000 arg_4        = dword ptr 8
.text:00541000 arg_8        = dword ptr 0Ch
.text:00541000 arg_C        = byte ptr 10h
.text:00541000
.text:00541000          mov     al, [esp+arg_C]
.text:00541004          mov     ecx, [esp+arg_8]
.text:00541008          push    esi
.text:00541009          mov     esi, [esp+4+arg_0]
.text:0054100D          test    al, al
.text:0054100F          mov     eax, [esp+4+arg_4]
.text:00541013          mov     dl, 1
.text:00541015          jz      short loc_54102B

```

DL is 1 here. Shift left it by arg_8 . For example, if arg_8 is 4, DL register value became 0x10 or 1000 in binary form.

```
.text:00541017      shl     dl, cl
.text:00541019      mov     cl, cube64[eax+esi*8]
```

Get bit from array and explicitly set one.

```
.text:00541020      or      cl, dl
```

Store it back:

```
.text:00541022      mov     cube64[eax+esi*8], cl
.text:00541029      pop     esi
.text:0054102A      retn
.text:0054102B ; -----
.text:0054102B
.text:0054102B loc_54102B:
.text:0054102B      shl     dl, cl
```

If arg_C is not zero...

```
.text:0054102D      mov     cl, cube64[eax+esi*8]
```

...invert DL. For example, if DL state after shift was 0x10 or 1000 in binary form, there will be 0xEF after NOT instruction or 11101111 in binary form.

```
.text:00541034      not     dl
```

This instruction clears bit, in other words, it saves all bits in CL which are also set in DL except those in DL which are cleared. This means that if DL is, for example, 11101111 in binary form, all bits will be saved except 5th (counting from lowest bit).

```
.text:00541036      and     cl, dl
```

Store it back:

```
.text:00541038      mov     cube64[eax+esi*8], cl
.text:0054103F      pop     esi
.text:00541040      retn
.text:00541040 set_bit      endp
```

It is almost the same as `get_bit()`, except, if `arg_C` is zero, the function clears specific bit in array, or sets it otherwise.

We also know that array size is 64. First two arguments both in `set_bit()` and `get_bit()` could be seen as 2D coordinates. Then array will be 8*8 matrix.

Here is C representation of what we already know:

```
#define IS_SET(flag, bit)      ((flag) & (bit))
#define SET_BIT(var, bit)     ((var) |= (bit))
#define REMOVE_BIT(var, bit)  ((var) &= ~(bit))

char cube[8][8];

void set_bit (int x, int y, int shift, int bit)
{
    if (bit)
        SET_BIT (cube[x][y], 1<<shift);
    else
        REMOVE_BIT (cube[x][y], 1<<shift);
};

int get_bit (int x, int y, int shift)
{
    if ((cube[x][y]>>shift)&1==1)
        return 1;
    return 0;
};
```

Now let's get back to rotate1/2/3 functions.

```
.text:00541070 rotate1      proc near
.text:00541070
```

Internal array allocation in local stack, its size 64 bytes:

```
.text:00541070 internal_array_64= byte ptr -40h
.text:00541070 arg_0          = dword ptr  4
.text:00541070
.text:00541070          sub     esp, 40h
.text:00541073          push    ebx
.text:00541074          push    ebp
.text:00541075          mov     ebp, [esp+48h+arg_0]
.text:00541079          push    esi
.text:0054107A          push    edi
.text:0054107B          xor     edi, edi      ; EDI is loop1 counter
```

EBX is a pointer to internal array:

```
.text:0054107D          lea     ebx, [esp+50h+internal_array_64]
.text:00541081
```

Two nested loops are here:

```
.text:00541081 first_loop1_begin:
.text:00541081          xor     esi, esi      ; ESI is loop 2 counter
.text:00541083
.text:00541083 first_loop2_begin:
.text:00541083          push    ebp          ; arg_0
.text:00541084          push    esi          ; loop 1 counter
.text:00541085          push    edi          ; loop 2 counter
.text:00541086          call   get_bit
.text:0054108B          add     esp, 0Ch
.text:0054108E          mov     [ebx+esi], al    ; store to internal array
.text:00541091          inc     esi          ; increment loop 1 counter
.text:00541092          cmp     esi, 8
.text:00541095          jnl     short first_loop2_begin
.text:00541097          inc     edi          ; increment loop 2 counter
.text:00541098          add     ebx, 8          ; increment internal array pointer by 8 at each loop 1 iteration
.text:0054109B          cmp     edi, 8
.text:0054109E          jnl     short first_loop1_begin
```

...we see that both loop counters are in range 0..7. Also, they are used as first and second arguments of get_bit(). Third argument of get_bit() is the only argument of rotate1(). What get_bit() returns, is being placed into internal array.

Prepare pointer to internal array again:

```
.text:005410A0          lea     ebx, [esp+50h+internal_array_64]
.text:005410A4          mov     edi, 7          ; EDI is loop 1 counter, initial state is 7
.text:005410A9
.text:005410A9 second_loop1_begin:
.text:005410A9          xor     esi, esi      ; ESI is loop 2 counter
.text:005410AB
.text:005410AB second_loop2_begin:
.text:005410AB          mov     al, [ebx+esi]    ; value from internal array
.text:005410AE          push    eax
.text:005410AF          push    ebp          ; arg_0
.text:005410B0          push    edi          ; loop 1 counter
.text:005410B1          push    esi          ; loop 2 counter
.text:005410B2          call   set_bit
.text:005410B7          add     esp, 10h
.text:005410BA          inc     esi          ; increment loop 2 counter
.text:005410BB          cmp     esi, 8
.text:005410BE          jnl     short second_loop2_begin
.text:005410C0          dec     edi          ; decrement loop 2 counter
.text:005410C1          add     ebx, 8          ; increment pointer in internal array
```

```

.text:005410C4    cmp     edi, 0FFFFFFFh
.text:005410C7    jg      short second_loop1_begin
.text:005410C9    pop     edi
.text:005410CA    pop     esi
.text:005410CB    pop     ebp
.text:005410CC    pop     ebx
.text:005410CD    add     esp, 40h
.text:005410D0    retn
.text:005410D0 rotate1      endp

```

...this code is placing contents from internal array to cube global array via `set_bit()` function, *but*, in different order! Now loop 1 counter is in range 7 to 0, decrementing at each iteration!

C code representation looks like:

```

void rotate1 (int v)
{
    bool tmp[8][8]; // internal array
    int i, j;

    for (i=0; i<8; i++)
        for (j=0; j<8; j++)
            tmp[i][j]=get_bit (i, j, v);

    for (i=0; i<8; i++)
        for (j=0; j<8; j++)
            set_bit (j, 7-i, v, tmp[x][y]);
};

```

Not very understandable, but if we will take a look at `rotate2()` function:

```

.text:005410E0 rotate2 proc near
.text:005410E0
.text:005410E0 internal_array_64 = byte ptr -40h
.text:005410E0 arg_0 = dword ptr 4
.text:005410E0
.text:005410E0    sub     esp, 40h
.text:005410E3    push    ebx
.text:005410E4    push    ebp
.text:005410E5    mov     ebp, [esp+48h+arg_0]
.text:005410E9    push    esi
.text:005410EA    push    edi
.text:005410EB    xor     edi, edi        ; loop 1 counter
.text:005410ED    lea     ebx, [esp+50h+internal_array_64]
.text:005410F1
.text:005410F1 loc_5410F1:
.text:005410F1    xor     esi, esi        ; loop 2 counter
.text:005410F3
.text:005410F3 loc_5410F3:
.text:005410F3    push    esi            ; loop 2 counter
.text:005410F4    push    edi            ; loop 1 counter
.text:005410F5    push    ebp            ; arg_0
.text:005410F6    call    get_bit
.text:005410FB    add     esp, 0Ch
.text:005410FE    mov     [ebx+esi], al    ; store to internal array
.text:00541101    inc     esi            ; increment loop 1 counter
.text:00541102    cmp     esi, 8
.text:00541105    jl      short loc_5410F3
.text:00541107    inc     edi            ; increment loop 2 counter
.text:00541108    add     ebx, 8
.text:0054110B    cmp     edi, 8
.text:0054110E    jl      short loc_5410F1
.text:00541110    lea     ebx, [esp+50h+internal_array_64]
.text:00541114    mov     edi, 7          ; loop 1 counter is initial state 7
.text:00541119
.text:00541119 loc_541119:
.text:00541119    xor     esi, esi        ; loop 2 counter
.text:0054111B

```

```
.text:0054111B loc_54111B:
.text:0054111B     mov     al, [ebx+esi]    ; get byte from internal array
.text:0054111E     push    eax
.text:0054111F     push    edi                ; loop 1 counter
.text:00541120     push    esi                ; loop 2 counter
.text:00541121     push    ebp                ; arg_0
.text:00541122     call    set_bit
.text:00541127     add     esp, 10h
.text:0054112A     inc     esi                ; increment loop 2 counter
.text:0054112B     cmp     esi, 8
.text:0054112E     jnl     short loc_54111B
.text:00541130     dec     edi                ; decrement loop 2 counter
.text:00541131     add     ebx, 8
.text:00541134     cmp     edi, 0FFFFFFFh
.text:00541137     jg     short loc_541119
.text:00541139     pop     edi
.text:0054113A     pop     esi
.text:0054113B     pop     ebp
.text:0054113C     pop     ebx
.text:0054113D     add     esp, 40h
.text:00541140     retn
.text:00541140 rotate2 endp
```

It is *almost* the same, except of different order of arguments of `get_bit()` and `set_bit()`. Let's rewrite it in C-like code:

```
void rotate2 (int v)
{
    bool tmp[8][8]; // internal array
    int i, j;

    for (i=0; i<8; i++)
        for (j=0; j<8; j++)
            tmp[i][j]=get_bit (v, i, j);

    for (i=0; i<8; i++)
        for (j=0; j<8; j++)
            set_bit (v, j, 7-i, tmp[i][j]);
};
```

Let's also rewrite `rotate3()` function:

```
void rotate3 (int v)
{
    bool tmp[8][8];
    int i, j;

    for (i=0; i<8; i++)
        for (j=0; j<8; j++)
            tmp[i][j]=get_bit (i, v, j);

    for (i=0; i<8; i++)
        for (j=0; j<8; j++)
            set_bit (7-j, v, i, tmp[i][j]);
};
```

Well, now things are simpler. If we consider `cube64` as 3D cube $8*8*8$, where each element is bit, `get_bit()` and `set_bit()` take just coordinates of bit on input.

`rotate1/2/3` functions are in fact rotating all bits in specific plane. Three functions are each for each cube side and `v` argument is setting plane in range 0..7.

Maybe, algorithm's author was thinking of $8*8*8$ [Rubik's cube](#)?!

Yes, indeed.

Let's get closer into `decrypt()` function, I rewrote it here:

```
void decrypt (BYTE *buf, int sz, char *pw)
{
```



```

    char *p=strdup (pw);
    strrev (p);
    int i=0;

    do
    {
        memcpy (cube, buf+i, 8*8);
        rotate_all (p, 3);
        memcpy (buf+i, cube, 8*8);
        i+=64;
    }
    while (i<sz);

    free (p);
};

```

It is almost the same except of `crypt()`, *but* password string is reversed by `strrev()` standard C function and `rotate_all()` is called with argument 3.

This means, that in case of decryption, each corresponding `rotate1/2/3` call will be performed thrice.

This is almost as in Rubik's cube! If you want to get back, do the same in reverse order and direction! If you need to undo effect of rotating one place in clockwise direction, rotate it thrice in counter-clockwise direction.

`rotate1()` is apparently for rotating "front" plane. `rotate2()` is apparently for rotating "top" plane. `rotate3()` is apparently for rotating "left" plane.

Let's get back to core of `rotate_all()` function:

```

q=c-'a';
if (q>24)
    q-=24;

int quotient=q/3; // in range 0..7
int remainder=q % 3;

switch (remainder)
{
    case 0: for (int i=0; i<v; i++) rotate1 (quotient); break; // front
    case 1: for (int i=0; i<v; i++) rotate2 (quotient); break; // top
    case 2: for (int i=0; i<v; i++) rotate3 (quotient); break; // left
};

```

Now it is much simpler to understand: each password character defines side (one of three) and plane (one of 8). $3 \times 8 = 24$, that's why two last characters of latin alphabet are remapped to fit an alphabet of exactly 24 elements.

The algorithm is clearly weak: in case of short passwords, one can see, that in crypted file there are some original bytes of original file in binary files editor.

Here is reconstructed whole source code:

```

#include <windows.h>

#include <stdio.h>
#include <assert.h>

#define IS_SET(flag, bit)      ((flag) & (bit))
#define SET_BIT(var, bit)      ((var) |= (bit))
#define REMOVE_BIT(var, bit)   ((var) &= ~(bit))

static BYTE cube[8][8];

void set_bit (int x, int y, int z, bool bit)
{
    if (bit)
        SET_BIT (cube[x][y], 1<<z);
    else
        REMOVE_BIT (cube[x][y], 1<<z);
};

```

```

bool get_bit (int x, int y, int z)
{
    if ((cube[x][y]>>z)&1==1)
        return true;
    return false;
};

void rotate_f (int row)
{
    bool tmp[8][8];
    int x, y;

    for (x=0; x<8; x++)
        for (y=0; y<8; y++)
            tmp[x][y]=get_bit (x, y, row);

    for (x=0; x<8; x++)
        for (y=0; y<8; y++)
            set_bit (y, 7-x, row, tmp[x][y]);
};

void rotate_t (int row)
{
    bool tmp[8][8];
    int y, z;

    for (y=0; y<8; y++)
        for (z=0; z<8; z++)
            tmp[y][z]=get_bit (row, y, z);

    for (y=0; y<8; y++)
        for (z=0; z<8; z++)
            set_bit (row, z, 7-y, tmp[y][z]);
};

void rotate_l (int row)
{
    bool tmp[8][8];
    int x, z;

    for (x=0; x<8; x++)
        for (z=0; z<8; z++)
            tmp[x][z]=get_bit (x, row, z);

    for (x=0; x<8; x++)
        for (z=0; z<8; z++)
            set_bit (7-z, row, x, tmp[x][z]);
};

void rotate_all (char *pwd, int v)
{
    char *p=pwd;

    while (*p)
    {
        char c=*p;
        int q;

        c=tolower (c);

        if (c>='a' && c<='z')
        {
            q=c-'a';
            if (q>24)
                q-=24;

```

```

        int quotient=q/3;
        int remainder=q % 3;

        switch (remainder)
        {
            case 0: for (int i=0; i<v; i++) rotate1 (quotient); break;
            case 1: for (int i=0; i<v; i++) rotate2 (quotient); break;
            case 2: for (int i=0; i<v; i++) rotate3 (quotient); break;
        };

    };

    p++;
};

void crypt (BYTE *buf, int sz, char *pw)
{
    int i=0;

    do
    {
        memcpy (cube, buf+i, 8*8);
        rotate_all (pw, 1);
        memcpy (buf+i, cube, 8*8);
        i+=64;
    }
    while (i<sz);
};

void decrypt (BYTE *buf, int sz, char *pw)
{
    char *p=strdup (pw);
    strrev (p);
    int i=0;

    do
    {
        memcpy (cube, buf+i, 8*8);
        rotate_all (p, 3);
        memcpy (buf+i, cube, 8*8);
        i+=64;
    }
    while (i<sz);

    free (p);
};

void crypt_file(char *fin, char* fout, char *pw)
{
    FILE *f;
    int flen, flen_aligned;
    BYTE *buf;

    f=fopen(fin, "rb");

    if (f==NULL)
    {
        printf ("Cannot open input file!\n");
        return;
    };

    fseek (f, 0, SEEK_END);
    flen=ftell (f);
    fseek (f, 0, SEEK_SET);

    flen_aligned=(flen&0xFFFFF0)+0x40;

```

```

    buf=(BYTE*)malloc (flen_aligned);
    memset (buf, 0, flen_aligned);

    fread (buf, flen, 1, f);

    fclose (f);

    crypt (buf, flen_aligned, pw);

    f=fopen(fout, "wb");

    fwrite ("QR9", 3, 1, f);
    fwrite (&flen, 4, 1, f);
    fwrite (buf, flen_aligned, 1, f);

    fclose (f);

    free (buf);
};

void decrypt_file(char *fin, char* fout, char *pw)
{
    FILE *f;
    int real_flen, flen;
    BYTE *buf;

    f=fopen(fin, "rb");

    if (f==NULL)
    {
        printf ("Cannot open input file!\n");
        return;
    };

    fseek (f, 0, SEEK_END);
    flen=ftell (f);
    fseek (f, 0, SEEK_SET);

    buf=(BYTE*)malloc (flen);

    fread (buf, flen, 1, f);

    fclose (f);

    if (memcmp (buf, "QR9", 3)!=0)
    {
        printf ("File is not crypted!\n");
        return;
    };

    memcpy (&real_flen, buf+3, 4);

    decrypt (buf+(3+4), flen-(3+4), pw);

    f=fopen(fout, "wb");

    fwrite (buf+(3+4), real_flen, 1, f);

    fclose (f);

    free (buf);
};

// run: input output 0/1 password

```

```
// 0 for encrypt, 1 for decrypt

int main(int argc, char *argv[])
{
    if (argc!=5)
    {
        printf ("Incorrect parameters!\n");
        return 1;
    };

    if (strcmp (argv[3], "0")==0)
        crypt_file (argv[1], argv[2], argv[4]);
    else
        if (strcmp (argv[3], "1")==0)
            decrypt_file (argv[1], argv[2], argv[4]);
        else
            printf ("Wrong param %s\n", argv[3]);

    return 0;
};
```

7.2 SAP

7.2.1 About SAP client network traffic compression

(Tracing connection between TDW_NOCOMPRESS SAPGUI³ environment variable to nagging pop-up window and actual data compression routine.)

It's known that network traffic between SAPGUI and SAP is not crypted by default, it's rather compressed (read [here](#) and [here](#)).

It's also known that by setting environment variable *TDW_NOCOMPRESS* to 1, it's possible to turn network packets compression off.

But you will see a nagging pop-up windows that cannot be closed:

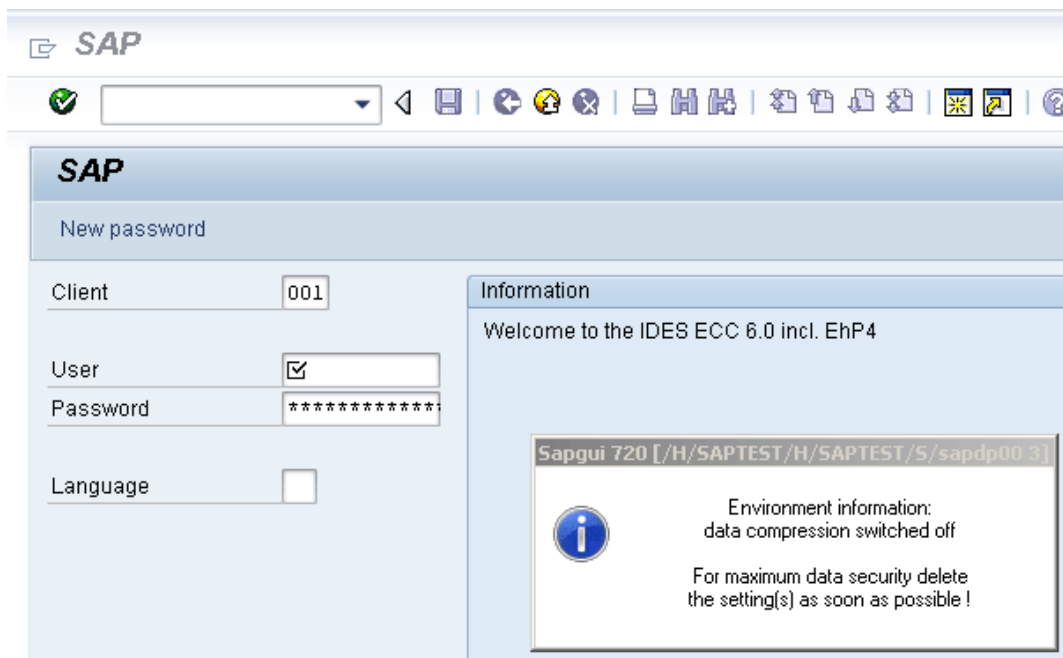


Figure 7.1: Screenshot

Let's see, if we can remove that window somehow.

³SAP GUI client

But before this, let's see what we already know. First: we know that environment variable *TDW_NOCOMPRESS* is checked somewhere inside of SAPGUI client. Second: string like "data compression switched off" must be present somewhere too. With the help of FAR file manager I found that both of these strings are stored in the SAPguilib.dll file.

So let's open SAPguilib.dll in IDA 5 and search for "TDW_NOCOMPRESS" string. Yes, it is present and there is only one reference to it.

We see the following fragment of code (all file offsets are valid for SAPGUI 720 win32, SAPguilib.dll file version 7200,1,0,9009):

```
.text:6440D51B      lea     eax, [ebp+2108h+var_211C]
.text:6440D51E      push    eax                ; int
.text:6440D51F      push    offset aTdw_nocompress ; "TDW_NOCOMPRESS"
.text:6440D524      mov     byte ptr [edi+15h], 0
.text:6440D528      call    chk_env
.text:6440D52D      pop     ecx
.text:6440D52E      pop     ecx
.text:6440D52F      push    offset byte_64443AF8
.text:6440D534      lea     ecx, [ebp+2108h+var_211C]

; demangled name: int ATL::CStringT::Compare(char const *)const
.text:6440D537      call    ds:mfc90_1603
.text:6440D53D      test    eax, eax
.text:6440D53F      jz      short loc_6440D55A
.text:6440D541      lea     ecx, [ebp+2108h+var_211C]

; demangled name: const char* ATL::CStringT::operator PCXSTR
.text:6440D544      call    ds:mfc90_910
.text:6440D54A      push    eax                ; Str
.text:6440D54B      call    ds:atoi
.text:6440D551      test    eax, eax
.text:6440D553      setnz   al
.text:6440D556      pop     ecx
.text:6440D557      mov     [edi+15h], al
```

String returned by *chk_env()* via second argument is then handled by MFC string functions and then *atoi()*⁴ is called. After that, numerical value is stored to *edi+15h*.

Also, take a look onto *chk_env()* function (I gave a name to it):

```
.text:64413F20 ; int __cdecl chk_env(char *VarName, int)
.text:64413F20 chk_env      proc near
.text:64413F20
.text:64413F20 DstSize      = dword ptr -0Ch
.text:64413F20 var_8          = dword ptr -8
.text:64413F20 DstBuf        = dword ptr -4
.text:64413F20 VarName        = dword ptr 8
.text:64413F20 arg_4          = dword ptr 0Ch
.text:64413F20
.text:64413F20      push    ebp
.text:64413F21      mov     ebp, esp
.text:64413F23      sub     esp, 0Ch
.text:64413F26      mov     [ebp+DstSize], 0
.text:64413F2D      mov     [ebp+DstBuf], 0
.text:64413F34      push    offset unk_6444C88C
.text:64413F39      mov     ecx, [ebp+arg_4]

; (demangled name) ATL::CStringT::operator=(char const *)
.text:64413F3C      call    ds:mfc90_820
.text:64413F42      mov     eax, [ebp+VarName]
.text:64413F45      push    eax                ; VarName
.text:64413F46      mov     ecx, [ebp+DstSize]
.text:64413F49      push    ecx                ; DstSize
.text:64413F4A      mov     edx, [ebp+DstBuf]
.text:64413F4D      push    edx                ; DstBuf
```

⁴standard C library function, converting number in string into number

```

.text:64413F4E      lea     eax, [ebp+DstSize]
.text:64413F51      push    eax                ; ReturnSize
.text:64413F52      call   ds:getenv_s
.text:64413F58      add     esp, 10h
.text:64413F5B      mov     [ebp+var_8], eax
.text:64413F5E      cmp     [ebp+var_8], 0
.text:64413F62      jz      short loc_64413F68
.text:64413F64      xor     eax, eax
.text:64413F66      jmp     short loc_64413FBC
.text:64413F68 ; -----
.text:64413F68      loc_64413F68:
.text:64413F68      cmp     [ebp+DstSize], 0
.text:64413F6C      jnz     short loc_64413F72
.text:64413F6E      xor     eax, eax
.text:64413F70      jmp     short loc_64413FBC
.text:64413F72 ; -----
.text:64413F72      loc_64413F72:
.text:64413F72      mov     ecx, [ebp+DstSize]
.text:64413F75      push    ecx
.text:64413F76      mov     ecx, [ebp+arg_4]

; demangled name: ATL::CStringT<char, 1>::Preallocate(int)
.text:64413F79      call   ds:mfc90_2691
.text:64413F7F      mov     [ebp+DstBuf], eax
.text:64413F82      mov     edx, [ebp+VarName]
.text:64413F85      push    edx                ; VarName
.text:64413F86      mov     eax, [ebp+DstSize]
.text:64413F89      push    eax                ; DstSize
.text:64413F8A      mov     ecx, [ebp+DstBuf]
.text:64413F8D      push    ecx                ; DstBuf
.text:64413F8E      lea     edx, [ebp+DstSize]
.text:64413F91      push    edx                ; ReturnSize
.text:64413F92      call   ds:getenv_s
.text:64413F98      add     esp, 10h
.text:64413F9B      mov     [ebp+var_8], eax
.text:64413F9E      push    0FFFFFFFh
.text:64413FA0      mov     ecx, [ebp+arg_4]

; demangled name: ATL::CStringT::ReleaseBuffer(int)
.text:64413FA3      call   ds:mfc90_5835
.text:64413FA9      cmp     [ebp+var_8], 0
.text:64413FAD      jz      short loc_64413FB3
.text:64413FAF      xor     eax, eax
.text:64413FB1      jmp     short loc_64413FBC
.text:64413FB3 ; -----
.text:64413FB3      loc_64413FB3:
.text:64413FB3      mov     ecx, [ebp+arg_4]

; demangled name: const char* ATL::CStringT::operator PCXSTR
.text:64413FB6      call   ds:mfc90_910
.text:64413FBC      loc_64413FBC:
.text:64413FBC      mov     esp, ebp
.text:64413FBE      pop     ebp
.text:64413FBF      retn
.text:64413FBF      chk_env      endp

```

Yes. `getenv_s()`⁵ function is Microsoft security-enhanced version of `getenv()`⁶.
There are also some MFC string manipulations.

⁵[http://msdn.microsoft.com/en-us/library/tb2sfw2z\(VS.80\).aspx](http://msdn.microsoft.com/en-us/library/tb2sfw2z(VS.80).aspx)

⁶Standard C library returning environment variable

Lots of other environment variables are checked as well. Here is a list of all variables being checked and what SAPGUI could write to trace log when logging is turned on:

DPTRACE	"GUI-OPTION: Trace set to %d"
TDW_HEXDUMP	"GUI-OPTION: Hexdump enabled"
TDW_WORKDIR	"GUI-OPTION: working directory '%s'"
TDW_SPLASHSRCEENOFF	"GUI-OPTION: Splash Screen Off" / "GUI-OPTION: Splash Screen On"
TDW_REPLYTIMEOUT	"GUI-OPTION: reply timeout %d milliseconds"
TDW_PLAYBACKTIMEOUT	"GUI-OPTION: PlaybackTimeout set to %d milliseconds"
TDW_NOCOMPRESS	"GUI-OPTION: no compression read"
TDW_EXPERT	"GUI-OPTION: expert mode"
TDW_PLAYBACKPROGRESS	"GUI-OPTION: PlaybackProgress"
TDW_PLAYBACKNETTRAFFIC	"GUI-OPTION: PlaybackNetTraffic"
TDW_PLAYLOG	"GUI-OPTION: /PlayLog is YES, file %s"
TDW_PLAYTIME	"GUI-OPTION: /PlayTime set to %d milliseconds"
TDW_LOGFILE	"GUI-OPTION: TDW_LOGFILE '%s'"
TDW_WAN	"GUI-OPTION: WAN - low speed connection enabled"
TDW_FULLMENU	"GUI-OPTION: FullMenu enabled"
SAP_CP / SAP_CODEPAGE	"GUI-OPTION: SAP_CODEPAGE '%d'"
UPDOWNLOAD_CP	"GUI-OPTION: UPDOWNLOAD_CP '%d'"
SNC_PARTNERNAME	"GUI-OPTION: SNC name '%s'"
SNC_QOP	"GUI-OPTION: SNC_QOP '%s'"
SNC_LIB	"GUI-OPTION: SNC is set to: %s"
SAPGUI_INPLACE	"GUI-OPTION: environment variable SAPGUI_INPLACE is on"

Settings for each variable are written to the array via pointer in EDI register. EDI is being set before that function call:

```
.text:6440EE00      lea     edi, [ebp+2884h+var_2884] ; options here like +0x15...
.text:6440EE03      lea     ecx, [esi+24h]
.text:6440EE06      call    load_command_line
.text:6440EE0B      mov     edi, eax
.text:6440EE0D      xor     ebx, ebx
.text:6440EE0F      cmp     edi, ebx
.text:6440EE11      jz      short loc_6440EE42
.text:6440EE13      push    edi
.text:6440EE14      push    offset aSapguiStoppedA ; "Sapgui stopped after commandline interp"...
.text:6440EE19      push    dword_644F93E8
.text:6440EE1F      call    FEWTraceError
```

Now, can we find *"data record mode switched on"* string? Yes, and here is the only reference in function CDwsGui::PrepareInfoWindow: How do I know class/method names? There is a lot of special debugging calls writing to log-files like:

```
.text:64405160      push    dword ptr [esi+2854h]
.text:64405166      push    offset aCdwsGuiPrepare ; "\nCDwsGui::PrepareInfoWindow: sapgui env
    "...
.text:6440516B      push    dword ptr [esi+2848h]
.text:64405171      call    dbg
.text:64405176      add     esp, 0Ch
```

...or:

```
.text:6440237A      push    eax
.text:6440237B      push    offset aCClientStart_6 ; "CClient::Start: set shortcut user to
    '\%' ...
.text:64402380      push    dword ptr [edi+4]
.text:64402383      call    dbg
.text:64402388      add     esp, 0Ch
```

It's **very** useful.

So let's see contents of that nagging pop-up window function:


```

.text:64404F4F CDwsGui__PrepareInfoWindow proc near
.text:64404F4F
.text:64404F4F pvParam          = byte ptr -3Ch
.text:64404F4F var_38            = dword ptr -38h
.text:64404F4F var_34            = dword ptr -34h
.text:64404F4F rc              = tagRECT ptr -2Ch
.text:64404F4F cy              = dword ptr -1Ch
.text:64404F4F h                = dword ptr -18h
.text:64404F4F var_14          = dword ptr -14h
.text:64404F4F var_10          = dword ptr -10h
.text:64404F4F var_4            = dword ptr -4
.text:64404F4F
.text:64404F4F                push    30h
.text:64404F51                mov     eax, offset loc_64438E00
.text:64404F56                call    __EH_prolog3
.text:64404F5B                mov     esi, ecx          ; ECX is pointer to object
.text:64404F5D                xor     ebx, ebx
.text:64404F5F                lea     ecx, [ebp+var_14]
.text:64404F62                mov     [ebp+var_10], ebx

; demangled name: ATL::CStringT(void)
.text:64404F65                call    ds:mfc90_316
.text:64404F6B                mov     [ebp+var_4], ebx
.text:64404F6E                lea     edi, [esi+2854h]
.text:64404F74                push    offset aEnvironmentInf ; "Environment information:\n"
.text:64404F79                mov     ecx, edi

; demangled name: ATL::CStringT::operator=(char const *)
.text:64404F7B                call    ds:mfc90_820
.text:64404F81                cmp     [esi+38h], ebx
.text:64404F84                mov     ebx, ds:mfc90_2539
.text:64404F8A                jbe     short loc_64404FA9
.text:64404F8C                push    dword ptr [esi+34h]
.text:64404F8F                lea     eax, [ebp+var_14]
.text:64404F92                push    offset aWorkingDirecto ; "working directory: '%s'\n"
.text:64404F97                push    eax

; demangled name: ATL::CStringT::Format(char const *,...)
.text:64404F98                call    ebx ; mfc90_2539
.text:64404F9A                add     esp, 0Ch
.text:64404F9D                lea     eax, [ebp+var_14]
.text:64404FA0                push    eax
.text:64404FA1                mov     ecx, edi

; demangled name: ATL::CStringT::operator+=(class ATL::CSimpleStringT<char, 1> const &)
.text:64404FA3                call    ds:mfc90_941
.text:64404FA9
.text:64404FA9 loc_64404FA9:
.text:64404FA9                mov     eax, [esi+38h]
.text:64404FAC                test    eax, eax
.text:64404FAE                jbe     short loc_64404FD3
.text:64404FB0                push    eax
.text:64404FB1                lea     eax, [ebp+var_14]
.text:64404FB4                push    offset aTraceLevelDAct ; "trace level %d activated\n"
.text:64404FB9                push    eax

; demangled name: ATL::CStringT::Format(char const *,...)
.text:64404FBA                call    ebx ; mfc90_2539
.text:64404FBC                add     esp, 0Ch
.text:64404FBF                lea     eax, [ebp+var_14]
.text:64404FC2                push    eax
.text:64404FC3                mov     ecx, edi

; demangled name: ATL::CStringT::operator+=(class ATL::CSimpleStringT<char, 1> const &)
.text:64404FC5                call    ds:mfc90_941

```

```

.text:64404FCB      xor     ebx, ebx
.text:64404FCD      inc     ebx
.text:64404FCE      mov     [ebp+var_10], ebx
.text:64404FD1      jmp     short loc_64404FD6
.text:64404FD3 ; -----
.text:64404FD3
.text:64404FD3 loc_64404FD3:
.text:64404FD3      xor     ebx, ebx
.text:64404FD5      inc     ebx
.text:64404FD6
.text:64404FD6 loc_64404FD6:
.text:64404FD6      cmp     [esi+38h], ebx
.text:64404FD9      jbe     short loc_64404FF1
.text:64404FDB      cmp     dword ptr [esi+2978h], 0
.text:64404FE2      jz      short loc_64404FF1
.text:64404FE4      push    offset aHexdumpInTrace ; "hexdump in trace activated\n"
.text:64404FE9      mov     ecx, edi

; demangled name: ATL::CStringT::operator+=(char const *)
.text:64404FEB      call    ds:mfc90_945
.text:64404FF1
.text:64404FF1 loc_64404FF1:
.text:64404FF1
.text:64404FF1      cmp     byte ptr [esi+78h], 0
.text:64404FF5      jz      short loc_64405007
.text:64404FF7      push    offset aLoggingActivat ; "logging activated\n"
.text:64404FFC      mov     ecx, edi

; demangled name: ATL::CStringT::operator+=(char const *)
.text:64404FFE      call    ds:mfc90_945
.text:64405004      mov     [ebp+var_10], ebx
.text:64405007
.text:64405007 loc_64405007:
.text:64405007      cmp     byte ptr [esi+3Dh], 0
.text:6440500B      jz      short bypass
.text:6440500D      push    offset aDataCompressio ; "data compression switched off\n"
.text:64405012      mov     ecx, edi

; demangled name: ATL::CStringT::operator+=(char const *)
.text:64405014      call    ds:mfc90_945
.text:6440501A      mov     [ebp+var_10], ebx
.text:6440501D
.text:6440501D bypass:
.text:6440501D      mov     eax, [esi+20h]
.text:64405020      test    eax, eax
.text:64405022      jz      short loc_6440503A
.text:64405024      cmp     dword ptr [eax+28h], 0
.text:64405028      jz      short loc_6440503A
.text:6440502A      push    offset aDataRecordMode ; "data record mode switched on\n"
.text:6440502F      mov     ecx, edi

; demangled name: ATL::CStringT::operator+=(char const *)
.text:64405031      call    ds:mfc90_945
.text:64405037      mov     [ebp+var_10], ebx
.text:6440503A
.text:6440503A loc_6440503A:
.text:6440503A
.text:6440503A      mov     ecx, edi
.text:6440503C      cmp     [ebp+var_10], ebx
.text:6440503F      jnz     loc_64405142
.text:64405045      push    offset aForMaximumData ; "\nFor maximum data security delete\nthe s
"..."
; demangled name: ATL::CStringT::operator+=(char const *)
.text:6440504A      call    ds:mfc90_945
.text:64405050      xor     edi, edi

```

```

.text:64405052      push     edi             ; fWinIni
.text:64405053      lea      eax, [ebp+pvParam]
.text:64405056      push     eax             ; pvParam
.text:64405057      push     edi             ; uiParam
.text:64405058      push     30h             ; uiAction
.text:6440505A      call     ds:SystemParametersInfoA
.text:64405060      mov      eax, [ebp+var_34]
.text:64405063      cmp      eax, 1600
.text:64405068      jle      short loc_64405072
.text:6440506A      cdq
.text:6440506B      sub      eax, edx
.text:6440506D      sar      eax, 1
.text:6440506F      mov      [ebp+var_34], eax
.text:64405072
.text:64405072 loc_64405072:
.text:64405072      push     edi             ; hWnd
.text:64405073      mov      [ebp+cy], 0A0h
.text:6440507A      call     ds:GetDC
.text:64405080      mov      [ebp+var_10], eax
.text:64405083      mov      ebx, 12Ch
.text:64405088      cmp      eax, edi
.text:6440508A      jz       loc_64405113
.text:64405090      push     11h             ; i
.text:64405092      call     ds:GetStockObject
.text:64405098      mov      edi, ds>SelectObject
.text:6440509E      push     eax             ; h
.text:6440509F      push     [ebp+var_10]     ; hdc
.text:644050A2      call     edi ; SelectObject
.text:644050A4      and      [ebp+rc.left], 0
.text:644050A8      and      [ebp+rc.top], 0
.text:644050AC      mov      [ebp+h], eax
.text:644050AF      push     401h            ; format
.text:644050B4      lea      eax, [ebp+rc]
.text:644050B7      push     eax             ; lprc
.text:644050B8      lea      ecx, [esi+2854h]
.text:644050BE      mov      [ebp+rc.right], ebx
.text:644050C1      mov      [ebp+rc.bottom], 0B4h

; demangled name: ATL::CSimpleStringT::GetLength(void)
.text:644050C8      call     ds:mfc90_3178
.text:644050CE      push     eax             ; cchText
.text:644050CF      lea      ecx, [esi+2854h]

; demangled name: const char* ATL::CSimpleStringT::operator PCXSTR
.text:644050D5      call     ds:mfc90_910
.text:644050DB      push     eax             ; lpchText
.text:644050DC      push     [ebp+var_10]     ; hdc
.text:644050DF      call     ds:DrawTextA
.text:644050E5      push     4               ; nIndex
.text:644050E7      call     ds:GetSystemMetrics
.text:644050ED      mov      ecx, [ebp+rc.bottom]
.text:644050F0      sub      ecx, [ebp+rc.top]
.text:644050F3      cmp      [ebp+h], 0
.text:644050F7      lea      eax, [eax+ecx+28h]
.text:644050FB      mov      [ebp+cy], eax
.text:644050FE      jz       short loc_64405108
.text:64405100      push     [ebp+h]         ; h
.text:64405103      push     [ebp+var_10]     ; hdc
.text:64405106      call     edi ; SelectObject
.text:64405108
.text:64405108 loc_64405108:
.text:64405108      push     [ebp+var_10]     ; hdc
.text:6440510B      push     0               ; hWnd
.text:6440510D      call     ds:ReleaseDC
.text:64405113
.text:64405113 loc_64405113:

```

```

.text:64405113      mov     eax, [ebp+var_38]
.text:64405116      push    80h          ; uFlags
.text:6440511B      push    [ebp+cy]      ; cy
.text:6440511E      inc     eax
.text:6440511F      push    ebx          ; cx
.text:64405120      push    eax          ; Y
.text:64405121      mov     eax, [ebp+var_34]
.text:64405124      add     eax, 0FFFFFFD4h
.text:64405129      cdq
.text:6440512A      sub     eax, edx
.text:6440512C      sar     eax, 1
.text:6440512E      push    eax          ; X
.text:6440512F      push    0            ; hWndInsertAfter
.text:64405131      push    dword ptr [esi+285Ch] ; hWnd
.text:64405137      call   ds:SetWindowPos
.text:6440513D      xor     ebx, ebx
.text:6440513F      inc     ebx
.text:64405140      jmp     short loc_6440514D
.text:64405142      ; -----
.text:64405142      loc_64405142:
.text:64405142      push    offset byte_64443AF8

; demangled name: ATL::CStringT::operator=(char const *)
.text:64405147      call   ds:mfc90_820
.text:6440514D      loc_6440514D:
.text:6440514D      cmp     dword_6450B970, ebx
.text:64405153      jl      short loc_64405188
.text:64405155      call   sub_6441C910
.text:6440515A      mov     dword_644F858C, ebx
.text:64405160      push    dword ptr [esi+2854h]
.text:64405166      push    offset aCdwsGuiPrepare ; "\nCDwsGui::PrepareInfoWindow: sapgui env
    "...
.text:6440516B      push    dword ptr [esi+2848h]
.text:64405171      call   dbg
.text:64405176      add     esp, 0Ch
.text:64405179      mov     dword_644F858C, 2
.text:64405183      call   sub_6441C920
.text:64405188      loc_64405188:
.text:64405188      or      [ebp+var_4], 0FFFFFFFh
.text:6440518C      lea     ecx, [ebp+var_14]

; demangled name: ATL::CStringT::~CStringT()
.text:6440518F      call   ds:mfc90_601
.text:64405195      call   __EH_epilog3
.text:6440519A      retn
.text:6440519A      CDwsGui__PrepareInfoWindow endp

```

ECX at function start gets pointer to object (because it is thiscall [2.5.4](#)-type of function). In our case, that object obviously has class type *CDwsGui*. Depends of option turned on in the object, specific message part will be concatenated to resulting message.

If value at this+0x3D address is not zero, compression is off:

```

.text:64405007      loc_64405007:
.text:64405007      cmp     byte ptr [esi+3Dh], 0
.text:6440500B      jz      short bypass
.text:6440500D      push    offset aDataCompressio ; "data compression switched off\n"
.text:64405012      mov     ecx, edi

; demangled name: ATL::CStringT::operator+=(char const *)
.text:64405014      call   ds:mfc90_945
.text:6440501A      mov     [ebp+var_10], ebx
.text:6440501D      bypass:

```

It is interesting, that finally, *var_10* variable state defines whether the message is to be shown at all:

```
.text:6440503C      cmp     [ebp+var_10], ebx
.text:6440503F      jnz     exit ; bypass drawing

; add strings "For maximum data security delete" / "the setting(s) as soon as possible !":

.text:64405045      push    offset aForMaximumData ; "\nFor maximum data security delete\nthe s
    "...
.text:6440504A      call    ds:mfc90_945 ; ATL::CStringT::operator+=(char const *)
.text:64405050      xor     edi, edi
.text:64405052      push    edi ; fWinIni
.text:64405053      lea     eax, [ebp+pvParam]
.text:64405056      push    eax ; pvParam
.text:64405057      push    edi ; uiParam
.text:64405058      push    30h ; uiAction
.text:6440505A      call    ds:SystemParametersInfoA
.text:64405060      mov     eax, [ebp+var_34]
.text:64405063      cmp     eax, 1600
.text:64405068      jle     short loc_64405072
.text:6440506A      cdq
.text:6440506B      sub     eax, edx
.text:6440506D      sar     eax, 1
.text:6440506F      mov     [ebp+var_34], eax
.text:64405072      loc_64405072:

start drawing:

.text:64405072      push    edi ; hWnd
.text:64405073      mov     [ebp+cy], 0A0h
.text:6440507A      call    ds:GetDC
```

Let's check our theory on practice.

JNZ at this line ...

```
.text:6440503F      jnz     exit ; bypass drawing
```

...replace it with just JMP, and get SAPGUI working without that nagging pop-up window appearing!

Now let's dig deeper and find connection between 0x15 offset in *load_command_line()* (I gave that name to that function) and *this+0x3D* variable in *CDwsGui::PrepareInfoWindow*. Are we sure that the value is the same?

I'm starting to search for all occurrences of 0x15 value in code. For some small programs like SAPGUI, it sometimes works. Here is the first occurrence I got:

```
.text:64404C19 sub_64404C19 proc near
.text:64404C19
.text:64404C19 arg_0 = dword ptr 4
.text:64404C19
.text:64404C19      push    ebx
.text:64404C1A      push    ebp
.text:64404C1B      push    esi
.text:64404C1C      push    edi
.text:64404C1D      mov     edi, [esp+10h+arg_0]
.text:64404C21      mov     eax, [edi]
.text:64404C23      mov     esi, ecx ; ESI/ECX are pointers to some unknown object.
.text:64404C25      mov     [esi], eax
.text:64404C27      mov     eax, [edi+4]
.text:64404C2A      mov     [esi+4], eax
.text:64404C2D      mov     eax, [edi+8]
.text:64404C30      mov     [esi+8], eax
.text:64404C33      lea     eax, [edi+0Ch]
.text:64404C36      push    eax
.text:64404C37      lea     ecx, [esi+0Ch]

; demangled name: ATL::CStringT::operator=(class ATL::CStringT ... &)
.text:64404C3A      call    ds:mfc90_817
```

```
.text:64404C40      mov     eax, [edi+10h]
.text:64404C43      mov     [esi+10h], eax
.text:64404C46      mov     al, [edi+14h]
.text:64404C49      mov     [esi+14h], al
.text:64404C4C      mov     al, [edi+15h] ; copy byte from 0x15 offset
.text:64404C4F      mov     [esi+15h], al ; to 0x15 offset in CDwsGui object
```

That function was called from the function named *CDwsGui::CopyOptions*! And thanks again for debugging information.

But the real answer in the function *CDwsGui::Init()*:

```
.text:6440B0BF loc_6440B0BF:
.text:6440B0BF      mov     eax, [ebp+arg_0]
.text:6440B0C2      push    [ebp+arg_4]
.text:6440B0C5      mov     [esi+2844h], eax
.text:6440B0CB      lea     eax, [esi+28h] ; ESI is pointer to CDwsGui object
.text:6440B0CE      push    eax
.text:6440B0CF      call    CDwsGui__CopyOptions
```

Finally, we understand: array filled in *load_command_line()* is actually placed in *CDwsGui* class but on this+0x28 address. 0x15 + 0x28 is exactly 0x3D. OK, we found the place where the value is copied to.

Let's also find other places where 0x3D offset is used. Here is one of them in *CDwsGui::SapguiRun* function (again, thanks to debugging calls):

```
.text:64409D58      cmp     [esi+3Dh], bl ; ESI is pointer to CDwsGui object
.text:64409D5B      lea     ecx, [esi+2B8h]
.text:64409D61      setz    al
.text:64409D64      push    eax ; arg_10 of CConnectionContext::CreateNetwork
.text:64409D65      push    dword ptr [esi+64h]

; demangled name: const char* ATL::CSimpleStringT::operator PCXSTR
.text:64409D68      call    ds:mfc90_910
.text:64409D68      ; no arguments
.text:64409D6E      push    eax
.text:64409D6F      lea     ecx, [esi+2BCh]

; demangled name: const char* ATL::CSimpleStringT::operator PCXSTR
.text:64409D75      call    ds:mfc90_910
.text:64409D75      ; no arguments
.text:64409D7B      push    eax
.text:64409D7C      push    esi
.text:64409D7D      lea     ecx, [esi+8]
.text:64409D80      call    CConnectionContext__CreateNetwork
```

Let's check our findings. Replace *setz al* here to *xor eax, eax / nop*, clear *TDW_NOCOMPRESS* environment variable and run *SAPGUI*. Wow! There is no nagging window (just as expected, because variable is not set), but in *Wireshark* we can see that the network packets are not compressed anymore! Obviously, this is the place where compression flag is to be set in *CConnectionContext* object.

So, compression flag is passed in the 5th argument of function *CConnectionContext::CreateNetwork*. Inside that function, another one is called:

```
...
.text:64403476      push    [ebp+compression]
.text:64403479      push    [ebp+arg_C]
.text:6440347C      push    [ebp+arg_8]
.text:6440347F      push    [ebp+arg_4]
.text:64403482      push    [ebp+arg_0]
.text:64403485      call    CNetwork__CNetwork
```

Compression flag is passing here in the 5th argument to *CNetwork::CNetwork* constructor.

And here is how *CNetwork* constructor sets some flag in *CNetwork* object according to the 5th argument **and** some another variable which probably could affect network packets compression too.

```
.text:64411DF1      cmp     [ebp+compression], esi
.text:64411DF7      jz      short set_EAX_to_0
```

```

.text:64411DF9      mov     al, [ebx+78h]    ; another value may affect compression?
.text:64411DFC      cmp     al, '3'
.text:64411DFE      jz      short set_EAX_to_1
.text:64411E00      cmp     al, '4'
.text:64411E02      jnz     short set_EAX_to_0
.text:64411E04      set_EAX_to_1:
.text:64411E04      xor     eax, eax
.text:64411E06      inc     eax             ; EAX -> 1
.text:64411E07      jmp     short loc_64411E0B
.text:64411E09      ; -----
.text:64411E09      set_EAX_to_0:
.text:64411E09
.text:64411E09      xor     eax, eax             ; EAX -> 0
.text:64411E0B
.text:64411E0B      loc_64411E0B:
.text:64411E0B      mov     [ebx+3A4h], eax ; EBX is pointer to CNetwork object

```

At this point we know that compression flag is stored in *CNetwork* class at *this+0x3A4* address.

Now let's dig across *SAPguilib.dll* for *0x3A4* value. And here is the second occurrence in *CDwsGui::OnClientMessageWrite* (endless thanks for debugging information):

```

.text:64406F76      loc_64406F76:
.text:64406F76      mov     ecx, [ebp+7728h+var_7794]
.text:64406F79      cmp     dword ptr [ecx+3A4h], 1
.text:64406F80      jnz     compression_flag_is_zero
.text:64406F86      mov     byte ptr [ebx+7], 1
.text:64406F8A      mov     eax, [esi+18h]
.text:64406F8D      mov     ecx, eax
.text:64406F8F      test    eax, eax
.text:64406F91      ja      short loc_64406FFF
.text:64406F93      mov     ecx, [esi+14h]
.text:64406F96      mov     eax, [esi+20h]
.text:64406F99      loc_64406F99:
.text:64406F99      push    dword ptr [edi+2868h] ; int
.text:64406F9F      lea     edx, [ebp+7728h+var_77A4]
.text:64406FA2      push    edx             ; int
.text:64406FA3      push    30000           ; int
.text:64406FA8      lea     edx, [ebp+7728h+Dst]
.text:64406FAB      push    edx             ; Dst
.text:64406FAC      push    ecx             ; int
.text:64406FAD      push    eax             ; Src
.text:64406FAE      push    dword ptr [edi+28C0h] ; int
.text:64406FB4      call    sub_644055C5     ; actual compression routine
.text:64406FB9      add     esp, 1Ch
.text:64406FBC      cmp     eax, 0FFFFFFFh
.text:64406FBF      jz      short loc_64407004
.text:64406FC1      cmp     eax, 1
.text:64406FC4      jz      loc_6440708C
.text:64406FCA      cmp     eax, 2
.text:64406FCD      jz      short loc_64407004
.text:64406FCF      push    eax
.text:64406FD0      push    offset aCompressionErr ; "compression error [rc = %d]- program wi
    "...
.text:64406FD5      push    offset aGui_err_compre ; "GUI_ERR_COMPRESS"
.text:64406FDA      push    dword ptr [edi+28D0h]
.text:64406FE0      call    SapPcTxtRead

```

Let's take a look into *sub_644055C5*. In it we can only see call to *memcpy()* and some other function named (by IDA 5) *sub_64417440*.

And, let's take a look inside *sub_64417440*. What we see is:

```

.text:6441747C      push    offset aErrorCsrrcompre ; "\nERROR: CsRCompress: invalid handle"
.text:64417481      call    eax ; dword_644F94C8

```

.text:64417483	add	esp, 4
----------------	-----	--------

Voilà! We've found the function which actually compresses data. As I [revealed in past](#), this function is used in SAP and also open-source MaxDB project. So it is available in sources.

Doing last check here:

.text:64406F79	cmp	dword ptr [ecx+3A4h], 1
.text:64406F80	jnz	compression_flag_is_zero

Replace JNZ here for unconditional JMP. Remove environment variable TDW_NOCOMPRESS. Voilà! In Wire-shark we see that client messages are not compressed. Server responses, however, are compressed.

So we found exact connection between environment variable and the point where data compression routine may be called or may be bypassed.

7.2.2 SAP 6.0 password checking functions

While returning again to my SAP 6.0 IDES installed in VMware box, I figured out I forgot the password for SAP* account, then it back to my memory, but now I got error message «*Password logon no longer possible - too many failed attempts*», because I've spent all these attempts in trying to recall it.

First extremely good news is that full *disp+work.pdb* file is supplied with SAP, it contain almost everything: function names, structures, types, local variable and argument names, etc. What a lavish gift!

I got TYPEINFODUMP⁷ utility for converting PDB files into something readable and grepable.

Here is an example of function information + its arguments + its local variables:

FUNCTION ThVmcSysEvent					
Address:	10143190	Size:	675 bytes	Index:	60483
Type:	int NEAR_C ThVmcSysEvent (unsigned int, unsigned char, unsigned short*)				
Flags:	0				
PARAMETER events					
Address:	Reg335+288	Size:	4 bytes	Index:	60488
Type:	unsigned int				
Flags:	d0				
PARAMETER opcode					
Address:	Reg335+296	Size:	1 bytes	Index:	60490
Type:	unsigned char				
Flags:	d0				
PARAMETER serverName					
Address:	Reg335+304	Size:	8 bytes	Index:	60492
Type:	unsigned short*				
Flags:	d0				
STATIC_LOCAL_VAR func					
Address:	12274af0	Size:	8 bytes	Index:	60495
Type:	wchar_t*				
Flags:	80				
LOCAL_VAR admhead					
Address:	Reg335+304	Size:	8 bytes	Index:	60498
Type:	unsigned char*				
Flags:	90				
LOCAL_VAR record					
Address:	Reg335+64	Size:	204 bytes	Index:	60501
Type:	AD_RECORD				
Flags:	90				
LOCAL_VAR adlen					
Address:	Reg335+296	Size:	4 bytes	Index:	60508
Type:	int				
Flags:	90				

And here is an example of some structure:

STRUCT DBSL_STMTID			
Size:	120	Variables:	4
Functions:	0	Base classes:	0
MEMBER moduletype			

⁷<http://www.debuginfo.com/tools/typeinfodump.html>


```

Type: DBSL_MODULETYPE
Offset:      0  Index:      3  TypeIndex:    38653
MEMBER module
Type: wchar_t module[40]
Offset:      4  Index:      3  TypeIndex:    831
MEMBER stmtnum
Type: long
Offset:     84  Index:      3  TypeIndex:    440
MEMBER timestamp
Type: wchar_t timestamp[15]
Offset:     88  Index:      3  TypeIndex:    6612

```

Wow!

Another good news is: *debugging* calls (there are plenty of them) are very useful.

Here you can also notice *ct_level* global variable⁸, reflecting current trace level.

There is a lot of such debugging inclusions in *disp+work.exe*:

```

cmp     cs:ct_level, 1
jl      short loc_1400375DA
call    DpLock
lea     rcx, aDpxxtool4_c ; "dpxxtool4.c"
mov     edx, 4Eh          ; line
call    CTrcSaveLocation
mov     r8, cs:func_48
mov     rcx, cs:hdl       ; hdl
lea     rdx, aSDpreadmemvalu ; "%s: DpReadMemValue (%d)"
mov     r9d, ebx
call    DpTrcErr
call    DpUnlock

```

If current trace level is bigger or equal to threshold defined in the code here, debugging message will be written to log files like *dev_w0*, *dev_disp*, and other *dev** files.

Let's do grepping on file I got with the help of TYPEINFODUMP utility:

```
cat "disp+work.pdb.d" | grep FUNCTION | grep -i password
```

I got:

```

FUNCTION rcui::AgiPassword::DiagISelection
FUNCTION ssf_password_encrypt
FUNCTION ssf_password_decrypt
FUNCTION password_logon_disabled
FUNCTION dySignSkipUserPassword
FUNCTION migrate_password_history
FUNCTION password_is_initial
FUNCTION rcui::AgiPassword::IsVisible
FUNCTION password_distance_ok
FUNCTION get_password_downwards_compatibility
FUNCTION dySignUnSkipUserPassword
FUNCTION rcui::AgiPassword::GetTypeNames
FUNCTION 'rcui::AgiPassword::AgiPassword'::'1'::dtor$2
FUNCTION 'rcui::AgiPassword::AgiPassword'::'1'::dtor$0
FUNCTION 'rcui::AgiPassword::AgiPassword'::'1'::dtor$1
FUNCTION usm_set_password
FUNCTION rcui::AgiPassword::TraceTo
FUNCTION days_since_last_password_change
FUNCTION rsecgrp_generate_random_password
FUNCTION rcui::AgiPassword::'scalar deleting destructor'
FUNCTION password_attempt_limit_exceeded
FUNCTION handle_incorrect_password
FUNCTION 'rcui::AgiPassword::'scalar deleting destructor'::'1'::dtor$1
FUNCTION calculate_new_password_hash
FUNCTION shift_password_to_history

```

⁸More about trace level: http://help.sap.com/saphelp_nwpi71/helpdata/en/46/962416a5a613e8e1000000a155369/content.htm

```

FUNCTION rcui::AgiPassword::GetType
FUNCTION found_password_in_history
FUNCTION 'rcui::AgiPassword::'scalar deleting destructor''::'1'::dtor$0
FUNCTION rcui::AgiObj::IsaPassword
FUNCTION password_idle_check
FUNCTION SlicHwPasswordForDay
FUNCTION rcui::AgiPassword::IsaPassword
FUNCTION rcui::AgiPassword::AgiPassword
FUNCTION delete_user_password
FUNCTION usm_set_user_password
FUNCTION Password_API
FUNCTION get_password_change_for_SSO
FUNCTION password_in_USR40
FUNCTION rsec_agrp_abap_generate_random_password

```

Let's also try to search for debug messages which contain words «password» and «locked». One of them is the string «user was locked by subsequently failed password logon attempts» referenced in function `password_attempt_limit_exceeded()`.

Other string this function I found may write to log file are: «password logon attempt will be rejected immediately (preventing dictionary attacks)», «failed-logon lock: expired (but not removed due to 'read-only' operation)», «failed-logon lock: expired => removed».

After playing for a little with this function, I quickly noticed that the problem is exactly in it. It is called from `chkpass()` function — one of the password checking functions.

First, I would like to be sure I'm at the correct point:

Run my *tracer* 5.0.1:

```
tracer64.exe -a:disp+work.exe bpf=disp+work.exe!chkpass,args:3,unicode
```

```

PID=2236|TID=2248|(0) disp+work.exe!chkpass (0x202c770, L"Brewered1", 0x41)
(called from 0x1402f1060 (disp+work.exe!usrexist+0x3c0))
PID=2236|TID=2248|(0) disp+work.exe!chkpass -> 0x35

```

Call path is: `syssigni()` -> `DylSigni()` -> `dychkusr()` -> `usrexist()` -> `chkpass()`.

Number 0x35 is error returning in `chkpass()` at that point:

```

.text:00000001402ED567 loc_1402ED567: ; CODE XREF: chkpass+B4
.text:00000001402ED567 mov rcx, rbx ; usr02
.text:00000001402ED56A call password_idle_check
.text:00000001402ED56F cmp eax, 33h
.text:00000001402ED572 jz loc_1402EDB4E
.text:00000001402ED578 cmp eax, 36h
.text:00000001402ED57B jz loc_1402EDB3D
.text:00000001402ED581 xor edx, edx ; usr02_readonly
.text:00000001402ED583 mov rcx, rbx ; usr02
.text:00000001402ED586 call password_attempt_limit_exceeded
.text:00000001402ED58B test al, al
.text:00000001402ED58D jz short loc_1402ED5A0
.text:00000001402ED58F mov eax, 35h
.text:00000001402ED594 add rsp, 60h
.text:00000001402ED598 pop r14
.text:00000001402ED59A pop r12
.text:00000001402ED59C pop rdi
.text:00000001402ED59D pop rsi
.text:00000001402ED59E pop rbx
.text:00000001402ED59F retn

```

Fine, let's check:

```
tracer64.exe -a:disp+work.exe bpf=disp+work.exe!password_attempt_limit_exceeded,args:4,unicode,rt:0
```

```

PID=2744|TID=360|(0) disp+work.exe!password_attempt_limit_exceeded (0x202c770, 0, 0x257758, 0) (called from
0x1402ed58b (disp+work.exe!chkpass+0xeb))
PID=2744|TID=360|(0) disp+work.exe!password_attempt_limit_exceeded -> 1
PID=2744|TID=360|We modify return value (EAX/RAX) of this function to 0

```

```
PID=2744|TID=360|(0) disp+work.exe!password_attempt_limit_exceeded (0x202c770, 0, 0, 0) (called from 0
x1402e9794 (disp+work.exe!chnypass+0xe4))
PID=2744|TID=360|(0) disp+work.exe!password_attempt_limit_exceeded -> 1
PID=2744|TID=360|We modify return value (EAX/RAX) of this function to 0
```

Excellent! I can successfully login now.

By the way, if I try to pretend I forgot the password, fixing *chkpass()* function return value at zero is enough to bypass check:

```
tracer64.exe -a:disp+work.exe bpf=disp+work.exe!chkpass,args:3,unicode,rt:0
```

```
PID=2744|TID=360|(0) disp+work.exe!chkpass (0x202c770, L"bogus", 0x41) (
called from 0x1402f1060 (disp+work.exe!usrexist+0x3c0))
PID=2744|TID=360|(0) disp+work.exe!chkpass -> 0x35
PID=2744|TID=360|We modify return value (EAX/RAX) of this function to 0
```

What also can be said while analyzing *password_attempt_limit_exceeded()* function is that at the very beginning of it, this call might be seen:

```
lea    rcx, aLoginFailed_us ; "login/failed_user_auto_unlock"
call   sapgparam
test   rax, rax
jz     short loc_1402E19DE
movzx  eax, word ptr [rax]
cmp    ax, 'N'
jz     short loc_1402E19D4
cmp    ax, 'n'
jz     short loc_1402E19D4
cmp    ax, '0'
jnz    short loc_1402E19DE
```

Obviously, function *sapgparam()* used to query value of some configuration parameter. This function can be called from 1768 different places. It seems, with the help of this information, we can easily find places in code, control flow of which can be affected by specific configuration parameters.

It is really sweet. Function names are very clear, much clearer than in Oracle RDBMS. It seems, *disp+work* process written in C++. It was apparently rewritten some time ago?

7.3 Oracle RDBMS

7.3.1 V\$VERSION table in Oracle RDBMS

Oracle RDBMS 11.2 is a huge program, main module *oracle.exe* contain approx. 124,000 functions. For comparison, Windows 7 x86 kernel (*ntoskrnl.exe*) — approx. 11,000 functions and Linux 3.9.8 kernel (with default drivers compiled) — 31,000 functions.

Let's start with an easy question. Where Oracle RDBMS get all this information, when we execute such simple statement in SQL*Plus:

```
SQL> select * from V$VERSION;
```

And we've got:

```
BANNER
-----
Oracle Database 11g Enterprise Edition Release 11.2.0.1.0 - Production
PL/SQL Release 11.2.0.1.0 - Production
CORE    11.2.0.1.0      Production
TNS for 32-bit Windows: Version 11.2.0.1.0 - Production
NLSRTL Version 11.2.0.1.0 - Production
```

Let's start. Where in Oracle RDBMS we may find a string *V\$VERSION*?

As of win32-version, `oracle.exe` file contain that string, that can be investigated easily. But we can also use object (.o) files from Linux version of Oracle RDBMS, because, unlike win32 version `oracle.exe`, function names (and global variables as well) are preserved there.

So, `kqf.o` file contain `V$VERSION` string. That object file is in main Oracle-library `libserver11.a`.

A reference to this text string we may find in `kqfviw` table stored in the same file, `kqf.o`:

Listing 7.1: `kqf.o`

```
.rodata:0800C4A0 kqfviw          dd 0Bh                ; DATA XREF: kqfchk:loc_8003A6D
.rodata:0800C4A0                ; kqfgbn+34
.rodata:0800C4A4                dd offset _2__STRING_10102_0 ; "GV$WAITSTAT"
.rodata:0800C4A8                dd 4
.rodata:0800C4AC                dd offset _2__STRING_10103_0 ; "NULL"
.rodata:0800C4B0                dd 3
.rodata:0800C4B4                dd 0
.rodata:0800C4B8                dd 195h
.rodata:0800C4BC                dd 4
.rodata:0800C4C0                dd 0
.rodata:0800C4C4                dd 0FFFFFFC1CBh
.rodata:0800C4C8                dd 3
.rodata:0800C4CC                dd 0
.rodata:0800C4D0                dd 0Ah
.rodata:0800C4D4                dd offset _2__STRING_10104_0 ; "V$WAITSTAT"
.rodata:0800C4D8                dd 4
.rodata:0800C4DC                dd offset _2__STRING_10103_0 ; "NULL"
.rodata:0800C4E0                dd 3
.rodata:0800C4E4                dd 0
.rodata:0800C4E8                dd 4Eh
.rodata:0800C4EC                dd 3
.rodata:0800C4F0                dd 0
.rodata:0800C4F4                dd 0FFFFFFC003h
.rodata:0800C4F8                dd 4
.rodata:0800C4FC                dd 0
.rodata:0800C500                dd 5
.rodata:0800C504                dd offset _2__STRING_10105_0 ; "GV$BH"
.rodata:0800C508                dd 4
.rodata:0800C50C                dd offset _2__STRING_10103_0 ; "NULL"
.rodata:0800C510                dd 3
.rodata:0800C514                dd 0
.rodata:0800C518                dd 269h
.rodata:0800C51C                dd 15h
.rodata:0800C520                dd 0
.rodata:0800C524                dd 0FFFFFFC1EDh
.rodata:0800C528                dd 8
.rodata:0800C52C                dd 0
.rodata:0800C530                dd 4
.rodata:0800C534                dd offset _2__STRING_10106_0 ; "V$BH"
.rodata:0800C538                dd 4
.rodata:0800C53C                dd offset _2__STRING_10103_0 ; "NULL"
.rodata:0800C540                dd 3
.rodata:0800C544                dd 0
.rodata:0800C548                dd 0F5h
.rodata:0800C54C                dd 14h
.rodata:0800C550                dd 0
.rodata:0800C554                dd 0FFFFFFC1EEh
.rodata:0800C558                dd 5
.rodata:0800C55C                dd 0
```

By the way, often, while analysing Oracle RDBMS internals, you may ask yourself, why functions and global variable names are so weird. Probably, because Oracle RDBMS is very old product and was developed in C in 1980-s. And that was a time when C standard guaranteed function names/variables support only up to 6 characters inclusive: «6 significant initial characters in an external identifier»⁹

Probably, the table `kqfviw` contain most (maybe even all) views prefixed with `V$`, these are *fixed views*, present

⁹Draft ANSI C Standard (ANSI X3J11/88-090) (May 13, 1988)

all the time. Superficially, by noticing cyclic recurrence of data, we can easily see that each `kqfviw` table element has 12 32-bit fields. It's very simple to create a 12-elements structure in IDA 5 and apply it to all table elements. As of Oracle RDBMS version 11.2, there are 1023 table elements, i.e., there are described 1023 of all possible *fixed views*. We will return to this number later.

As we can see, there are not much information in these numbers in fields. The very first number is always equals to name of view (without terminating zero. This is correct for each element. But this information is not very useful.

We also know that information about all fixed views can be retrieved from *fixed view* named `V$FIXED_VIEW_DEFINITION` (by the way, the information for this view is also taken from `kqfviw` and `kqfvip` tables.) By the way, there are 1023 elements too.

```
SQL> select * from V$FIXED_VIEW_DEFINITION where view_name='V$VERSION';
```

```
VIEW_NAME
```

```
VIEW_DEFINITION
```

```
V$VERSION
```

```
select BANNER from GV$VERSION where inst_id = USERENV('Instance')
```

So, `V$VERSION` is some kind of *thunk view* for another view, named `GV$VERSION`, which is, in turn:

```
SQL> select * from V$FIXED_VIEW_DEFINITION where view_name='GV$VERSION';
```

```
VIEW_NAME
```

```
VIEW_DEFINITION
```

```
GV$VERSION
```

```
select inst_id, banner from x$version
```

Tables prefixed as `X$` in Oracle RDBMS– is service tables too, undocumented, cannot be changed by user and refreshed dynamically.

Let's also try to search the text `select BANNER from GV$VERSION where inst_id = USERENV('Instance')` in `kqf.o` file and we find it in `kqfvip` table:

Listing 7.2: `kqf.o`

```
rodata:080185A0 kqfvip          dd offset _2__STRING_11126_0 ; DATA XREF: kqfgvcn+18
.rodata:080185A0                                ; kqfgvt+F
.rodata:080185A0                                ; "select inst_id,decode(indx,1,'data bloc"...
.rodata:080185A4          dd offset kqfv459_c_0
.rodata:080185A8          dd 0
.rodata:080185AC          dd 0

...

.rodata:08019570          dd offset _2__STRING_11378_0 ; "select  BANNER from GV$VERSION where in"...
.rodata:08019574          dd offset kqfv133_c_0
.rodata:08019578          dd 0
.rodata:0801957C          dd 0
.rodata:08019580          dd offset _2__STRING_11379_0 ; "select  inst_id,decode(bitand(cfflg,1),0)"...
.rodata:08019584          dd offset kqfv403_c_0
.rodata:08019588          dd 0
.rodata:0801958C          dd 0
.rodata:08019590          dd offset _2__STRING_11380_0 ; "select  STATUS , NAME, IS_RECOVERY_DEST"...
.rodata:08019594          dd offset kqfv199_c_0
```

The table appear to have 4 fields in each element. By the way, there are 1023 elements too. The second field pointing to another table, containing table fields for this *fixed view*. As of `V$VERSION`, this table contain only two elements, first is 6 and second is `BANNER` string (the number (6) is this string length) and after, *terminating* element contain 0 and *null* C-string:

Listing 7.3: kqf.o

```
.rodata:080BBAC4 kqfv133_c_0      dd 6                      ; DATA XREF: .rodata:08019574
.rodata:080BBAC8                  dd offset _2__STRING_5017_0 ; "BANNER"
.rodata:080BBACC                  dd 0
.rodata:080BBAD0                  dd offset _2__STRING_0_0
```

By joining data from both `kqfviw` and `kqfvip` tables, we may get SQL-statements which are executed when user wants to query information from specific *fixed view*.

So I wrote an oracle tables¹⁰ program, so to gather all this information from Oracle RDBMS for Linux object files. For `V$VERSION`, we may find this:

Listing 7.4: Result of oracle tables

```
kqfviw_element.viewname: [V$VERSION] ?: 0x3 0x43 0x1 0xffffc085 0x4
kqfvip_element.statement: [select BANNER from GV$VERSION where inst_id = USERENV('Instance')]
kqfvip_element.params:
[BANNER]
```

And:

Listing 7.5: Result of oracle tables

```
kqfviw_element.viewname: [GV$VERSION] ?: 0x3 0x26 0x2 0xffffc192 0x1
kqfvip_element.statement: [select inst_id, banner from x$version]
kqfvip_element.params:
[INST_ID] [BANNER]
```

`GV$VERSION` *fixed view* is distinct from `V$VERSION` in only that way that it contains one more field with *instance* identifier. Anyway, we stuck at the table `X$VERSION`. Just like any other `X$`-tables, it's undocumented, however, we can query it:

```
SQL> select * from x$version;

ADDR          INDX    INST_ID
-----
BANNER
-----

ODBAF574          0          1
Oracle Database 11g Enterprise Edition Release 11.2.0.1.0 - Production
...
```

This table has additional fields like `ADDR` and `INDX`.

While scrolling `kqf.o` in IDA 5 we may spot another table containing pointer to `X$VERSION` string, this is `kqf` tab:

Listing 7.6: kqf.o

```
.rodata:0803CAC0                  dd 9                      ; element number 0x1f6
.rodata:0803CAC4                  dd offset _2__STRING_13113_0 ; "X$VERSION"
.rodata:0803CAC8                  dd 4
.rodata:0803CACC                  dd offset _2__STRING_13114_0 ; "kqvt"
.rodata:0803CAD0                  dd 4
.rodata:0803CAD4                  dd 4
.rodata:0803CAD8                  dd 0
.rodata:0803CADC                  dd 4
.rodata:0803CAE0                  dd 0Ch
.rodata:0803CAE4                  dd 0FFFFFFC075h
.rodata:0803CAE8                  dd 3
.rodata:0803CAEC                  dd 0
.rodata:0803CAF0                  dd 7
.rodata:0803CAF4                  dd offset _2__STRING_13115_0 ; "X$KQFSZ"
.rodata:0803CAF8                  dd 5
.rodata:0803CAFC                  dd offset _2__STRING_13116_0 ; "kqfsz"
.rodata:0803CB00                  dd 1
```

¹⁰<http://yurichev.com/oracle-tables.html>

```
.rodata:0803CB04      dd 38h
.rodata:0803CB08      dd 0
.rodata:0803CB0C      dd 7
.rodata:0803CB10      dd 0
.rodata:0803CB14      dd 0FFFFFFC09Dh
.rodata:0803CB18      dd 2
.rodata:0803CB1C      dd 0
```

There are a lot of references to X\$-table names, apparently, to all Oracle RDBMS 11.2 X\$-tables. But again, we have not enough information. I have no idea, what kqvt string mean. kq prefix may mean *kernal* and *query*. v, apparently, mean *version* and t — *type*? Frankly speaking, I don't know.

The table named similarly can be found in kqf.o:

Listing 7.7: kqf.o

```
.rodata:0808C360 kqvt_c_0      kqftap_param <4, offset _2__STRING_19_0, 917h, 0, 0, 0, 4, 0, 0>
.rodata:0808C360      ; DATA XREF: .rodata:08042680
.rodata:0808C360      ; "ADDR"
.rodata:0808C384      kqftap_param <4, offset _2__STRING_20_0, 0B02h, 0, 0, 0, 4, 0, 0> ; "INDX"
.rodata:0808C3A8      kqftap_param <7, offset _2__STRING_21_0, 0B02h, 0, 0, 0, 4, 0, 0> ; "
INST_ID"
.rodata:0808C3CC      kqftap_param <6, offset _2__STRING_5017_0, 601h, 0, 0, 0, 50h, 0, 0> ; "
BANNER"
.rodata:0808C3F0      kqftap_param <0, offset _2__STRING_0_0, 0, 0, 0, 0, 0, 0, 0, 0>
```

It contain information about all fields in the X\$VERSION table. The only reference to this table present in kqftap table:

Listing 7.8: kqf.o

```
.rodata:08042680      kqftap_element <0, offset kqvt_c_0, offset kqvrow, 0> ; element 0x1f6
```

It's interesting that this element here is *0x1f6th* (502nd), just as a pointer to X\$VERSION string in kqftab table. Probably, kqftap and kqftab tables are complement each other, just like kqfvip and kqfviv. We also see a pointer to kqvrow() function. Something useful at last!

So I added these tables to my oracle tables¹¹ utility too. For X\$VERSION I've got:

Listing 7.9: Result of oracle tables

```
kqftab_element.name: [X$VERSION] ?: [kqvt] 0x4 0x4 0x4 0xc 0xffffc075 0x3
kqftap_param.name=[ADDR] ?: 0x917 0x0 0x0 0x0 0x4 0x0 0x0
kqftap_param.name=[INDX] ?: 0xb02 0x0 0x0 0x0 0x4 0x0 0x0
kqftap_param.name=[INST_ID] ?: 0xb02 0x0 0x0 0x0 0x4 0x0 0x0
kqftap_param.name=[BANNER] ?: 0x601 0x0 0x0 0x0 0x50 0x0 0x0
kqftap_element.fn1=kqvrow
kqftap_element.fn2=NULL
```

Using *tracer* 5.0.1, it's easy to check that this function called 6 times in row (from the qerfxFetch() function) while querying X\$VERSION table.

Let's run *tracer* 5.0.1 in cc mode (it will comment each executed instruction):

```
tracer -a:oracle.exe bpf=oracle.exe!_kqvrow,trace:cc
```

```
_kqvrow_      proc near
var_7C        = byte ptr -7Ch
var_18        = dword ptr -18h
var_14        = dword ptr -14h
Dest          = dword ptr -10h
var_C         = dword ptr -0Ch
var_8         = dword ptr -8
var_4         = dword ptr -4
arg_8         = dword ptr 10h
arg_C         = dword ptr 14h
```

¹¹http://yurichev.com/oracle_tables.html

```
arg_14      = dword ptr 1Ch
arg_18      = dword ptr 20h
```

```
; FUNCTION CHUNK AT .text1:056C11A0 SIZE 00000049 BYTES
```

```
push    ebp
mov     ebp, esp
sub     esp, 7Ch
mov     eax, [ebp+arg_14] ; [EBP+1Ch]=1
mov     ecx, TlsIndex    ; [69AEB08h]=0
mov     edx, large fs:2Ch
mov     edx, [edx+ecx*4] ; [EDX+ECX*4]=0xc98c938
cmp     eax, 2           ; EAX=1
mov     eax, [ebp+arg_8] ; [EBP+10h]=0xcdfe554
jz      loc_2CE1288
mov     ecx, [eax]       ; [EAX]=0..5
mov     [ebp+var_4], edi ; EDI=0xc98c938
```

```
loc_2CE10F6: ; CODE XREF: _kqvrow_+10A
```

```
; _kqvrow_+1A9
cmp     ecx, 5           ; ECX=0..5
ja      loc_56C11C7
mov     edi, [ebp+arg_18] ; [EBP+20h]=0
mov     [ebp+var_14], edx ; EDX=0xc98c938
mov     [ebp+var_8], ebx ; EBX=0
mov     ebx, eax         ; EAX=0xcdfe554
mov     [ebp+var_C], esi ; ESI=0xcdfe248
```

```
loc_2CE110D: ; CODE XREF: _kqvrow_+29E00E6
```

```
mov     edx, ds:off_628B09C[ecx*4] ; [ECX*4+628B09Ch]=0x2ce1116, 0x2ce11ac, 0x2ce11db, 0
x2ce11f6, 0x2ce1236, 0x2ce127a
jmp     edx              ; EDX=0x2ce1116, 0x2ce11ac, 0x2ce11db, 0x2ce11f6, 0x2ce1236, 0
x2ce127a
```

```
; -----
```

```
loc_2CE1116: ; DATA XREF: .rdata:off_628B09C
```

```
push    offset aXKqvvsBuffer ; "x$kqvvs buffer"
mov     ecx, [ebp+arg_C] ; [EBP+14h]=0x8a172b4
xor     edx, edx
mov     esi, [ebp+var_14] ; [EBP-14h]=0xc98c938
push    edx              ; EDX=0
push    edx              ; EDX=0
push    50h
push    ecx              ; ECX=0x8a172b4
push    dword ptr [esi+10494h] ; [ESI+10494h]=0xc98cd58
call    _kgghalf         ; tracing nested maximum level (1) reached, skipping this CALL
mov     esi, ds:__imp__vsnum ; [59771A8h]=0x61bc49e0
mov     [ebp+Dest], eax ; EAX=0xce2ffb0
mov     [ebx+8], eax     ; EAX=0xce2ffb0
mov     [ebx+4], eax     ; EAX=0xce2ffb0
mov     edi, [esi]       ; [ESI]=0xb200100
mov     esi, ds:__imp__vsenstr ; [597D6D4h]=0x65852148, "- Production"
push    esi              ; ESI=0x65852148, "- Production"
mov     ebx, edi         ; EDI=0xb200100
shr     ebx, 18h         ; EBX=0xb200100
mov     ecx, edi         ; EDI=0xb200100
shr     ecx, 14h         ; ECX=0xb200100
and     ecx, 0Fh         ; ECX=0xb2
mov     edx, edi         ; EDI=0xb200100
shr     edx, 0Ch         ; EDX=0xb200100
movzx   edx, dl          ; DL=0
mov     eax, edi         ; EDI=0xb200100
shr     eax, 8           ; EAX=0xb200100
and     eax, 0Fh         ; EAX=0xb2001
and     edi, 0FFh        ; EDI=0xb200100
push    edi              ; EDI=0
```



```

        mov     edi, [ebp+arg_18] ; [EBP+20h]=0
        push    eax                ; EAX=1
        mov     eax, ds:__imp__vsnb ; [597D6D8h]=0x65852100, "Oracle Database 11g Enterprise
Edition Release %d.%d.%d.%d.%d %s"
        push    edx                ; EDX=0
        push    ecx                ; ECX=2
        push    ebx                ; EBX=0xb
        mov     ebx, [ebp+arg_8] ; [EBP+10h]=0xcdfe554
        push    eax                ; EAX=0x65852100, "Oracle Database 11g Enterprise Edition Release %d
.%d.%d.%d.%d %s"
        mov     eax, [ebp+Dest] ; [EBP-10h]=0xce2ffb0
        push    eax                ; EAX=0xce2ffb0
        call    ds:__imp__sprintf ; op1=MSVCR80.dll!sprintf tracing nested maximum level (1) reached
, skipping this CALL
        add     esp, 38h
        mov     dword ptr [ebx], 1

loc_2CE1192:                                ; CODE XREF: _kqvrow_+FB
                                                ; _kqvrow_+128 ...
        test    edi, edi            ; EDI=0
        jnz     __VInfreq__kqvrow
        mov     esi, [ebp+var_C] ; [EBP-0Ch]=0xcdfe248
        mov     edi, [ebp+var_4] ; [EBP-4]=0xc98c938
        mov     eax, ebx            ; EBX=0xcdfe554
        mov     ebx, [ebp+var_8] ; [EBP-8]=0
        lea     eax, [eax+4] ; [EAX+4]=0xce2ffb0, "NLSRTL Version 11.2.0.1.0 - Production", "
Oracle Database 11g Enterprise Edition Release 11.2.0.1.0 - Production", "PL/SQL Release 11.2.0.1.0 -
Production", "TNS for 32-bit Windows: Version 11.2.0.1.0 - Production"

loc_2CE11A8:                                ; CODE XREF: _kqvrow_+29E00F6
        mov     esp, ebp
        pop     ebp
        retn                                ; EAX=0xcdfe558
; -----

loc_2CE11AC:                                ; DATA XREF: .rdata:0628B0A0
        mov     edx, [ebx+8] ; [EBX+8]=0xce2ffb0, "Oracle Database 11g Enterprise Edition Release
11.2.0.1.0 - Production"
        mov     dword ptr [ebx], 2
        mov     [ebx+4], edx ; EDX=0xce2ffb0, "Oracle Database 11g Enterprise Edition Release
11.2.0.1.0 - Production"
        push    edx                ; EDX=0xce2ffb0, "Oracle Database 11g Enterprise Edition Release
11.2.0.1.0 - Production"
        call    _kxvsn             ; tracing nested maximum level (1) reached, skipping this CALL
        pop     ecx
        mov     edx, [ebx+4] ; [EBX+4]=0xce2ffb0, "PL/SQL Release 11.2.0.1.0 - Production"
        movzx   ecx, byte ptr [edx] ; [EDX]=0x50
        test    ecx, ecx            ; ECX=0x50
        jnz     short loc_2CE1192
        mov     edx, [ebp+var_14]
        mov     esi, [ebp+var_C]
        mov     eax, ebx
        mov     ebx, [ebp+var_8]
        mov     ecx, [eax]
        jmp     loc_2CE10F6
; -----

loc_2CE11DB:                                ; DATA XREF: .rdata:0628B0A4
        push    0
        push    50h
        mov     edx, [ebx+8] ; [EBX+8]=0xce2ffb0, "PL/SQL Release 11.2.0.1.0 - Production"
        mov     [ebx+4], edx ; EDX=0xce2ffb0, "PL/SQL Release 11.2.0.1.0 - Production"
        push    edx                ; EDX=0xce2ffb0, "PL/SQL Release 11.2.0.1.0 - Production"
        call    _lmxver            ; tracing nested maximum level (1) reached, skipping this CALL
        add     esp, 0Ch
        mov     dword ptr [ebx], 3

```

```

        jmp     short loc_2CE1192
; -----
loc_2CE11F6:
        ; DATA XREF: .rdata:0628B0A8
        mov     edx, [ebx+8] ; [EBX+8]=0xce2ffb0
        mov     [ebp+var_18], 50h
        mov     [ebx+4], edx ; EDX=0xce2ffb0
        push    0
        call    _npinli ; tracing nested maximum level (1) reached, skipping this CALL
        pop     ecx
        test    eax, eax ; EAX=0
        jnz     loc_56C11DA
        mov     ecx, [ebp+var_14] ; [EBP-14h]=0xc98c938
        lea     edx, [ebp+var_18] ; [EBP-18h]=0x50
        push    edx ; EDX=0xd76c93c
        push    dword ptr [ebx+8] ; [EBX+8]=0xce2ffb0
        push    dword ptr [ecx+13278h] ; [ECX+13278h]=0xacce190
        call    _nrtnsvrs ; tracing nested maximum level (1) reached, skipping this CALL
        add     esp, 0Ch

loc_2CE122B:
        ; CODE XREF: _kqvrow_+29E0118
        mov     dword ptr [ebx], 4
        jmp     loc_2CE1192
; -----
loc_2CE1236:
        ; DATA XREF: .rdata:0628B0AC
        lea     edx, [ebp+var_7C] ; [EBP-7Ch]=1
        push    edx ; EDX=0xd76c8d8
        push    0
        mov     esi, [ebx+8] ; [EBX+8]=0xce2ffb0, "TNS for 32-bit Windows: Version 11.2.0.1.0 -
Production"
        mov     [ebx+4], esi ; ESI=0xce2ffb0, "TNS for 32-bit Windows: Version 11.2.0.1.0 -
Production"
        mov     ecx, 50h
        mov     [ebp+var_18], ecx ; ECX=0x50
        push    ecx ; ECX=0x50
        push    esi ; ESI=0xce2ffb0, "TNS for 32-bit Windows: Version 11.2.0.1.0 -
Production"
        call    _lxvers ; tracing nested maximum level (1) reached, skipping this CALL
        add     esp, 10h
        mov     edx, [ebp+var_18] ; [EBP-18h]=0x50
        mov     dword ptr [ebx], 5
        test    edx, edx ; EDX=0x50
        jnz     loc_2CE1192
        mov     edx, [ebp+var_14]
        mov     esi, [ebp+var_C]
        mov     eax, ebx
        mov     ebx, [ebp+var_8]
        mov     ecx, 5
        jmp     loc_2CE10F6
; -----
loc_2CE127A:
        ; DATA XREF: .rdata:0628B0B0
        mov     edx, [ebp+var_14] ; [EBP-14h]=0xc98c938
        mov     esi, [ebp+var_C] ; [EBP-0Ch]=0xcdfe248
        mov     edi, [ebp+var_4] ; [EBP-4]=0xc98c938
        mov     eax, ebx ; EBX=0xcdfe554
        mov     ebx, [ebp+var_8] ; [EBP-8]=0

loc_2CE1288:
        ; CODE XREF: _kqvrow_+1F
        mov     eax, [eax+8] ; [EAX+8]=0xce2ffb0, "NLSRTL Version 11.2.0.1.0 - Production"
        test    eax, eax ; EAX=0xce2ffb0, "NLSRTL Version 11.2.0.1.0 - Production"
        jz      short loc_2CE12A7
        push    offset aXKqvvsnBuffer ; "x$kqvvsn buffer"
        push    eax ; EAX=0xce2ffb0, "NLSRTL Version 11.2.0.1.0 - Production"
        mov     eax, [ebp+arg_C] ; [EBP+14h]=0x8a172b4

```

```

        push    eax                ; EAX=0x8a172b4
        push    dword ptr [edx+10494h] ; [EDX+10494h]=0xc98cd58
        call    _kghfrf            ; tracing nested maximum level (1) reached, skipping this CALL
        add     esp, 10h

loc_2CE12A7:
        xor     eax, eax            ; CODE XREF: _kqvrow_+1C1
        mov     esp, ebp
        pop     ebp
        retn    4                  ; EAX=0
_kqvrow_    endp

```

Now it's easy to see that row number is passed from outside of function. The function returning the string constructing it as follows:

String 1	Using vsnstr, vsnnum, vsnban global variables. Calling sprintf().
String 2	Calling kkvsn().
String 3	Calling lmxver().
String 4	Calling npinli(), nrtnsvrs().
String 5	Calling lxvers().

That's how corresponding functions are called for determining each module's version.

7.3.2 X\$KSMLRU table in Oracle RDBMS

There is a mention of some special table in the *Diagnosing and Resolving Error ORA-04031 on the Shared Pool or Other Memory Pools [Video] [ID 146599.1]* note:

There is a fixed table called X\$KSMLRU that tracks allocations in the shared pool that cause other objects in the shared pool to be aged out. This fixed table can be used to identify what is causing the large allocation.

If many objects are being periodically flushed from the shared pool then this will cause response time problems and will likely cause library cache latch contention problems when the objects are reloaded into the shared pool.

One unusual thing about the X\$KSMLRU fixed table is that the contents of the fixed table are erased whenever someone selects from the fixed table. This is done since the fixed table stores only the largest allocations that have occurred. The values are reset after being selected so that subsequent large allocations can be noted even if they were not quite as large as others that occurred previously. Because of this resetting, the output of selecting from this table should be carefully kept since it cannot be retrieved back after the query is issued.

However, as it can be easily checked, this table's contents is cleared each time table querying. Are we able to find why? Let's back to tables we already know: kqftab and kqftap which were generated with oracle tables¹² help, containing all information about X\$-tables, now we can see here, the ksmlrs() function is called to prepare this table's elements:

Listing 7.10: Result of oracle tables

```

kqftab_element.name: [X$KSMLRU] ?: [ksmlr] 0x4 0x64 0x11 0xc 0xffffc0bb 0x5
kqftap_param.name=[ADDR] ?: 0x917 0x0 0x0 0x0 0x4 0x0 0x0
kqftap_param.name=[INDX] ?: 0xb02 0x0 0x0 0x0 0x4 0x0 0x0
kqftap_param.name=[INST_ID] ?: 0xb02 0x0 0x0 0x0 0x4 0x0 0x0
kqftap_param.name=[KSMLRIDX] ?: 0xb02 0x0 0x0 0x0 0x4 0x0 0x0
kqftap_param.name=[KSMLRDUR] ?: 0xb02 0x0 0x0 0x0 0x4 0x4 0x0
kqftap_param.name=[KSMLRSHRPOOL] ?: 0xb02 0x0 0x0 0x0 0x4 0x8 0x0

```

¹²<http://yurichev.com/oracle-tables.html>

```

kqftap_param.name=[KSMLRCOM] ?: 0x501 0x0 0x0 0x0 0x14 0xc 0x0
kqftap_param.name=[KSMLRSIZ] ?: 0x2 0x0 0x0 0x0 0x4 0x20 0x0
kqftap_param.name=[KSMLRNUM] ?: 0x2 0x0 0x0 0x0 0x4 0x24 0x0
kqftap_param.name=[KSMLRHON] ?: 0x501 0x0 0x0 0x0 0x20 0x28 0x0
kqftap_param.name=[KSMLROHV] ?: 0xb02 0x0 0x0 0x0 0x4 0x48 0x0
kqftap_param.name=[KSMLRSES] ?: 0x17 0x0 0x0 0x0 0x4 0x4c 0x0
kqftap_param.name=[KSMLRADU] ?: 0x2 0x0 0x0 0x0 0x4 0x50 0x0
kqftap_param.name=[KSMLRNID] ?: 0x2 0x0 0x0 0x0 0x4 0x54 0x0
kqftap_param.name=[KSMLRNSD] ?: 0x2 0x0 0x0 0x0 0x4 0x58 0x0
kqftap_param.name=[KSMLRNCD] ?: 0x2 0x0 0x0 0x0 0x4 0x5c 0x0
kqftap_param.name=[KSMLRNED] ?: 0x2 0x0 0x0 0x0 0x4 0x60 0x0
kqftap_element.fn1=ksmlrs
kqftap_element.fn2=NULL

```

Indeed, with the *tracer* 5.0.1 help it's easy to see that this function is called each time we query the X\$KSMLRU table.

Here we see a references to `ksmsplu_sp()` and `ksmsplu_jp()` functions, each of them call `ksmsplu()` finally. At the end of `ksmsplu()` function we see a call to `memset()`:

Listing 7.11: ksm.o

```

...
.text:00434C50 loc_434C50:                                ; DATA XREF: .rdata:off_5E50EA8
.text:00434C50      mov     edx, [ebp-4]
.text:00434C53      mov     [eax], esi
.text:00434C55      mov     esi, [edi]
.text:00434C57      mov     [eax+4], esi
.text:00434C5A      mov     [edi], eax
.text:00434C5C      add     edx, 1
.text:00434C5F      mov     [ebp-4], edx
.text:00434C62      jnz     loc_434B7D
.text:00434C68      mov     ecx, [ebp+14h]
.text:00434C6B      mov     ebx, [ebp-10h]
.text:00434C6E      mov     esi, [ebp-0Ch]
.text:00434C71      mov     edi, [ebp-8]
.text:00434C74      lea     eax, [ecx+8Ch]
.text:00434C7A      push    370h                ; Size
.text:00434C7F      push    0                   ; Val
.text:00434C81      push    eax                 ; Dst
.text:00434C82      call    __intel_fast_memset
.text:00434C87      add     esp, 0Ch
.text:00434C8A      mov     esp, ebp
.text:00434C8C      pop     ebp
.text:00434C8D      retn
.text:00434C8D _ksmsplu      endp

```

Constructions like `memset (block, 0, size)` are often used just to zero memory block. What if we would take a risk, block `memset()` call and see what will happen?

Let's run *tracer* 5.0.1 with the following options: set breakpoint at 0x434C7A (the point where `memset()` arguments are to be passed), thus, that *tracer* 5.0.1 set program counter EIP at this place to the place where passed to `memset()` arguments are to be cleared (at 0x434C8A) It can be said, we just simulate an unconditional jump from the address 0x434C7A to 0x434C8A.

```
tracer -a:oracle.exe bpx=oracle.exe!0x00434C7A,set(eip,0x00434C8A)
```

(Important: all these addresses are valid only for win32-version of Oracle RDBMS 11.2)

Indeed, now we can query X\$KSMLRU table as many times as we want and it's not cleared anymore!

Don't try this at home ("MythBusters") Don't try this on your production servers.

It's probably not a very useful or desired system behaviour, but as an experiment of locating piece of code we need, that's perfectly suit our needs!

7.3.3 V\$TIMER table in Oracle RDBMS

V\$TIMER is another *fixed view*, reflecting some often changed value:

V\$TIMER displays the elapsed time in hundredths of a second. Time is measured since the beginning of the epoch, which is operating system specific, and wraps around to 0 again whenever the value overflows four bytes (roughly 497 days).

(From Oracle RDBMS documentation ¹³)

It's interesting that periods are different for Oracle for win32 and for Linux. Will we able to find a function generating this value?

As we can see, this information is finally taken from X\$KSUTM table.

```
SQL> select * from V$FIXED_VIEW_DEFINITION where view_name='V$TIMER';

VIEW_NAME
-----
VIEW_DEFINITION
-----

V$TIMER
select  HSECS from GV$TIMER where inst_id = USERENV('Instance')

SQL> select * from V$FIXED_VIEW_DEFINITION where view_name='GV$TIMER';

VIEW_NAME
-----
VIEW_DEFINITION
-----

GV$TIMER
select inst_id,ksutmtim from x$ksutm
```

Now we stuck in problem, there are no references to value generating function(s) in the tables kqftab/kqftap:

Listing 7.12: Result of oracle tables

```
kqftab_element.name: [X$KSUTM] ?: [ksutm] 0x1 0x4 0x4 0x0 0xffffc09b 0x3
kqftap_param.name=[ADDR] ?: 0x10917 0x0 0x0 0x0 0x4 0x0 0x0
kqftap_param.name=[INDX] ?: 0x20b02 0x0 0x0 0x0 0x4 0x0 0x0
kqftap_param.name=[INST_ID] ?: 0xb02 0x0 0x0 0x0 0x4 0x0 0x0
kqftap_param.name=[KSUTMTIM] ?: 0x1302 0x0 0x0 0x0 0x4 0x0 0x1e
kqftap_element.fn1=NULL
kqftap_element.fn2=NULL
```

Let's try to find a string KSUTMTIM, and we find it in this function:

```
kqfd_DRN_ksutm_c proc near                                ; DATA XREF: .rodata:0805B4E8

arg_0              = dword ptr  8
arg_8              = dword ptr 10h
arg_C              = dword ptr 14h

                push    ebp
                mov     ebp, esp
                push    [ebp+arg_C]
                push    offset ksugtm
                push    offset _2__STRING_1263_0 ; "KSUTMTIM"
                push    [ebp+arg_8]
                push    [ebp+arg_0]
```

¹³http://docs.oracle.com/cd/B28359_01/server.111/b28320/dynviews_3104.htm

```

        call    kqfd_cfui_drain
        add     esp, 14h
        mov     esp, ebp
        pop     ebp
        retn
kqfd_DRN_ksutm_c endp

```

The function `kqfd_DRN_ksutm_c()` is mentioned in `kqfd_tab_registry_0` table:

```

dd offset _2__STRING_62_0 ; "X$KSUTM"
dd offset kqfd_OPN_ksutm_c
dd offset kqfd_tabl_fetch
dd 0
dd 0
dd offset kqfd_DRN_ksutm_c

```

There are also some function `ksugtm()` referenced here. Let's see what's in it (Linux x86):

Listing 7.13: `ksu.o`

```

ksugtm      proc near
var_1C      = byte ptr -1Ch
arg_4       = dword ptr  0Ch

        push    ebp
        mov     ebp, esp
        sub     esp, 1Ch
        lea     eax, [ebp+var_1C]
        push    eax
        call    slgcs
        pop     ecx
        mov     edx, [ebp+arg_4]
        mov     [edx], eax
        mov     eax, 4
        mov     esp, ebp
        pop     ebp
        retn
ksugtm      endp

```

Almost the same code in win32-version.

Is this the function we are looking for? Let's see:

```

tracer -a:oracle.exe bpf=oracle.exe!_ksugtm,args:2,dump_args:0x4

```

Let's try again:

```

SQL> select * from V$TIMER;

HSECS
-----
27294929

SQL> select * from V$TIMER;

HSECS
-----
27295006

SQL> select * from V$TIMER;

HSECS
-----
27295167

```

Output of *tracer* 5.0.1:

```

TID=2428|(0) oracle.exe!_ksugtm (0x0, 0xd76c5f0) (called from oracle.exe!__VInfreq__qerfxFetch+0xfad (0
x56bb6d5))
Argument 2/2
0D76C5F0: 38 C9 "8. "
TID=2428|(0) oracle.exe!_ksugtm () -> 0x4 (0x4)
Argument 2/2 difference
00000000: D1 7C A0 01 ".|.. "
TID=2428|(0) oracle.exe!_ksugtm (0x0, 0xd76c5f0) (called from oracle.exe!__VInfreq__qerfxFetch+0xfad (0
x56bb6d5))
Argument 2/2
0D76C5F0: 38 C9 "8. "
TID=2428|(0) oracle.exe!_ksugtm () -> 0x4 (0x4)
Argument 2/2 difference
00000000: 1E 7D A0 01 ".}.. "
TID=2428|(0) oracle.exe!_ksugtm (0x0, 0xd76c5f0) (called from oracle.exe!__VInfreq__qerfxFetch+0xfad (0
x56bb6d5))
Argument 2/2
0D76C5F0: 38 C9 "8. "
TID=2428|(0) oracle.exe!_ksugtm () -> 0x4 (0x4)
Argument 2/2 difference
00000000: BF 7D A0 01 ".}.. "

```

Indeed — the value is the same we see in SQL*Plus and it's returning via second argument.

Let's see what is in `slgcs()` (Linux x86):

```

slgcs      proc near
var_4      = dword ptr -4
arg_0      = dword ptr 8

        push    ebp
        mov     ebp, esp
        push    esi
        mov     [ebp+var_4], ebx
        mov     eax, [ebp+arg_0]
        call    $+5
        pop     ebx
        nop     ; PIC mode
        mov     ebx, offset _GLOBAL_OFFSET_TABLE_
        mov     dword ptr [eax], 0
        call    sltrgtime64 ; PIC mode
        push    0
        push    0Ah
        push    edx
        push    eax
        call    __udivdi3 ; PIC mode
        mov     ebx, [ebp+var_4]
        add     esp, 10h
        mov     esp, ebp
        pop     ebp
        retn
slgcs      endp

```

(it's just a call to `sltrgtime64()` and division of its result by 10 [1.12](#))

And win32-version:

```

_slgcs     proc near ; CODE XREF: _dbgefghTtElResetCount+15
; _dbgerRunActions+1528
        db      66h
        nop
        push    ebp
        mov     ebp, esp
        mov     eax, [ebp+8]
        mov     dword ptr [eax], 0
        call    ds:__imp__GetTickCount@0 ; GetTickCount()
        mov     edx, eax

```

```

        mov     eax, 0CCCCCDh
        mul     edx
        shr     edx, 3
        mov     eax, edx
        mov     esp, ebp
        pop     ebp
        retn
_slgcs   endp

```

It's just result of `GetTickCount()`¹⁴ divided by 10^{1.12}.

Voilà! That's why win32-version and Linux x86 version show different results, just because they are generated by different operation system functions.

Drain apparently mean *connecting* specific table column to specific function.

I added the table `kqfd_tab_registry_0` to oracle tables¹⁵, now we can see, how table column's variables are *connected* to specific functions:

```

[X$KSUTM] [kqfd_OPN_ksutm_c] [kqfd_tabl_fetch] [NULL] [NULL] [kqfd_DRN_ksutm_c]
[X$KSUSGIF] [kqfd_OPN_ksusg_c] [kqfd_tabl_fetch] [NULL] [NULL] [kqfd_DRN_ksusg_c]

```

OPN, apparently, *open*, and *DRN*, apparently, meaning *drain*.

¹⁴[http://msdn.microsoft.com/en-us/library/windows/desktop/ms724408\(v=vs.85\).aspx](http://msdn.microsoft.com/en-us/library/windows/desktop/ms724408(v=vs.85).aspx)

¹⁵http://yurichev.com/oracle_tables.html

Chapter 8

Other things

8.1 Compiler's anomalies

Intel C++ 10.1, which was used for Oracle RDBMS 11.2 Linux86 compilation, may emit two JZ in row, and there are no references to the second JZ. Second JZ is thus senseless.

Listing 8.1: kdli.o from libserver11.a

```
.text:08114CF1          loc_8114CF1:                                ; CODE XREF: __PGOSF539_kdlimemSer
+89A
.text:08114CF1                                ; __PGOSF539_kdlimemSer+3994
.text:08114CF1 8B 45 08          mov     eax, [ebp+arg_0]
.text:08114CF4 0F B6 50 14      movzx   edx, byte ptr [eax+14h]
.text:08114CF8 F6 C2 01          test    dl, 1
.text:08114CFB 0F 85 17 08 00 00 jnz     loc_8115518
.text:08114D01 85 C9            test    ecx, ecx
.text:08114D03 0F 84 8A 00 00 00 jz      loc_8114D93
.text:08114D09 0F 84 09 08 00 00 jz      loc_8115518
.text:08114D0F 8B 53 08          mov     edx, [ebx+8]
.text:08114D12 89 55 FC          mov     [ebp+var_4], edx
.text:08114D15 31 C0            xor     eax, eax
.text:08114D17 89 45 F4          mov     [ebp+var_C], eax
.text:08114D1A 50              push    eax
.text:08114D1B 52              push    edx
.text:08114D1C E8 03 54 00 00      call    len2nbytes
.text:08114D21 83 C4 08          add     esp, 8
```

Listing 8.2: from the same code

```
.text:0811A2A5          loc_811A2A5:                                ; CODE XREF: kdliSerLengths+11C
.text:0811A2A5                                ; kdliSerLengths+1C1
.text:0811A2A5 8B 7D 08          mov     edi, [ebp+arg_0]
.text:0811A2A8 8B 7F 10          mov     edi, [edi+10h]
.text:0811A2AB 0F B6 57 14      movzx   edx, byte ptr [edi+14h]
.text:0811A2AF F6 C2 01          test    dl, 1
.text:0811A2B2 75 3E            jnz     short loc_811A2F2
.text:0811A2B4 83 E0 01          and     eax, 1
.text:0811A2B7 74 1F            jz      short loc_811A2D8
.text:0811A2B9 74 37            jz      short loc_811A2F2
.text:0811A2BB 6A 00            push    0
.text:0811A2BD FF 71 08          push    dword ptr [ecx+8]
.text:0811A2C0 E8 5F FE FF FF      call    len2nbytes
```

It's probably code generator bug wasn't found by tests, because, resulting code is working correctly anyway.

Another compiler anomaly I described here [1.15.2](#).

I'm showing such cases here, so to understand that such compilers errors are possible and sometimes one should not to rack one's brain and think why compiler generated such strange code.

Chapter 9

Tasks solutions

9.1 Easy level

9.1.1 Task 1.1

Solution: toupper().

C source code:

```
char toupper ( char c )
{
    if( c >= 'a' && c <= 'z' ) {
        c = c - 'a' + 'A';
    }
    return( c );
}
```

9.1.2 Task 1.2

Solution: atoi()

C source code:

```
#include <stdio.h>
#include <string.h>
#include <ctype.h>

int atoi ( const *p ) /* convert ASCII string to integer */
{
    int i;
    char s;

    while( isspace ( *p ) )
        ++p;
    s = *p;
    if( s == '+' || s == '-' )
        ++p;
    i = 0;
    while( isdigit(*p) ) {
        i = i * 10 + *p - '0';
        ++p;
    }
    if( s == '-' )
        i = - i;
    return( i );
}
```

9.1.3 Task 1.3

Solution: srand() / rand().

C source code:

```
static unsigned int v;

void srand (unsigned int s)
{
    v = s;
}

int rand ()
{
    return( ((v = v * 214013L
            + 2531011L) >> 16) & 0x7fff );
}
```

9.1.4 Task 1.4

Solution: strstr().

C source code:

```
char * strstr (
    const char * str1,
    const char * str2
)
{
    char *cp = (char *) str1;
    char *s1, *s2;

    if ( !*str2 )
        return((char *)str1);

    while (*cp)
    {
        s1 = cp;
        s2 = (char *) str2;

        while ( *s1 && *s2 && !(*s1-*s2) )
            s1++, s2++;

        if (!*s2)
            return(cp);

        cp++;
    }

    return(NULL);
}
```

9.1.5 Task 1.5

Hint #1: Keep in mind that __v — global variable.

Hint #2: That function is called in startup code, before main() execution.

Solution: early Pentium CPU FDIV bug checking¹.

C source code:

```
unsigned _v; // _v

enum e {
    PROB_P5_DIV = 0x0001
};
```

¹http://en.wikipedia.org/wiki/Pentium_FDIV_bug

```

void f( void ) // __verify_pentium_fdiv_bug
{
    /*
     * Verify we have got the Pentium FDIV problem.
     * The volatiles are to scare the optimizer away.
     */
    volatile double    v1    = 4195835;
    volatile double    v2    = 3145727;

    if( (v1 - (v1/v2)*v2) > 1.0e-8 ) {
        _v |= PROB_P5_DIV;
    }
}

```

9.1.6 Task 1.6

Hint: it might be helpful to google a constant used here.

Solution: TEA encryption algorithm².

C source code (taken from http://en.wikipedia.org/wiki/Tiny_Encryption_Algorithm):

```

void f (unsigned int* v, unsigned int* k) {
    unsigned int v0=v[0], v1=v[1], sum=0, i;           /* set up */
    unsigned int delta=0x9e3779b9;                     /* a key schedule constant */
    unsigned int k0=k[0], k1=k[1], k2=k[2], k3=k[3];   /* cache key */
    for (i=0; i < 32; i++) {                           /* basic cycle start */
        sum += delta;
        v0 += ((v1<<4) + k0) ^ (v1 + sum) ^ ((v1>>5) + k1);
        v1 += ((v0<<4) + k2) ^ (v0 + sum) ^ ((v0>>5) + k3);
    }                                                    /* end cycle */
    v[0]=v0; v[1]=v1;
}

```

9.1.7 Task 1.7

Hint: the table contain pre-calculated values. It's possible to implement the function without it, but it will work slower, though.

Solution: this function reverse all bits in input 32-bit integer. It's lib/bitrev.c from Linux kernel.

C source code:

```

const unsigned char byte_rev_table[256] = {
    0x00, 0x80, 0x40, 0xc0, 0x20, 0xa0, 0x60, 0xe0,
    0x10, 0x90, 0x50, 0xd0, 0x30, 0xb0, 0x70, 0xf0,
    0x08, 0x88, 0x48, 0xc8, 0x28, 0xa8, 0x68, 0xe8,
    0x18, 0x98, 0x58, 0xd8, 0x38, 0xb8, 0x78, 0xf8,
    0x04, 0x84, 0x44, 0xc4, 0x24, 0xa4, 0x64, 0xe4,
    0x14, 0x94, 0x54, 0xd4, 0x34, 0xb4, 0x74, 0xf4,
    0x0c, 0x8c, 0x4c, 0xcc, 0x2c, 0xac, 0x6c, 0xec,
    0x1c, 0x9c, 0x5c, 0xdc, 0x3c, 0xbc, 0x7c, 0xfc,
    0x02, 0x82, 0x42, 0xc2, 0x22, 0xa2, 0x62, 0xe2,
    0x12, 0x92, 0x52, 0xd2, 0x32, 0xb2, 0x72, 0xf2,
    0x0a, 0x8a, 0x4a, 0xca, 0x2a, 0xaa, 0x6a, 0xea,
    0x1a, 0x9a, 0x5a, 0xda, 0x3a, 0xba, 0x7a, 0xfa,
    0x06, 0x86, 0x46, 0xc6, 0x26, 0xa6, 0x66, 0xe6,
    0x16, 0x96, 0x56, 0xd6, 0x36, 0xb6, 0x76, 0xf6,
    0x0e, 0x8e, 0x4e, 0xce, 0x2e, 0xae, 0x6e, 0xee,
    0x1e, 0x9e, 0x5e, 0xde, 0x3e, 0xbe, 0x7e, 0xfe,
    0x01, 0x81, 0x41, 0xc1, 0x21, 0xa1, 0x61, 0xe1,
    0x11, 0x91, 0x51, 0xd1, 0x31, 0xb1, 0x71, 0xf1,
    0x09, 0x89, 0x49, 0xc9, 0x29, 0xa9, 0x69, 0xe9,
    0x19, 0x99, 0x59, 0xd9, 0x39, 0xb9, 0x79, 0xf9,
}

```

²Tiny Encryption Algorithm

```

    0x05, 0x85, 0x45, 0xc5, 0x25, 0xa5, 0x65, 0xe5,
    0x15, 0x95, 0x55, 0xd5, 0x35, 0xb5, 0x75, 0xf5,
    0x0d, 0x8d, 0x4d, 0xcd, 0x2d, 0xad, 0x6d, 0xed,
    0x1d, 0x9d, 0x5d, 0xdd, 0x3d, 0xbd, 0x7d, 0xfd,
    0x03, 0x83, 0x43, 0xc3, 0x23, 0xa3, 0x63, 0xe3,
    0x13, 0x93, 0x53, 0xd3, 0x33, 0xb3, 0x73, 0xf3,
    0x0b, 0x8b, 0x4b, 0xcb, 0x2b, 0xab, 0x6b, 0xeb,
    0x1b, 0x9b, 0x5b, 0xdb, 0x3b, 0xbb, 0x7b, 0xfb,
    0x07, 0x87, 0x47, 0xc7, 0x27, 0xa7, 0x67, 0xe7,
    0x17, 0x97, 0x57, 0xd7, 0x37, 0xb7, 0x77, 0xf7,
    0x0f, 0x8f, 0x4f, 0xcf, 0x2f, 0xaf, 0x6f, 0xef,
    0x1f, 0x9f, 0x5f, 0xdf, 0x3f, 0xbf, 0x7f, 0xff,
};

unsigned char bitrev8(unsigned char byte)
{
    return byte_rev_table[byte];
}

unsigned short bitrev16(unsigned short x)
{
    return (bitrev8(x & 0xff) << 8) | bitrev8(x >> 8);
}

/**
 * bitrev32 - reverse the order of bits in a unsigned int value
 * @x: value to be bit-reversed
 */
unsigned int bitrev32(unsigned int x)
{
    return (bitrev16(x & 0xffff) << 16) | bitrev16(x >> 16);
}

```

9.1.8 Task 1.8

Solution: two 100*200 matrices of *double* type addition.

C/C++ source code:

```

#define M    100
#define N    200

void s(double *a, double *b, double *c)
{
    for(int i=0;i<N;i++)
        for(int j=0;j<M;j++)
            *(c+i*M+j)=*(a+i*M+j) + *(b+i*M+j);
};

```

9.1.9 Task 1.9

Solution: two matrices (one is 100*200, second is 100*300) of *double* type multiplication, result: 100*300 matrix.

C/C++ source code:

```

#define M    100
#define N    200
#define P    300

void m(double *a, double *b, double *c)
{
    for(int i=0;i<M;i++)
        for(int j=0;j<P;j++)
        {

```

```

        *(c+i*M+j)=0;
        for (int k=0;k<N;k++) *(c+i*M+j)+=*(a+i*M+j) * *(b+i*M+j);
    }
};

```

9.2 Middle level

9.2.1 Task 2.1

Hint #1: The code has one characteristic thing, considering it, it may help narrowing search of right function among glibc functions.

Solution: characteristic — is callback-function calling [1.19](#), pointer to which is passed in 4th argument. It's quicksort().

C source code:

```

/* Copyright (C) 1991,1992,1996,1997,1999,2004 Free Software Foundation, Inc.
   This file is part of the GNU C Library.
   Written by Douglas C. Schmidt (schmidt@ics.uci.edu).

   The GNU C Library is free software; you can redistribute it and/or
   modify it under the terms of the GNU Lesser General Public
   License as published by the Free Software Foundation; either
   version 2.1 of the License, or (at your option) any later version.

   The GNU C Library is distributed in the hope that it will be useful,
   but WITHOUT ANY WARRANTY; without even the implied warranty of
   MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU
   Lesser General Public License for more details.

   You should have received a copy of the GNU Lesser General Public
   License along with the GNU C Library; if not, write to the Free
   Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA
   02111-1307 USA. */

/* If you consider tuning this algorithm, you should consult first:
   Engineering a sort function; Jon Bentley and M. Douglas McIlroy;
   Software - Practice and Experience; Vol. 23 (11), 1249-1265, 1993. */

#include <alloca.h>
#include <limits.h>
#include <stdlib.h>
#include <string.h>

typedef int (*__compar_d_fn_t) (__const void *, __const void *, void *);

/* Byte-wise swap two items of size SIZE. */
#define SWAP(a, b, size) \
do \
{ \
    register size_t __size = (size); \
    register char *__a = (a), *__b = (b); \
    do \
    { \
        char __tmp = *__a; \
        *__a++ = *__b; \
        *__b++ = __tmp; \
    } while (--__size > 0); \
} while (0)

/* Discontinue quicksort algorithm when partition gets below this size.
   This particular magic number was chosen to work best on a Sun 4/260. */
#define MAX_THRESH 4

```

```

/* Stack node declarations used to store unfulfilled partition obligations. */
typedef struct
{
    char *lo;
    char *hi;
} stack_node;

/* The next 4 #defines implement a very fast in-line stack abstraction. */
/* The stack needs log (total_elements) entries (we could even subtract
log(MAX_THRESH)). Since total_elements has type size_t, we get as
upper bound for log (total_elements):
bits per byte (CHAR_BIT) * sizeof(size_t). */
#define STACK_SIZE      (CHAR_BIT * sizeof(size_t))
#define PUSH(low, high) ((void) ((top->lo = (low)), (top->hi = (high)), ++top))
#define POP(low, high)  ((void) (--top, (low = top->lo), (high = top->hi)))
#define STACK_NOT_EMPTY (stack < top)

/* Order size using quicksort. This implementation incorporates
four optimizations discussed in Sedgewick:

1. Non-recursive, using an explicit stack of pointer that store the
next array partition to sort. To save time, this maximum amount
of space required to store an array of SIZE_MAX is allocated on the
stack. Assuming a 32-bit (64 bit) integer for size_t, this needs
only 32 * sizeof(stack_node) == 256 bytes (for 64 bit: 1024 bytes).
Pretty cheap, actually.

2. Chose the pivot element using a median-of-three decision tree.
This reduces the probability of selecting a bad pivot value and
eliminates certain extraneous comparisons.

3. Only quicksorts TOTAL_ELEMS / MAX_THRESH partitions, leaving
insertion sort to order the MAX_THRESH items within each partition.
This is a big win, since insertion sort is faster for small, mostly
sorted array segments.

4. The larger of the two sub-partitions is always pushed onto the
stack first, with the algorithm then concentrating on the
smaller partition. This *guarantees* no more than log (total_elems)
stack size is needed (actually O(1) in this case)! */

void
_quicksort (void *const pbase, size_t total_elems, size_t size,
            __compar_d_fn_t cmp, void *arg)
{
    register char *base_ptr = (char *) pbase;

    const size_t max_thresh = MAX_THRESH * size;

    if (total_elems == 0)
        /* Avoid lossage with unsigned arithmetic below. */
        return;

    if (total_elems > MAX_THRESH)
    {
        char *lo = base_ptr;
        char *hi = &lo[size * (total_elems - 1)];
        stack_node stack[STACK_SIZE];
        stack_node *top = stack;

        PUSH (NULL, NULL);

        while (STACK_NOT_EMPTY)
        {
            char *left_ptr;

```

```

char *right_ptr;

/* Select median value from among LO, MID, and HI. Rearrange
   LO and HI so the three values are sorted. This lowers the
   probability of picking a pathological pivot value and
   skips a comparison for both the LEFT_PTR and RIGHT_PTR in
   the while loops. */

char *mid = lo + size * ((hi - lo) / size >> 1);

if ((*cmp) ((void *) mid, (void *) lo, arg) < 0)
    SWAP (mid, lo, size);
if ((*cmp) ((void *) hi, (void *) mid, arg) < 0)
    SWAP (mid, hi, size);
else
    goto jump_over;
if ((*cmp) ((void *) mid, (void *) lo, arg) < 0)
    SWAP (mid, lo, size);
jump_over::

left_ptr  = lo + size;
right_ptr = hi - size;

/* Here's the famous "collapse the walls" section of quicksort.
   Gotta like those tight inner loops!  They are the main reason
   that this algorithm runs much faster than others. */
do
{
    while ((*cmp) ((void *) left_ptr, (void *) mid, arg) < 0)
        left_ptr += size;

    while ((*cmp) ((void *) mid, (void *) right_ptr, arg) < 0)
        right_ptr -= size;

    if (left_ptr < right_ptr)
    {
        SWAP (left_ptr, right_ptr, size);
        if (mid == left_ptr)
            mid = right_ptr;
        else if (mid == right_ptr)
            mid = left_ptr;
        left_ptr += size;
        right_ptr -= size;
    }
    else if (left_ptr == right_ptr)
    {
        left_ptr += size;
        right_ptr -= size;
        break;
    }
}
while (left_ptr <= right_ptr);

/* Set up pointers for next iteration.  First determine whether
   left and right partitions are below the threshold size.  If so,
   ignore one or both.  Otherwise, push the larger partition's
   bounds on the stack and continue sorting the smaller one. */

if ((size_t) (right_ptr - lo) <= max_thresh)
{
    if ((size_t) (hi - left_ptr) <= max_thresh)
        /* Ignore both small partitions. */
        POP (lo, hi);
    else
        /* Ignore small left partition. */
        lo = left_ptr;
}

```



```

    }
    else if ((size_t) (hi - left_ptr) <= max_thresh)
        /* Ignore small right partition. */
        hi = right_ptr;
    else if ((right_ptr - lo) > (hi - left_ptr))
    {
        /* Push larger left partition indices. */
        PUSH (lo, right_ptr);
        lo = left_ptr;
    }
    else
    {
        /* Push larger right partition indices. */
        PUSH (left_ptr, hi);
        hi = right_ptr;
    }
}

/* Once the BASE_PTR array is partially sorted by quicksort the rest
is completely sorted using insertion sort, since this is efficient
for partitions below MAX_THRESH size. BASE_PTR points to the beginning
of the array to sort, and END_PTR points at the very last element in
the array (*not* one beyond it!). */

#define min(x, y) ((x) < (y) ? (x) : (y))

{
    char *const end_ptr = &base_ptr[size * (total_elems - 1)];
    char *tmp_ptr = base_ptr;
    char *thresh = min(end_ptr, base_ptr + max_thresh);
    register char *run_ptr;

    /* Find smallest element in first threshold and place it at the
    array's beginning. This is the smallest array element,
    and the operation speeds up insertion sort's inner loop. */

    for (run_ptr = tmp_ptr + size; run_ptr <= thresh; run_ptr += size)
        if ((*cmp) ((void *) run_ptr, (void *) tmp_ptr, arg) < 0)
            tmp_ptr = run_ptr;

    if (tmp_ptr != base_ptr)
        SWAP (tmp_ptr, base_ptr, size);

    /* Insertion sort, running from left-hand-side up to right-hand-side. */

    run_ptr = base_ptr + size;
    while ((run_ptr += size) <= end_ptr)
    {
        tmp_ptr = run_ptr - size;
        while ((*cmp) ((void *) run_ptr, (void *) tmp_ptr, arg) < 0)
            tmp_ptr -= size;

        tmp_ptr += size;
        if (tmp_ptr != run_ptr)
        {
            char *trav;

            trav = run_ptr + size;
            while (--trav >= run_ptr)
            {
                char c = *trav;
                char *hi, *lo;

                for (hi = lo = trav; (lo -= size) >= tmp_ptr; hi = lo)
                    *hi = *lo;
            }
        }
    }
}

```

```
        *hi = c;  
    }  
}  
}
```

Afterword

9.3 Donate

This book is free, available freely and available in source code form³ (LaTeX), and it will be so forever.

If you want to support my work, so that I could continue to add things to it regularly, you may consider donation.

You may donate by sending small (or not small) donation to bitcoin⁴
1HRGTRdFNH1cE81zxWQg6jTtkLzAiGU9Lp



Figure 9.1: bitcoin address

Other ways to donate are available on the page: <http://yurichev.com/donate.html>
Major benefactors will be mentioned right here.

9.4 Questions?

Do not hesitate to mail any questions to the author: <dennis@yurichev.com>

Please, also do not hesitate to send me any corrections (including grammar ones (you see how horrible my English is?)), etc.

³<https://github.com/dennis714/RE-for-beginners>

⁴<http://en.wikipedia.org/wiki/Bitcoin>

Bibliography

- [App10] Apple. iOS ABI Function Call Guide. 2010. Also available as <http://developer.apple.com/library/ios/documentation/Xcode/Conceptual/iPhoneOSABIReference/iPhoneOSABIReference.pdf>.
- [Cli] Marshall Cline. C++ faq. Also available as <http://www.parashift.com/c++-faq-lite/index.html>.
- [ISO07] ISO. ISO/IEC 9899:TC3 (C C99 standard). 2007. Also available as <http://www.open-std.org/jtc1/sc22/WG14/www/docs/n1256.pdf>.
- [ISO13] ISO. ISO/IEC 14882:2011 (C++ 11 standard). 2013. Also available as <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2013/n3690.pdf>.
- [Ker88] Brian W. Kernighan. The C Programming Language. Prentice Hall Professional Technical Reference, 2nd edition, 1988.
- [Knu98] Donald E. Knuth. The Art of Computer Programming Volumes 1-3 Boxed Set. Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA, 2nd edition, 1998.
- [Loh10] Eugene Loh. The ideal hpc programming language. Queue, 8(6):30:30–30:38, June 2010.
- [Ltd94] Advanced RISC Machines Ltd. The ARM Cookbook. 1994. Also available as [http://yurichev.com/ref/ARM%20Cookbook%20\(1994\)](http://yurichev.com/ref/ARM%20Cookbook%20(1994)).
- [Ray03] Eric S. Raymond. The Art of UNIX Programming. Pearson Education, 2003. Also available as <http://catb.org/esr/writings/taoup/html/>.
- [Rit86] Dennis M. Ritchie. Where did ++ come from? (net.lang.c). http://yurichev.com/mirrors/c_dmr_postincrement.txt, 1986. [Online; accessed 2013].
- [Rit93] Dennis M. Ritchie. The development of the c language. SIGPLAN Not., 28(3):201–208, March 1993. Also available as <http://yurichev.com/mirrors/dmr-The%20Development%20of%20the%20C%20Language-1993.pdf>.
- [War02] Henry S. Warren. Hacker's Delight. Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA, 2002.

Index

C language elements

- Pointers, [18](#), [20](#), [29](#), [133](#), [147](#)
- Post-decrement, [53](#)
- Post-increment, [53](#)
- Pre-decrement, [53](#)
- Pre-increment, [53](#)
- alloca(), [10](#), [79](#)
- assert(), [169](#)
- C99
 - bool, [82](#)
 - restrict, [154](#)
 - variable length arrays, [79](#)
- calloc(), [204](#)
- const, [2](#), [22](#)
- for, [44](#), [93](#)
- if, [31](#), [36](#)
- longjmp(), [37](#)
- malloc(), [98](#)
- memcpy(), [169](#)
- memcpy(), [18](#)
- memset(), [244](#)
- qsort(), [133](#)
- restrict, [154](#)
- return, [3](#), [23](#), [28](#)
- scanf, [18](#)
- strlen(), [49](#), [143](#)
- switch, [35–37](#)
- tolower(), [210](#)
- while, [49](#)

Compiler’s anomalies, [88](#), [249](#)

grep usage, [68](#), [168](#), [170](#), [172](#), [233](#)

Dynamically loaded libraries, [7](#)

Global variables, [20](#)

Linker, [22](#), [115](#)

RISC pipeline, [34](#)

Non-a-numbers (NaNs), [65](#)

Buffer overflow, [75](#)

Out-of-order execution, [16](#)

position-independent code, [4](#), [165](#)

RAM, [22](#)

ROM, [22](#)

Recursion, [8](#), [159](#)

- Tail recursion, [159](#)

Stack, [8](#), [25](#), [36](#)

Stack overflow, [8](#)

Stack frame, [19](#)

Syntactic Sugar, [36](#), [102](#)

iPod/iPhone/iPad, [4](#)

8080, [52](#)

Angry Birds, [68](#)

ARM, [52](#)

ARM mode, [4](#)

Pipeline, [42](#)

Mode switching, [27](#), [43](#)

Adressing modes, [53](#)

mode switching, [7](#)

Instructions

ADD, [6](#), [33](#), [47](#), [55](#), [91](#)

ADDAL, [33](#)

ADDCC, [41](#)

ADDS, [27](#), [37](#)

ADR, [4](#), [33](#)

ADREQ, [33](#), [37](#)

ADRGT, [33](#)

ADRHI, [33](#)

ADRNE, [37](#)

ASRS, [56](#), [88](#)

B, [13](#), [33](#), [34](#)

BCS, [34](#), [69](#)

BEQ, [24](#), [37](#)

BGE, [34](#)

BIC, [87](#)

BL, [5](#), [7](#), [33](#)

BLE, [34](#)

BLEQ, [33](#)

BLGT, [33](#)

BLHI, [33](#)

BLS, [34](#)

BLT, [47](#)

BLX, [7](#)

BNE, [34](#)

BX, [27](#), [43](#)

CMP, [24](#), [33](#), [37](#), [41](#), [47](#), [91](#)

IDIV, [54](#)

IT, [67](#), [78](#)

LDMCSFD, [33](#)

LDMEA, [8](#)

- LDMED, [8](#)
- LDMFA, [8](#)
- LDMFD, [5](#), [8](#), [33](#)
- LDMGEFD, [33](#)
- LDR, [15](#), [20](#), [72](#)
- LDR.W, [81](#)
- LDRB, [106](#)
- LDRB.W, [53](#)
- LDRSB, [52](#)
- LSL, [91](#)
- LSL.W, [91](#)
- LSLS, [73](#)
- MLA, [27](#)
- MOV, [5](#), [55](#), [91](#)
- MOVT, [6](#), [55](#)
- MOVT.W, [7](#)
- MOVW, [7](#)
- MULS, [27](#)
- MVNS, [53](#)
- ORR, [87](#)
- POP, [4](#), [5](#), [8](#), [9](#)
- PUSH, [8](#), [9](#)
- RSB, [81](#), [91](#)
- SMMUL, [55](#)
- STMEA, [8](#)
- STMED, [8](#)
- STMFA, [8](#), [16](#)
- STMFD, [4](#), [8](#)
- STMIA, [15](#)
- STMIB, [16](#)
- STR, [14](#), [72](#)
- SUB, [14](#), [81](#), [91](#)
- SUBEQ, [53](#)
- SXTB, [106](#)
- TEST, [50](#)
- TST, [85](#), [91](#)
- VADD, [59](#)
- VDIV, [59](#)
- VLDR, [59](#)
- VMOV, [59](#), [67](#)
- VMOVGT, [67](#)
- VMRS, [67](#)
- VMUL, [59](#)
- Registers
 - APSR, [67](#)
 - FPSCR, [67](#)
 - Link Register, [5](#), [9](#), [13](#), [43](#)
 - R0, [28](#)
 - scratch registers, [52](#)
 - Z, [24](#)
- thumb mode, [4](#), [34](#), [43](#)
- thumb-2 mode, [4](#), [43](#), [67](#), [68](#)
- armel, [60](#)
- armhf, [60](#)
- Condition codes, [33](#)
- D-registers, [59](#)
- Data processing instructions, [55](#)
- DCB, [5](#)
- hard float, [60](#)
- if-then block, [67](#)
- Leaf function, [9](#)
- Optional operators
 - ASR, [55](#), [91](#)
 - LSL, [72](#), [81](#), [91](#)
 - LSR, [55](#), [91](#)
 - ROR, [91](#)
 - RRX, [91](#)
- S-registers, [59](#)
- soft float, [60](#)
- BASIC
 - POKE, [172](#)
- C++, [235](#)
 - References, [30](#)
- Callbacks, [133](#)
- Canary, [77](#)
- cdecl, [12](#), [163](#)
- column-major order, [79](#)
- Compiler intrinsic, [10](#)
- CRC32, [91](#)
- DES, [137](#), [147](#)
- DosBox, [172](#)
- double, [57](#), [164](#)
- ELF, [21](#)
- Error messages, [169](#)
- fastcall, [83](#), [163](#)
- float, [57](#), [111](#), [164](#)
- FORTRAN, [79](#), [154](#)
- Function epilogue, [13](#), [15](#), [33](#), [106](#), [159](#), [171](#)
- Function prologue, [3](#), [9](#), [14](#), [77](#), [159](#), [171](#)
- Fused multiply-add, [27](#)
- GDB, [76](#)
- IDA
 - var_?, [15](#), [20](#)
- IEEE 754, [57](#), [111](#), [131](#)
- Inline code, [48](#), [87](#), [120](#)
- Intel C++, [2](#), [137](#), [249](#)
- jumptable, [39](#), [43](#)
- Keil, [4](#)
- Linux, [235](#)
 - libc.so.6, [83](#), [136](#)
- LLVM, [4](#)

- long double, [57](#)
- Loop unwinding, [46](#)
- MD5, [169](#)
- MIDI, [169](#)
- Name mangling, [115](#)
- objdump, [166](#)
- Oracle RDBMS, [2](#), [137](#), [169](#), [235](#), [243](#), [245](#), [249](#)
- Page (memory), [145](#)
- PDB, [168](#), [232](#)
- PDP-11, [53](#)
- puts() instead of printf(), [6](#), [20](#), [32](#)
- Raspberry Pi, [60](#)
- Register allocation, [147](#)
- Relocation, [7](#)
- row-major order, [79](#)
- RTTI, [130](#)
- SAP, [168](#), [232](#)
- Signed numbers, [32](#), [162](#)
- stdcall, [163](#)
- this, [115](#)
- thiscall, [116](#), [164](#)
- ThumbTwoMode, [7](#)
- thunk-functions, [7](#)
- Unrolled loop, [48](#), [78](#)
- Windows
 - KERNEL32.DLL, [82](#)
 - MSVCR80.DLL, [134](#)
 - ntoskrnl.exe, [235](#)
 - Structured Exception Handling, [11](#)
- x86
 - Instructions
 - ADD, [2](#), [12](#), [26](#)
 - AND, [3](#), [82](#), [86](#), [88](#), [110](#)
 - BSF, [146](#)
 - CALL, [2](#), [8](#)
 - CMOVcc, [34](#)
 - CMP, [23](#), [24](#)
 - CMPSB, [169](#)
 - CPUID, [107](#)
 - DEC, [51](#)
 - DIVSD, [171](#)
 - FADDP, [58](#), [59](#)
 - FCOM, [64](#), [65](#)
 - FCOMP, [63](#)
 - FDIV, [58](#), [170](#)
 - FDIVP, [58](#)
 - FDIVR, [59](#)
 - FLD, [61](#), [63](#)
 - FMUL, [58](#)
 - FNSTSW, [63](#), [66](#)
 - FSTP, [61](#)
 - FUCOM, [65](#)
 - FUCOMPP, [65](#)
 - IMUL, [26](#)
 - INC, [51](#)
 - JA, [32](#), [162](#)
 - JAE, [32](#)
 - JB, [32](#), [162](#)
 - JBE, [32](#)
 - JE, [36](#)
 - JG, [32](#), [162](#)
 - JGE, [32](#)
 - JL, [32](#), [162](#)
 - JLE, [32](#)
 - JMP, [8](#), [13](#)
 - JNBE, [66](#)
 - JNE, [23](#), [24](#), [32](#)
 - JP, [63](#)
 - JZ, [24](#), [36](#), [249](#)
 - LEA, [19](#), [94](#), [100](#), [158](#)
 - LEAVE, [3](#)
 - LOOP, [44](#), [171](#)
 - MOV, [3](#)
 - MOVDQA, [140](#)
 - MOVDQU, [140](#)
 - MOVSD, [209](#)
 - MOVSX, [49](#), [52](#), [106](#)
 - MOVZX, [50](#), [98](#)
 - NOP, [94](#), [160](#)
 - NOT, [51](#), [53](#), [213](#)
 - OR, [86](#)
 - PADDD, [140](#)
 - PCMPEQB, [145](#)
 - PLMULHW, [137](#)
 - PLMULLD, [137](#)
 - PMOVMASKB, [145](#)
 - POP, [2](#), [8](#)
 - PUSH, [2](#), [3](#), [8](#), [19](#)
 - PXOR, [145](#)
 - RCL, [171](#)
 - RET, [3](#), [8](#), [77](#), [116](#)
 - SAHF, [65](#)
 - SETcc, [66](#)
 - SETNBE, [66](#)
 - SETNZ, [50](#)
 - SHL, [71](#), [88](#), [90](#)
 - SHR, [90](#), [110](#)
 - SUB, [3](#), [24](#), [36](#)
 - TEST, [49](#), [82](#), [85](#)
 - XOR, [3](#), [23](#), [51](#)
 - Registers

- Flags, [24](#)
- Parity flag, [63](#)
- EAX, [23](#), [28](#)
- EBP, [19](#), [25](#)
- ECX, [115](#)
- ESP, [12](#), [19](#)
- JMP, [40](#)
- RIP, [166](#)
- ZF, [24](#), [82](#)
- 8086, [52](#), [87](#)
- 80386, [87](#)
- 80486, [57](#)
- AVX, [137](#)
- MMX, [137](#)
- SSE, [137](#)
- SSE2, [137](#)
- x86-64, [18](#), [147](#), [166](#)
- Xcode, [4](#)